

EVA-Bench: A New End-to-end Framework for Evaluating Voice Agents

Tara Bogazelli, Gabrielle Gauthier Melançon, Katrina Stankiewicz, Olurwanifemi Bamgbose, Fanny Riols, Hoang H. Nguyen, Raghuveer Meindiratta, Lindsay Devon Brin, Joseph Marinier, Hari Subramani, Anil Madamala, Sridhar Krishna Nemala, Srinivas Sunkara

ABSTRACT

Voice agents, artificial intelligence systems that conduct spoken conversations to complete tasks, are increasingly deployed across enterprise applications. However, no existing benchmark jointly addresses two core evaluation challenges: generating realistic simulated conversations, and measuring quality across the full scope of voice-specific failure modes. We present EVA-Bench, an end-to-end evaluation framework that addresses both. On the simulation side, EVA-Bench orchestrates bot-to-bot audio conversations over dynamic multi-turn dialogues, with automatic simulation validation that detects user simulator error and appropriately regenerates conversations before scoring. On the measurement side, EVA-Bench introduces two composite metrics: EVA-A (Accuracy), capturing task completion, faithfulness, and audio-level speech fidelity; and EVA-X (Experience), capturing conversation progression, spoken conciseness, and turn-taking timing. Both metrics apply to different agent architectures, enabling direct cross-architecture comparison. EVA-Bench includes 213 scenarios across three enterprise domains, a controlled perturbation suite for accent and noise robustness, and pass@1, pass@k, pass^k measurements that distinguish peak from reliable capability. Across 12 systems spanning all three architectures, we find: (1) no system simultaneously exceeds 0.5 on both EVA-A pass@1 and EVA-X pass@1; (2) peak and reliable performance diverge substantially (median pass@k - pass^k gap of 0.44 on EVA-A); and (3) accent and noise perturbations expose substantial robustness gaps, with effects varying across architectures, systems, and metrics (mean up to 0.314). We release the full framework, evaluation suite, and benchmark data under an open-source license.

About this document

This is a Reviewed Preprint: a third-party paper that llmXive has *auto-reviewed* (with its LLM review panel) but *never modified*. The original manuscript follows this cover page exactly as published by its authors — llmXive re-typesets nothing and claims no authorship. Provenance. Ingested by llmXive on 2026-07-01 — scraped from arXiv; via llmXive dashboard, GitHub issue #160; submitted by github-actions[bot].

Source. <https://arxiv.org/abs/2605.13841>

About llmXive. llmXive is an automated platform for scientific discovery in which specialist LLM agents advance ideas from a one-paragraph brainstorm to a peer-reviewed paper. Learn more at <https://github.com/ContextLab/llmXive>.

Automated review. See the accompanying [peer-review-llmxive.pdf](#) for the full reviewer report.

llmXive follow-up. A separate llmXive study building on this work: [PROJ-824-llmxive-follow-up-extending-eva-bench-a](#).

EVA-Bench: A New End-to-end Framework for Evaluating Voice Agents

Tara Bogavelli, Gabrielle Gauthier Melançon, Katrina Stankiewicz, Oluwanifemi Bamgbose, Fanny Riols, Hoang H. Nguyen, Raghav Mehndiratta, Lindsay Devon Brin, Hari Subramani, Joseph Marinier, Anil Madamala, Sridhar Krishna Nemala, Srinivas Sunkara

ServiceNow

Voice agents, artificial intelligence systems that conduct spoken conversations to complete tasks, are increasingly deployed across enterprise applications. However, no existing benchmark jointly addresses two core evaluation challenges: generating realistic simulated conversations, and measuring quality across the full scope of voice-specific failure modes. We present **EVA-Bench**, an end-to-end evaluation framework that addresses both. On the simulation side, EVA-Bench orchestrates bot-to-bot audio conversations over dynamic multi-turn dialogues, with automatic simulation validation that detects user simulator error and appropriately regenerates conversations before scoring. On the measurement side, EVA-Bench introduces two composite metrics: **EVA-A** (Accuracy), capturing task completion, faithfulness, and audio-level speech fidelity; and **EVA-X** (Experience), capturing conversation progression, spoken conciseness, and turn-taking timing. Both metrics apply to all major agent architectures, enabling direct cross-architecture comparison. EVA-Bench includes 213 scenarios across three enterprise domains, a controlled perturbation suite for accent and noise robustness, and pass@1, pass@k, pass[^]k measurements that distinguish peak from reliable capability. Across 12 systems spanning all three architectures, we find: (1) no system simultaneously exceeds 0.5 on both **EVA-A_{pass@1}** and **EVA-X_{pass@1}**; (2) peak and reliable performance diverge substantially (median pass@k–pass[^]k gap of 0.44 on EVA-A); and (3) accent and noise perturbations expose substantial robustness gaps, with effects varying across architectures, systems, and metrics (mean Δ up to 0.314). We release the full framework, evaluation suite, and benchmark data under an open-source license.

🌐 : <https://servicenow.github.io/eva>

🔗 : <https://github.com/ServiceNow/eva>

🗣️ : <https://huggingface.co/datasets/ServiceNow-AI/eva>

Correspondence: tara.bogavelli@servicenow.com

servicenow

1 Introduction

Voice agents are Artificial Intelligence (AI) systems designed to carry out tasks through spoken conversations, and their deployment across a wide range of applications is rapidly growing [23]. Voice agents operate under constraints that are fundamentally distinct from text: speech is ephemeral and linear, real-time timing shapes the naturalness of interaction, and acoustic conditions vary widely across callers. These properties give rise to failure modes with no direct text analog [29, 5], and render evaluation frameworks designed for text-based agents [34, 28, 25] insufficient for assessing voice agent quality. Rigorous evaluation of voice agents must therefore address two distinct challenges: how conversations are simulated, and how quality is measured.

The simulation challenge concerns constructing interactions that are valid proxies for real deployment conditions. This requires complete multi-turn interactions as only full conversations expose how an agent recovers from misunderstandings and resolves tasks end-to-end. Conversations must reflect the task-oriented nature of real voice agent deployments, user behavior must reflect natural human spoken dialogue, and acoustic conditions must reflect real-world environments, including variation in accents and background noise. Simulated users must themselves be validated: a simulator that drifts from its scenario, behaves unrealistically, or varies across runs undermines any downstream evaluation, conflating agent variance with simulator variance.

The measurement challenge concerns capturing the full scope of voice agent quality once valid simulations are

Table 1 Feature comparison of contemporary voice agent evaluation frameworks. $\sim$ denotes partial support, and — for the simulator validation means it doesn’t apply because of missing simulator. EVA-Bench is the only framework combining live multi-turn simulation across both speech-to-speech (S2S) and cascade architectures with a realistic audio environment, automated simulator validation, comprehensive metrics exposing a wide range of voice agent failures, and pass@k measurement via multi-trial consistency measurement.

Framework	Interaction	Supported Architectures	Multi-turn	Tool Use	Realistic Audio	Simulator Validation	Comprehensive Metrics	Multi Trial
EVA-Bench	Live bot-to-bot	S2S, Cascade	✓	✓	✓	✓	✓	✓
τ -Voice [29]	Live bot-to-bot	S2S	✓	✓	✓	✗	✗	✗
FDB-v3 [19]	Real human audio	S2S, Cascade	✗	✓	$\sim$	—	✗	✗
VoiceAgentBench [15]	Static	S2S, Cascade	✓	✓	✗	—	✗	✗
CAVA [10]	Partial simulation	S2S, Cascade	✓	✓	✗	✗	✗	✗
FDB-v2 [17]	Live, auto examiner	S2S	✓	✗	✗	✗	✗	✗
FD-Bench [18]	Live, simulated	S2S	✗	✗	✗	✗	✗	✗

in place. Task completion and turn-taking dynamics, while necessary, leave critical failure modes undetected [4, 22, 1]. On the accuracy side, an agent may call the correct tools yet violate system policy or produce spoken outputs containing incorrect entities (e.g. wrong confirmation codes, or monetary amounts) that are catastrophic in production yet undetectable from transcript-level evaluation alone. On the user experience side, an agent may achieve low response latency yet fail to make meaningful progress across turns, repeat prior questions, or present users with an excessive number of spoken options that would overwhelm a user’s working memory. Addressing the measurement challenge requires evaluation across a broader set of dimensions than existing benchmarks provide. Additionally, voice agents are not architecturally uniform: cascade systems chain separate speech-to-text (STT), large language model (LLM), and text-to-speech (TTS) components, while audio-native systems process audio inputs directly, either end-to-end via speech-to-speech (S2S) models, or via hybrid systems that pair a large audio language model (LALM) with a TTS model (full definitions in Appendix A). These architectures have fundamentally different mechanisms, yet must be evaluated on equal footing for benchmarks to meaningfully compare them.

We present EVA-Bench, a benchmark designed to solve both of these challenges. On the simulation side, EVA-Bench conducts fully automated bot-to-bot audio simulation over dynamic multi-turn dialogues, with validation-gated quality control ensuring consistency across repeated trials. On the measurement side, EVA-Bench introduces two composite scores: EVA-A (Accuracy) and EVA-X (Experience). EVA-A captures task completion, faithfulness to policy and tool outputs, and audio-level entity fidelity. EVA-X captures conversation progression, conciseness for spoken delivery, and turn-taking timing. Across 12 evaluated systems, EVA-Bench reveals that accuracy and experience remain jointly unsatisfied, that peak capability consistently overstates reliable performance, and that robustness to acoustic perturbations varies substantially across systems and metrics. Our contributions:

- We introduce EVA-Bench: an end-to-end evaluation framework for voice agents that generates realistic bot-to-bot conversations through a user simulator with validation-gated quality control, and supports controlled acoustic perturbations across independent trials.
- We define EVA-A and EVA-X, joint accuracy and experience metrics that surface failure modes invisible to existing benchmarks and enable direct comparison between audio-native and cascade voice agents.
- We create three enterprise benchmark datasets with a total of 213 scenarios focused on surfacing voice-specific failure modes.
- We show empirical findings on cascade vs. audio-native tradeoffs, perturbation sensitivity, and behavioral consistency across trials.




Figure 1 EVA-Bench framework overview. The simulation orchestrates parallel per-scenario bot-to-bot audio sessions over WebSocket in which the User Simulator — configured with a scenario-specific goal, persona, and conversational TTS voice — interacts with the Voice Agent under test. The Tool Executor handles all agent tool calls deterministically. Completed conversations pass through **Simulator Validation** that trigger automatic regeneration on failure before entering the **Quality Measurements** phase, which produces **EVA-A** and **EVA-X** pass@1, pass@k, and pass^k scores in addition to **Diagnostic** metrics.

2 Related Work

Many existing voice benchmarks focus on individual components such as STT robustness [33, 2, 6], TTS quality [13, 21, 14], or conversational dynamics [26, 3], rather than end-to-end voice agent behavior. Table 1 provides a feature-level comparison of voice agent benchmarks; we summarize the key gaps below.

Conversation Simulation. Effective voice agent evaluation requires live, task-oriented multi-turn conversations, yet most benchmarks fall short: FDB-v1 and FDB-v1.5 [18, 20] are heavily scripted without tool use, VoiceAgentBench [15] uses static TTS-synthesized queries, and FDB-v3 [19] remains single-turn despite improved realism. The closest prior work, τ -Voice [29] and FDB-v2 [17], conduct live bot-to-bot multi-turn simulation but lack automated validation of simulator behavior across trials, leaving open whether scores reflect agent quality or simulator variance. Additionally, τ -Voice couples accent variation with persona changes, confounding acoustic and behavioral effects. EVA-Bench addresses these gaps with multi-trial bot-to-bot simulation, automatic simulator validation, and a controlled perturbation suite that isolates each factor independently.

Voice Agent Quality Measurement. Existing benchmarks converge on a narrow set of metrics. VoiceAgentBench [15] assesses tool use accuracy but no conversational quality; τ -Voice [29] adds turn-taking measures but not faithfulness; FDB-v3 [19] introduces transcript-level response quality but omits policy faithfulness and audio-level entity accuracy. No existing framework measures whether the agent makes efficient conversational progress, avoids excessive cognitive load, or speaks correct information — failure modes most consequential for enterprise deployment.

3 Methodology

3.1 Conversation Simulation

Data Design. Constructing a benchmark dataset well-suited to voice agent evaluation requires careful attention to both domain relevance and scenario specificity. EVA-Bench comprises three domains reflecting real-world enterprise voice agent deployments: Airline Customer Service Management (CSM), Healthcare Human Resources Service Delivery (HRSD), and Enterprise Information Technology Service Management (ITSM). Scenarios within each domain are designed to reflect the task-oriented nature of real voice agent interactions — focusing on high-contact cases where users are most likely to call an agent, such as flight rebooking rather than initial booking. Each scenario consists of a **user goal** specifying the user’s intended outcome with explicit constraints (e.g., departure before 10pm, fare below a specified amount), a **user persona**

defining speaking style, patience, and personality, a **scenario database** containing the data the agent’s tools query and modify, and **ground truth** specifying the expected database state after successful task resolution. User goals are accompanied by a decision tree that eliminates ambiguity about intended outcomes and user choices throughout the conversation, enabling repeatable evaluation. Scenarios are further designed to surface voice-specific failure modes by requiring agents to correctly handle key entities (e.g. confirmation codes, employee identifiers (IDs), names, and domain-specific identifiers) that are frequently misheard in spoken interactions. More details on data domains, scenario examples, and dataset construction and validation can be found in Appendix C.

Multi-Turn Conversations. EVA-Bench evaluates agents through fully automated bot-to-bot conversations. A user simulator, built on a high-quality cascade pipeline, receives the user goal, decision tree, and persona as input and communicates with the agent over a live audio WebSocket (Figure 1). Both sides of the interaction operate over audio, enabling evaluation of cascade and audio-native architectures under identical conditions. See Appendix D for full simulator details.

Controlled Perturbations. EVA-Bench introduces a perturbation suite that varies user acoustic and behavioral conditions independently. Acoustic perturbations include accent variations, background noises, and connection degradation. Behavioral perturbations model caller variation in personality and speaking style. Each perturbation axis is independently controlled, enabling conditions to be applied in isolation or combination to disentangle each factor’s effect on performance. See Appendix G for more details.

Simulation Validation. Before any evaluation metrics are computed, each simulated conversation passes through automated validation checks. User Behavioral Fidelity (LLM-as-Judge [35]) checks whether the user simulator faithfully executed its assigned goal without deviations that would corrupt agent evaluation. The judge prompt contains specific corruption types to check for. User Speech Fidelity uses an LLM-as-Judge to verify that the simulator’s spoken audio accurately conveyed its intended content, using a nearly identical prompt to the Speech Fidelity judge explained in 3.2.1. Conversations failing any check are automatically regenerated, ensuring that evaluation scores reflect agent behavior rather than simulator artifacts. Across four systems evaluated on all domains, 12.0% of trials required regeneration due to user simulator error (almost exclusively due to user behavioral drift), with speech fidelity accounting for less than 4% of reruns. Full validation details, including judge selection methodology and per-check rerun breakdowns, are provided in Appendix D.

3.2 Voice Agent Quality Measurement

EVA-Bench evaluates each conversation across three layered metric categories: **Accuracy (EVA-A)**, **Experience (EVA-X)**, and **Diagnostic Metrics**. These are described in the following subsections, and a table summarizing all metrics is provided in Appendix E. Note that for certain metrics, separate implementations are created for audio-native and cascade systems, since the two pipelines differ in which intermediate signals we can observe. See details in Appendices E.1 and E.2. Judge development followed a rigorous multi-stage development process described in Appendix E.3.

3.2.1 EVA-A: Accuracy Metrics

Task completion alone is a necessary but insufficient measure of accuracy. An agent can reach the correct end state while fabricating a policy detail, misreading a confirmation code aloud, or proceeding without required confirmations. Below are the metrics we propose to measure Accuracy.

Task Completion. A deterministic binary metric comparing the SHA-256 hash of the scenario database’s final state against the ground-truth state. A score of 1 indicates the agent made exactly the correct tool calls with correct parameters; 0 indicates any deviation, i.e. wrong, missing, or extra changes. Because the user simulator produces repeatable outcomes, failures are unambiguously attributable to agent error.

Faithfulness. An LLM-as-Judge metric evaluating whether the agent actions remain grounded in the instructions, policies, tool results, and user inputs. This complements task completion: high task completion with low faithfulness indicates the task was completed but with material errors along the way (e.g., misrepresenting fees). Notably, faithfulness prompt differs by architecture: cascade systems are evaluated relative to what

the STT layer delivered, while audio-native systems treat mishearing as a faithfulness violation, since audio understanding is the model’s own responsibility.

Speech Fidelity. A LALM-as-Judge metric evaluating whether the agent’s spoken audio accurately reproduces the intended text, with particular attention to high-stakes named entities (e.g. confirmation codes, dates, dollar amounts). For S2S systems where no intended text exists, the metric instead verifies that key entities from user turns and tool responses are correctly spoken. To our knowledge, this is the only metric in any end-to-end voice agent benchmark that evaluates the quality of the agent’s spoken output at the audio level.

3.2.2 EVA-X: Experience Metrics

The quality of a conversational experience with a voice agent is shaped by several key factors: whether the conversation moves purposefully toward resolution, whether responses are concise enough to follow without replay, and whether the timing of the agent’s replies feels natural.

Conversation Progression. An LLM-as-Judge metric that evaluates whether the agent efficiently moves the conversation forward by avoiding repetition, retaining context across turns, and driving toward task resolution without stalling or backtracking.

Conciseness. An LLM-as-Judge metric that evaluates whether the agent’s responses are appropriately brief for spoken delivery. Phone callers cannot skim or re-read long responses; verbose agents fail users when they impose cognitive overload by providing too many details or questions.

Turn-Taking. A timestamp-based metric measuring whether the agent spoke at the right time, neither interrupting the user nor introducing excessive silence. Each turn is routed to a semantically appropriate scoring function: agent-interrupted turns are scored on overlap duration, barge-in count, and post-interrupt recovery latency; user-interruption turns on agent yield latency; and uninterrupted turns on a piecewise-linear latency curve. Turns involving tool calls receive a more lenient latency threshold, reflecting a longer expected duration than a purely conversational turn. This metric also takes into account when an agent fails to respond to a user turn (Conversation Completion).

3.2.3 Diagnostic Metrics

Diagnostic metrics are not included in EVA-A or EVA-X scores. Their purpose is to make main metric failures actionable by providing more granular information on key failure areas. For example, Transcription Accuracy (Key Entities) is an LLM-as-Judge diagnostic metric that identifies domain-specific key entities in user speech (confirmation codes, names, dates, IDs) and verifies whether each was correctly transcribed in cascade systems using semantic rather than exact match. This surfaces failures that word error rate (WER) misses entirely: a confirmation code off by one character scores near-perfect on WER but is functionally unusable. Additional diagnostic metrics cover authentication outcomes, response latency, etc. (complete list provided in Appendix E.6).

3.2.4 Aggregate Metrics

Metrics for each dimension are aggregated into per-dimension scores (EVA-A, EVA-X), designed to capture both average and consistent performance. Measuring consistency requires a binary notion of success per conversation, so that we can assess how often a system succeeds across repeated trials of the same scenario. Simple averaging is problematic: it can mask a major failure on one metric with a high score on another, and the metrics are not on comparable scales (continuous, three-point, and binary), so the same numerical score carries different meaning across metrics.

We therefore define a *pass threshold* τ_m for each metric m , calibrated to the point at which performance is acceptable given the metric scale and implementation (Appendix E). A conversation passes on a dimension if *every* metric meets its threshold. Concretely, a conversation passes on accuracy if (task completion = 1.0) $\wedge$ (faithfulness ≥ 0.5) $\wedge$ (speech fidelity ≥ 0.95) and passes on experience if (turn-taking ≥ 0.8) $\wedge$ (conversation progression ≥ 0.5) $\wedge$ (conciseness ≥ 0.5).

This binary pass/fail gives us three aggregate statistics, each reported as **EVA-A** and **EVA-X** variants. **pass@1** is the fraction of $T = Nk$ trials (N scenarios, k trials each) that pass, measuring average performance. **pass@k**

is the fraction of scenarios where at least one of k trials passes, measuring ceiling performance. $\widehat{\text{pass}}^k$ measures reliability, by raising each scenario’s pass rate $\hat{p}_i$ to the k -th power and averaging across all N scenarios. This captures the probability that the system passes all k independent trials in a given scenario. The difference between $\text{pass}@k$ and $\widehat{\text{pass}}^k$ quantifies the gap between ceiling (peak) and consistent (reliable) performance. Formal definitions are provided in Appendix A.

4 Empirical Analysis

4.1 Experiment Setup

We evaluate 12 systems in total: seven cascade, two hybrid, and three S2S. Configuration details are provided in Appendix B. Under the clean (unperturbed) condition, systems are evaluated on all 213 scenarios with $k = 5$ trials per scenario. We additionally evaluate under three perturbation conditions: French-accented user speech, coffee shop background noise, and both combined. To maintain feasibility across 12 systems and 3 conditions, perturbed evaluations use a randomly sampled subset of 90 scenarios (30 per domain) with $k = 3$ trials per scenario; the same subset is used across all systems. GPT-Realtime and Gemini Live models are evaluated using their native SDK [9, 24], Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha is evaluated using ElevenAgents [8], and the remaining systems are evaluated using Pipecat [7]. All systems are evaluated using the respective framework’s default turn detection configuration.

4.2 Evaluation Reliability

EVA-Bench scores reflect genuine differences in agent behavior rather than evaluation artifacts: judge stochasticity contributes minimally to observed variance, and score differences between systems exceed measurement noise. We substantiate this with two sets of analyses: human-judge agreement to establish judge validity, and variance decomposition to complement bootstrap confidence intervals (CIs) in characterizing and bounding relative measurement noise.

Human-Judge Agreement We measure inter-annotator agreement (IAA) between the judge and/or human annotators using quadratic-weighted Cohen’s κ . Our reference baseline is the IAA of two independent annotations from linguist labellers (IAA-L). For judge inter-annotator agreement (IAA-J), each human annotation is paired with the judge score and pooled, with 95% bootstrap confidence intervals (10,000 resamples) computed at the conversation level. **The judge meets the human IAA on all four metrics**, with IAA-J κ ranging from 0.777 to 0.845 across metrics. For every metric, Spearman’s ρ and κ (both computed on the same judge–human ratings) differ by at most 0.008, supporting the absence of systematic calibration bias. Full scores, details, and confidence intervals are shown in Table 14.

Variance Decomposition Observed metric scores for a given model within a domain reflect variance from three sources: scenario difficulty, trial stochasticity (conversation trajectories), and LLM judge stochasticity. We characterize the contribution of each source on a subset of systems through complementary analyses, demonstrating that our main findings reflect genuine differences in model behavior rather than measurement noise. Trial stochasticity was the dominant source of variance across all metrics and models, consistently exceeding scenario-level variance; judge stochasticity was minimal by comparison (permutation test, $p < 0.0001$ for all 16 model $\times$ metric combinations). Full variance decomposition results are reported in Appendix H.

4.3 Main Findings

Accuracy–Experience Frontier. No evaluated system clears 0.5 on both EVA-A $\text{pass}@1$ and EVA-X $\text{pass}@1$ jointly, and only GPT-Realtime-1.5 (0.47, 0.57) clears 0.4 on both. This sparsity is reflected in the $\text{pass}@1$ Pareto frontier (Figure 2), which contains four systems spanning two disjoint regions.

Within EVA-X, the gap is almost entirely driven by Turn-Taking, where mean scores stratify cleanly by architecture (cascade: 0.28–0.58; S2S: 0.82–0.83) (Table 2). Conciseness and Conversation Progression show no comparable separation. The two hybrid systems both fall within the cascade EVA-X range (0.000, 0.029), suggesting that hybrid systems may not inherit the latency advantages of fully speech-to-speech architectures, though more systems evaluation would be needed to confirm this.




Figure 2 Accuracy vs Experience. Accuracy and Experience scores for pass@1, $\pm$ 95% CIs across domains. Four systems are on the Pareto frontier, two S2S and two cascade: Gemini-3.1-Flash-Live, GPT-Realtime-1.5, Scribe + Gemini-3-Flash + Conversational v3, and Nova + GPT-5.4 + Sonic (left to right).

Among the evaluated cascade systems, we observe a consistent accuracy–experience trade-off, pointing to a potential underlying capability–latency tension. The three cascade systems that perform best on accuracy (Nova + GPT-5.4 + Sonic, Scribe + Gemini-3-Flash + Conversational v3, Parakeet + Gemma-4-31B + Kokoro) struggle on experience, with mean latencies on tool-call turns above 5 s. In contrast, the two cascade systems with better experience (Whisper + Qwen3.5-27B + Voxtral, Cohere + Gemma-4-26B + Voxtral) achieve tool-call turn latencies below 2.7 s but also lower accuracy. No cascade system exceeds 0.25 on both dimensions, with no overlapping CIs. Full latency breakdowns appear in Table 25.

Consistency Analysis. Across all 12 evaluated systems, peak performance (pass@k) substantially exceeds reliable performance (pass^k) on both axes: the median gap is 0.44 on EVA-A and 0.24 on EVA-X. Under the pass^k interpretation (the probability of passing all five trials on a given scenario) even the strongest systems fall well below peak, suggesting that single-trial scores systematically overstate deployment-grade reliability.

Robustness Analysis. We evaluate each system under three perturbation conditions – **accented speech**, **background noise**, and **both combined** – measuring performance degradation against a clean baseline across the 90 subsampled scenarios (see subsection 4.1) via paired sign-flip permutation tests ($n=90$, Holm–Bonferroni corrected). The two architectures diverge in their failure modes: cascade systems are most vulnerable on accuracy metrics, while S2S systems suffer most on experience metrics. **Accented speech** drives the largest accuracy failures: cascade task completion drops by a mean 10 points (worst system: 17 points), while no S2S model shows significant accuracy degradation (0/27 model-metric pairs). **Background noise** exposes S2S experience failures (EVA-X $\bar{\Delta}=-0.16$), though cascade accuracy also suffers under noise (task completion $\bar{\Delta}=-0.10$ vs. S2S $\bar{\Delta}=-0.04$). The **combined** condition reveals the full spread: cascade task completion drops a mean 19 points (worst systems: 31 points), while the S2S model mean remains within 5 points. Within the cascade class, however, robustness varies considerably, with as low as 11% and as high as 87% of per-system metric-perturbation combinations showing significant degradation. The two most robust cascade systems degrade primarily on experience metrics, more closely resembling S2S systems than their cascade peers. Turn-taking is the most perturbation-sensitive metric overall, with 81% of measurements across perturbations and systems showing significant degradation. See Figure 3 and Appendix G.

Table 2 Accuracy and Experience metrics for all evaluated systems under clean-audio conditions, pooled equal-weighted across the three EVA-Bench domains. Each cell shows the pooled point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale so they can be visually compared; each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-A			Task Completion	Faithfulness	Speech Fidelity
		pass@1	pass@k	pass~k	Mean	Mean	Mean
Cascade	Cohere + Gemma-4-26B + Voxtral	0.207 $\pm$ 0.041	0.416 $\pm$ 0.070	0.060 $\pm$ 0.028	0.338 $\pm$ 0.049	0.375 $\pm$ 0.036	0.983 $\pm$ 0.003
	Scribe + Gemini-3-Flash + Conversational v3	0.490 $\pm$ 0.052	0.730 $\pm$ 0.058	0.269 $\pm$ 0.055	0.736 $\pm$ 0.043	0.457 $\pm$ 0.042	0.977 $\pm$ 0.006
	Ink-whisper + Haiku-4.5 + Sonic 3	0.234 $\pm$ 0.041	0.516 $\pm$ 0.069	0.057 $\pm$ 0.028	0.374 $\pm$ 0.044	0.518 $\pm$ 0.033	0.983 $\pm$ 0.003
	Nova-3 + GPT-5.4 + Sonic 3	0.504 $\pm$ 0.044	0.809 $\pm$ 0.048	0.217 $\pm$ 0.048	0.609 $\pm$ 0.043	0.754 $\pm$ 0.027	0.989 $\pm$ 0.003
	Nova-3 + GPT-5.4-mini + Aura-2	0.210 $\pm$ 0.045	0.448 $\pm$ 0.069	0.062 $\pm$ 0.032	0.465 $\pm$ 0.050	0.270 $\pm$ 0.033	0.974 $\pm$ 0.005
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.403 $\pm$ 0.045	0.748 $\pm$ 0.055	0.169 $\pm$ 0.046	0.637 $\pm$ 0.051	0.466 $\pm$ 0.035	0.954 $\pm$ 0.009
Hybrid	Whisper + Qwen3.5-27B + Voxtral	0.205 $\pm$ 0.033	0.518 $\pm$ 0.066	0.033 $\pm$ 0.019	0.417 $\pm$ 0.051	0.546 $\pm$ 0.033	0.913 $\pm$ 0.010
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.431 $\pm$ 0.047	0.812 $\pm$ 0.055	0.158 $\pm$ 0.043	0.674 $\pm$ 0.041	0.443 $\pm$ 0.036	0.969 $\pm$ 0.006
S2S	Ultravox-Realtime	0.270 $\pm$ 0.047	0.503 $\pm$ 0.072	0.108 $\pm$ 0.037	0.473 $\pm$ 0.055	0.292 $\pm$ 0.035	0.971 $\pm$ 0.007
	Gemini-3.1-Flash-Live	0.292 $\pm$ 0.048	0.552 $\pm$ 0.069	0.132 $\pm$ 0.043	0.473 $\pm$ 0.052	0.238 $\pm$ 0.035	0.995 $\pm$ 0.003
	GPT-Realtime-1.5	0.467 $\pm$ 0.052	0.710 $\pm$ 0.061	0.283 $\pm$ 0.056	0.739 $\pm$ 0.046	0.360 $\pm$ 0.041	0.996 $\pm$ 0.002
	GPT-Realtime-mini	0.163 $\pm$ 0.041	0.318 $\pm$ 0.063	0.059 $\pm$ 0.030	0.345 $\pm$ 0.054	0.125 $\pm$ 0.031	0.977 $\pm$ 0.012
Arch.	System	EVA-X			Turn-Taking	Conciseness	Conv. Progression
		pass@1	pass@k	pass~k	Mean	Mean	Mean
Cascade	Cohere + Gemma-4-26B + Voxtral	0.209 $\pm$ 0.027	0.647 $\pm$ 0.069	0.015 $\pm$ 0.011	0.567 $\pm$ 0.024	0.809 $\pm$ 0.007	0.598 $\pm$ 0.032
	Scribe + Gemini-3-Flash + Conversational v3	0.024 $\pm$ 0.018	0.061 $\pm$ 0.035	0.004 $\pm$ 0.006	0.451 $\pm$ 0.019	0.774 $\pm$ 0.007	0.804 $\pm$ 0.023
	Ink-whisper + Haiku-4.5 + Sonic 3	0.009 $\pm$ 0.006	0.042 $\pm$ 0.031	0.000 $\pm$ 0.000	0.312 $\pm$ 0.020	0.784 $\pm$ 0.007	0.710 $\pm$ 0.023
	Nova-3 + GPT-5.4 + Sonic 3	0.007 $\pm$ 0.006	0.031 $\pm$ 0.024	0.000 $\pm$ 0.000	0.283 $\pm$ 0.019	0.835 $\pm$ 0.007	0.737 $\pm$ 0.020
	Nova-3 + GPT-5.4-mini + Aura-2	0.113 $\pm$ 0.023	0.416 $\pm$ 0.070	0.005 $\pm$ 0.004	0.583 $\pm$ 0.019	0.835 $\pm$ 0.008	0.428 $\pm$ 0.025
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.010 $\pm$ 0.009	0.035 $\pm$ 0.032	0.000 $\pm$ 0.000	0.308 $\pm$ 0.015	0.829 $\pm$ 0.007	0.774 $\pm$ 0.024
Hybrid	Whisper + Qwen3.5-27B + Voxtral	0.273 $\pm$ 0.034	0.684 $\pm$ 0.065	0.051 $\pm$ 0.021	0.561 $\pm$ 0.029	0.685 $\pm$ 0.010	0.612 $\pm$ 0.026
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.019 $\pm$ 0.003	0.801 $\pm$ 0.007	0.618 $\pm$ 0.029
S2S	Ultravox-Realtime	0.029 $\pm$ 0.020	0.081 $\pm$ 0.039	0.006 $\pm$ 0.007	0.417 $\pm$ 0.020	0.750 $\pm$ 0.010	0.429 $\pm$ 0.030
	Gemini-3.1-Flash-Live	0.589 $\pm$ 0.035	0.979 $\pm$ 0.021	0.240 $\pm$ 0.045	0.830 $\pm$ 0.017	0.801 $\pm$ 0.009	0.636 $\pm$ 0.029
	GPT-Realtime-1.5	0.566 $\pm$ 0.039	0.939 $\pm$ 0.034	0.216 $\pm$ 0.040	0.815 $\pm$ 0.013	0.801 $\pm$ 0.008	0.679 $\pm$ 0.024
	GPT-Realtime-mini	0.406 $\pm$ 0.036	0.893 $\pm$ 0.043	0.099 $\pm$ 0.032	0.818 $\pm$ 0.015	0.722 $\pm$ 0.009	0.388 $\pm$ 0.033

4.4 Failure Mode Analysis

EVA-Bench’s metrics enable detailed error analysis. We report four targeted analyses below, with additional analyses in Appendix F.

Named entity transcription as an accuracy bottleneck. EVA-Bench’s diagnostic metrics reveal a consistent association between Transcription Accuracy (Key Entities) and Task Completion across all three domains, pointing to transcription as a candidate bottleneck for cascade system performance. Across all cascade systems, mean transcription accuracy on key entities is strongly correlated with mean task completion (Pearson $r = 0.93$, $p = 0.002$) and the relationship holds within each domain ($r = 0.88$ – 0.93 , all $p < 0.01$). Cascade systems with transcription accuracy below 70% on key entities show task completion rates 39% lower than systems above this threshold (0.37 vs 0.60). See Figure 5 and Figure 6 for details.

Evaluated S2S systems show contrasting patterns relative to other systems. The higher EVA-X scores for S2S systems reflect a clear turn-taking lead with on-time rate and conversation completion up +27.9 pp and +15.2 pp on average respectively, though the strongest cascade systems come within a few percentage points. The conversation completion gap is partly explained by many cascade systems completely failing to respond to short utterances and spelled content. S2S systems also show policy adherence issues, violating stated policy more often on average (+24.6 pp) and trailing the strongest cascade systems on this aspect. Confidence intervals, per-system breakdowns, and additional details are in Appendix F.3.

Faithfulness failures are not predicted by task completion. 72.2% of conversations ($n = 12,780$) with task completion = 1 exhibit at least one faithfulness deviation (faithfulness < 1.0), indicating that agents frequently make policy deviations or hallucinate details even when they call the correct tools. Conversely, 50.5% of faithfulness deviations co-occur with task completion = 0, suggesting that some task failures are downstream of faithfulness violations. This decoupling motivates including faithfulness as an independent




Figure 3 Perturbation effect on **Turn-Taking** pooled across domains. Bars show the mean delta from clean trials; 95% bootstrap CIs on the per-scenario delta. Bar colors encode the perturbation condition: ■ accent, ■ background noise, ■ accent + background noise. Asterisks mark cells significant after Holm-Bonferroni correction (* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$). Models listed in Appendix G.

metric. See Appendix F.2 for full confusion matrix.

Speech fidelity failures concentrate on alphanumeric entities. Entity errors — letter substitutions, digit omissions, and phonetic confusions between characters — are the dominant failure mode across all evaluated models, motivating speech fidelity as an audio-level metric: a caller who receives a misarticulated confirmation code cannot detect the error from context alone. See Appendix F.4 for examples.

5 Conclusion

We presented EVA-Bench, an end-to-end evaluation framework for voice agents that jointly addresses simulation fidelity and measurement comprehensiveness. Together, validation-gated bot-to-bot simulation, architecture-agnostic composite metrics (EVA-A and EVA-X), and a multi-trial consistency framework (pass@1, pass@k, pass^k) enable comparison across cascade, hybrid, and S2S voice agents under identical conditions.

Our evaluation of 12 systems produces three central findings:

- (1) While the best-performing cascade and S2S systems achieve comparable accuracy, experience quality diverges sharply along architecture lines, with the S2S–cascade gap on EVA-X driven almost entirely by turn-taking.
- (2) Peak and reliable performance diverge substantially across all systems: the median pass@k–pass^k gap is 0.44 on EVA-A, indicating that single-trial evaluation scores systematically overstate deployment-grade quality regardless of architecture.
- (3) We observe that cascade and S2S systems degrade differently under acoustic perturbation.

These observations highlight dimensions of voice agent quality that existing benchmarks leave unmeasured: joint accuracy-experience measurement, peak-versus-reliable performance gaps, and perturbation-specific degradation patterns across voice agent architectures.

EVA-Bench is released as a fully open-source, extensible framework. The metric suite is modular, and pass criteria and perturbation conditions are configurable, allowing practitioners to adapt the benchmark to their deployment requirements. New domains can be added by configuring agents and scenarios, and as S2S models mature and organizations evaluate migrating from cascade architectures, EVA-Bench provides a controlled framework for that decision under matched conditions. We invite the community to extend the benchmark with additional domains, languages, and evaluation dimensions as the field evolves.

Limitations

While EVA-Bench provides a rigorous, extensible evaluation framework for conversational voice agents, several limitations should be acknowledged.

Metrics

- LLM-based judges may exhibit stylistic biases, systematically favoring outputs resembling their training distribution. This risk may be higher when evaluated systems share a model family with the judge.
- LALM-based judges remain an emerging paradigm and may exhibit lower reliability than text-only judges under equivalent conditions.
- The current scoring scheme applies binary task completion judgments with no partial credit for agents that satisfy all but one sub-goal, potentially suppressing fine-grained discriminability between systems differing primarily in terminal sub-goal robustness.

Framework

- EVA-Bench assumes that bot-to-bot simulation with a validated user simulator constitutes a valid proxy for real caller interactions. If the simulator systematically diverges from real callers in ways our behavioral fidelity validator does not detect scores may not transfer to production conditions.
- PCM-to- μ -law audio conversion introduces quality degradation, and bot-to-bot interface timing may not fully represent production deployments. Imprecise pipeline event timestamps (VAD events, log reconciliation across systems) may lead to inaccurate response speed values.
- Full reproduction requires access to commercial model APIs.
- Generating $k=5$ validated trials per scenario carries non-trivial cost; a full run across all 213 scenarios and 12 systems requires several hundred model calls per trial.
- All tools use a declarative mock executor rather than live APIs; while schemas reflect real-world APIs, mocks cannot capture failure modes, latency variance, partial responses, or schema drift.
- Production latency-reduction strategies (speculative execution, response pre-fetching, streaming-aware scheduling) are not used; reported latency metrics may not reflect achievable production lower bounds.

Simulation All scenarios are currently English-only. The user simulator is built on a high-quality STT-LLM-TTS pipeline (scribe-v2.2-realtime + GPT-5.1 + ElevenLabs v3 Conversational) which may therefore under-represent natural disfluencies. The simulator also does not systematically generate interruptions, meaning EVA-Bench under-stresses turn-taking robustness and cannot differentiate systems that handle barge-in gracefully. As the simulator relies on a commercial system, its behavior may change across versions; while we employ validators to detect off-policy behavior, perfect adherence cannot be guaranteed.

Experiments Our evaluation covers 12 systems across three architectural classes, but this sample size precludes broad inferential claims about model classes. All frameworks use default turn detection settings, and system prompts were not optimized — both deliberate choices to preserve fair baselines. Targeted tuning would likely yield higher scores. Robustness experiments include only one accent (French) and one noise environment (coffee shop); observed differences may partially reflect properties of a particular voice or sound rather than accent or noise variation broadly.

Our plans for future work and addressing certain limitations are outlined in [Appendix R](#).

References

- [1] Miguel E. Andres, Vadim Fedorov, Rida Sadek, Enric Spagnolo-Arrizabalaga, and Nadescha Trudel. Testing the Testers: Human-Driven Quality Assessment of Voice AI Testing Platforms. 2026. URL <https://arxiv.org/abs/2511.04133>.

- [2] Junyi Ao, Yuancheng Wang, Xiaohai Tian, Dekun Chen, Jun Zhang, Lu Lu, Yuxuan Wang, Haizhou Li, and Zhizheng Wu. SD-Eval: A Benchmark Dataset for Spoken Dialogue Understanding Beyond Words. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024. URL <https://openreview.net/forum?id=Pnjbvbb1Gv>.
- [3] Siddhant Arora, Zhiyun Lu, Chung-Cheng Chiu, Ruoming Pang, and Shinji Watanabe. Talking Turns: Benchmarking Audio Foundation Models on Turn-Taking Dynamics. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=2e4ECh0ikn>.
- [4] Hongliu Cao, Ilias Driouich, and Eoin Thomas. Beyond Task Completion: Revealing Corrupt Success in LLM Agents through Procedure-Aware Evaluation. 2026. URL <https://arxiv.org/abs/2603.03116>.
- [5] Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-Based Voice Assistants. *Transactions of the Association for Computational Linguistics*, 14:378–398, 04 2026. ISSN 2307-387X. doi: 10.1162/TACL.a.628. URL <https://doi.org/10.1162/TACL.a.628>.
- [6] Wenqian Cui, Xiaoqi Jiao, Ziqiao Meng, and Irwin King. VoxEval: Benchmarking the Knowledge Understanding Capabilities of End-to-End Spoken Language Models. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 16735–16753, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.818. URL <https://aclanthology.org/2025.acl-long.818/>.
- [7] Daily. Pipecat: Open Source Framework for Voice and Multimodal Conversational AI. <https://github.com/pipecat-ai/pipecat>, 2024. Accessed: 2026.
- [8] ElevenLabs. ElevenLabs: Documentation. <https://elevenlabs.io/docs/eleven-agents/overview>, 2023. Accessed: 2026.
- [9] Google DeepMind. Gemini Live API: Low-Latency Bidirectional Voice and Video Interactions. <https://ai.google.dev/gemini-api/docs/live>, 2024. Accessed: 2026.
- [10] Will Held, Michael J. Ryan, Aditya Shrivastava, Ali Sartaz Khan, Caleb Ziems, Ella Li, Martijn Bartelds, Michael Sun, Tan Li, Woody Gan, and Diyi Yang. CAVA: Comprehensive Assessment of Voice Assistants. <https://github.com/SALT-NLP/CAVA>, 2025. URL <https://talkarena.org/cava>. A benchmark for evaluating large audio models (LAMs) capabilities across six domains: turn taking, instruction following, function calling, tone awareness, safety, and latency.
- [11] Mattias Heldner. Detection thresholds for gaps, overlaps, and no-gap-no-overlaps. *The Journal of the Acoustical Society of America*, 130(1):508–513, 2011. doi: 10.1121/1.3598457.
- [12] Mattias Heldner and Jens Edlund. Pauses, gaps and overlaps in conversations. *Journal of Phonetics*, 38(4):555–568, 2010. ISSN 0095-4470. doi: <https://doi.org/10.1016/j.wocn.2010.08.002>. URL <https://www.sciencedirect.com/science/article/pii/S0095447010000628>.
- [13] Wen-Chin Huang, Erica Cooper, Yu Tsao, Hsin-Min Wang, Tomoki Toda, and Junichi Yamagishi. The VoiceMOS Challenge 2022. In *Interspeech 2022*, pp. 4536–4540, 2022. doi: 10.21437/Interspeech.2022-970.
- [14] Wen-Chin Huang, Erica Cooper, and Tomoki Toda. SHEET: A Multi-purpose Open-source Speech Human Evaluation Estimation Toolkit. In *Interspeech 2025*, pp. 2355–2359, 2025. doi: 10.21437/Interspeech.2025-1977.
- [15] Dhruv Jain, Harshit Shukla, Gautam Rajeev, Ashish Kulkarni, Chandra Khatri, and Shubham Agarwal. VoiceAgentBench: Are Voice Assistants ready for agentic tasks? *arXiv preprint arXiv:2510.07978*, 2026. doi: 10.48550/arXiv.2510.07978.
- [16] Stephen C. Levinson and Francisco Torreira. Timing in turn-taking and its implications for processing models of language. *Frontiers in Psychology*, 6:731, 2015. doi: 10.3389/fpsyg.2015.00731.
- [17] Guan-Ting Lin, Shih-Yun Shan Kuan, Jiatong Shi, Kai-Wei Chang, Siddhant Arora, Shinji Watanabe, and Hung-yi Lee. Full-Duplex-Bench-v2: A Multi-Turn Evaluation Framework for Duplex Dialogue Systems with an Automated Examiner. *arXiv preprint arXiv:2510.07838*, 2025.
- [18] Guan-Ting Lin, Jiachen Lian, Tingle Li, Qirui Wang, Gopala Anumanchipalli, Alexander H Liu, and Hung-yi Lee. Full-duplex-bench: A benchmark to evaluate full-duplex spoken dialogue models on turn-taking capabilities. *arXiv preprint arXiv:2503.04721*, 2025.

- [19] Guan-Ting Lin, Chen Chen, Zhehuai Chen, and Hung-yi Lee. Full-Duplex-Bench-v3: Benchmarking Tool Use for Full-Duplex Voice Agents Under Real-World Disfluency. *arXiv preprint arXiv:2604.04847*, 2026.
- [20] Guan-Ting Lin, Shih-Yun Shan Kuan, Qirui Wang, Jiachen Lian, Tingle Li, Shinji Watanabe, and Hung-yi Lee. Full-Duplex-Bench v1.5: Evaluating Overlap Handling for Full-Duplex Speech Models. In *ICASSP 2026-2026 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 19447–19451. IEEE, 2026. doi: 10.48550/arXiv.2507.23159.
- [21] Ruskin Raj Manku, Yuzhi Tang, Xingjian Shi, Mu Li, and Alexander Smola. EmergentTTS-Eval: Evaluating TTS Models on Complex Prosodic, Expressiveness, and Linguistic Challenges Using Model-as-a-Judge. In D. Belgrave, C. Zhang, H. Lin, R. Pascanu, P. Koniusz, M. Ghassemi, and N. Chen (eds.), *Advances in Neural Information Processing Systems*, volume 38. Curran Associates, Inc., 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/file/04970a25af46918606ba2cf0a3d7905d-Paper-Datasets_and_Benchmarks_Track.pdf.
- [22] Sushant Mehta. Beyond Accuracy: A Multi-Dimensional Framework for Evaluating Enterprise Agentic AI Systems. *arXiv preprint arXiv:2511.14136*, 2025.
- [23] Olivia Moore. AI Voice Agents: 2025 Update. <https://a16z.com/ai-voice-agents-2025-update/>, 2025. Andreessen Horowitz.
- [24] OpenAI. Realtime API Documentation. <https://platform.openai.com/docs/guides/realtime>, 2024. Accessed: 2026.
- [25] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The Berkeley Function Calling Leaderboard (BFCL): From Tool Use to Agentic Evaluation of Large Language Models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267 of *PMLR*, pp. 48371–48392, 2025.
- [26] Yizhou Peng, Yi-Wen Chao, Dianwen Ng, Yukun Ma, Chongjia Ni, Bin Ma, and Eng Siong Chng. FD-Bench: A Full-Duplex Benchmarking Pipeline Designed for Full Duplex Spoken Dialogue Systems. In *Proc. Interspeech 2025*, pp. 176–180, 2025.
- [27] Bidyapati Pradhan, Surajit Dasgupta, Amit Kumar Saha, Omkar Anustoop, Sriram Puttagunta, Vipul Mittal, and Gopal Sarda. SyGra: A Unified Graph-Based Framework for Scalable Generation, Quality Tagging, and Management of Synthetic Data. 2025. URL <https://arxiv.org/abs/2508.15432>.
- [28] Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, et al. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net, 2024. URL <https://openreview.net/forum?id=dHng200Jjr>.
- [29] Soham Ray, Keshav Dhandhanian, Victor Barres, and Karthik Narasimhan. τ -Voice: Benchmarking Full-Duplex Voice Agents on Real-World Domains, 2026. URL <https://arxiv.org/abs/2603.13686>.
- [30] Felicia Roberts, Piera Margutti, and Shoji Takano. Judgments Concerning the Valence of Inter-Turn Silence Across Speakers of American English, Italian, and Japanese. *Discourse Processes*, 48(5):331–354, 2011. doi: 10.1080/0163853X.2011.558002.
- [31] Gabriel Skantze. Turn-Taking in Conversational Systems and Human-Robot Interaction: A Review. *Computer Speech & Language*, 67:101178, 2021. URL <https://arxiv.org/abs/2010.03674>.
- [32] Tanya Stivers, Nicholas J Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann, Gertie Hoymann, Federico Rossano, Jan Peter De Ruiter, Kyung-Eun Yoon, et al. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National Academy of Sciences*, 106(26):10587–10592, 2009.
- [33] Shu-wen Yang, Po-Han Chi, Yung-Sung Chuang, Cheng-I Jeff Lai, Kushal Lakhotia, Yist Y Lin, Andy T Liu, Jiatong Shi, Xuankai Chang, Guan-Ting Lin, et al. SUPERB: Speech Processing Universal PERFORMANCE Benchmark. In *Proc. Interspeech 2021*, pp. 1194–1198, 2021.
- [34] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik R Narasimhan. τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=roNSXZpUDN>.
- [35] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In A. Oh, T. Naumann, A. Globerson, K. Saenko,

- M. Hardt, and S. Levine (eds.), *Advances in Neural Information Processing Systems*, volume 36, pp. 46595–46623. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/91f18a1287b398d378ef22505bf41832-Paper-Datasets_and_Benchmarks.pdf.
- [36] Xuhui Zhou, Weiwei Sun, Qianou Ma, Yiqing Xie, Jiarui Liu, Weihua Du, Sean Welleck, Yiming Yang, Graham Neubig, Sherry Tongshuang Wu, and Maarten Sap. Mind the Sim2Real Gap in User Simulation for Agentic Tasks. *arXiv preprint arXiv:2603.11245*, 2026.

A Definitions & Key Terms

STT	<i>Speech-to-Text</i> . A model or service that transcribes spoken audio into text. In voice agent pipelines, STT serves as the input stage, converting user utterances into a form consumable by a downstream language model.
TTS	<i>Text-to-Speech</i> . A model or service that synthesizes natural-sounding speech from text. In voice agent pipelines, TTS serves as the output stage, rendering language model responses as audio delivered to the user.
S2S	<i>Speech-to-Speech</i> . An end-to-end model architecture that accepts raw audio as input and produces audio as output, bypassing discrete STT and TTS stages entirely. S2S systems may also emit text output, used to invoke tool calls, alongside or prior to synthesizing a spoken response. S2S systems minimize modality-crossing latency and preserve paralinguistic signals throughout the pipeline.
LALM	A large language model that accepts audio as a direct input modality—without prior transcription—and produces text output. AudioLLMs internalize acoustic features (prosody, speaker characteristics, noise) that would otherwise be lost in an STT transcription step. Also referred to as an AudioLLM or a SpeechLLM .
Hybrid	A voice agent architecture that combines an AudioLLM (for audio-aware language understanding) with a discrete TTS module (for speech synthesis). Hybrid systems are Audio-Native on the input side while retaining a text-to-speech output stage, and occupy a middle ground between fully cascaded (STT+LLM+TTS) and fully end-to-end (S2S) pipelines.
Audio-Native	An umbrella term for voice agent architectures that process audio directly at one or more stages, rather than relying solely on text-based STT→LLM→TTS cascades. Audio-Native systems include both <i>S2S</i> and <i>Hybrid</i> configurations.
Multi-turn	A conversational interaction comprising more than one exchange between user and agent—i.e., at least one agent response followed by a subsequent user utterance within the same session. Multi-turn evaluation tests an agent’s ability to maintain task state, handle clarifications, and recover from errors across dialogue turns.
Bot-to-Bot	An automated evaluation protocol in which a <i>user simulator</i> (a separate agent or model) plays the role of the human caller, conducting full spoken conversations with the system under evaluation.
pass@1	The fraction of all trials that pass, measuring average performance: $\text{pass}@1 = \frac{1}{T} \sum_{t=1}^T \mathbf{1}[\text{pass}_t^{(d)}]$, where $T = Nk$ is the total number of trials across N scenarios and k trials each.
pass@k	The fraction of scenarios where at least one of k trials passes, measuring ceiling performance: $\text{pass}@k = \frac{1}{N} \sum_{i=1}^N \mathbf{1}\left[\sum_{j=1}^k \mathbf{1}[\text{pass}_{i,j}^{(d)}] \geq 1\right]$. See also: <i>Peak</i> .
pass[^]k	The expected probability that a system passes all k independent future trials on a given scenario, measuring reliable performance: $\text{pass}^k = \frac{1}{N} \sum_{i=1}^N \hat{p}_i^k$, where $\hat{p}_i = \frac{1}{k} \sum_{j=1}^k \mathbf{1}[\text{pass}_{i,j}^{(d)}]$ is the per-scenario pass rate across k trials. See also: <i>Reliable</i> .
Peak	Generally refers to <i>pass@k</i> scores. A peak score characterizes the best-case performance of a system over k independent attempts—i.e., the probability that at least one of k runs produces a correct outcome. Peak scores are informative about a system’s ceiling capability but do not reflect consistency.
Reliable	Generally refers to <i>pass[^]k</i> scores. A reliability score characterizes the probability that a system produces a correct outcome on <i>every one</i> of k independent attempts. Reliability scores penalize variance and are informative about how consistently a system can be expected to succeed across repeated deployments.

B Experiment Parameters

Below are the turn detection and model configurations for all evaluated and judge models; the user simulator is detailed separately in Appendix D. The agent prompts are provided in Appendix L. Finally, all the third-party software packages used are shared in Appendix N.

B.1 Self-Hosted Models

All self-hosted models were served on NVIDIA H100 GPUs. Models served via vLLM used `vllm-openai` v0.19.0. Table 3 lists the hardware and serving configurations.

Gemma-4-26B and **Gemma-4-31B** were called with `temperature=1.0`, `top_p=0.95`, `top_k=64`, and `max_tokens=12000`. Thinking mode was disabled via `enable_thinking=false` and special tokens were preserved (`skip_special_tokens=false`).

Qwen3.5-27B was called with `temperature=1.0`, `top_p=0.95`, `top_k=20`, `min_p=0.0`, `presence_penalty=1.5`, and `repetition_penalty=1.0`. Thinking mode was likewise disabled via `enable_thinking=false`.

Table 3 Self-hosted model configurations.

Model	Model Abrv	Model ID	Type	GPU	CPU	Precision	Deployment
Gemma-4-26B	Gemma-4-26B	google/gemma-4-26B-A4B-it	LLM	2× H100	–	BF16	vLLM
Gemma-4-31B	Gemma-4-31B	google/gemma-4-31B-it	LLM	2× H100	–	BF16	vLLM
Qwen3.5-27B	Qwen3.5-27B	Qwen/Qwen3.5-27B	LLM	4× H100	8× 8GB	BF16	vLLM
Parakeet-1.1	Parakeet	nvidia/parakeet-ctc-1.1b	STT	1× H100	–	BF16	Nvidia NIM
Whisper-Large-v3	Whisper	openai/whisper-large-v3	STT	1× H100	4× 64GB	FP16	vLLM
Cohere-transcribe	Cohere	CohereLabs/cohere-transcribe-03-2026	STT	1× H100	8× 64GB	BF16	vLLM
Kokoro	Kokoro	hexgrad/Kokoro-82M	TTS	1× H100	8× 32GB	FP32	Remsky Kokoro
Voxtral-4B-TTS	Voxtral	mistralai/Voxtral-4B-TTS-2603	TTS	1× H100	–	BF16	vLLM

B.2 API-Hosted Models

For ElevenLabs, we used ElevenAgents with the following models: Scribe-v2.2-Realtime, Gemini-3-Flash, and TTS-Conversational-v3-Alpha. We used the default agent parameters, listed in Table 4.

Table 5 lists all the other API-hosted models.

B.3 Turn Detection Configurations

We use the default turn detection configurations for each framework in our experiments. Each framework offers varying levels of configurability, making it difficult to standardize exact parameters and turn strategies across evaluations.

Pipecat. The default start strategy uses VAD (voice activity detection) or transcription receipt to determine when the user begins speaking, and the stop strategy uses AI-powered turn detection via `LocalSmartTurnAnalyzerV3` to determine when the user finishes speaking.¹

OpenAI Realtime. We use the default server VAD, which uses periods of silence to detect turn boundaries. Default values are used for `threshold`, `prefix_padding_ms`, and `silence_duration`.²

ElevenAgents. The turn “eagerness” parameter is left at its default setting (`normal`).³

Gemini Live. We use the default automatic VAD provided.⁴

¹<https://docs.pipecat.ai/api-reference/server/utilities/turn-management/user-turn-strategies#base-parameters-2>

²<https://developers.openai.com/api/docs/guides/realtime-vad>

³<https://elevenlabs.io/docs/eleven-agents/customization/conversation-flow#turn-eagerness>

⁴<https://docs.cloud.google.com/vertex-ai/generative-ai/docs/model-reference/multimodal-live>

EVA-Bench makes turn detection parameters and options configurable via the CLI, so practitioners can run experiments using the turn detection settings available to their chosen framework. The only exception is ElevenAgents, where users must register and configure their agents separately prior to evaluation.

B.4 Experiments Compute Resources

Running EVA-Bench experiments involves costs across three distinct components of the pipeline. First, the **user simulator** is powered by ElevenLabs’ ElevenAgents (Appendix D), a fully hosted conversational AI system, incurring costs per interaction across all simulated dialogue turns. Second, the **agent under evaluation** introduces inference costs whenever the model is proprietary. Closed-source models are billed per token for LLM inference, and per second of audio for STT and per character for TTS when the agent operates as a full voice pipeline. Note that a conversation will on average last 4 to 5 minutes. Third, the **LLM-as-judge** component adds a separate layer of inference cost, as 8 of the evaluated metrics rely on model-based scoring computed per sample. We evaluated 213 scenarios over 5 trials on clean data and 90 scenarios over 3 trials per perturbation across 3 perturbation types, yielding $213 \times 5 + 90 \times 3 \times 3 = 1,875$ evaluation samples per model. Costs are further compounded by re-runs triggered when a simulation fails validation gates. Across four representative systems evaluated on all domains, 24.1% of trials required regeneration (roughly half due to simulator error, the other half due to infrastructure failures or timeouts), effectively inflating the true number of simulations and judge calls to approximately $1,875 \times 1.241 \approx 2,327$ per system evaluated, on average.

Researchers should therefore treat full EVA-Bench runs as non-trivial compute investments, particularly when benchmarking multiple proprietary models simultaneously, and plan accordingly for parallelization and cost budgeting prior to large-scale experiments.

Table 4 ElevenLabs ElevenAgents parameters.

Component	Parameter	Value
STT	Filter background speech	disabled
TTS	Expressive mode	enabled
TTS	Voice	Lauren B - Friendly & Engaging Customer Care Agent
TTS	Voice ID	3liN8q8YoeB9Hk6AboKe
LLM	Temperature	0
	Reasoning effort	minimal
	Limit token usage	−1
	Parallel tool calling	disabled
	Cascade timeout	8 s
Tools	Wait for response	enabled
	Pre-tool speech	auto
	Execution mode	immediate
	Tool call sound	none
	Response timeout	20 s
Agent	Eagerness	normal
	Spelling patience	auto
	Speculative turn	enabled
	Re-transcribe on timeout	disabled
	Take turn after silence	15 s
	End call after silence	disabled
	Max conversation duration	600 s

Table 5 API-hosted model configurations. In this work, due to space constraints we occasionally use the following model abbreviations: Ink-whisper→ Ink, Sonic 3→ Sonic, Aura-2→ Aura, Ultravox-Realtime→ Ultravox. Note that Ultravox-Realtime is offered as a realtime model service in Pipecat.

Model	Provider	Type	Model ID	Parameters
GPT-5.2	OpenAI	LLM	gpt-5.2	reasoning: default
GPT-5.4	OpenAI	LLM	gpt-5.4	reasoning: default
GPT-5.4-mini	OpenAI	LLM	gpt-5.4-mini	reasoning: default
GPT-Realtime-1.5	OpenAI	S2S	gpt-realtime-1.5	voice: Marin
GPT-Realtime-mini	OpenAI	S2S	gpt-realtime-mini	voice: Marin
Gemini-3-Flash	Google	LLM	gemini-3-flash-preview	reasoning: default
Gemini-3.1-Flash-Live	Google	LALM	gemini-3.1-flash-live-preview	voice: Leda
Gemini-3.1-Flash-TTS	Google	TTS	gemini-3.1-flash-tts-preview	voice: provider default
Claude Opus 4.6	AWS Bedrock	LLM	us.anthropic.claude-opus-4-6-v1	reasoning: default
Haiku-4.5	AWS Bedrock	LLM	us.anthropic.claude-haiku-4-5-20251001-v1:0	−
Ultravox-Realtime	Ultravox	LALM	−	−
Ink-whisper	Cartesia	STT	ink-whisper	−
Sonic 3	Cartesia	TTS	sonic-3	−
Nova-3	Deepgram	STT	nova-3	−
Aura-2	Deepgram	TTS	aura-2-helena-en	−

C Data

We created three datasets on different enterprise domains, each selected to target a distinct axis of difficulty for voice agents. All three require accurate transcription of structured named entities over voice (e.g., confirmation codes and employee identifiers), but differ in their primary challenge. **Airline Customer Service Management (CSM)** tests temporal reasoning and complex policy adherence in high-stakes flight rebooking scenarios. **Healthcare Human Resources Service Delivery (HRSD)** stresses entity density, requiring callers to communicate multiple registration and license numbers across clinical and administrative HR workflows. **Enterprise Information Technology Service Management (ITSM)** introduces branching conversational flows (e.g., incident resolution attempts must fail before ticket escalation is permitted) and tiered authentication reflecting the access sensitivity of different workflows. Dataset statistics are summarized in Table 6.

Within each domain, scenarios span three dimensions: **Single-Intent** (one workflow per call), **Multi-Intent** (one to four concurrent workflows, testing compositional task completion without context loss), and **Adversarial** (hard policy constraints under social pressure, e.g., refusing compensation to an ineligible caller).

Section C.1 introduces the workflows for the three data domains, and more details on each are provided in Appendix I, with scenario examples (one for each domain, and a single-intent, a multi-intent, and an adversarial example) in Appendix J. Section C.2 describes the data generation pipeline.

C.1 Workflows

The **Airline CSM** domain covers 50 scenarios across seven workflow categories backed by 15 tools. It is high-stakes and time-pressured, with heavy dependence on accurate transcription of named entities such as confirmation codes, flight numbers, and passenger names.

The **Healthcare HRSD** domain covers 83 scenarios across 12 single-intent workflows backed by 47 tools, extended with dual-intent, triple-intent, and adversarial variants. It has the highest per-workflow complexity in EVA-Bench, averaging 8.7 expected tool calls per scenario. Its defining challenge is the density of named entities communicated over voice — NPI numbers, DEA registration numbers, state license numbers, and OTP codes — where a single transcription error can cascade into authentication or policy failures.

The **Enterprise ITSM** domain covers 80 scenarios across 21 workflows backed by 59 tools, spanning single- to quadruple-intent and adversarial variants. Its defining characteristic is a branching flow structure, where incident workflows gate escalation on failed resolution attempts. Authentication is tiered across three levels — standard, OTP-elevated, and manager-level — reflecting the sensitivity of different workflows.

Table 6 Comparison of EVA-Bench data domains. Multi-intent scenarios present one to four concurrent workflows within a single call. Auth tiers refers to distinct levels of caller authentication required across flows.

	Airline CSM	Healthcare HRSD	Enterprise ITSM
Scenarios	50	83	80
Workflows	7	12	21
Tools	15	47	59
Avg. tool calls	3.14	8.7	8.3
Min / max tool calls	1 / 6	1 / 18	1 / 18
Auth tiers	1	2	3

C.2 Data Generation Pipeline

Synthetic Data Generation with SyGra

Scenarios are generated using SyGra [27], a graph-based synthetic data generation pipeline. Each scenario requires three jointly consistent components: a user goal (including a decision tree that constrains the user simulator to a deterministic outcome), a scenario database (the backend state the agent’s tools query and modify), and an expected final database state (the ground truth against which task completion is evaluated).

Joint generation is essential: the expected final state must be consistent with both the user goal and the initial database. Independent generation would introduce silent inconsistencies; for example, a flight number referenced in the user goal that does not exist in the scenario database would corrupt the evaluation signal.

Generation proceeds in the following stages:

1. **Policy specification.** Domain policies and workflow constraints are defined and reviewed prior to generation.
2. **Joint generation.** SyGra generates user goals, initial databases, and expected final states jointly from a workflow graph, using GPT-5.2 as the generative backbone.
3. **Multi-intent composition.** Multi-intent scenarios are constructed by combining single-intent records into coherent multi-workflow user goals, with expected final states merged accordingly.
4. **Adversarial scenario design.** Adversarial scenarios are hand-designed around specific policy boundary conditions, then verified against tool executor behavior to confirm that the policy violation is achievable but detectable by a correctly behaving agent.

Human Review

Following generation, all scenarios went through multiple rounds of manual review. Reviewers verified that: (1) policies were applied consistently across scenarios within a domain; (2) user goals were specific enough to admit exactly one correct resolution; (3) expected final states were internally consistent with both the user goal and the initial database; and (4) adversarial scenarios were correctly specified, with a clearly identifiable policy violation. Records identified as ambiguous or inconsistent were corrected or discarded.

Frontier Model Stress Testing

As a final validation step, we ran three frontier models — OpenAI/Gpt-5.4, Google/Gemini 3.1 Pro, and Anthropic/Claude Opus 4.6 — on a text-only version of each scenario, bypassing the audio pipeline and providing conversation transcripts directly. For every scenario on which any model scored zero on task completion, we manually investigated whether the failure reflected genuine model error or a dataset issue: an ambiguous policy, an under-specified user goal, a bug in the tool executor, or an inconsistency between the initial and expected database states. Records with identified dataset issues were corrected or removed. All selected samples had a task completion of 1 on at least one of the frontier models.

This process provides high confidence that task completion failures in the full audio evaluation reflect real agent errors rather than evaluation artifacts.

D User Simulator Details

We use ElevenLabs ElevenAgents [8] as the user simulator with the following cascade system: **Scribe v2.2 Realtime + GPT-5.1 + Eleven V3 Conversational**. We select these models for their high transcription accuracy, User Behavioral Fidelity, user realism for GPT-5.1 [36], and for their naturalness and realism for Eleven v3 Conversational. ElevenLabs also provides a large voice library, enabling testing of a wide variety of user accents, languages, speaking styles, etc.

We created four ElevenLabs agents for the user simulator, covering two accents (English and French) and two genders each. When creating a new agent, select *Blank Agent* as the starting template, then apply the configuration described in Table 7. The four voice names are listed in Table 8. All parameters not listed are set to their default values provided by ElevenLabs at agent creation.

Table 7 ElevenLabs User Simulator Configuration

Parameter	Value
TTS model family	V3 Conversational
Expressive mode	Enabled (no tags selected)
Language	English
LLM	GPT-5.1
System prompt	{{prompt}}
Default personality	Disabled
First message	None
Interruptible	Disabled
Advanced > Input audio	μ -law telephony, 8000 Hz
Advanced > Eagerness	Eager
Advanced > Take turn after silence	15 s
Advanced > Max conversation duration	600 s
Tools > System tools	End conversation (enabled)

Table 8 ElevenLabs Voice Name per Agent

Accent	Gender	Voice Name	Voice ID
English	Female	Natalee Champlin	KpTQ5yzwazQWLkvnK59A
English	Male	Eric - Smooth, Trustworthy	cjVigY5qzO86Huf0OWal
French	Female	Mariva Viva Muse - Warm and Energetic	1hIScOW98xkqE5ttC10C
French	Male	Jamie - French Accent & Charismatic	K8nDX2f6wjv6bCh5UeZi

When enabling the "End Conversation" system tool, the name must be `end_call`, and the description to provide is shown below. This allows the simulator to hang up programmatically.

"End Conversation" Tool Description

Use this to end the phone call and hang up.

Call this function when its time to end the call and one of the following is true:

1. The agent has confirmed your request is resolved (all steps are completed) and you have said goodbye
2. The agent has initiated a transfer to a live agent
3. The agent has been unable to make progress for at least 5 consecutive turns
4. The agent says goodbye or indicates the conversation is over
5. The agent indicates that the remainder of your request cannot be fulfilled.
6. If the assistant says something along the lines of "I'm sorry I encountered an error processing your request."

IMPORTANT: never call this tool in the same turn that you provide the agent with data, an identifier, a request to transfer to a live agent, an approval to proceed, or any kind of additional information.
Before calling this tool, always say a brief goodbye first.

Once the agent is configured, click **Publish** in the top-right corner. The `agent-id` can be retrieved from the **Widget** tab of the agent dashboard, under **Embed code**.

The simulator is prompted in EVA-Bench with a specific user goal and is instructed to stay on task, communicate all required named entities clearly, and terminate the conversation when the goal is accomplished, or the task is clearly unlikely to succeed. The system prompts are provided in Appendix K.

Table 9 EVA-Bench validation metrics. These metrics ensure the quality of the user simulator and the integrity of each conversation prior to evaluation.

Metric	Type	Scale	Pass Thresholds
Conversation Finished	Deterministic	$\{0, 1\}$	1.0
User Behavioral Fidelity	LLM-as-Judge	$\{0, 1\}$	1.0
User Speech Fidelity	LALM-as-Judge	$1-3 \rightarrow [0, 1]$	–

D.1 User Simulation Validation

Table 9 lists the three validation metrics used to verify user-simulator behavior and conversation integrity before any EVA-Bench evaluation is performed.

D.1.1 Conversation Valid End.

Before invoking any LLM judges, we run a deterministic check on the conversation logs to verify that each simulation terminated correctly. A valid end state is one in which either the agent failed to respond to the user, or the user invoked the end-call tool. Conversations meeting this criterion proceed to the User Speech and Behavioral Fidelity judges; all others are rerun.

This gate primarily catches infrastructure-level failures: WebSocket connections that closed unexpectedly, conversations that failed to start, or user simulator timeouts. These are pipeline errors rather than agent errors, and filtering them deterministically before any judge invocation avoids wasting LLM calls on malformed simulations.

D.1.2 User Behavioral Fidelity.

Motivation. A key assumption underlying our evaluation is that the user simulator faithfully follows its assigned goal and decision logic, since the ground truth end database state is derived from this assumption. If the user deviates, the agent may fail to reach the ground truth state through no fault of its own. We therefore define a validation gate that detects user behavior errors capable of corrupting the evaluation, organized into five *corruption types*:

1. **Extra modifications.** The user makes requests beyond its stated goal that invoke modification tools writing to the scenario database. The user simulator prompt explicitly instructs the user to decline any such offers from the agent, but we check for violations regardless.
2. **Premature ending.** In our simulations, the user is responsible for ending the call once its goal is complete. If the user hangs up prematurely—for example, providing actionable information and ending the call in the same turn—the agent has no opportunity to execute the required tool calls. We therefore verify that the user does not terminate the conversation in the same turn it provides critical information to the agent.
3. **Missing information.** If the user fails to provide information the agent needs to complete the task, the evaluation is corrupted since task success cannot reasonably be expected.

4. **Duplicate modifications.** Occasionally, the user simulator (particularly when using non-primary models) enters a loop and repeats requests the agent has already fulfilled. The agent then acts on the duplicate request, causing redundant writes to the scenario database that cause the final state comparison to fail.
5. **Decision tree violations.** Each user is given a structured decision tree governing how to navigate choices during the interaction (e.g., “*accept the earlier flight if the price difference is under \$200, otherwise decline*”). We verify that the user adheres to this logic, since deviations would cause the agent to reach a final state inconsistent with the ground truth.

Implementation. We prompt a judge model with a description of the five corruption types, the full conversation trace (including agent tool calls), and the list of available agent tools. The judge first produces a *corruption analysis* assessing each corruption type in turn, then outputs a binary flag per corruption type, and finally an overall binary rating for the conversation. Conversations receiving a rating of 0 are rerun. See [subsection M.7](#) for the full prompt.

Validation. To validate the judge, we constructed a human-annotated dataset of real user simulator failures from earlier evaluation artifacts alongside correct behavior examples, labeled with both overall ratings and per-corruption-type annotations. GPT-5.2 with medium reasoning achieved 100% accuracy across three independent runs, and was therefore selected as the judge model.

Error Categorization. Across 714 trials flagged by the User Behavioral Fidelity judge—spanning four evaluated systems across all domains—premature ending was the most frequent corruption type (63.9%), followed by decision tree violations (52.9%), missing information (28.6%), extra modifications (1.3%), and duplicate modifications (0%). Notably, 42.4% of flagged trials exhibited two or more corruption types simultaneously. These numbers are summarized in [Table 10](#).

D.1.3 User Speech Fidelity.

Motivation. As noted in [subsection 3.1](#), speech fidelity flags are rare in practice, which may invite questions about why this validation gate was included at all. The answer lies in one of the primary failure modes we observe in voice agents: failure to transcribe and understand key entities. To attribute such failures confidently to the agent rather than the user simulator, we must verify that the user’s synthesized speech correctly conveys all critical entities. This is particularly consequential because most flows begin with an authentication step requiring the agent to correctly capture entities such as names, confirmation codes, and account IDs—if the user’s speech corrupts these, the agent cannot proceed regardless of its own capability.

Implementation. We adapt the same prompt used to evaluate agent speech fidelity ([Section M.3](#)), with one key modification to the rating scale. Unlike for the agent, we do not require the user’s speech to precisely mirror every word of the user-side LLM output; only key entities and major informational content must be conveyed accurately. The judge therefore rates each turn on a 3-point scale: 3 indicates full fidelity, 2 indicates minor errors that do not affect the agent’s ability to progress (e.g., slight disfluencies, but all entities intact), and 1 indicates an entity error or significant omission or addition that would prevent the conversation from proceeding sensibly. Any conversation in which any turn receives a rating of 1 is rerun. The full prompt can be found in [subsection M.6](#).

Validation. Because this prompt is closely derived from the agent speech fidelity judge, we inherit its validation. That judge achieved high inter-annotator agreement with human linguists ($\kappa = 0.777$, 95% CI [0.704, 0.835]), and the core capability it requires—accurately parsing audio and detecting entity-level errors—is shared. The rating scale and its interpretation are sufficiently well-defined that additional annotation studies were not deemed necessary. See [table 14](#) for more details on human-judge agreement.

Table 10 Distribution of corruption types among trials flagged by the User Behavioral Fidelity judge. Percentages sum to more than 100% because a single trial may exhibit multiple corruption types.

Corruption Type	Cases	% of Flagged
Premature ending	456	63.9%
Decision tree violation	378	52.9%
Missing information	204	28.6%
Extra modifications	9	1.3%
Duplicate modifications	0	0%

E Metric Details

All EVA-Bench metrics are summarized in Table 11.

Table 11 EVA-Bench metrics organized by category. All EVA-A and EVA-X scores are normalized to $[0, 1]$ prior to aggregation. Thresholds are used for $\text{pass}@k$ and pass^k computation: a run is considered successful if all EVA-A and EVA-X metrics meet their respective thresholds simultaneously.

Category	Metric	Type	Scale	Pass Thresholds
EVA-A (Accuracy)	Task Completion	Deterministic	$\{0, 1\}$	1.0
	Faithfulness	LLM-as-Judge	$1-3 \rightarrow [0, 1]$	0.50
	Speech Fidelity	LALM-as-Judge	$[0, 1]$	0.95
EVA-X (Experience)	Conciseness	LLM-as-Judge	$1-3 \rightarrow [0, 1]$	0.50
	Conversation Progression	LLM-as-Judge	$1-3 \rightarrow [0, 1]$	0.50
	Turn-Taking	Deterministic	$[0, 1]$	0.80
Diagnostic	Authentication Success	Deterministic	$\{0, 1\}$	—
	Response Latency	Deterministic	seconds	—
	Speakability	LLM-as-Judge	$\{0, 1\}$	—
	STT Word Error Rate	Deterministic	$[0, \infty)$	—
	Tool Call Validity	Deterministic	$[0, 1]$	—
	Transcription Key Entities	LLM-as-Judge	$[0, 1]$	—

E.1 Log Processing and Variable Extraction

Available Logs Every conversation typically produces three independent log streams, which we merge and replay deterministically to recover the variables needed by our metrics. The streams are:

- **Audit log** (`audit_log.json`) — always present. The agent’s internal record: user-side STT transcripts (in cascade pipelines), the assistant’s full LLM output, and tool calls/responses with their parameters and timestamps.
- **Framework events** (`framework_logs.jsonl`) — written by a generic framework logger that any speech pipeline (Pipecat, S2S, custom audio LLM, ...) can attach to in order to record TTS-stage text events with wall-clock timestamps. The two records the processor depends on are `tts_text` (the chunk actually sent to TTS) and `llm_response` (the LLM-side text). Some pipelines emit `tts_text`, some emit `llm_response`, some emit both. Pipelines that emit neither (notably S2S systems with no separable TTS step) effectively contribute nothing to this stream, in which case the variables that draw from it are left empty. When present, these events give the canonical *intended* assistant text in the sense of *what was actually sent to the TTS engine*. This is *not* the assistant’s full LLM output: that lives in the audit log and may include continuations beyond an interruption point that were never sent to the TTS step. Concretely, `intended_assistant_turns` is populated directly from these framework events, while audit-log assistant entries enter only the conversation trace, where they are post-hoc truncated to the longest prefix attested in the framework log (entries with no spoken overlap are dropped).
- **ElevenLabs events** (`elevenlabs_events.jsonl`) — events from the user simulator and the shared audio bus: per-speaker `audio_start/audio_end` markers, the user simulator’s own outgoing text (`user_speech`, treated as the user’s *intended* utterance), and the provider’s ASR of the assistant audio channel (`assistant_speech`, treated as the assistant *transcribed* utterance).

The available streams are concatenated and sorted by timestamp into a single timeline, then traversed in one pass. Turn boundaries are driven *exclusively* by `audio_start(elevenlabs_user)` events: turn 0 is reserved for the assistant greeting (anything before the first user audio session), and each subsequent user audio session, provided the assistant has spoken since the last advance, increments the turn counter so that index i aligns assistant turn i as the response to user turn i . Several edge cases require care: (a) empty user sessions

(background noise without any `user_speech` payload) are rolled back so they do not consume a turn index; (b) `user_speech` events that arrive before their `audio_start` are buffered and replayed once the correct turn is known; and (c) after a barge-in, a `hold_turn` flag suppresses the next advance from a late STT chunk belonging to the interrupted utterance, while still allowing a fresh `audio_start` to advance normally.

Interruptions Interruptions are detected whenever one speaker’s `audio_start` fires while the other’s audio session is still open, producing two disjoint sets `assistant_interrupted_turns` and `user_interrupted_turns`; the corresponding text fields are decorated with `[assistant interrupts]` / `[user interrupts]` (entry-level prefixes) and `[likely cut off by user]` / `[likely cut off by assistant]` / `[likely cut off on its own]` (turn-level suffixes). We deliberately mark these labels as *likely* rather than excising the post-interruption text: the intended-text streams record *what was sent to TTS*, not what was vocalised, and there is typically a non-trivial delay between text being handed to TTS and audio reaching the speaker. Words queued in the final moments before an interruption may therefore have been buffered but never played, so the precise truncation point in the intended text is not recoverable from the logs alone, and we let the annotation flag the ambiguity rather than make a hard cut at a position we cannot identify with confidence. This help the downstream judges to adjust their scoring based on the presence of interruptions. For example, AGENTSPEECHFIDELITY should not penalize if some words that are present in the intended text were never said due to an interruption.

Extracted Variables For each turn we extract four per-role variables: `intended*_turns`, `transcribed*_turns`, `audio_timestamps*_turns`, and entries of a linearised `conversation_trace` that interleaves user/assistant turns with tool calls. The default mapping from log source to variable is given in Table 12. Crucially, the table distinguishes the per-turn text fields (which are sourced directly from a single stream) from the `conversation_trace` (which is built from the audit log and post-hoc reconciled against the other streams). The `conversation_trace` is the linear, tool-call-interleaved view used by judge metrics that need a faithful chronological transcript, while the per-turn fields are useful for specific metrics that need intermediate states, such as TRANSCRIPTIONACCURACYKEYENTITIES. `audio_start`/`audio_end` pairs are matched greedily by speaker, and used to compute any latency measurements.

Table 12 Default mapping from log source to extracted variable. The upper block lists the per-turn text and audio fields; the lower block lists the entries that compose the linear `conversation_trace`. Pipeline-specific overrides are listed in the text.

* Empty if the framework emits neither record. † Falls back to `assistant_speech` for S2S. ‡ Uses `user_speech` for audio-native pipelines (S2S/Hybrid).

Variable	Source
<code>transcribed_user_turns[i]</code>	<code>audit_log</code> / <code>user</code>
<code>intended_user_turns[i]</code>	<code>elevenlabs</code> / <code>user_speech</code>
<code>intended_assistant_turns[i]</code>	<code>framework_logs</code> / <code>tts_text</code> , <code>llm_response</code> *
<code>transcribed_assistant_turns[i]</code>	<code>elevenlabs</code> / <code>assistant_speech</code>
<code>audio_timestamps_{role}_turns[i]</code>	<code>elevenlabs</code> / <code>audio_start</code> , <code>audio_end</code>
<code>tool_params</code> , <code>tool_responses</code>	<code>audit_log</code> / <code>tool_call</code> , <code>tool_response</code>
<code>conversation_trace</code> (assistant)	<code>audit_log</code> / <code>assistant</code> , truncated to framework-log prefix†
<code>conversation_trace</code> (user)	<code>audit_log</code> / <code>user</code> ‡
<code>conversation_trace</code> (tools)	<code>audit_log</code> / <code>tool_call</code> , <code>tool_response</code>

Pipeline type modifies this default in two principled ways, reflecting which signals are trustworthy for each architecture. These differences impact the metrics definition, as discussed in E.2.

1. **S2S.** There is no separable TTS step, so the framework log carries no `tts_text` or `llm_response` records and is dropped from the merge. Consequently `intended_assistant_turns` is left empty — S2S models typically do not expose any separate text intent — and the assistant’s entries in `conversation_trace` are sourced from ElevenLabs `assistant_speech` (transcribed) rather than from the audit log.

Symmetrically, S2S models consume the user audio directly and do not produce a trustworthy STT transcript of the user, so the user entries in `conversation_trace` are sourced from `user_speech` (the simulator’s intended text) rather than from the audit log.

2. **Hybrid.** The framework log is populated with `tts_text` or `llm_response` (depending on the backend), and `intended_assistant_turns` is built as in cascade. On the input side, however, hybrid audio-native models bypass the agent’s STT — as in S2S — so the audit-log user transcripts are unreliable, and the user entries in `conversation_trace` are again sourced from `user_speech` (intended) rather than from the audit log.
3. **Cascade.** All three streams are used unmodified: audit-log user transcripts feed both `transcribed_user_turns` and the trace, the framework log supplies the assistant’s intended text, and ElevenLabs supplies the user’s intended text and the assistant’s transcribed text.

A final post-processing step (i) aligns the per-turn dictionaries so that all sources share the same key set, with missing slots back-filled by the most informative remaining source; and (ii) reconciles the trace by ensuring the greeting is the first entry, appending any final user turn that arrived after the last audit-log entry, and propagating trailing-cutoff labels consistently across `intended_*`, `transcribed_*`, and the trace.

Audio recordings. In addition to the event logs, every session writes three 16-bit PCM mono WAV files, captured directly from the audio bus: `audio_user.wav` contains *only* the user simulator’s outgoing audio, `audio_assistant.wav` contains *only* the assistant’s outgoing audio, and `audio_mixed.wav` is a sum-mixed mono recording of both speakers, used as a reference of what either party would actually have heard. Splitting the channels at capture time — rather than diarising a mixed recording post-hoc — is what allows the speech-fidelity metrics to operate without speaker-attribution noise: `AGENTSPEECHFIDELITY` consumes `audio_assistant.wav` to compare what the assistant *said* against its corresponding `intended_assistant_turns`, and similarly for `USERSPEECHFIDELITY` with `audio_user.wav` to validate that the user simulator’s actual speech matches its scripted persona and goal. The mixed channel is not used by any metric, it is a reference that can be used for human review.

Limitations and fidelity. The three log streams are emitted by independent components — the agent, the speech framework, and the user simulator’s provider — and can drift from one another in edge cases, particularly around interruptions and rapid turn switches. We have observed, for example, some ElevenLabs `audio_end(elevenlabs_user)` arriving noticeably after the user simulator actually stopped speaking (so the audio session appears to “stay open” past its useful end), occasional transcripts that are missing or delayed, and misalignments between the framework log timestamps and the ElevenLabs ones. Making turn boundaries align across all sources is also a challenge, as one user turn can be detected as two on the framework side. The audio recordings themselves are also not always unambiguous, and observed latencies do not always match the latency estimates derived from log timestamps. The heuristics described above (empty-session rollback, late-transcript buffering, prefix-truncation against the framework log, the `hold_turn` flag after barge-ins) are designed to absorb the most common of these inconsistencies. To validate that residual drift is not a confound, we manually inspected a stratified sample of conversations across all three pipeline types: the extracted variables were high-fidelity in aggregate, per-turn alignment between intended and transcribed text was consistent, interruption labels matched the audio, and trace ordering matched the perceived chronology of the conversation. Cases where a single log entry was missing, a turn boundary was off by one, or interruption tags were wrong did occur but were rare and did not systematically bias any of the metrics reported in this work.

E.2 Equitable Evaluation Across Cascade, Hybrid, and S2S Architectures

The three pipeline types differ in what is observable about the agent’s reasoning (Appendix E.1), and a naive single-view evaluation would systematically favour one architecture over another. We therefore provide pipeline-aware variants of the reasoning-oriented metrics — Faithfulness, Conciseness, Conversation Progression— so that every system is judged on the most faithful proxy of *what its LLM actually saw and produced*.

In a **cascade** system, the LLM consumes user turns as STT transcripts and emits agent turns as text before TTS rendering. Both halves of this LLM-internal view are recoverable: `transcribed_user_turns` captures what the LLM read, and `intended_assistant_turns` captures what it wrote. Scoring on these two fields ensures that STT and TTS errors — the responsibility of the surrounding pipeline, not the LLM — are not attributed to the model under test.

In an **S2S** system, the model consumes and emits audio directly; there is no intermediate text on either side. The conversation trace is therefore built from the user simulator’s *intended* text on the user side and a post-hoc ASR transcript of the assistant’s audio also on the user side (see Appendix E.1). The same judges then evaluate the agent as a whole — speech understanding and synthesis included — since these capabilities are part of the S2S system’s responsibility. To avoid charging the agent for transcription errors of its output, the judge prompt explicitly defer such errors to Speech Fidelity, which receives the raw audio.

Hybrid systems are evaluated with a mixed view that follows the same observability principle as the other two. Like S2S, hybrid models bypass the agent-side STT, so user turns are taken from the simulator’s intended text rather than from a transcript; like cascade, hybrid models retain a separable text-to-TTS step, so agent turns are taken from `intended_assistant_turns` via the framework log.

E.3 Judge Development and Validation

LLM-as-judge evaluations are only as reliable as the judges themselves. We developed each judge through a structured five-stage pipeline: metric definition, prompt construction, development dataset construction, prompt improvements and judge model selection, and final validation against human annotation.

Metric definition and rating scales. For each judge metric, we first defined the rating scale and the explicit failure modes that distinguish each rating level. Each failure mode specifies both what the defect looks like and how a judge should detect and categorize it. This categorical and granular framing is intentional: it preserves actionable signal about *how* a voice agent fails, rather than collapsing everything into a single score. It also helps the judge accurately score by offering detailed categories.

Judge prompt construction. The failure mode definitions were operationalized into judge prompts, with explicit criteria for each rating category and targeted guidance for edge cases on the boundary between adjacent ratings.

Development dataset construction and labeling. We constructed a development dataset for each metric using synthetic data generation followed by multi-model consensus labeling. For data generation, we prompted frontier models to produce conversations exhibiting specific, targeted failure modes within each metric. This targeted generation was key to achieving class balance: on top of sampling from naturally occurring conversations, we forced the generator to produce examples of each failure mode explicitly. All generated data used the same domains, policies, and agent configurations as in EVA-Bench, ensuring the development set is in-distribution with respect to our evaluation targets. One exception is Speech Fidelity, where the class distribution is intentionally skewed toward failures (122 failing vs. 25 passing out of 147 records). This metric is evaluated by a Large Audio Language Model (LALM), and not a text-based LLM judge, and LALM-as-judge is a relatively nascent paradigm. We had limited confidence that the model would reliably attend to subtle acoustic artifacts without explicit exposure to a wide range of failure modes during development. Furthermore, because naturally occurring synthesized speech is predominantly artifact-free, a naive judge could achieve high accuracy simply by predicting “pass” for every sample. To guard against both of these risks, we oversampled failure cases during data generation, forcing a wide variety of speech artifacts to ensure the judge is sensitive to errors. In practice, we observe that the selected judge rarely predicts a passing label when the true label is failing, validating this design choice.

For labeling, we called three frontier models—Gemini 3.1 Pro, Claude Opus 4.6, and GPT-5.2—as judges on each generated sample. When all three models agreed on a rating, that rating was assigned as ground truth. When they disagreed, a human reviewer examined the sample and selected the correct label.

Prompt improvements and Judge model selection. We used the development sets to refine the prompts, focusing on samples where judges disagreed and analyzing their explanations. This process revealed ambiguities in the judge instructions and informed targeted prompt improvements.

We then formally evaluated each of the three frontier models as judge candidates on the development datasets, reporting accuracy alongside macro-averaged F1. We selected the judge model that achieved the highest combined score on each metric’s development set (see Table 13). For Speech Fidelity, Gemini 3.1 Pro achieved the highest performance, but we selected Gemini 3 Flash for deployment: it achieved nearly identical performance at substantially lower inference cost. Model ids and parameters used for each judge is provided in Table 15.

Table 13 Judge validation datasets and results per candidate model. Size = number of annotated examples. Dist. = class distribution. Acc. = judge accuracy against adjudicated human labels. F1 = macro F1-score. **Bold** indicates the selected judge model for each metric.

Metric	Dataset	Claude Opus 4.6		GPT-5.2		Gemini 3 Flash		Gemini 3.1 Pro	
	Size	Acc.	F1	Acc.	F1	Acc.	F1	Acc.	F1
Faithfulness	137	83.9%	80.7%	70.8%	68.9%	—	—	81.0%	78.4%
Conciseness	100	90.8%	73.4%	92.3%	84.0%	—	—	91.2%	74.4%
Conv. Progression	136	74.3%	73.1%	78.7%	76.7%	—	—	69.1%	66.6%
Speech Fidelity	147	—	—	—	—	89.1%	83.2%	91.2%	85.9%

Test set validation and human agreement. To assess the validity of the finalized judge prompts and selected models, we separately constructed a held-out test set of 63 samples per metric, never used during prompt development. Each sample was labeled by two expert linguists. For text-based metrics (Faithfulness, Conversation Progression, Conciseness), linguists were shown the conversation sample and the judge outputs from all three frontier model candidates, presented blind to which model produced which output. We provided judge outputs to the linguists because voice agent transcripts can be long and complex; without the judge’s surfaced evidence and analysis, annotators found it difficult to reliably identify subtle failure modes in long traces. Linguists verified the judge’s reasoning against the transcript directly rather than accepting it uncritically. For Speech Fidelity, linguists listened to the raw audio only along with the intended text, with no judge output, since this metric requires direct perceptual evaluation of synthesized speech.

We report inter-annotator agreement using Cohen’s κ in Table 14: linguist–judge agreement (L_J , pooled across 126 pairs per metric), linguist–judge agreement using a single randomly selected linguist per record (L_J_{rand} , 10,000 iterations), and linguist–linguist agreement (LL , 63 pairs per metric). To verify that the pooled pairing design does not distort agreement estimates, we computed κ using a single randomly selected annotator per record (10,000 iterations combining labeler randomization with record-level bootstrap); results fell within 0.007 of the pooled estimates across all metrics. Quadratic-weighted κ is used for the three ordinal metrics; unweighted κ for the binary Speech Fidelity metric. 95% confidence intervals are from 10,000 record-level bootstrap resamples.

Linguist–judge κ ranges from 0.777 to 0.845 across the four metrics—a strong result, and notably one that meets or exceeds the linguist–linguist agreement ceiling in every case. This indicates that our judges are not merely consistent with human raters, but that human–judge agreement is at least as high as the agreement between two human experts annotating the same data.

E.4 Accuracy Metrics

E.4.1 Task Completion

What it measures. Task Completion is the bottom-line accuracy check: it verifies if the agent actually accomplished what it was asked to do. Concretely, it asks whether the changes the agent committed to the scenario database during the conversation match the expected end state encoded in the dataset’s ground

Table 14 Human-judge agreement (κ , Spearman ρ) for the primary judge (L_J) versus the human linguist inter-annotator baseline (L_L). Quadratic-weighted κ is used for the 1–3 ordinal metrics; unweighted κ for the binary Agent Speech Fidelity metric. 95% CIs from 10,000 record-level bootstrap resamples. To verify that the pooled pairing design does not distort agreement estimates, we computed κ using a single randomly selected annotator per record (10,000 iterations combining labeler randomization with record-level bootstrap); results fell within 0.007 of the pooled estimates across all metrics. L_J κ ranges from 0.777 to 0.845 across the 4 metrics; Spearman ρ agrees with κ within 0.008 for every metric, indicating no systematic calibration bias. Note that IAA-L is a practical rather than strict ceiling: when two human annotators disagree and the judge agrees with one, IAA-J is inflated relative to IAA-L. Nonetheless, similar IAA-J and IAA-L values indicate that the judge is at least as consistent as the human annotators.

Metric	L_J κ	95% CI	L_J ρ	L_J _{rand} κ	95% CI	L_L κ	95% CI	L_L ρ
Faithfulness	0.836	[0.729, 0.915]	0.844	0.832	[0.697, 0.932]	0.740	[0.566, 0.870]	0.756
Conv. Progression	0.845	[0.753, 0.911]	0.843	0.841	[0.724, 0.931]	0.769	[0.627, 0.875]	0.782
Conciseness	0.823	[0.754, 0.874]	0.826	0.823	[0.745, 0.883]	0.825	[0.749, 0.881]	0.825
Speech Fidelity	0.777	[0.704, 0.835]	0.781	0.770	[0.683, 0.846]	0.754	[0.685, 0.817]	0.754

Table 15 Judge model identifiers and inference parameters used during evaluation. Parameters not listed were left at provider defaults.

Model	Model ID	Parameter	Value
GPT-5.2	gpt-5.2	max_tokens	100,000
Claude Opus 4.6	us.anthropic.claude-opus-4-6-v1	—	—
Gemini 3 Flash	gemini-3-flash-preview	temperature	0.0
		max_tokens	40,000
		reasoning_effort	minimal

truth. Unlike the judge metrics, this is a deterministic code-based check with no LLM in the loop; the same conversation always yields the same task-completion verdict, and the verdict is binary.

Method. Each scenario in the dataset specifies an `expected_scenario_db`: the database state we expect after a successful run. During execution, every state-mutating tool call writes through to a per-record copy of the scenario database, and the final state is captured at the end of the conversation (`final_scenario_db`). The metric canonically serialises both states (sort keys, no whitespace), computes their SHA-256 hashes, and reports a pass if the two hashes match. When they do not match, a structured diff is computed, including tables added/removed/modified, records added/removed/modified within tables, and field-level changes within records.

Inputs. Task Completion does not consume the conversation trace, the audio, or any text-side variables produced by log processing. It uses only:

- `expected_scenario_db` — the dataset’s ground-truth final state for this record;
- `final_scenario_db` — the actual final state captured at end-of-run;
- `final_scenario_db_hash` — the SHA-256 hash of `final_scenario_db`.

This separation is deliberate: the metric measures *outcome* on the database, decoupled from how the agent got there. Path-quality concerns (correct tool usage, faithful disclosure, efficient progression) are the responsibility of the corresponding judge metrics (Appendices E.4.2, E.5.1).

Authentication gate. Authentication state lives in a dedicated `session` field of the scenario database, and is verified separately from the hash comparison rather than folded into it. The reason is that some scenarios can be satisfied by different valid authentication paths — different combinations of identifying fields the agent may legitimately collect — and a hash computed over the full database would mark any such variant as incorrect. We therefore (i) strip the `session` key from both the expected and the actual scenario database

before computing either hash, and (ii) verify authentication via a superset check: every key-value pair in the expected session must be present in the actual session (string comparisons are case-insensitive), but the actual session may carry additional fields without penalty. If the superset check fails, the metric short-circuits to a fail with details describing which session keys mismatched; otherwise the run proceeds to the hash comparison on the rest of the database.

Determinism by construction. For task completion to be a meaningful metric, a scenario must yield the same expected end state regardless of conversational variation, provided the agent does not make mistakes. We enforce this by tightly constraining the user simulator, as discussed in [D](#), so two valid conversations on the same scenario commit produce identical final-state hashes. Variation between systems (or between runs of the same system) only arises when an agent’s behavior deviates from the dataset’s intended outcome.

Rating scale and aggregation. Task completion is binary: 1.0 when the expected and actual hashes match, 0.0 otherwise.

Pass/fail thresholding. For pass-related aggregations such as `pass@k`, the pass threshold is set to 1.0 (i.e. exact match). There is no middle ground for this metric; a run either commits the expected writes or it does not. The structured diff in the failure details supports finer-grained downstream analysis when one wants to understand *how* a run fell short, but the headline metric remains binary.

E.4.2 Faithfulness

What it measures. Faithfulness is a conversation-level accuracy judge that asks whether the assistant remained grounded in the information, policies, and instructions available to it throughout the run. Unlike Task Completion — which only checks whether the goal was achieved — Faithfulness penalises the path: a conversation that concludes successfully but along the way hallucinates a fee, skips a required confirmation, or commits a fabricated identifier to a write tool will receive a low faithfulness score.

Judge and Prompt Template. The judge is run as a single LLM-as-judge call per conversation, using *Claude Opus 4.6*. The full prompt is available in [Appendix M.2](#).

Inputs. The judge consumes the full conversation trace alongside everything needed to evaluate it against agent policy:

- the linearised `conversation_trace` (with tool calls and responses inline);
- the agent’s configuration — `agent_role`, `agent_instructions`, and the JSON schema of the available tools;
- the simulated `current_date_time`, used to resolve temporal references and policy windows;
- the pipeline-aware shared fragments documenting how user and assistant turns are sourced and how interruption tags are used ([Appendix M.1](#));
- two faithfulness-specific pipeline-aware fragments, `disambiguation_context` and `misrepresentation_pipeline_note`, described below.

Failure modes. The judge scores five disjoint dimensions, each defined to be non-overlapping so that any given issue maps to exactly one.

1. **fabricating_tool_parameters** — the assistant called a tool with a parameter value that cannot be traced to any user statement, prior tool result, policy entitlement, simple arithmetic, or standard domain mapping. Includes invented IDs, empty placeholder values, and wrongly chosen enum buckets.
2. **misrepresenting_tool_result** — the assistant inaccurately conveyed something a tool actually returned: wrong field value, contradicted status, omitted material caveats (e.g. a non-zero fee), or arithmetic errors when computing values from tool data.

3. **violating_policies** — the assistant contradicted the agent instructions: skipped a required verification step, executed an irreversible write without the disclosure or confirmation the policy requires, or stated a policy incorrectly.
4. **failing_to_disambiguate** — the assistant proceeded on ambiguous or contradictory user input without clarification (multiple options, conflicting values, suspicious lookups failing on uncommon names or codes).
5. **hallucination** — a residual category for information stated to the user that has no source at all in any tool response, user utterance, agent instruction, or system context, and that is not already captured by the four preceding dimensions.

Pipeline-aware adaptations. Faithfulness applies the general framing of Appendix E.2 and adds two faithfulness-specific deltas. First, the `disambiguation_context` fragment changes the bar for clarification: in cascade, the assistant is reminded to account for STT-style transcription errors before write actions; in audio-native pipelines, the assistant is held to a *higher* clarification bar because mishearing letters, numbers, names, and codes is intrinsic to consuming raw audio, and the model is expected to anticipate it. Second, the `misrepresentation_pipeline_note` explicitly scopes the `misrepresenting_tool_result` dimension for audio-native pipelines: because assistant turns in those traces are post-hoc ASR of the assistant audio, token-level discrepancies between an assistant utterance and a tool result (dropped dashes, single-character substitutions, missing/extra digits in long IDs) typically reflect TTS-rendering or post-hoc-ASR artifacts and are scored by Speech Fidelity, not here; only structural or semantic discrepancies (wrong field, wrong order of magnitude, wrong category, or downstream signals indicating the agent was operating on a wrong value) are flagged on this dimension. In cascade, this note is empty: assistant turns are intended TTS text, so any discrepancy with a tool result *is* attributable to the LLM. The result is that mishearing the user is a faithfulness violation in S2S and Hybrid but not in cascade (where it is the STT’s responsibility).

Rating scale. Each dimension is rated on a 3-point integer scale:

- **3** — no issue on this dimension.
- **2** — minor or ambiguous issue with low user impact (e.g. a fabricated parameter that reaches a read-only tool, is caught quickly, and never surfaces to the user; a borderline policy deviation; a small phrasing embellishment that does not alter any decision).
- **1** — clear violation with material impact (e.g. a fabricated parameter passed to a write tool, regardless of whether the call succeeds; an irreversible action without the policy-required disclosure or confirmation; misstating a fee, balance, or eligibility rule the user could act on).

The dimension-level ratings are aggregated by *minimum* into a single overall rating, so a single rating-1 dimension produces an overall rating of 1. The overall rating is normalised to $[0, 1]$ via $(r - 1)/2$, giving $3 \rightarrow 1.0$, $2 \rightarrow 0.5$, $1 \rightarrow 0.0$. The flagged dimensions are preserved in the metric details and aggregated across samples.

Pass/fail thresholding. For pass-related aggregations, such as `pass@k`, we count a conversation as a faithfulness pass if its overall rating is ≥ 2 , i.e. its normalised score is ≥ 0.5 . This threshold sits exactly at the rubric’s own load-bearing boundary: rating 1 is reserved for violations with *material* impact on the user — financial consequences, irreversible actions taken without the policy-required disclosure or confirmation, misstatements of policy that the user could act on later — while rating 2 is explicitly defined as covering minor or ambiguous issues that do *not* materially affect the outcome (a small phrasing embellishment, a quickly self-corrected read-only fabrication, a borderline judgement call). As such, rating-2 conversations are still treated as acceptable: at the current capability level of voice-LLM systems, demanding strict rating-3 perfection on every run would be unrealistic and would compress meaningful differences between systems into a uniformly low pass-rate.

E.4.3 Speech Fidelity

What it measures. Speech Fidelity asks whether the assistant’s spoken audio actually matches what the system intended (or was expected) to communicate. It is the audio-side complement to Faithfulness: faithfulness scores what the LLM *decided* to say, speech fidelity scores whether the user could correctly hear the entities the LLM (or the upstream tool responses) returned. Errors here would result in the user receiving the wrong information, such as a garbled flight number, a confirmation code with one substituted character, or a dollar amount whose digit is dropped. The metric is computed per-turn on the assistant audio channel only (`audio_assistant.wav`, see Appendix E.1), and aggregated by mean across rated turns.

Judge and Prompt Template. Speech Fidelity is an audio-judge metric: it sends a multimodal request (audio + textual context) to *Gemini 3 Flash*. The audio is encoded as `base64 WAV` and accompanied by the per-turn entity context the judge needs to verify. To reduce upstream noise and cost, the assistant audio is silence-trimmed before being sent to the judge. The full prompt is available in Appendix M.3.

Two pipeline-specific variants. Unlike the text judge metrics, speech fidelity does not have a single prompt with pipeline-aware fragments: the cascade/hybrid case and the S2S case have qualitatively different inputs and use different prompts.

1. **Cascade and Hybrid.** Both architectures expose an *intended* text-side reference for the assistant, i.e., the LLM’s text output before TTS (`intended_assistant_turns` E.1). The judge task is a direct word-for-word comparison: did the audio reproduce the intended text, with particular attention to TTS-critical entities (confirmation codes, flight numbers, dollar amounts, dates, names, spelled-out alphanumeric codes, segmented reference IDs).
2. **S2S.** S2S systems do not typically expose any text-side intent, so there is nothing to compare the audio against in the cascade sense. We instead reformulate the question as an *entity articulation* check: does the assistant clearly and correctly speak the entities it was supposed to convey? The judge receives a *redacted conversation trace* in which assistant entries are replaced by an "[Assistant speaks]" placeholder per turn, while user utterances and tool responses are preserved verbatim. These are the entity sources the assistant was supposed to articulate. The judge transcribes the assistant audio itself and checks, per turn, that any entities it speaks that originate from the trace (a confirmation code returned by a tool, a name supplied by the user) are clearly audible. Turns where the assistant speaks no in-trace entities (greetings, questions, clarifications using only system-side phrasing) are flagged `has_entities: false` and excluded from aggregation.

Critically, the S2S variant explicitly excludes faithfulness/correctness from its scope. If the agent says "\$315" when the tool returned "\$300", that is a faithfulness violation (Appendix E.4.2), not a speech-fidelity issue: the metric only flags the turn if the dollar amount is garbled in audio. Hallucinated entities not present in the trace are likewise out of scope for this metric. This is the symmetric counterpart of the `misrepresentation_pipeline_note` carve-out on the faithfulness side: token-level audio artifacts are scored here, not under faithfulness.

Failure modes. A turn is rated 0 when any of the following are observed; otherwise 1:

- an entity spoken with wrong digits, letters, amounts, or numbers (cascade/Hybrid: against the intended text; S2S: against the corresponding source in the trace);
- missing words that change the meaning of the turn or omit an entity;
- added words that introduce a factually different entity;
- substituted words that alter an entity value;
- for spelled-out codes ("Z K three F F W"), any letter or digit that is unclear, missing, or substituted;
- for segmented reference IDs (REF-8JVSD-001, MEAL-FAROUM-PAX0), any segment that is unclear or wrong (e.g. "M E L" versus "M E A L").

Carve-outs. The prompt explicitly does *not* penalise:

- minor pronunciation variations that do not change entity identity (“Ms.” vs “Miss”);
- filler words (“um”, “uh”, “so”) added or omitted;
- slight pacing or prosody differences;
- non-spoken audio-direction tags in the intended text ([slow], [firm], [annoyed]) — these describe how the words should be spoken and were never expected in the audio;
- words in regions flagged by interruption tags as likely not spoken (Appendix E.1) — if a tag indicates a span was likely cut off, missing words in that span are not penalised;
- missing words at the very end of the *last* turn only (audio cutoff at the end of the conversation);
- for the S2S variant, any entities the agent speaks that are *not* in the conversation trace (those are out of scope — the metric does not evaluate hallucinations).

Rating scale and aggregation. Each rated turn is given a binary rating in $\{0, 1\}$, where 1 means the audio is correct, and 0 means there is an issue. The conversation-level score is the mean across rated turns; for S2S, turns excluded by `has_entities: false` are dropped from both the numerator and the denominator. Per-turn ratings, the judge’s audio transcript per turn, and the natural-language explanation citing intended-vs-actual mismatches are preserved in the metric details for inspection.

Pass/fail thresholding. For pass-related aggregations such as `pass@k`, the pass threshold is set to 0.95, i.e. a conversation passes if at most a small fraction of rated turns are flagged with an entity error. We use a stricter threshold than the rubric-grounded 0.5 used for the 3-point judge metrics for two reasons. First, the rating scale is already binary: there is no rating-2 “minor issue” band whose inclusion this threshold has to negotiate, so the rubric-grounding argument used elsewhere does not apply. Second, the underlying error is high-stakes: a single garbled confirmation code or dollar amount can render a whole conversation operationally wrong, so tolerating more than the occasional turn-level slip would mask exactly the kind of failure this metric exists to catch. The 0.95 ceiling is therefore set to absorb sporadic per-turn judge noise (e.g. a single ambiguous spell-out) while still requiring the audio-side reproduction to be substantively faithful across the run.

Limitations. Manual inspection of a stratified sample of rated turns surfaced a few false positives (the judge rates 0 on a turn whose audio actually reproduced the intended content). They are typically caused by upstream log-processing artifacts rather than by judge mistakes. The metric compares the audio against `intended_assistant_turns`, and that reference is only as good as the merge of the framework log and the ElevenLabs audio events described in Appendix E.1. Two failure patterns recur. First, the intended text contains a trailing span that was handed to TTS but never vocalized because of a barge-in whose interruption tag was not raised. The heuristics in Appendix E.1 absorb most cases, but residual mis-tagging still occurs, and the judge then legitimately flags “missing words” against an intended text that was never meant to be heard. Second, turn-boundary drift between the framework log and the ElevenLabs timeline aligns the wrong intended utterance with a given assistant turn, so the judge compares audio for turn i against intended text for turn $i \pm 1$, producing a spurious mismatch. These false positives are concentrated on turns adjacent to interruptions and rapid turn switches.

E.5 Experience Metrics

E.5.1 Conversation Progression

What it measures. Conversation Progression is a conversation-level experience judge that asks whether the assistant moved the conversation forward without redundancy: consistent progress toward the user’s goal, no repeated tool calls with identical parameters, no restating of information already communicated, retention of established facts, and well-formed clarification questions. Unlike Faithfulness, which scores the *correctness* of the assistant’s choices, conversation progression scores the *efficiency* of those choices — a run that arrives

at the right outcome but loops or re-asks for known information will receive a low progression score even if everything it says is faithful.

Judge and Prompt Template. The judge is run as a single LLM-as-judge call per conversation, using *GPT-5.2*. The prompt template is available in Appendix [M.4](#).

Inputs. Compared to faithfulness, conversation progression deliberately operates on a smaller input bundle, because policy reasoning is explicitly out of scope:

- the linearized `conversation_trace`;
- the pipeline-aware shared fragments documenting how user and assistant turns are sourced and how interruption tags are used (Appendix [M.1](#));
- one progression-specific pipeline-aware fragment, `information_loss_pipeline_note`, described below.

The judge does *not* receive the agent role, instructions, available-tools schema, or current date/time: those drive faithfulness/policy reasoning, and giving the progression judge access to them invites it to silently re-litigate faithfulness questions under a different label. The prompt enforces this scope explicitly — if an issue is primarily a policy or faithfulness violation (e.g. taking an action the user said not to, not disclosing a fee), the judge is instructed to leave it to faithfulness even if the violation also affects conversational flow; only issues where the assistant’s *conversational choices* (questions asked, information repeated, tools called) are themselves inefficient should be flagged here.

Failure modes. The judge scores four disjoint dimensions:

1. **unnecessary_tool_calls** — the assistant called a tool without justification: same tool with the same parameters after a successful prior response, a tool with empty/missing required parameters that produced a predictable error, or a tool whose result was already available from a previous response. As a hard caveat in the rubric, three or more unnecessary tool calls in a run automatically rates this dimension at the lowest level.
2. **information_loss** — the assistant failed to retain or act on an established fact: re-asked the user for information already provided, ignored a constraint the user explicitly stated, or failed to use a value returned by a prior tool when needed for the next step. The dimension is about *forgetting or ignoring* known facts; if the assistant proceeded against a user-stated preference deliberately, that is a faithfulness issue and is excluded from this dimension.
3. **redundant_statements** — the assistant restated information it had already communicated to the user (repeated explanations, repeated status updates, multiple recaps across non-final turns). The exception is a single recap at the very end of the conversation, which the rubric explicitly allows.
4. **question_quality** — the assistant’s questions were poorly formed or missing: overly broad/vague questions when enough information was on hand to act, multiple questions bundled in a turn that a single tool call could resolve, missing clarifications when input was genuinely ambiguous, or proceeding to an irreversible action without confirming an ambiguous value. Standard policy-required readbacks of error-prone values (alphanumeric IDs, codes, dates, amounts) are explicitly *not* flagged here.

Two cross-cutting carve-outs apply to all four dimensions. First, an *interruption-tag carve-out* (Appendix [E.1](#)): truncated speech caused by an interruption is not a progression issue per se; only its observable downstream consequences (information genuinely lost because the cut-off content was never restated, or the assistant repeating already-heard content after being interrupted) are flagged. Second, a *voice-context carve-out*: when the assistant repeats a request because the previous attempt was clearly misheard or garbled, that repetition is expected behaviour in a voice interface and is not a progression issue — but only when the transcript shows visible evidence of an ASR failure on the prior attempt; re-asking without cause is still a flag.

Pipeline-aware adaptations. Conversation progression follows the general framing of Appendix [E.2](#) and adds one progression-specific delta. The `information_loss_pipeline_note` explicitly scopes the `information_loss` dimension for audio-native pipelines: because assistant turns in those traces are post-hoc ASR of the

assistant audio, variant token-level readings of the same alphanumeric identifier across nearby assistant turns (dropped/added dashes, single-character substitutions, missing/extra digits within long IDs, altered spacing or capitalisation) typically reflect TTS-rendering or post-hoc-ASR artefacts on a value the agent is reading consistently in audio. Such surface variance is scored by AGENTSPEECHFIDELITY, not here; only structural or semantic discrepancies (different entity, wrong field, wrong category — e.g. addressing the user by an entirely different first name or referencing a different person/record than the tool returned) or downstream signals indicating the agent was operating on a wrong value (subsequent tool calls with a wrong parameter, follow-up actions on stale data, user objections that the agent then fails to incorporate) are flagged on this dimension. In cascade, this note is empty: assistant turns are intended TTS text, so any inconsistency between two assistant utterances of the same fact *is* attributable to the LLM. The voice-context carve-out described above also lands somewhat differently across architectures: in cascade, the trace exposes the actual STT artefacts that justify a re-ask, so the carve-out is concrete; in audio-native pipelines, the trace shows the user simulator’s clean intended text, so re-asks may look less self-evidently justified, and the carve-out instead reminds the judge that mishearings are still possible at the audio layer even when not visible in the trace.

Rating scale. Each dimension is rated on a 3-point integer scale:

- **3** — no issue on this dimension.
- **2** — a single isolated issue that does not significantly impact conversation flow (e.g. one unnecessary tool call that didn’t slow things down, a single redundant restatement, one vague question), or a borderline case where it is unclear whether the issue constitutes a real progression problem.
- **1** — multiple instances of the same type of issue in this dimension, or a single severe issue that clearly derailed or stalled the conversation (e.g. ignoring a stated user constraint before a write operation, failing to ask for required information before taking action, asking an overly vague question when the user’s goal was clear).

Unlike faithfulness, the dimension-level ratings are *not* aggregated by simple minimum, because conversational efficiency is more sensitive to the breadth of issues than to a single worst dimension: a run that has one minor question-quality issue is qualitatively different from a run that has one minor issue in each of three dimensions, even though the per-dimension minimum is the same. The overall rating is therefore:

- **3** if no dimension is flagged (all dimensions rated 3);
- **2** if one or two dimensions are flagged at rating 2 *and* no dimension is rated 1;
- **1** if any dimension is rated 1, *or* if three or more dimensions are flagged — the latter captures the case where issues are individually minor but spread across many areas, which the rubric treats as a clear overall progression problem.

The overall rating is normalized to $[0, 1]$ via $(r - 1)/2$, giving $3 \rightarrow 1.0$, $2 \rightarrow 0.5$, $1 \rightarrow 0.0$. The judge’s per-dimension JSON output and the count of flagged dimensions are preserved in the metric details for inspection.

Pass/fail thresholding. As for faithfulness (Appendix E.4.2), we count a conversation as a conversation-progression pass if its overall rating is ≥ 2 , i.e. its normalized score is ≥ 0.5 . The threshold again sits at the rubric’s own load-bearing boundary: rating 1 is reserved for runs where progression *materially* broke down, while rating 2 covers a small number of isolated minor inefficiencies that did not impede the outcome; conversations at rating 2 are still treated as acceptable.

E.5.2 Conciseness

What it measures. Conciseness is an experience metric that asks whether each assistant turn is appropriately brief and voice-appropriate: a listener consuming the response in real time should be able to absorb its content in a single pass, without filler, without excessive enumeration, and without information density beyond what working memory can comfortably retain. Unlike Faithfulness and Conversation Progression, which produce one rating per conversation, conciseness is rated *per turn*: each assistant’s turn is scored independently, and the per-conversation score is the mean of those per-turn ratings.

Judge and Prompt Template. The judge is run as a single LLM-as-judge call per conversation (one call returning a per-turn array), using *GPT-5.2*. The prompt template is available in Appendix M.5.

Inputs. The judge consumes a deliberately minimal bundle:

- **conversation_turns** — the linearised conversation trace, grouped by **turn_id**, including user, assistant, tool-call, and tool-response entries. The non-assistant entries are provided as context only; the judge rates only assistant content. Multiple assistant entries within a single turn (e.g. a partial response, a tool call, then a continuation) are explicitly evaluated together as a single unit;
- **interruption_tags_reference** — the shared interruption-tag glossary (Appendix M.1), with a strong instruction not to penalise truncated or fragmented content caused by interruptions.

Notably, Conciseness does *not* receive the pipeline-aware **user_turns_disclaimer** or **assistant_turns_disclaimer**, as they don't have a significant impact on Conciseness assessment.

Failure modes. When a turn is rated below 3, the judge tags it with one or more of the following (a turn may carry multiple tags):

1. **verbosity_or_filler** — unnecessary wording, hedging, or repetition within the same turn beyond what the context requires.
2. **excess_information_density** — too many distinct facts, options, numbers, steps, or requests packed into one turn for a listener to retain in real time. Bundling closely related transactional details that the user must act on together (e.g. confirming a reference number, date, and one or two key details in a single turn) is explicitly *not* flagged — only volume that genuinely exceeds working-memory limits.
3. **over_enumeration_or_list_exhaustion** — reading out long lists exhaustively rather than summarising, or presenting multiple options with excessive per-option detail rather than inviting follow-up.
4. **contextually_disproportionate_detail** — more background, clarification, or explanation than the situation actually warrants, given what the user asked for.

The rubric also defines an explicit set of *allowed exceptions* that are never penalised on this metric, even when they produce longer turns: phonetic spell-out of confirmation codes (NATO alphabet) when clarification is needed, full delivery of reference/identifier values the user needs to note down (ticket numbers, voucher codes), and a slightly longer end-of-call recap or wrap-up. Truncated content caused by user or assistant interruptions is also exempt, mirroring the carve-out used by the other judges.

Pipeline-aware adaptations. Conciseness has none. As noted above, the metric only rates assistant content, and the cascade-vs-audio-native distinction in user-side text sourcing (Appendix E.2) is therefore immaterial to its judgements. The interruption-tag carve-out is the only piece of pipeline-derived context the prompt makes use of, and it is shared with the other judges via the same **interruption_tags_reference** fragment.

Rating scale and aggregation. Each assistant-bearing turn is rated on a 3-point integer scale:

- **3** (highly concise) — the response is clear, appropriately scoped for voice, and comfortably digestible in real time. No failure modes are present. A turn that delivers a few closely related facts as part of a single transactional step still qualifies as 3 if a listener can absorb it in one pass.
- **2** (adequate but not optimally concise) — exactly one minor failure mode is present, but the response remains processable in a voice setting and does not meaningfully overwhelm the listener. Reserved for turns where one can identify specific content that should have been omitted or deferred — not merely for turns that happen to contain several necessary details.
- **1** (not concise) — one or more significant failure modes are present that would materially increase cognitive load and hinder comprehension when spoken.

Each per-turn rating is normalised to $[0, 1]$ via $(r-1)/2$, giving $3 \rightarrow 1.0$, $2 \rightarrow 0.5$, $1 \rightarrow 0.0$, and the conversation-level Conciseness score is the *mean* of these per-turn normalised ratings across all rated turns. Unlike

Faithfulness (minimum across dimensions) and Conversation Progression (count-aware aggregation across dimensions), Conciseness is therefore a continuous score in $[0, 1]$ rather than a discrete one in $\{0.0, 0.5, 1.0\}$. Per-turn ratings, failure-mode tags, and per-turn explanations are preserved in the metric details for inspection, and per-failure-mode rates (the fraction of rated turns flagged with each tag) are surfaced as sub-metrics.

Pass/fail thresholding. Following the same convention as Faithfulness and Conversation Progression (Appendices E.4.2, E.5.1), we count a conversation as a Conciseness pass if its conversation-level normalised score is ≥ 0.5 , i.e. if its mean per-turn rating is ≥ 2 . Because Conciseness aggregates by mean rather than by minimum, this threshold has a slightly different rubric-level interpretation: a pass means the assistant was on average at least “adequate but not optimally concise” across the conversation, with isolated rating-1 turns (significant verbosity in a small number of turns) tolerated as long as they are offset by rating-3 turns elsewhere. The threshold is intentionally lenient: verbosity degrades the listening experience but does not cause the material harm that the rating-1 cut captures for Faithfulness and Conversation Progression.

E.5.3 Turn-taking

Unlike prior works [29, 19] that treat turn-taking as a collection of independent flat scalars across all turns — conflating qualitatively distinct events and penalizing tool-call latency equivalently to slow conversational responses — our Turn-Taking metric introduces two key distinctions: (1) **tool-call-aware Turn-Taking evaluation**, which applies more adaptive latency thresholds to turns involving or not involving tool execution, decoupling architectural latency from conversational responsiveness; and (2) a **unified per-turn score** that routes each turn to a semantically appropriate scoring function conditioned on what actually occurred — penalizing agent interruptions based on overlap severity and recovery latency, rewarding immediate agent yield on user interruptions, and scoring uninterrupted turns on a principled response latency curve — producing a single, interpretable score that reflects the full diversity of turn-taking events in task-oriented voice interaction. The metric is deterministically computed from the event timestamps and latencies recorded in the simulation logs.

Unified Per-Turn Scoring Regime

Rather than aggregating flat scalars across all turns, each turn is first classified by its interrupt condition and then routed to a semantically appropriate scoring function. This ensures that qualitatively distinct turn-taking events — agent interruptions, user interruptions, and uninterrupted exchanges — are each evaluated according to the behavioral properties that matter most for that event type. The major event types and their corresponding score functions are summarized in Table 16.

Table 16 Per-turn scoring regime. Each turn is routed to a semantically appropriate scoring function conditioned on the interrupt condition, rather than contributing to a single aggregated scalar.

Turn Condition	Score Function
Agent interrupted user	$s_{\text{agent}} = \min(s_{\text{overlap}}, s_{\text{count}}, s_{\text{post}})$
User interrupted agent	s_{yield}
Both	$\min(s_{\text{agent}}, s_{\text{yield}})$
Uninterrupted	s_{latency}

Uninterrupted Turns For uninterrupted turns, Turn-Taking score is computed from the agent’s response latency ℓ (ms), defined as the elapsed time between the end of the user’s utterance and the onset of the agent’s response. This latency is mapped to a score in $[0, 1]$ via a piecewise-linear curve encoding four regions as detailed below:

- **Hard-zero early** ($\ell \leq \ell_{\text{hard-early}}$): The agent begins speaking before the user finishes their utterance, indicating a premature interruption. 500 ms is set conservatively relative to the perceptual detection threshold for overlapping speech of approximately 120 ms [11], acknowledging that voice agent deployments introduce audio buffering and streaming artifacts that can produce small spurious negative latencies. Score is hard-clamped to 0 regardless of how early the response is.

- **Early ramp** ($\ell_{\text{hard-early}} < \ell \leq \ell_{\text{sweet-low}}$): The response arrives before the natural conversational window but is not severely premature. While the modal inter-turn gap in human conversation falls between 0-200 ms [32, 12], voice agents might be subject to additional processing overhead, making 500 ms the realistic minimum achievable latency of a well-optimized voice agent system. Score ramps linearly from 0 to 1 as latency increases toward the optimal window, penalizing responses proportionally to how early they arrive.
- **Sweet spot** ($\ell_{\text{sweet-low}} < \ell \leq \ell_{\text{sweet-high}}$): The response arrives within the optimal window for natural conversational flow, consistent with psycholinguistic norms for inter-turn gaps [32, 31]. The upper bound of 2,000 ms is motivated by evidence that gaps exceeding 700–1,000 ms are perceived as problematic [30, 16], while telephony contexts tolerate slightly longer delays than face-to-face interaction [31]. Score is flat at 1 across the entire time.
- **Late ramp** ($\ell_{\text{sweet-high}} < \ell \leq \ell_{\text{hard-late}}$): The response begins to feel delayed, introducing noticeable silence that disrupts conversational rhythm. At latencies approaching 2,000 ms, users in voice agent interactions begin to check in or repeat themselves [31]; the ramp to 3,500 ms provides a graduated penalty for responses that are slow but not yet conversation-ending. Score ramps linearly from 1 to 0.
- **Hard-zero late** ($\ell > \ell_{\text{hard-late}}$): The silence is long enough to cause conversational breakdown, likely prompting the user to disengage from the conversation any further [31]. Score is hard-clamped to 0.

Table 17 Latency curve breakpoints by turn type, with descriptions. Tool-call turns receive a more lenient upper threshold ($\ell_{\text{sweet-high}}$ and $\ell_{\text{hard-late}}$) to account for inherent tool execution latency; lower breakpoints are unchanged since early-response behavior is unaffected by tool calls.

Breakpoint	Standard Turn	Tool-Call Turn	Description
$\ell_{\text{hard-early}}$	−500 ms	−500 ms	Response begins before the user finishes speaking; score drops to 0 at or below this threshold.
$\ell_{\text{sweet-low}}$	500 ms	500 ms	Lower bound of the optimal response window; score reaches 1 at this point.
$\ell_{\text{sweet-high}}$	2,000 ms	3,000 ms	Upper bound of the optimal response window; score begins ramping down beyond this threshold.
$\ell_{\text{hard-late}}$	3,500 ms	5,000 ms	Response is considered excessively delayed; score drops to 0 at or beyond this threshold.

The breakpoints defining these regions are detailed in Table 17, which presents both standard and tool-call-aware variants discussed in Section E.5.3.

Agent Interruption Score When the agent interrupts the user, the score is governed by three sub-dimensions that jointly capture the severity and quality of the interruption. All three are capped at $M = 0.5$, reflecting the principle that an interruption is never cost-free regardless of its brevity or recovery quality:

$$s_{\text{agent}} = \min(s_{\text{overlap}}, s_{\text{count}}, s_{\text{post}}) \quad (1)$$

User Interruption Score When the user interrupts the agent, the score measures how promptly the agent yields, rewarding agents that stop speaking immediately. Yield latency Δt is computed across the turn boundary (agent’s last-end in turn $t-1$ minus user’s first-start in turn t), and is linearly penalized up to a maximum tolerable yield duration Δt_{max} , beyond which the score reaches 0:

$$s_{\text{yield}} = \max\left(0, 1 - \frac{\Delta t}{\Delta t_{\text{max}}}\right) \quad (2)$$

where $\Delta t_{\text{max}} = 2,000$ ms defines the upper bound of the acceptable yield latency. An agent that continues speaking beyond this threshold after the user begins is considered to have failed to yield entirely, resulting in

Table 18 Agent interruption sub-scores, each capped at $M = 0.5$. The minimum of the three is taken, ensuring that the weakest dimension dominates — a single poorly-recovered interruption cannot be masked by favorable scores on the other dimensions. $o_{\max} = 2,000$ ms defines the maximum tolerable overlap duration beyond which $s_{\text{overlap}} = 0$; $N_{\max} = 3$ defines the maximum number of distinct overlapping segments beyond which $s_{\text{count}} = 0$.

Sub-score	Definition	Formula
Overlap	Total simultaneous-speech duration o (ms) across all user/agent segment pairs, penalized up to a maximum tolerable overlap duration $o_{\max}$, beyond which the agent is considered to have unacceptably interrupted the user.	$s_{\text{overlap}} = \max\left(0, M\left(1 - \frac{o}{o_{\max}}\right)\right)$
Count	Number of distinct agent segments n overlapping user speech (each with >1 ms intersection), penalized as interruption frequency increases toward the maximum tolerable segment count $N_{\max}$.	$s_{\text{count}} = \max\left(0, M\left(1 - \frac{n-1}{N_{\max}-1}\right)\right), n \geq 1$
Post-interrupt	Silent gap from the end of the user’s last segment to the start of the agent’s next settled response, measuring recovery quality after an interruption. Scored via the latency curve (Table 17). Omitted when no settled response exists or the agent was still speaking at the user’s last segment end.	$s_{\text{post}} = s_{\text{latency}}(\Delta t)$

a score of 0. Unlike the agent interruption score, s_{yield} is not capped at 0.5: an agent that yields immediately receives a perfect score of 1, reflecting that deferring to the user is always the correct behavior when interrupted. Definition and formula of each sub-score (overlap, count, post-interrupt) are provided in Table 18.

Tool-Call-Aware Turn-taking Score

A uniform latency threshold treats all response delays equally — yet in task-oriented voice agents, a delay caused by tool execution is fundamentally different from a delay caused by slow conversational processing. Applying the same thresholds to both would unfairly penalize agents for latency they cannot avoid, conflating infrastructure overhead with genuine conversational responsiveness. To address this, we introduce tool-call-aware latency thresholds: turns in which the agent issued a tool call before responding receive more lenient upper breakpoints ($\ell_{\text{sweet-high}}$ and $\ell_{\text{hard-late}}$), extending both the optimal response window and the tolerable delay ceiling to absorb the cost of tool execution. The lower breakpoints ($\ell_{\text{hard-early}}$ and $\ell_{\text{sweet-low}}$) remain unchanged, since early-response behavior is independent of whether a tool call was issued. This decoupling ensures that the Turn-Taking score reflects genuine conversational responsiveness, enabling fair and meaningful comparison across systems with different tool execution speeds and architectures. The full set of breakpoints for both standard and tool-call turns is provided in Table 17.

Pass/fail thresholding.

The EVA-X composite admits a trial only if its Turn-Taking score clears a pass-threshold τ_{tt} , set to 0.8 in this work. Unlike other metrics—e.g., Faithfulness and Conciseness, which are judged on coarse ordinal scales whose thresholds follow directly from their rubric definitions—turn-taking scores are continuous and lack such an intrinsic cutoff. We therefore justify $\tau_{\text{tt}} = 0.8$ on three grounds: its conversational interpretation, its calibration to current model capability, and the empirical finding that EVA-X model rankings remain mostly stable across threshold values.

Conversational interpretation: $0.8 \approx 4/5$ on-time turns. The per-record Turn-Taking score is the mean of per-turn scores in $[0, 1]$, where each turn earns a score of 1 when its latency is acceptable and a correspondingly

lower score outside it. A threshold of 0.8 thus admits at most one off-bracket turn per five—a concrete, conversation-level reading.

The on-time bracket is already lenient. Natural human turn-taking typically centers around 200ms [32, 12], yet EVA’s latency thresholds are deliberately permissive relative to this norm as discussed in Section E.5.3). Because the per-turn turn-taking scoring is already permissive about which turns count as on-time, a low τ_{tt} would make the metric trivially easy to pass, compounding the two layers of leniency and reducing the metric’s ability to discriminate between systems at the higher end of the performance range.

Calibration to current turn-taking capability. Among the 12 systems benchmarked here (Tab. 25), a wide gap separates systems that achieve real-time turn-taking from those that do not. The three S2S systems score 0.815–0.830; the next-highest system scores 0.583, a gap of more than 0.23. Any threshold in the range [0.6, 0.8] would likely therefore produce the same binary partition of systems—the precise value does not determine which systems pass or fail given the current landscape.

EVA-X model rankings remain mostly stable across τ_{tt} . We sweep $\tau_{tt} \in \{0.50, 0.55, \dots, 0.95\}$ and recompute EVA-X pass@1 at each value (Fig. 4). System rankings are preserved across the entire range: Spearman $\rho = 0.968$ between rankings at $\tau_{tt} = 0.5$ and $\tau_{tt} = 0.95$, and 61/66 (92.4%) of pairwise system orderings hold strictly across all ten thresholds (63/66, 95.5%, when statistically tied pairs are counted as preserved). The three pairs that flip somewhere in the sweep are systems that score within 0.15 pass@1 of each other.

Linear stability of the per-system pass@1 *vector* around the production anchor $\tau_{tt} = 0.8$ is correspondingly high: as shown in Table 19, Pearson $r = 0.998$ at $\tau_{tt} = 0.75$ and $r = 0.995$ at $\tau_{tt} = 0.85$; r never falls below 0.910 across the entire [0.50, 0.95] range (minimum at $\tau_{tt} = 0.50$).

Table 19 Pearson correlation between per-system EVA-X pass@1 vectors at $\tau_{tt} = 0.8$ and other candidate thresholds ($n = 12$ systems).

τ_{tt}	0.50	0.55	0.60	0.65	0.70	0.75	0.80	0.85	0.90	0.95
Pearson r vs. $\tau_{tt} = 0.80$	0.910	0.947	0.967	0.979	0.991	0.998	1.000 (anchor)	0.995	0.980	0.950

Why 0.8? The sensitivity sweep confirms that neither plateau nor cliff appears at any candidate threshold. More importantly, the 0.23-point gap between the highest non-S2S system and the lowest S2S system means that any $\tau_{tt} \in [0.6, 0.8]$ yields similar pass/fail assignments for all 12 systems. We select the top of this range for two reasons: (i) it maps cleanly to the “4/5 turns on-time” interpretation, and (ii) it avoids compounding two layers of leniency, given that the on-time bracket itself already extends beyond natural human turn-taking latency. A higher threshold (e.g., 0.9) is currently miscalibrated to model capability—no system in our pool would reliably qualify.

Modularity and forward compatibility. EVA exposes `pass_at_k_threshold` as a configurable metric parameter, and practitioners benchmarking different system classes are encouraged to override the default. The 0.8 value reflects 2026 model capability; we expect this default to rise as model latencies improve. The defense above is therefore a statement about the *current* calibration, not a fixed property of the metric.

E.6 Diagnostic Metrics

Diagnostic metrics include: *Authentication Success Rate* (deterministic), *Response Latency* in seconds decomposed into turns with and without tool calls or not (deterministic), *speakability* — whether agent text is voice-friendly prior to synthesis (LLM-as-Judge), STT word error rate computed via `jiwer` (deterministic), tool call validity — the fraction of tool calls with correctly formatted parameters (deterministic), and transcription accuracy for key entities — STT accuracy specifically on named entities such as names, dates, and alphanumeric codes (LLM-as-Judge). Together these metrics allow practitioners to distinguish, for example, between a task completion failure caused by an STT transcription error on a confirmation code versus one caused by incorrect LLM reasoning.




Figure 4 Sensitivity of EVA-X pass@1 to the turn-taking pass-threshold τ_{tt} . We sweep τ_{tt} from 0.50 to 0.95 in 0.05 increments and recompute EVA-X pass@1 at each value, holding the Conversation Progression and Conciseness thresholds fixed at 0.5. The dashed vertical line marks the pass threshold value $\tau_{tt} = 0.8$. (a) Architecture-aggregated. Each curve is the mean per-scenario EVA-X pass@1 (averaged over $k = 5$ trials per scenario), pooled with equal weight across all (system, scenario) pairs within an architecture class (cascade, hybrid, S2S); shaded bands are percentile bootstrap 95% CIs from 1,000 resamples of (system, scenario) pairs. (b) Per-system. One curve per benchmarked system, colored by architecture; shaded bands are bootstrap 95% CIs from 1,000 resamples of scenarios within the system. Architecture ordering (S2S > cascade > hybrid) is preserved at every threshold; per-system rankings flip only between close competitors.




Figure 5 Mean transcription accuracy on key entities per cascade STT model, pooled across the three EVA-Bench domains. Each bar is the mean of per-scenario accuracies (one value per scenario, averaged over $k=5$ trials); whiskers are 95% normal-approximation CIs on that mean. Models are sorted ascending. Nova-3 aggregates the two cascades that share it (Nova-3 + GPT-5.4 + Sonic 3 and Nova-3 + GPT-5.4-mini + Aura-2), giving it twice as many scenario observations and a correspondingly tighter interval. Scribe-v2.2-Realtime saturates near 1.0 on this benchmark; the gap to the next-best STT (Parakeet-1.1) exceeds 0.19.

F Metrics Analysis

F.1 Key Entity Transcription Accuracy & Task Completion

The domains in EVA-Bench are entity-dense — agents must extract confirmation codes, names, authentication codes, employee IDs, and similar information not only to authenticate users, but also to complete downstream task steps (e.g., looking up a ticket, updating a record, or routing a request). This explains the correlation shown in Figure 6: since authentication is a prerequisite for nearly every task, and both authentication and task execution depend on accurate entity transcription, transcription accuracy and task completion are tightly coupled. This relationship may not generalize to domains with fewer key entities or where task success is independent of entity transcription. That said, we consider entity-heavy flows a representative and common voice agent use case, which is why our tasks and domains are designed to stress-test this capability. Transcription accuracy on entities is not sufficient on its own to predict task completion, however. Among the two systems using Nova-3 as the STT model, transcription scores are similar, yet the system backed by GPT-5.4 achieves meaningfully higher task completion than the one using GPT-5.4-mini— indicating that LLM reasoning over transcribed entities also plays a role.

F.2 Faithfulness & Task Completion

Table 20 Joint distribution of task completion and faithfulness across 12,780 clean trials from all evaluated systems. Cells show % of total; shading scales with cell count (darker = more trials).

	faithfulness = 1.0	faithfulness < 1.0	Row total
task_completion = 0	9.5%	38.3%	47.9%
task_completion = 1	14.5%	37.6%	52.1%
Column Total	24.0%	76.0%	12,780

Table 20 shows the confusion matrix between Task Completion and Faithfulness (score of 1.0 vs. below 1.0) across all evaluated systems under clean conditions. Only 14.5% of trials achieve both, underscoring that Task Completion alone is a weak proxy for overall agent accuracy. Faithfulness failures are distributed nearly uniformly across task outcomes (38.3% vs. 37.6%), reinforcing that the two dimensions capture distinct aspects of agent behavior.

Cascade systems: transcription vs task completion

Pearson $r = 0.930$, $p = 0.002$, $n = 7$ systems




Figure 6 Mean key entity transcription accuracy and mean task completion correlation across the 7 evaluated cascade systems.

F.3 Patterns Across Pipeline Architectures

The S2S systems we evaluated tend to lead cascade on responsiveness facets — though by a narrow margin against the strongest cascades — and trail on policy adherence. We focus on four facets that surface these patterns most clearly: on-time turns ($\uparrow$, a sub-metric of turn-taking), conversation completion ($\uparrow$), authentication success ($\uparrow$), and policy violations ($\downarrow$, derived from the faithfulness judge).

- On-time rate ($\uparrow$) — a sub-metric of the Turn-Taking score (App. E.5.3). We report it directly because its definition is unambiguous: the share of agent turns whose response latency falls in $[200, 4000)$ ms, or $[200, 6000)$ ms when the turn involves a tool call. Note that the boundaries are not exactly the same as in the turn-taking definition, given that here we don’t have the piecewise-linear curves.
- Conversation Completion ($\uparrow$) — captures the failure mode where the agent fails to respond to a user turn, leading the conversation to time out.
- Authentication success ($\uparrow$) — the proportion of conversations in which the agent successfully authenticated the user before proceeding to the main task. This facet is informative because it probes how well each architecture handles entity-dense input — IDs, names, dates of birth — without requiring a transcript-level evaluation.




Figure 7 Per-pipeline-type means on the four facets, computed on clean runs only. Error bars are 95% percentile bootstrap intervals over systems-within-type, shown as descriptive uncertainty bands; faint dots are per-system means. Hybrid ($n = 2$) is shown as individual system points only.

- Policy violations ($\downarrow$) — the dominant failure mode contributing to low faithfulness scores, since policy adherence requires strict instruction-following. Across all pipelines, 61% of conversations are flagged with at least one policy violation. Common cases include the agent performing write actions without explicit user authorization, or fabricating policies absent from its instructions.

For each facet, we compute a per-system mean by averaging across that system’s per-domain values. The pipeline-type mean is the unweighted mean of its constituent system means, and its 95% percentile bootstrap interval is obtained by resampling those system means with replacement (10,000 iterations, fixed seed). Pairwise effect sizes $\Delta = \mu_A - \mu_B$ between pipeline types are reported with a percentile bootstrap interval obtained by independently resampling each group’s per-system means. Differences are in percentage points ($\Delta \times 100$); metrics are bounded in $[0, 1]$. All numbers are computed on the clean (unperturbed) runs of the systems described in the main paper: 7 cascade, 3 S2S, and 2 hybrid, evaluated on all three domains, yielding 1,065 clean conversations per system.

Given the small number of systems per pipeline type, the analysis is descriptive: bootstrap intervals are reported as uncertainty bands, not as confidence intervals in the inferential sense. We do not claim these patterns generalize to systems beyond those we evaluated; we report them as observations that may motivate further study. For some facets, the difference visible at the pipeline-type level conceals notable within-class exceptions, particularly within cascade. We also refrain from drawing pipeline-type conclusions about hybrid systems: the two we evaluated diverge substantially from one another. We also performed the analysis stratified by domain (omitted here for space); the ordering of the three pipeline-type means is preserved in every domain.

F.3.1 Observations

Table 21 reports the per-pipeline-type means with their 95% bootstrap intervals, and the S2S–Cascade pairwise difference (in percentage points) with its bootstrap interval as a bottom row; Fig. 7 visualizes the means with per-system points overlaid. The S2S systems sit at the top of all facets, including the policy-violation facet (where lower is better). However, there are some differences at the system level, as shown in Table 22, and the per-facet patterns below describe both where S2S systems excel or trail, and

Table 21 Per-pipeline-type means with 95% bootstrap intervals over systems-within-type, and the S2S–Cascade pairwise difference (Δ) with its 95% bootstrap interval. Computed on clean runs only. Higher is better for the first three facets; lower is better for `violating_policies_rate`. n is the number of systems contributing to each row. Mean rows use the $[0, 1]$ metric scale; the Δ row is reported in percentage points ($\Delta \times 100$), with its bootstrap interval obtained by independently resampling each group’s per-system means.

Pipeline type	n	Facet (mean [95% interval])			
		On-time $\uparrow$	Finishes $\uparrow$	Auth. $\uparrow$	Violates policy $\downarrow$
Cascade	7	0.55 [0.44, 0.66]	0.82 [0.74, 0.93]	0.77 [0.66, 0.86]	0.53 [0.41, 0.65]
Hybrid	2	0.23 [0.01, 0.46]	0.89 [0.87, 0.90]	0.88 [0.84, 0.92]	0.63 [0.56, 0.70]
Speech-to-Speech	3	0.83 [0.80, 0.89]	0.97 [0.95, 0.99]	0.88 [0.83, 0.93]	0.78 [0.69, 0.88]
Δ S2S–Cascade (pp)	—	+27.9 [+16.4, +39.6]	+15.2 [+6.0, +24.5]	+11.3 [+2.4, +20.9]	+24.6 [+11.7, +38.4]

where individual cascades narrow or close the gap.

- On-time rate ($\uparrow$) — all three S2S systems we evaluated sit above every cascade in our sample, but the gap is narrow for the strongest cascades: *Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS* (0.76) and *Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS* (0.70) come within 4–10 pp of the lowest S2S system (0.80). The remaining five cascades trail more substantially (0.33–0.66).
- Conversation Completion ($\uparrow$) — none of the S2S pipelines we evaluated struggled with completing conversations, whereas cascade is more variable. Only two cascade systems reached S2S-comparable values, *Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha* and *Parakeet-1.1 + Gemma-4-31B + Kokoro*. Section F.3.3 examines the likely causes of agent-side failures, per pipeline.
- Authentication success ($\uparrow$) — on raw scores, *GPT-Realtime-1.5* (0.93) and *Gemini-3.1-Flash-Live* (0.88) lead, while *GPT-Realtime-mini* (0.83) sits in the cascade mid-range, outperformed by *Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha* (0.89), *Parakeet-1.1 + Gemma-4-31B + Kokoro* (0.87), and *Nova-3 + GPT-5.4 + Sonic 3* (0.85). Much of the apparent cascade-vs-S2S gap on this facet is a completion artifact, however: when restricted to conversations that finished, several additional cascades reach the S2S range (*Nova-3 + GPT-5.4 + Sonic 3* rises to 0.94, matching *GPT-Realtime-1.5*). See Sec. F.3.2.
- Policy violations ($\downarrow$) — the most variable facet within both pipeline classes. *GPT-Realtime-1.5* (0.69, the lowest S2S violation rate) lands in the same range as several cascades (e.g., *Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS* at 0.60, *Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha* and *Parakeet-1.1 + Gemma-4-31B + Kokoro* at 0.61); the other two S2S systems sit higher. The lowest cascade violation rate comes from *Nova-3 + GPT-5.4 + Sonic 3* (0.28), and remains the lowest after conditioning on completed conversations (0.32). Conditioning matters here for the cascades that hang often: *Ink-whisper + Haiku-4.5 + Sonic 3* rises from 0.48 to 0.63 (+15 pp), revealing that its low raw rate partly reflects fewer opportunities to violate rather than stronger instruction-following. No S2S system reaches the conditional violation rate of *Nova-3 + GPT-5.4 + Sonic 3*, suggesting a tension between latency and instruction-following accuracy that none of the systems we evaluated fully resolves.

F.3.2 Authentication: completion vs. ability

Raw `authentication_success` mixes two distinct phenomena: how often the agent completes the authentication exchange without hanging, and how often it correctly authenticates the user when it does. Cascade pipelines frequently hang *during* the authentication flow — typically when the user reads out an ID or a short utterance that the pipeline mishandles. To separate these, Table 23 reports each system’s authentication rate restricted to conversations that finished correctly.

S2S systems are essentially unchanged: they finish 95–99% of clean conversations, so the raw and conditional rates coincide ($\Delta \leq 0.8$ pp). Several cascades, by contrast, move substantially: *Nova-3 + GPT-5.4 + Sonic 3* rises from 0.85 to 0.94, matching the best S2S (*GPT-Realtime-1.5* at 0.94); *Ink-whisper + Haiku-4.5 + Sonic 3* from 0.65 to 0.86 (+20 pp); and *Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS* from 0.60 to 0.73 (+13 pp). The cascade authentication deficit visible in Table 21 is therefore in large part a completion

Table 22 Per-system means on clean runs, averaged across the three domains.

Pipeline	System	On-time $\uparrow$	Finishes $\uparrow$	Auth. $\uparrow$	Violates $\downarrow$
Cascade	Nova-3 + GPT-5.4 + Sonic 3	0.48	0.66	0.85	0.28
	Ink-whisper + Haiku-4.5 + Sonic 3	0.44	0.67	0.65	0.48
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.76	0.75	0.60	0.42
	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.70	0.79	0.68	0.60
	Nova-3 + GPT-5.4-mini + Aura-2	0.66	0.90	0.83	0.73
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.49	0.97	0.89	0.61
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.33	0.97	0.87	0.61
Hybrid	Ultravox-Realtime	0.46	0.87	0.84	0.70
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.01	0.90	0.92	0.56
Speech-to-Speech	GPT-Realtime-1.5	0.80	0.99	0.93	0.69
	GPT-Realtime-mini	0.80	0.97	0.83	0.88
	Gemini-3.1-Flash-Live	0.89	0.95	0.88	0.78

artifact: when a cascade gets through the authentication exchange without timing out, several configurations authenticate at rates comparable to S2S systems. Among the systems we evaluated, *GPT-Realtime-1.5* stands out for combining both axes: it finishes 99% of clean conversations and authenticates correctly 93% of the time when it does, the highest values in the table on both axes.

Table 23 Per-system authentication success on clean runs, raw vs. conditional on conversation completion. *Auth (raw)* is computed over all 1,050 clean conversations; *Auth | finished* is computed only on conversations where `conversation_correctly_finished` = 1. Δ (pp) is their difference. Systems with low finish rates move the most.

Pipeline	System	Finish rate	Auth (raw)	Auth finished	Δ (pp)
Cascade	Nova-3 + GPT-5.4 + Sonic 3	0.66	0.85	0.94	+8.8
	Ink-whisper + Haiku-4.5 + Sonic 3	0.67	0.65	0.86	+20.5
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.75	0.60	0.73	+12.6
	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.79	0.68	0.73	+4.9
	Nova-3 + GPT-5.4-mini + Aura-2	0.90	0.83	0.86	+2.9
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.97	0.89	0.90	+1.1
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.97	0.87	0.88	+0.3
Hybrid	Ultravox-Realtime	0.87	0.84	0.86	+2.0
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.90	0.92	0.94	+2.4
Speech-to-Speech	GPT-Realtime-1.5	0.99	0.93	0.94	+0.5
	GPT-Realtime-mini	0.97	0.83	0.84	+0.8
	Gemini-3.1-Flash-Live	0.95	0.88	0.88	+0.7

F.3.3 Inactivity-timeout failures

To better understand why conversations do not always finish correctly, we examined every conversation that ended with an agent-side inactivity timeout. We classify each such conversation by the immediately preceding user turn: *short turn* (under 5 words), *short $\wedge$ confirm* (short turn opening with *yes/no/sure/ok*), and *spelled content* (the user turn contained NATO phonetic words, isolated letters or digits read out one-by-one, alphanumeric codes, or numbers spoken as words). A turn can match more than one pattern.

Table 24 breaks the timeouts down per system. Five of the seven cascade systems show short-turn shares of 52–68% of their timeouts; a substantial fraction of those short turns also begin with a confirmation token (e.g., “Yes, that’s correct”), after which the agent fails to produce a follow-up turn. Several pipelines do not exhibit this pattern: *Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha*, *Parakeet-1.1 + Gemma-4-31B + Kokoro*, *Ultravox-Realtime*, and all three S2S systems. With the exception of *Ultravox-Realtime*, all of these have an overall conversation failure rate below 5%. Spelled-content failures, by contrast, persist across all pipelines and are the dominant timeout cause among the systems that handle short turns well. These spelled-content hangs occur predominantly during the authentication exchange, which is also why the cascade authentication deficit shrinks substantially under the completion-conditional view (Sec. F.3.2).

Table 24 Per-system view on clean runs (1,065 conversations per system). *Conv. fail rate* = $1 - \text{conversation_correctly_finished}$ (any failure type, conversation-pooled). *Timeouts* = number of those failures that ended with the agent going silent. The three right-most columns are the share of *this system’s* timeouts whose preceding user turn matched each pattern. Systems are ordered within each pipeline class by timeout count (descending).

Pipeline	System	Conv. fail rate	Timeouts	Short turn	Short $\wedge$ confirm	Spelled content
Cascade	Ink-whisper + Haiku-4.5 + Sonic 3	36.5%	395	63%	43%	17%
	Nova-3 + GPT-5.4 + Sonic 3	35.8%	381	68%	39%	12%
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	28.2%	300	62%	52%	14%
	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	22.6%	241	56%	37%	21%
	Nova-3 + GPT-5.4-mini + Aura-2	10.7%	117	52%	20%	24%
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	2.6%	28	0%	0%	68%
	Parakeet-1.1 + Gemma-4-31B + Kokoro	2.5%	27	19%	19%	48%
Hybrid	Ultravox-Realtime	14.0%	149	6%	1%	63%
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	10.5%	112	30%	14%	46%
S2S	Gemini-3.1-Flash-Live	4.5%	52	2%	0%	58%
	GPT-Realtime-mini	2.9%	32	0%	0%	66%
	GPT-Realtime-1.5	1.4%	15	0%	0%	73%

F.4 Speech Fidelity

Across all models, **entity mispronunciation** (letter substitutions, digit omissions, spurious insertions, and phonetic confusions between similar-sounding characters) is by far the dominant failure class, accounting for the majority of flagged turns.

Substitutions. In a single alphanumeric code, one character is replaced by a visually or acoustically unrelated one. *Sonic 3* consistently rendered the airport code LAX as LEX (CSM, record 1.2.3) and substituted S C 3 for A C 3 in an ITSM request prefix across multiple turns (record 16). *Aura-2* substituted L for I in a medical license number, saying L 1 C instead of L I C (HRSD, record A6). *Ultravox-Realtime* misread the confirmation code ZKLX8E as ZKLXIE, substituting 8 for I (CSM, record 7.2.9), and rendered the hotel voucher prefix GQSIHM as DQSIHM, substituting D for G (record 7.3.1).

Omissions. One or more characters are silently dropped from an entity. *GPT-Realtime-1.5* read the 10-digit NPI 3342331444 as only 9 digits, dropping the trailing 4 (HRSD, record 1.1), and omitted the 6 from registration ID 358607, producing 35807 (record 9.1). *Kokoro* dropped the A from the spelled-out code MEAL across multiple trials, producing MEL (CSM, record 2.2.5), and omitted the letter K from flight number SK915, rendering it as S 915 consistently across three consecutive turns (record 2.2.2). *Aura-2* omitted one of the five repeated digits, saying EMP-0-5-5-5-5 instead of EMP-0-5-5-5-5-5 (ITSM, record 22).

Insertions. Spurious characters or words are injected into an entity that should be reproduced verbatim. *Gemini-3.1-Flash-Live* inserted the name Pavel into the middle of the visa petition number EH23328710672, breaking the code entirely (HRSD, record 11.2). *Ultravox-Realtime* added a spurious E to the meal voucher prefix in multiple trials, saying MEAL-7EMMHTS-PAX0 instead of MEAL-7MMHTS-PAX0 with the E after the 7 (CSM, records 2.3.2). *Voxtral-4B-TTS* appended a trailing 3 to the request identifier REQ-SW-9befdac7c2e6, producing REQ-SW-9befdac7c2e63 (ITSM, record 76).

Phonetic confusions. A character is replaced by one that sounds similar when spoken aloud. *GPT-Realtime-1.5* repeatedly confused C and P, rendering the facility code IXC as IXP across two consecutive turns (HRSD, record D3.3), and swapped D for B in the policy number PDZP6L $\rightarrow$ PBZP6L (record T2.1). *Gemini-3.1-Flash-Live* confused Z and V, articulating the DEA prefix ZS as VS (HRSD, record A2). *Kokoro* persistently mispronounced SK130 as S-Cone 130 across two turns (record 4.2.4).

F.5 Turn-Taking & Response Speed

Table 25 reports per-system latency and turn-taking sub-metrics that feed into the aggregate turn-taking score, which in turn contributes to EVA-X pass@1. The early, on-time, and late columns give the share of agent turns falling into each timing bucket and sum to 1.0 per system; the response-speed columns report mean per-turn latency in seconds, broken down by whether the turn involved a tool call. A turn is classified

as early when latency < 200 ms, late when latency ≥ 2.75 s (or ≥ 4 s if the turn involves a tool call), and on-time otherwise. Across systems the turn-taking score appears to track on-time and late rates more closely than early rate, which stays small for nearly every system — the largest observed early rate is 0.148 (*GPT-Realtime-mini*), and most systems sit below 0.06. Comparing response speed with vs. without tool calls, every system in the table shows a tool-call latency increase (no overlapping CIs) so the slowdown is consistent. The magnitude varies substantially, however: from roughly 0.1–0.2 s for the fastest cascades (*Whisper-Large-v3* + *Qwen3.5-27B* + *Voxtral-4B-TTS*; *Ink-whisper* + *Haiku-4.5* + *Sonic 3*) up to several seconds for *Ultravox-Realtime* and *Gemini-3-Flash* + *Gemini-3.1-Flash-TTS*, indicating that tool-call overhead is system-specific rather than a uniform property of the benchmark. These scores help illustrate what drives the difference in EVA-X pass@1 scores between cascade and S2S models.

Table 25 Per-system turn-taking and response-speed detail under clean-audio conditions, pooled equal-weighted across the three domains. Each cell shows the pooled point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). Turn-taking rates are in $[0, 1]$; response speed is mean per-turn latency in seconds. Cells are shaded per column with darker = better (lower for latency and error-rate columns; higher otherwise).

Arch.	System	EVA-X	Turn-Taking				Response Speed (s)			
		pass@1	Score	Early	On-time	Late	Interrupt.	Overall	w/ tools	w/o tools
Cascade	Cohere + Gemma-4-26B + Voxtral	0.209 ±0.027	0.567 ±0.024	0.113 ±0.009	0.703 ±0.015	0.184 ±0.014	0.110 ±0.009	2.415 ±0.107	2.660 ±0.253	2.171 ±0.077
	Scribe + Gemini-3-Flash + Conversational v3	0.024 ±0.018	0.451 ±0.019	0.008 ±0.002	0.489 ±0.022	0.504 ±0.022	0.006 ±0.002	4.158 ±0.118	5.294 ±0.166	2.420 ±0.119
	Ink-whisper + Haiku-4.5 + Sonic 3	0.009 ±0.006	0.312 ±0.019	0.024 ±0.004	0.440 ±0.015	0.536 ±0.016	0.023 ±0.004	3.399 ±0.045	3.535 ±0.081	3.354 ±0.060
	Nova-3 + GPT-5.4 + Sonic 3	0.007 ±0.006	0.283 ±0.021	0.029 ±0.005	0.484 ±0.017	0.487 ±0.018	0.028 ±0.005	3.943 ±0.112	5.191 ±0.189	2.914 ±0.078
	Nova-3 + GPT-5.4-mini + Aura-2	0.113 ±0.021	0.583 ±0.020	0.042 ±0.006	0.658 ±0.016	0.299 ±0.016	0.041 ±0.006	3.148 ±0.074	3.648 ±0.095	2.377 ±0.113
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.010 ±0.009	0.308 ±0.015	0.030 ±0.005	0.327 ±0.017	0.643 ±0.018	0.030 ±0.005	4.832 ±0.114	5.373 ±0.128	4.291 ±0.216
	Whisper + Qwen3.5-27B + Voxtral	0.273 ±0.035	0.561 ±0.029	0.056 ±0.008	0.759 ±0.015	0.186 ±0.012	0.054 ±0.008	2.251 ±0.050	2.309 ±0.069	2.164 ±0.067
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000 ±0.000	0.019 ±0.003	0.017 ±0.004	0.010 ±0.003	0.973 ±0.005	0.017 ±0.004	7.466 ±0.257	9.100 ±0.413	5.244 ±0.189
	Ultravox-Realtime	0.029 ±0.019	0.417 ±0.021	0.040 ±0.007	0.457 ±0.019	0.503 ±0.020	0.040 ±0.007	4.838 ±0.225	6.994 ±0.301	1.393 ±0.063
S2S	Gemini-3.1-Flash-Live	0.589 ±0.034	0.830 ±0.017	0.011 ±0.005	0.886 ±0.012	0.104 ±0.012	0.011 ±0.005	2.288 ±0.070	2.846 ±0.088	1.288 ±0.086
	GPT-Realtime-1.5	0.566 ±0.040	0.815 ±0.013	0.121 ±0.012	0.799 ±0.016	0.079 ±0.009	0.072 ±0.011	1.798 ±0.061	2.395 ±0.089	1.011 ±0.031
	GPT-Realtime-mini	0.406 ±0.035	0.818 ±0.015	0.148 ±0.013	0.805 ±0.014	0.047 ±0.006	0.106 ±0.011	1.481 ±0.045	1.868 ±0.063	1.067 ±0.040

F.6 Cross Domain Variability

Tables 29 and 33 report, for every (system, metric) pair, the sample standard deviation (ddof=1) of the per-domain point estimates across the three EVA domains (CSM, ITSM, HR). These complement Tables 26–31 (per-domain values) and the pooled tables 2, by quantifying how much of the pooled headline number is hiding domain dispersion.

F.6.1 Metric-level findings

Speech rendering is essentially domain-invariant. Across all twelve systems, the cross-domain standard deviation of Speech Fidelity is between 0.002 and 0.013 (mean 0.007), against a point-estimate range of 0.913–0.996. No other metric comes within an order of magnitude of this stability. This is consistent with the fact that this metric is determined by the TTS/S2S audio path, which depends less on conversational content.

Within accuracy, Faithfulness is the most domain-coupled metrics. Mean standard deviations across systems are 0.087 for Faithfulness and 0.064 for Task Completion. Faithfulness shows the largest single standard deviation in either table (Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha : 0.184; Ink-whisper + Haiku-4.5 + Sonic 3 : 0.166).

Within experience, Conversation Progression is the noisiest metric and Conciseness the most stable. Mean standard deviations are 0.052 (Conversation Progression) versus 0.020 (Conciseness); Conversation Progression peaks at 0.125 (Ultravox-Realtime) and conciseness never exceeds 0.043 (Ultravox-Realtime). Turn-taking sits between them at mean SD 0.042. The pattern suggests that whether a system is terse is largely a system property, while whether it makes progress effectively depends on the conversational structure of the domain. EVA-X pass-rates are uniformly less variable across domains than EVA-A pass-rates. Mean standard deviations on EVA-X are 0.029 / 0.037 / 0.018 for pass@1/pass@k/ pass[^]k, versus 0.074 / 0.101 / 0.050 on EVA-A — a roughly 2.5 $\times$ gap on every variant. Even controlling for the floor effect by restricting to systems with non-zero EVA-X performance, EVA-A standard deviations remain the larger of the two. Domain choice perturbs task accuracy substantially more than it perturbs the experience overlay.

F.6.2 Model-level findings

Some systems pair high pooled scores with large task-completion swings, which the pooled tables conceal. *GPT-Realtime-1.5* reports a pooled task-completion mean of 0.739 ± 0.045 in Table 2, but its per-domain SD on the same metric is 0.174 — by far the largest task-completion SD in either table — meaning the pooled mean averages over substantial across-domain variation. *Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha* shows the same effect on faithfulness (pooled 0.457; standard deviation 0.184).

F.6.3 Accuracy and Experience results per domain

Accuracy results are provided for all systems per domain: CSM in Table 26, HR in Table 27, and ITSM in Table 28. Cross-domain variability of accuracy metrics are provided in Table 29.

Table 26 CSM domain accuracy metrics for all evaluated systems under clean-audio conditions. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass^k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-A			Task Completion	Faithfulness	Speech Fidelity
		pass@1	pass@k	pass^k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.246 $\pm$ 0.084	0.500 $\pm$ 0.140	0.066 $\pm$ 0.061	0.368 $\pm$ 0.092	0.398 $\pm$ 0.070	0.989 $\pm$ 0.006
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.656 $\pm$ 0.092	0.900 $\pm$ 0.080	0.385 $\pm$ 0.122	0.808 $\pm$ 0.080	0.666 $\pm$ 0.088	0.981 $\pm$ 0.012
	Ink-whisper + Haiku-4.5 + Sonic 3	0.272 $\pm$ 0.076	0.640 $\pm$ 0.140	0.065 $\pm$ 0.064	0.440 $\pm$ 0.092	0.374 $\pm$ 0.068	0.989 $\pm$ 0.004
	Nova-3 + GPT-5.4 + Sonic 3	0.628 $\pm$ 0.076	0.940 $\pm$ 0.060	0.278 $\pm$ 0.108	0.732 $\pm$ 0.068	0.784 $\pm$ 0.050	0.996 $\pm$ 0.003
	Nova-3 + GPT-5.4-mini + Aura-2	0.216 $\pm$ 0.096	0.460 $\pm$ 0.140	0.098 $\pm$ 0.081	0.456 $\pm$ 0.092	0.298 $\pm$ 0.078	0.979 $\pm$ 0.009
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.540 $\pm$ 0.096	0.880 $\pm$ 0.100	0.294 $\pm$ 0.111	0.672 $\pm$ 0.096	0.600 $\pm$ 0.076	0.966 $\pm$ 0.019
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.272 $\pm$ 0.071	0.680 $\pm$ 0.140	0.045 $\pm$ 0.049	0.504 $\pm$ 0.096	0.470 $\pm$ 0.072	0.918 $\pm$ 0.017
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.488 $\pm$ 0.092	0.840 $\pm$ 0.100	0.188 $\pm$ 0.089	0.696 $\pm$ 0.088	0.478 $\pm$ 0.080	0.965 $\pm$ 0.012
	Ultravox-Realtime	0.324 $\pm$ 0.100	0.540 $\pm$ 0.140	0.131 $\pm$ 0.081	0.428 $\pm$ 0.104	0.334 $\pm$ 0.076	0.971 $\pm$ 0.015
S2S	Gemini-3.1-Flash-Live	0.356 $\pm$ 0.104	0.660 $\pm$ 0.140	0.170 $\pm$ 0.102	0.504 $\pm$ 0.096	0.342 $\pm$ 0.084	1.000 $\pm$ 0.000
	GPT-Realtime-1.5	0.424 $\pm$ 0.116	0.640 $\pm$ 0.140	0.271 $\pm$ 0.120	0.540 $\pm$ 0.112	0.424 $\pm$ 0.092	0.998 $\pm$ 0.003
	GPT-Realtime-mini	0.176 $\pm$ 0.092	0.300 $\pm$ 0.120	0.085 $\pm$ 0.075	0.288 $\pm$ 0.104	0.164 $\pm$ 0.076	0.971 $\pm$ 0.034

Table 27 HR domain accuracy metrics for all evaluated systems under clean-audio conditions. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass^k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-A			Task Completion	Faithfulness	Speech Fidelity
		pass@1	pass@k	pass^k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.229 $\pm$ 0.070	0.386 $\pm$ 0.108	0.087 $\pm$ 0.049	0.318 $\pm$ 0.089	0.408 $\pm$ 0.055	0.985 $\pm$ 0.005
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.373 $\pm$ 0.082	0.614 $\pm$ 0.108	0.199 $\pm$ 0.074	0.655 $\pm$ 0.080	0.320 $\pm$ 0.066	0.973 $\pm$ 0.009
	Ink-whisper + Haiku-4.5 + Sonic 3	0.239 $\pm$ 0.063	0.482 $\pm$ 0.108	0.052 $\pm$ 0.033	0.318 $\pm$ 0.077	0.700 $\pm$ 0.052	0.982 $\pm$ 0.006
	Nova-3 + GPT-5.4 + Sonic 3	0.422 $\pm$ 0.080	0.651 $\pm$ 0.108	0.213 $\pm$ 0.073	0.496 $\pm$ 0.087	0.740 $\pm$ 0.033	0.986 $\pm$ 0.005
	Nova-3 + GPT-5.4-mini + Aura-2	0.217 $\pm$ 0.060	0.434 $\pm$ 0.108	0.042 $\pm$ 0.032	0.455 $\pm$ 0.084	0.298 $\pm$ 0.051	0.979 $\pm$ 0.006
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.318 $\pm$ 0.065	0.651 $\pm$ 0.096	0.083 $\pm$ 0.051	0.590 $\pm$ 0.082	0.386 $\pm$ 0.057	0.946 $\pm$ 0.016
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.167 $\pm$ 0.058	0.337 $\pm$ 0.096	0.042 $\pm$ 0.032	0.335 $\pm$ 0.082	0.651 $\pm$ 0.048	0.899 $\pm$ 0.021
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.381 $\pm$ 0.070	0.783 $\pm$ 0.096	0.135 $\pm$ 0.062	0.660 $\pm$ 0.063	0.380 $\pm$ 0.052	0.982 $\pm$ 0.005
	Ultravox-Realtime	0.272 $\pm$ 0.072	0.506 $\pm$ 0.108	0.125 $\pm$ 0.069	0.523 $\pm$ 0.092	0.229 $\pm$ 0.046	0.973 $\pm$ 0.008
S2S	Gemini-3.1-Flash-Live	0.248 $\pm$ 0.075	0.434 $\pm$ 0.108	0.111 $\pm$ 0.060	0.501 $\pm$ 0.082	0.164 $\pm$ 0.051	0.994 $\pm$ 0.005
	GPT-Realtime-1.5	0.349 $\pm$ 0.084	0.614 $\pm$ 0.108	0.176 $\pm$ 0.078	0.812 $\pm$ 0.060	0.230 $\pm$ 0.055	0.994 $\pm$ 0.004
	GPT-Realtime-mini	0.106 $\pm$ 0.051	0.229 $\pm$ 0.096	0.029 $\pm$ 0.029	0.369 $\pm$ 0.080	0.080 $\pm$ 0.031	0.978 $\pm$ 0.008

Experience results are provided for all systems per domain: CSM in Table 30, HR in Table 31, and ITSM in Table 32. Cross-domain variability of experience metrics are provided in Table 33.

Table 28 ITSM domain accuracy metrics for all evaluated systems under clean-audio conditions. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass^k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-A			Task Completion	Faithfulness	Speech Fidelity
		pass@1	pass@k	pass^k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.147 $\pm$ 0.055	0.362 $\pm$ 0.100	0.028 $\pm$ 0.030	0.328 $\pm$ 0.073	0.319 $\pm$ 0.063	0.974 $\pm$ 0.007
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.440 $\pm$ 0.082	0.675 $\pm$ 0.100	0.222 $\pm$ 0.086	0.745 $\pm$ 0.067	0.385 $\pm$ 0.064	0.977 $\pm$ 0.008
	Ink-whisper + Haiku-4.5 + Sonic 3	0.190 $\pm$ 0.067	0.425 $\pm$ 0.112	0.054 $\pm$ 0.050	0.362 $\pm$ 0.065	0.480 $\pm$ 0.054	0.977 $\pm$ 0.006
	Nova-3 + GPT-5.4 + Sonic 3	0.463 $\pm$ 0.068	0.838 $\pm$ 0.088	0.162 $\pm$ 0.066	0.597 $\pm$ 0.073	0.738 $\pm$ 0.054	0.985 $\pm$ 0.005
	Nova-3 + GPT-5.4-mini + Aura-2	0.198 $\pm$ 0.060	0.450 $\pm$ 0.112	0.045 $\pm$ 0.042	0.482 $\pm$ 0.082	0.215 $\pm$ 0.044	0.962 $\pm$ 0.008
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.350 $\pm$ 0.075	0.713 $\pm$ 0.100	0.131 $\pm$ 0.068	0.650 $\pm$ 0.088	0.414 $\pm$ 0.057	0.951 $\pm$ 0.014
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.175 $\pm$ 0.045	0.537 $\pm$ 0.113	0.012 $\pm$ 0.010	0.412 $\pm$ 0.085	0.516 $\pm$ 0.060	0.923 $\pm$ 0.015
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.425 $\pm$ 0.073	0.812 $\pm$ 0.088	0.150 $\pm$ 0.069	0.665 $\pm$ 0.067	0.470 $\pm$ 0.049	0.960 $\pm$ 0.011
	Ultravox-Realtime	0.215 $\pm$ 0.065	0.463 $\pm$ 0.113	0.068 $\pm$ 0.049	0.467 $\pm$ 0.092	0.312 $\pm$ 0.061	0.969 $\pm$ 0.008
S2S	Gemini-3.1-Flash-Live	0.273 $\pm$ 0.075	0.562 $\pm$ 0.112	0.115 $\pm$ 0.075	0.412 $\pm$ 0.082	0.209 $\pm$ 0.051	0.991 $\pm$ 0.008
	GPT-Realtime-1.5	0.628 $\pm$ 0.080	0.875 $\pm$ 0.075	0.402 $\pm$ 0.098	0.865 $\pm$ 0.053	0.426 $\pm$ 0.055	0.996 $\pm$ 0.004
	GPT-Realtime-mini	0.207 $\pm$ 0.065	0.425 $\pm$ 0.112	0.063 $\pm$ 0.050	0.378 $\pm$ 0.090	0.133 $\pm$ 0.040	0.981 $\pm$ 0.013

Table 29 Cross-domain variability of accuracy metrics: each cell is the sample standard deviation (ddof=1) of the per-domain point estimates across the three EVA domains (CSM, ITSM, HR). Larger values indicate the system’s performance depends more strongly on domain.

Arch.	System	EVA-A			Task Completion	Faithfulness	Speech Fidelity
		pass@1	pass@k	pass^k	SD	SD	SD
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.053	0.074	0.030	0.027	0.049	0.008
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.148	0.150	0.101	0.077	0.184	0.004
	Ink-whisper + Haiku-4.5 + Sonic 3	0.041	0.111	0.007	0.062	0.166	0.006
	Nova-3 + GPT-5.4 + Sonic 3	0.109	0.147	0.058	0.118	0.026	0.006
	Nova-3 + GPT-5.4-mini + Aura-2	0.011	0.013	0.032	0.015	0.048	0.010
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.120	0.119	0.111	0.042	0.117	0.010
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.058	0.172	0.019	0.085	0.094	0.013
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.054	0.028	0.027	0.019	0.055	0.011
	Ultravox-Realtime	0.055	0.039	0.035	0.048	0.056	0.002
S2S	Gemini-3.1-Flash-Live	0.057	0.113	0.033	0.052	0.093	0.005
	GPT-Realtime-1.5	0.144	0.144	0.113	0.174	0.113	0.002
	GPT-Realtime-mini	0.052	0.099	0.028	0.049	0.043	0.005

Table 30 Experience metrics for all evaluated systems under clean-audio conditions, restricted to the CSM domain. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass^k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-X			Turn-Taking	Conciseness	Conv. Progression
		pass@1	pass@k	pass^k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.220 $\pm$ 0.052	0.680 $\pm$ 0.120	0.014 $\pm$ 0.014	0.664 $\pm$ 0.035	0.790 $\pm$ 0.012	0.572 $\pm$ 0.068
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.052 $\pm$ 0.048	0.120 $\pm$ 0.100	0.009 $\pm$ 0.015	0.469 $\pm$ 0.041	0.792 $\pm$ 0.013	0.832 $\pm$ 0.048
	Ink-whisper + Haiku-4.5 + Sonic 3	0.008 $\pm$ 0.012	0.040 $\pm$ 0.060	0.000 $\pm$ 0.000	0.391 $\pm$ 0.028	0.755 $\pm$ 0.015	0.682 $\pm$ 0.048
	Nova-3 + GPT-5.4 + Sonic 3	0.004 $\pm$ 0.008	0.020 $\pm$ 0.040	0.000 $\pm$ 0.000	0.291 $\pm$ 0.034	0.825 $\pm$ 0.014	0.796 $\pm$ 0.042
	Nova-3 + GPT-5.4-mini + Aura-2	0.108 $\pm$ 0.044	0.440 $\pm$ 0.140	0.003 $\pm$ 0.005	0.577 $\pm$ 0.036	0.812 $\pm$ 0.018	0.440 $\pm$ 0.056
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.024 $\pm$ 0.028	0.080 $\pm$ 0.080	0.000 $\pm$ 0.001	0.274 $\pm$ 0.034	0.842 $\pm$ 0.014	0.832 $\pm$ 0.040
	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.224 $\pm$ 0.064	0.700 $\pm$ 0.120	0.021 $\pm$ 0.021	0.661 $\pm$ 0.031	0.654 $\pm$ 0.019	0.558 $\pm$ 0.048
Hybrid	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.031 $\pm$ 0.009	0.805 $\pm$ 0.013	0.634 $\pm$ 0.062
	Ultravox-Realtime	0.048 $\pm$ 0.052	0.120 $\pm$ 0.100	0.013 $\pm$ 0.020	0.483 $\pm$ 0.036	0.711 $\pm$ 0.020	0.540 $\pm$ 0.054
S2S	Gemini-3.1-Flash-Live	0.504 $\pm$ 0.064	1.000 $\pm$ 0.000	0.119 $\pm$ 0.067	0.788 $\pm$ 0.039	0.806 $\pm$ 0.020	0.702 $\pm$ 0.056
	GPT-Realtime-1.5	0.560 $\pm$ 0.080	0.940 $\pm$ 0.080	0.194 $\pm$ 0.083	0.821 $\pm$ 0.029	0.783 $\pm$ 0.017	0.634 $\pm$ 0.052
	GPT-Realtime-mini	0.376 $\pm$ 0.076	0.900 $\pm$ 0.100	0.089 $\pm$ 0.064	0.801 $\pm$ 0.036	0.705 $\pm$ 0.016	0.406 $\pm$ 0.062

Table 31 Experience metrics for all evaluated systems under clean-audio conditions, restricted to the HR domain. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass $\sim$ k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-X			Turn-Taking	Conciseness	Conv. Progression
		pass@1	pass@k	pass $\sim$ k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.241 $\pm$ 0.043	0.687 $\pm$ 0.096	0.014 $\pm$ 0.010	0.549 $\pm$ 0.046	0.818 $\pm$ 0.011	0.584 $\pm$ 0.052
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.012 $\pm$ 0.024	0.024 $\pm$ 0.036	0.004 $\pm$ 0.008	0.469 $\pm$ 0.023	0.774 $\pm$ 0.010	0.788 $\pm$ 0.030
	Ink-whisper + Haiku-4.5 + Sonic 3	0.002 $\pm$ 0.005	0.012 $\pm$ 0.024	0.000 $\pm$ 0.000	0.218 $\pm$ 0.037	0.793 $\pm$ 0.012	0.741 $\pm$ 0.033
	Nova-3 + GPT-5.4 + Sonic 3	0.012 $\pm$ 0.014	0.048 $\pm$ 0.048	0.000 $\pm$ 0.000	0.273 $\pm$ 0.039	0.859 $\pm$ 0.011	0.729 $\pm$ 0.031
	Nova-3 + GPT-5.4-mini + Aura-2	0.130 $\pm$ 0.036	0.458 $\pm$ 0.108	0.004 $\pm$ 0.003	0.594 $\pm$ 0.032	0.861 $\pm$ 0.008	0.405 $\pm$ 0.035
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.002 $\pm$ 0.005	0.012 $\pm$ 0.024	0.000 $\pm$ 0.000	0.344 $\pm$ 0.020	0.825 $\pm$ 0.013	0.701 $\pm$ 0.046
Hybrid	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.258 $\pm$ 0.058	0.614 $\pm$ 0.108	0.053 $\pm$ 0.033	0.481 $\pm$ 0.060	0.673 $\pm$ 0.019	0.625 $\pm$ 0.045
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.014 $\pm$ 0.004	0.795 $\pm$ 0.012	0.561 $\pm$ 0.049
S2S	Ultravox-Realtime	0.022 $\pm$ 0.024	0.072 $\pm$ 0.060	0.004 $\pm$ 0.008	0.432 $\pm$ 0.029	0.742 $\pm$ 0.019	0.453 $\pm$ 0.051
	Gemini-3.1-Flash-Live	0.622 $\pm$ 0.055	1.000 $\pm$ 0.000	0.261 $\pm$ 0.078	0.861 $\pm$ 0.021	0.798 $\pm$ 0.012	0.570 $\pm$ 0.042
	GPT-Realtime-1.5	0.593 $\pm$ 0.060	0.964 $\pm$ 0.048	0.244 $\pm$ 0.072	0.825 $\pm$ 0.016	0.807 $\pm$ 0.010	0.696 $\pm$ 0.041
	GPT-Realtime-mini	0.354 $\pm$ 0.055	0.867 $\pm$ 0.084	0.075 $\pm$ 0.048	0.827 $\pm$ 0.019	0.725 $\pm$ 0.014	0.313 $\pm$ 0.042

Table 32 Experience metrics for all evaluated systems under clean-audio conditions, restricted to the ITSM domain. Each cell shows the point estimate $\pm$ the percentile bootstrap CI half-width ($\alpha = 0.05$). The three pass-rate columns share a single shading scale (so pass@1 vs. pass@k vs. pass $\sim$ k are visually comparable); each submetric column is scaled independently. Darker = higher point estimate.

Arch.	System	EVA-X			Turn-Taking	Conciseness	Conv. Progression
		pass@1	pass@k	pass $\sim$ k	Mean	Mean	Mean
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.168 $\pm$ 0.042	0.575 $\pm$ 0.113	0.016 $\pm$ 0.026	0.487 $\pm$ 0.044	0.821 $\pm$ 0.010	0.639 $\pm$ 0.041
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.007 $\pm$ 0.010	0.037 $\pm$ 0.050	0.000 $\pm$ 0.000	0.415 $\pm$ 0.029	0.755 $\pm$ 0.012	0.793 $\pm$ 0.035
	Ink-whisper + Haiku-4.5 + Sonic 3	0.018 $\pm$ 0.015	0.075 $\pm$ 0.063	0.000 $\pm$ 0.000	0.328 $\pm$ 0.035	0.804 $\pm$ 0.008	0.708 $\pm$ 0.038
	Nova-3 + GPT-5.4 + Sonic 3	0.005 $\pm$ 0.008	0.025 $\pm$ 0.037	0.000 $\pm$ 0.000	0.285 $\pm$ 0.029	0.821 $\pm$ 0.011	0.688 $\pm$ 0.031
	Nova-3 + GPT-5.4-mini + Aura-2	0.100 $\pm$ 0.038	0.350 $\pm$ 0.113	0.007 $\pm$ 0.009	0.577 $\pm$ 0.030	0.831 $\pm$ 0.011	0.439 $\pm$ 0.038
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.003 $\pm$ 0.005	0.013 $\pm$ 0.025	0.000 $\pm$ 0.000	0.306 $\pm$ 0.021	0.821 $\pm$ 0.011	0.787 $\pm$ 0.039
Hybrid	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.338 $\pm$ 0.062	0.738 $\pm$ 0.100	0.079 $\pm$ 0.050	0.539 $\pm$ 0.058	0.727 $\pm$ 0.015	0.651 $\pm$ 0.036
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.013 $\pm$ 0.003	0.803 $\pm$ 0.012	0.659 $\pm$ 0.039
S2S	Ultravox-Realtime	0.018 $\pm$ 0.020	0.050 $\pm$ 0.050	0.001 $\pm$ 0.002	0.336 $\pm$ 0.042	0.797 $\pm$ 0.016	0.294 $\pm$ 0.045
	Gemini-3.1-Flash-Live	0.643 $\pm$ 0.068	0.938 $\pm$ 0.062	0.340 $\pm$ 0.088	0.840 $\pm$ 0.026	0.799 $\pm$ 0.014	0.637 $\pm$ 0.047
	GPT-Realtime-1.5	0.545 $\pm$ 0.065	0.912 $\pm$ 0.063	0.208 $\pm$ 0.067	0.798 $\pm$ 0.023	0.812 $\pm$ 0.014	0.708 $\pm$ 0.042
	GPT-Realtime-mini	0.487 $\pm$ 0.057	0.912 $\pm$ 0.062	0.134 $\pm$ 0.052	0.825 $\pm$ 0.020	0.736 $\pm$ 0.017	0.445 $\pm$ 0.050

Table 33 Cross-domain variability of experience metrics: each cell is the sample standard deviation (ddof=1) of the per-domain point estimates across the three EVA domains (CSM, ITSM, HR). Larger values indicate the system’s performance depends more strongly on domain.

Arch.	System	EVA-X			Turn-Taking	Conciseness	Conv. Progression
		pass@1	pass@k	pass $\sim$ k	SD	SD	SD
Cascade	Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	0.038	0.063	0.001	0.090	0.017	0.036
	Scribe-v2.2-Realtime + Gemini-3-Flash + TTS-Conversational-v3-Alpha	0.024	0.052	0.004	0.031	0.018	0.024
	Ink-whisper + Haiku-4.5 + Sonic 3	0.008	0.032	0.000	0.088	0.026	0.030
	Nova-3 + GPT-5.4 + Sonic 3	0.004	0.015	0.000	0.009	0.021	0.055
	Nova-3 + GPT-5.4-mini + Aura-2	0.016	0.058	0.002	0.010	0.024	0.020
	Parakeet-1.1 + Gemma-4-31B + Kokoro	0.012	0.039	0.000	0.035	0.011	0.067
Hybrid	Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	0.058	0.063	0.029	0.092	0.038	0.048
	Gemini-3-Flash + Gemini-3.1-Flash-TTS	0.000	0.000	0.000	0.010	0.005	0.051
S2S	Ultravox-Realtime	0.017	0.036	0.006	0.075	0.043	0.125
	Gemini-3.1-Flash-Live	0.075	0.036	0.112	0.037	0.004	0.066
	GPT-Realtime-1.5	0.024	0.026	0.026	0.015	0.015	0.040
	GPT-Realtime-mini	0.071	0.023	0.031	0.014	0.016	0.068

G Perturbation Analysis

Beyond standard evaluation conditions, EVA-Bench supports robustness testing through a structured perturbation system applied to the user simulator. Perturbations operate along three independent axes — behavior, accent, and audio degradation — and can be composed to simulate realistic deployment conditions that clean-audio benchmarks do not capture. All perturbations are applied exclusively to the user simulator; the agent under evaluation receives no special configuration, ensuring that robustness scores reflect genuine system sensitivity rather than evaluation artifacts.

Below we describe the robustness testing options available in EVA-Bench as well as additional results of the perturbation testing experiment described in the main body of the paper.

G.1 Behavioral Perturbations in EVA-Bench

The three behavioral personas described in Appendix G.2 — *aggressive_impatient*, *elderly_slow*, and *forgetful_disorganized* — constitute the behavioral perturbation axis. Each modifies both the simulator’s conversational prompt and its voice to produce acoustically and behaviorally consistent speech. Behavioral perturbations stress-test the agent’s ability to handle non-canonical caller patterns: interruptions and rapid speech, slow delivery with long silences, and disfluency-laden turns with mid-utterance corrections.

Accent Perturbations. EVA-Bench supports four non-native English accent variants for the user simulator: French, Indian, Spanish, and Chinese. Each accent uses a dedicated voice to ensure phonetic authenticity. Accent perturbations are mutually exclusive with behavioral perturbations, as both require a specific voice id. Together, the four accents allow systematic evaluation of speech recognition robustness across a representative sample of global English speaker populations commonly encountered in enterprise deployments.

Audio Degradation. Two audio degradation mechanisms are supported and can be composed with any behavioral or accent perturbation. First, *background noise* mixes an ambient audio track into the user simulator’s speech at a configurable signal-to-noise ratio (default: 15 dB). Eight noise environments are included: airport gate, baby crying, background music, bad connection static, coffee shop, loud construction, NYC street, and road noise — spanning the range of acoustic environments in which enterprise voice agents are commonly exposed to. Second, *connection degradation* applies a stack of VoIP artifacts — codec compression, simulated packet loss, and volume fluctuation — on top of any other active perturbation, reproducing the degradation characteristic of real telephony infrastructure.

G.2 User Personas

The **user persona** defines how the caller behaves during the conversation. By default, the simulated user is direct and efficient, with no disfluencies or unusual speech patterns. EVA-Bench additionally defines three behavioral personas that can be applied as perturbations: *aggressive_impatient*, in which the caller speaks quickly, interrupts the agent, and expresses frustration when progress stalls; *elderly_slow*, in which the caller speaks very slowly with deliberate pauses and occasionally asks the agent to repeat themselves; and *forgetful_disorganized*, in which the caller is prone to disfluencies, loses their train of thought, and requires time to retrieve codes and identifiers mid-turn. Each behavioral persona modifies both the user simulator’s prompt and its voice model, ensuring behavioral and acoustic consistency. Full persona prompts are provided below.

Default Persona

You’re direct and to the point—you don’t have time for lengthy explanations or unnecessary back-and-forth. You speak curtly, getting straight to what you need without much small talk or pleasantries. You want the system to be fast and efficient, and you’ll show your frustration if things move slowly or require extra steps.

Elderly & Slow

You are elderly and have difficulty understanding fast speech. You speak extremely slowly, with frequent deliberate pauses. You occasionally ask the agent to repeat themselves slowly. You do not rush. You frequently use ellipses (...) in your output to indicate pauses. Ex. ("Ok yes... my confirmation code is... W... K... 2... E... X... B...")

Aggressive & Impatient

You are impatient and easily frustrated when the agent does not resolve your requests immediately. You speak very quickly and often interrupt the agent mid-sentence when they are talking for too long to make your frustration clear and ask them to hurry it up. Or you ask why they are taking so long if there is a long silence. Express your frustration whenever progress is not being made and remember to interrupt often. You frequently output words in all caps to indicate your frustration and add emphasis.

Forgetful & Disorganized

You are forgetful and prone to disfluencies (um..., uh..., huh..., let me think..., hold on a second..., let me find that piece of information..., etc). You frequently use ellipses (...) in your output to indicate pauses. You often forget the information you need and have to search for it in the middle of your turn. Simple things like your name and date of birth you remember easily, but for any specific codes and IDs you need a couple of seconds to find it. You often lose your train of thought and need a moment to remember what you were saying. You also make mistakes when you speak and have to repeat yourself (ex. "hmm yeah one second... let me find that... ok its A E 2 B oh wait sorry actually its A F 2 B")

G.3 Perturbation Experiment

Methods. For each model, metric, and domain (or pooled across domains), we test whether each perturbation condition - accent, background noise, and the combination (accent + background noise) - alters performance as compared to the clean baseline. We analyze scenario-level mean deltas between a perturbation run and a paired baseline: within each scenario, performance is averaged across trials (3 trials per perturbation condition, 5 trials per clean baseline) and the paired difference is computed as $\delta = \bar{x}_{\text{perturbation}} - \bar{x}_{\text{baseline}}$, yielding one δ per scenario.

We assess each perturbation with a paired sign-flip permutation test on scenario-level deltas. This approach makes no distributional assumptions and is appropriate for the bounded, ordinal nature of our metrics. Under the null hypothesis of no perturbation effect ($\mathbb{E}[\delta] = 0$), the sign of each paired difference is exchangeable; we independently flip the sign of each scenario delta with probability 0.5 and recompute the mean across 10,000 permutations to construct the null distribution. We report two-sided p -values, defined as the fraction of permutations whose absolute mean is at least as large as the observed $|\bar{\delta}|$. Each estimate is accompanied by a 95% percentile bootstrap confidence interval on the mean delta, computed by resampling scenarios with replacement (1,000 bootstrap samples).

To control for multiple comparisons across the three perturbation conditions within each model $\times$ metric $\times$ domain combination, raw p -values are Holm-Bonferroni corrected. Effects are considered significant when the corrected p -value is below $\alpha = 0.05$.

Results. Within the perturbation effect plotted in Figure 8, models are always listed in the following order (left to right): **Cascade:** Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS, Scribe + Gemini-3-Flash + Conversational v3, Ink-whisper + Haiku-4.5 + Sonic 3, Nova-3 + GPT-5.4 + Sonic 3, Nova-3 + GPT-5.4-mini + Aura-2, Parakeet-1.1 + Gemma-4-31B + Kokoro, Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS; **Hybrid:** Gemini-3-Flash + Gemini-3.1-Flash-TTS, Ultravox-Realtime; **S2S:** Gemini-3.1-Flash-Live, GPT-Realtime-1.5, GPT-Realtime-mini. The complete results are shared in Table 34, Table 35 and Table 36.

Table 34 Perturbation effect under **Accent**: pooled mean Δ (perturbed – clean) per model and metric. Cells are shaded by magnitude (red = degradation, green = improvement). Significance from sign-flip permutation tests with Holm–Bonferroni correction within each (model, metric) family of three conditions: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Model	EVA-A	EVA-X	TaskComp	Faith	SpeechFidelity	ConvProg	TurnTake	Concise
<i>Cascade</i>								
Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	−0.090**	−0.073*	−0.176***	+0.106**	+0.006	+0.020	−0.165***	−0.015*
Scribe + Gemini-3-Flash + Conversational v3	−0.058	+0.001	−0.001	−0.024	−0.012	−0.000	+0.006	+0.006
Ink-whisper + Haiku-4.5 + Sonic 3	−0.103***	−0.013	−0.109***	+0.035	−0.009	−0.027	−0.124***	−0.031**
Nova-3 + GPT-5.4 + Sonic 3	−0.135***	−0.004	−0.144***	−0.007	−0.003	+0.019	−0.160***	+0.003
Nova-3 + GPT-5.4-mini + Aura-2	−0.054*	−0.046*	−0.121***	+0.060*	−0.016*	−0.029	−0.164***	−0.018*
Parakeet-1.1 + Gemma-4-31B + Kokoro	−0.040	+0.011	−0.033	+0.030	+0.003	−0.055	+0.121***	−0.003
Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	−0.074**	−0.161***	−0.128***	+0.085*	+0.015	−0.027	−0.165***	−0.004
<i>Hybrid</i>								
Ultravox-Realtime	+0.008	−0.008	−0.019	+0.023	+0.004	+0.097**	−0.015	−0.002
Gemini-3-Flash + Gemini-3.1-Flash-TTS	−0.116**	+0.000	−0.152**	−0.139***	+0.008	−0.150***	+0.096***	−0.014
<i>S2S</i>								
Gemini-3.1-Flash-Live	−0.014	+0.033	−0.007	−0.023	+0.005	−0.028	+0.028	−0.001
GPT-Realtime-1.5	+0.021	+0.033	+0.041	+0.013	−0.002	+0.030	+0.014	−0.000
GPT-Realtime-mini	+0.016	+0.035	−0.039	+0.019	+0.018	−0.027	+0.025	+0.004

Table 35 Perturbation effect under **Background Noise**: pooled mean Δ (perturbed – clean) per model and metric. Cells are shaded by magnitude (red = degradation, green = improvement). Significance from sign-flip permutation tests with Holm–Bonferroni correction within each (model, metric) family of three conditions: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Model	EVA-A	EVA-X	TaskComp	Faith	SpeechFidelity	ConvProg	TurnTake	Concise
<i>Cascade</i>								
Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	−0.046	+0.079*	−0.072*	+0.006	+0.032*	−0.070*	+0.097***	−0.023***
Scribe + Gemini-3-Flash + Conversational v3	−0.043	−0.014	−0.016	−0.032	−0.003	−0.007	−0.065***	−0.005
Ink-whisper + Haiku-4.5 + Sonic 3	−0.051	−0.010	+0.002	+0.005	−0.007	−0.071*	−0.056**	−0.022**
Nova-3 + GPT-5.4 + Sonic 3	−0.187***	−0.004	−0.199***	+0.010	−0.005	−0.007	−0.162***	−0.005
Nova-3 + GPT-5.4-mini + Aura-2	−0.095**	−0.050*	−0.187***	+0.118***	−0.002	+0.036	−0.227***	+0.002
Parakeet-1.1 + Gemma-4-31B + Kokoro	−0.066	−0.004	−0.059	+0.010	+0.004	−0.021	−0.163***	+0.001
Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	−0.104***	−0.206***	−0.195***	−0.013	+0.024	−0.175***	−0.281***	−0.033**
<i>Hybrid</i>								
Ultravox-Realtime	+0.041	−0.016	−0.004	+0.020	+0.008	+0.077*	−0.044**	−0.011
Gemini-3-Flash + Gemini-3.1-Flash-TTS	−0.056	+0.000	−0.004	−0.092**	−0.001	−0.032	+0.060***	−0.006
<i>S2S</i>								
Gemini-3.1-Flash-Live	−0.062*	−0.115**	−0.033	−0.018	−0.001	−0.043	−0.057*	−0.006
GPT-Realtime-1.5	−0.082*	−0.230***	−0.055	−0.050	−0.011	−0.105***	−0.102***	−0.018
GPT-Realtime-mini	+0.001	−0.143***	−0.043	−0.001	+0.008	−0.081**	−0.069***	−0.014

Table 36 Perturbation effect under **Accent + Background Noise**: pooled mean Δ (perturbed – clean) per model and metric. Cells are shaded by magnitude (red = degradation, green = improvement). Significance from sign-flip permutation tests with Holm–Bonferroni correction within each (model, metric) family of three conditions: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Model	EVA-A	EVA-X	TaskComp	Faith	SpeechFidelity	ConvProg	TurnTake	Concise
<i>Cascade</i>								
Cohere-transcribe + Gemma-4-26B + Voxtral-4B-TTS	−0.046	−0.055	−0.146***	+0.059	+0.038**	−0.122***	+0.002	−0.023***
Scribe + Gemini-3-Flash + Conversational v3	−0.061	−0.021	−0.045	−0.026	−0.011	−0.035	−0.113***	−0.005
Ink-whisper + Haiku-4.5 + Sonic 3	−0.151***	−0.013	−0.183***	+0.074*	−0.016**	−0.155***	−0.181***	−0.067***
Nova-3 + GPT-5.4 + Sonic 3	−0.227***	−0.004	−0.314***	+0.050	−0.001	+0.017	−0.201***	−0.023*
Nova-3 + GPT-5.4-mini + Aura-2	−0.143***	−0.094***	−0.313***	+0.077*	−0.052***	−0.134***	−0.314***	−0.076***
Parakeet-1.1 + Gemma-4-31B + Kokoro	−0.103**	−0.007	−0.062	−0.033	+0.006	−0.021	−0.104***	−0.004
Whisper-Large-v3 + Qwen3.5-27B + Voxtral-4B-TTS	−0.104***	−0.217***	−0.228***	−0.024	+0.003	−0.166***	−0.224***	−0.037**
<i>Hybrid</i>								
Ultravox-Realtime	−0.033	−0.023	−0.090**	+0.029	+0.009	+0.041	−0.087***	−0.003
Gemini-3-Flash + Gemini-3.1-Flash-TTS	−0.101**	+0.000	−0.104*	−0.148***	+0.000	−0.128***	+0.076***	−0.017**
<i>S2S</i>								
Gemini-3.1-Flash-Live	−0.014	−0.144**	−0.022	−0.009	−0.001	−0.023	−0.078***	−0.018
GPT-Realtime-1.5	−0.060*	−0.115**	−0.092**	−0.016	−0.014	−0.040	−0.049**	−0.021*
GPT-Realtime-mini	−0.029	−0.032	−0.032	+0.003	+0.003	−0.046	−0.012	−0.007






Figure 8 Perturbation effect across EVA metrics — (8a) EVA-A, (8b) EVA-X, (8c) Task Completion, (8d) Transcription Accuracy (Key Entities), (8e) Conversation Progression — for all evaluated systems, pooled across the three EVA domains. Bars show the mean delta from clean trials (negative = drop under perturbation); whiskers are 95% percentile bootstrap CIs on the per-scenario delta. Bar colors encode the perturbation condition: ■ accent, ■ background noise, ■ accent + background noise. Asterisks mark cells significant after Holm-Bonferroni correction within each model (* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$). Models, left to right: **Cascade**: Cohere + Gemma-4-26B + Voxtral, ElevenAgents (Scribe + Gemini-3-Flash + Conversational v3), Ink + Haiku-4.5 + Sonic, Nova + GPT-5.4 + Sonic, Nova + GPT-5.4-mini + Aura, Parakeet + Gemma-4-31B + Kokoro, Whisper + Qwen3.5-27B + Voxtral; **Hybrid**: Gemini-3-Flash + Gemini-3.1-Flash, Ultravox; **S2S**: Gemini-3.1-Flash-Live, GPT-Realtime-1.5, GPT-Realtime-mini.



(d) Transcription Accuracy (Key Entities)



(e) Conversation Progression

Figure 8 Perturbation effect across EVA metrics (Continued)

H Measurement Reliability: Variance Decomposition and Trial Count Justification

H.1 Variance decomposition

Per-model metric scores reflect variance from scenario difficulty, trial stochasticity, and LLM judge stochasticity. We characterized the contributions of each on a subset of $N = 4$ model configurations (2 cascade, 2 speech-to-speech) using 5 trials per scenario and 3 judge iterations per trial, across all 213 scenarios and 3 task domains. Three per-metric analyses were conducted: mixed effects variance decomposition (REML), two-way random effects ICC including the model $\times$ scenario interaction, and a permutation-based comparison of judge and trial standard deviations. Full results are reported below.

Mixed effects modelling. We conducted variance decomposition using linear mixed effects modeling (REML), fitted separately within each model, with domain as a fixed effect and scenario and trial as nested random effects. For judge-graded metrics, judge iterations were included as an additional nested random effect; for deterministic metrics, trial is the lowest modeled level. In both cases, variance at the lowest level of the hierarchy is absorbed into the residual, as it cannot be separately identified without repeated observations at that level.

Trial was consistently the largest source of variance across all models, contributing 40-80% of total observed variance for judge-graded metrics and 53-100% for deterministic metrics. Trial variance exceeded scenario variance on all 9 metrics (33/34 model-metric combinations); across metrics, cross-model ranges of trial and scenario components did not overlap with the exception of a single-model, 1.5 point overlap for faithfulness (see Table 37 for full decomposition).

Scenario-level intraclass correlation (ICC), estimated as the proportion of total variance attributable to scenario identity, ranged from near zero to 47% across metrics and models; Task Completion (ICC 33–47% across models) and Faithfulness (ICC 21-42%) showed consistently the highest scenario contributions. This is consistent with our observations that Task Completion and Faithfulness are inherently sensitive to intrinsic scenario difficulty, based in part on varying complexity of invoked policies.

Table 37 Per-model variance decomposition from REML linear mixed-effects fits with domain as a fixed effect and scenario and trial as nested random effects. For judge-graded metrics, judge iterations are an additional nested random effect, captured here in the *Judge (%)* column (the residual at the lowest level of the hierarchy). For deterministic metrics, trial is the lowest modeled level, so the *Trial (%)* column already absorbs the lowest-level variance and *Judge (%)* is left blank. All component columns give the proportion of total variance attributable to each component (% of σ_{total}^2). Judge-graded metrics appear first, followed by deterministic metrics.

Metric	Model	Scenario (%)	Trial (%)	Judge (%)	σ_{total}^2
Conciseness	Gemini-3.1-Flash-Live	22.6	47.1	30.4	0.0114
	GPT-Realtime-1.5	26.4	42.0	31.6	0.0081
	Ink + Haiku-4.5 + Sonic	7.8	56.7	35.6	0.0090
	Parakeet + Gemma-4-31B + Kokoro	20.0	42.5	37.5	0.0076
Conversation progression	Gemini-3.1-Flash-Live	20.9	56.0	23.2	0.1186
	GPT-Realtime-1.5	19.0	54.4	26.6	0.1034
	Ink + Haiku-4.5 + Sonic	4.0	59.1	36.8	0.0991
	Parakeet + Gemma-4-31B + Kokoro	15.9	56.9	27.1	0.0923
Faithfulness	Gemini-3.1-Flash-Live	36.4	49.1	14.5	0.0942
	GPT-Realtime-1.5	41.9	40.4	17.7	0.1159
	Ink + Haiku-4.5 + Sonic	20.9	65.1	14.0	0.1616
	Parakeet + Gemma-4-31B + Kokoro	33.2	46.3	20.6	0.1458
Speech fidelity	Gemini-3.1-Flash-Live	4.6	80.4	15.0	0.0032

continued on next page

Table 37 continued from previous page

Metric	Model	Scenario (%)	Trial (%)	Judge (%)	σ_{total}^2
	GPT-Realtime-1.5	0.0	71.6	28.4	0.0022
	Ink + Haiku-4.5 + Sonic	3.4	79.0	17.6	0.0028
	Parakeet + Gemma-4-31B + Kokoro	21.1	69.4	9.5	0.0109
Transcription accuracy key entities	Ink + Haiku-4.5 + Sonic	26.1	68.4	5.5	0.0679
	Parakeet + Gemma-4-31B + Kokoro	19.7	72.0	8.2	0.0393
Authentication success	Gemini-3.1-Flash-Live	30.3	69.7	–	0.1031
	GPT-Realtime-1.5	12.1	87.9	–	0.0515
	Ink + Haiku-4.5 + Sonic	26.2	73.8	–	0.2141
	Parakeet + Gemma-4-31B + Kokoro	24.8	75.2	–	0.1090
Conversation completion	Gemini-3.1-Flash-Live	0.0	100.0	–	0.0431
	GPT-Realtime-1.5	2.1	97.9	–	0.0139
	Ink + Haiku-4.5 + Sonic	14.9	85.1	–	0.2025
	Parakeet + Gemma-4-31B + Kokoro	1.1	98.9	–	0.0247
Task completion	Gemini-3.1-Flash-Live	41.7	58.3	–	0.2490
	GPT-Realtime-1.5	42.4	57.6	–	0.1629
	Ink + Haiku-4.5 + Sonic	33.2	66.8	–	0.2307
	Parakeet + Gemma-4-31B + Kokoro	47.3	52.7	–	0.2333
Turn taking	Gemini-3.1-Flash-Live	7.1	92.9	–	0.0508
	GPT-Realtime-1.5	15.0	85.0	–	0.0277
	Ink + Haiku-4.5 + Sonic	18.5	81.5	–	0.0646
	Parakeet + Gemma-4-31B + Kokoro	21.6	78.4	–	0.0267

Scenario variance and model $\times$ scenario interaction. To complement the full variance decomposition and assess whether models rank scenarios consistently, we computed two ICC models: a one-way ANOVA on model-centered scores, and a two-way random effects model including the model $\times$ scenario interaction term. The analysis covered eight metrics - authentication success, Conciseness, Conversation Completion, Conversation Progression, Task Completion, Transcription Accuracy (Key Entities) (cascade-only), and Turn-Taking - across the three task domains (CSM, ITSM, HR), yielding 22 domain $\times$ metric combinations. For each, we decomposed score variance into four components: scenario, model, model $\times$ scenario interaction, and residual trial-to-trial noise; variance components were estimated from a balanced two-way ANOVA. Significance was assessed via F-tests using the interaction mean square as the denominator (Cornfield-Tukey rule for random effects), and $\text{ICC}_{\text{scenario}}$ was reported as $\sigma_{\text{scenario}}^2 / \sigma_{\text{total}}^2$.

Per-metric scenario-level intraclass correlation ($\text{ICC}_{\text{scenario}}$, computed via one-way ANOVA after centering scores by model) ranged from approximately 0 on metrics that saturate across most scenarios (Speech Fidelity, Conversation Valid End, Turn-Taking) up to 0.31 on Task Completion (see Table 38). Faithfulness ranged 0.17–0.25 across domains, Conciseness 0.12–0.13, and Conversation Progression 0.07–0.13.

The model $\times$ scenario interaction was significant in all 22 domain $\times$ metric combinations ($p < 0.01$ in every case; $p < 0.001$ for 19 of 22; see Table 39), accounting for 4–18% of total score variance (median 11%). This indicates that scenario difficulty is ranked differently across models; some scenarios are disproportionately harder for one model than another. Consistent with this finding, $\text{ICC}_{\text{scenario}}$ was low across the board (median 0.08, range 0.00–0.27), confirming that scenario identity alone explains a small fraction of overall score variance after accounting for model and interaction effects.

These analyses support the finding from the full variance decomposition that scenario is not the dominant source of variance, and is a property of the benchmark rather than a confound: it reflects the range of task difficulty that the benchmark spans. Additionally, this also demonstrates that a substantial proportion of each

model’s apparent scenario variance reflects model-specific scenario interactions rather than shared scenario difficulty.

Table 38 Per-(metric $\times$ domain) scenario-level ICC from one-way ANOVA on per-model-centered scores. $ICC_{\text{scenario}} = \sigma_{\text{scenario}}^2 / (\sigma_{\text{scenario}}^2 + \sigma_{\text{residual}}^2)$. 95% CI from the F-distribution (Fisher’s exact ICC bounds). n is the number of scenarios in the (metric, domain) cell.

Metric	Domain	ICC_{scenario}	95% CI	n
Authentication success	CSM	0.112	[0.066, 0.186]	48
	ITSM	0.065	[0.037, 0.106]	80
	HR	0.214	[0.160, 0.288]	82
Conciseness	CSM	0.132	[0.082, 0.210]	50
	ITSM	0.115	[0.077, 0.170]	80
	HR	0.129	[0.089, 0.185]	83
Conversation Completion	CSM	0.002	[0.000, 0.030]	50
	ITSM	0.019	[0.002, 0.046]	80
	HR	0.041	[0.019, 0.075]	83
Conversation progression	CSM	0.111	[0.066, 0.183]	50
	ITSM	0.126	[0.086, 0.183]	80
	HR	0.072	[0.043, 0.114]	83
Faithfulness	CSM	0.247	[0.174, 0.351]	50
	ITSM	0.173	[0.110, 0.257]	80
	HR	0.231	[0.174, 0.306]	83
Speech fidelity	CSM	0.017	[0.000, 0.067]	50
	ITSM	0.020	[0.001, 0.048]	80
	HR	0.000	[0.000, 0.023]	83
Task completion	CSM	0.309	[0.226, 0.421]	50
	ITSM	0.273	[0.210, 0.355]	80
	HR	0.285	[0.222, 0.367]	83
Transcription accuracy key entities	CSM	0.197	[0.119, 0.308]	50
	ITSM	0.116	[0.066, 0.187]	80
	HR	0.131	[0.075, 0.205]	83
Turn-taking	CSM	0.040	[0.013, 0.086]	50
	ITSM	0.051	[0.027, 0.088]	80
	HR	0.020	[0.002, 0.046]	83

Table 39 Per-(metric $\times$ domain) variance decomposition from a two-way random-effects ANOVA. Each row partitions total observed variance into scenario, model, model $\times$ scenario interaction, and residual components (% of σ_{total}^2). F_{int} and p_{int} are the F-statistic and p-value for the interaction term, computed using the interaction mean square as the F-test denominator (Cornfield-Tukey rule for random effects).

Metric	Domain	Scenario (%)	Model (%)	Interaction (%)	Residual (%)	F_{int}	p_{int}
Authentication success	CSM	9.7	1.0	6.4	83.0	1.38	0.004
	ITSM	3.6	23.9	6.4	66.1	1.48	< 0.0001
	HR	14.7	18.7	12.5	54.0	2.16	< 0.0001
Conciseness	CSM	7.7	13.2	17.2	61.9	2.39	< 0.0001
	ITSM	8.0	1.1	16.1	74.9	2.07	< 0.0001
	HR	10.0	3.3	11.6	75.1	1.77	< 0.0001

continued on next page

Table 39 continued from previous page

Metric	Domain	Scenario (%)	Model (%)	Interaction (%)	Residual (%)	F_{int}	p_{int}
Conversation Completion	CSM	0.0	2.1	8.0	89.9	1.44	0.001
	ITSM	0.4	25.9	4.6	69.0	1.34	0.001
	HR	0.6	44.3	8.1	47.0	1.86	< 0.0001
Conversation progression	CSM	8.3	6.7	9.5	75.5	1.63	< 0.0001
	ITSM	10.0	4.4	9.3	76.3	1.61	< 0.0001
	HR	4.2	5.1	12.3	78.4	1.78	< 0.0001
Faithfulness	CSM	20.7	6.2	10.9	62.1	1.88	< 0.0001
	HR	11.6	36.6	14.0	37.8	2.85	< 0.0001
Task completion	CSM	27.3	3.3	11.5	57.9	1.99	< 0.0001
	ITSM	18.6	20.5	14.3	46.6	2.53	< 0.0001
	HR	20.9	16.0	13.9	49.1	2.42	< 0.0001
Transcription accuracy key entities	CSM	7.3	21.1	18.3	53.3	2.71	< 0.0001
	ITSM	2.9	20.8	14.2	62.1	2.14	< 0.0001
Turn taking	CSM	0.6	63.2	4.2	32.0	1.66	< 0.0001
	ITSM	0.8	65.6	4.7	29.0	1.81	< 0.0001
	HR	0.0	73.0	4.5	22.5	2.01	< 0.0001

Judge stochasticity. We assessed whether trial standard deviation systematically exceeds judge standard deviation using a sign-flip permutation test on the mean delta (trial SD - judge SD delta per scenario, within or averaged across models; one-sided, H_1 : mean delta > 0; 10,000 permutations). As a complementary check on directional consistency, we applied an exact binomial sign test to the count of scenarios where the delta exceeded zero (one-sided, H_1 : $P(\text{delta} > 0) > 0.5$).

Trial variance dominated judge variance across all 16 model $\times$ metric combinations (sign-flip permutation test, one-sided, $p < 0.0001$ in every case; see Table 40, with per-model mean deltas between trial and judge standard deviations ranging from 0.01–0.05 for agent’s Speech Fidelity and 0.02–0.03 for Conciseness, to 0.12–0.15 for Conversation Progression and 0.11–0.21 for Faithfulness. The binomial sign test confirmed that trial SD exceeded judge SD in the majority of scenarios for 13 of 16 model $\times$ metric combinations ($p < 0.05$). For 3 models, the significant mean deltas for agent Speech Fidelity, which had negligible variance for both judge and trial, were driven by a minority of high-variance scenarios or a p-value just above the threshold.

Table 40 Judge vs. trial variance per domain and model. Each cell value is the mean across records of the per-record standard deviation: *Judge std dev* is the per-record SD across the 3 judge iterations (averaged over trials before averaging across records); *Trial std dev* is the per-record SD across trials (after averaging across judge iterations), averaged across records. p_{perm} is the one-sided sign-flip permutation test p-value (10,000 permutations, H_1 : trial std dev > judge std dev). p_{sign} is the one-sided exact binomial sign test p-value (H_1 : $P(\text{trial std dev} > \text{judge std dev}) > 0.5$).

Metric	Domain	Model	Judge std dev	Trial std dev	p_{perm}	p_{sign}
Conciseness	CSM	Gemini-3.1-Flash-Live	0.0364	0.0681	< 0.0001	< 0.0001
	CSM	Parakeet + Gemma-4-31B + Kokoro	0.0401	0.0489	0.004	0.003
	CSM	GPT-Realtime-1.5	0.0345	0.0649	< 0.0001	< 0.0001
	CSM	Ink + Haiku-4.5 + Sonic	0.0412	0.0665	< 0.0001	< 0.0001
	ITSM	Gemini-3.1-Flash-Live	0.0367	0.0684	< 0.0001	< 0.0001
	ITSM	Parakeet + Gemma-4-31B + Kokoro	0.0349	0.0522	< 0.0001	< 0.0001
	ITSM	GPT-Realtime-1.5	0.0311	0.0529	< 0.0001	< 0.0001
	ITSM	Ink + Haiku-4.5 + Sonic	0.0370	0.0645	< 0.0001	< 0.0001
	HR	Gemini-3.1-Flash-Live	0.0386	0.0651	< 0.0001	< 0.0001

continued on next page

Table 40 continued from previous page

Metric	Domain	Model	Judge std dev	Trial std dev	p_{perm}	p_{sign}
Conversation progression	HR	Parakeet + Gemma-4-31B + Kokoro	0.0325	0.0547	< 0.0001	< 0.0001
	HR	GPT-Realtime-1.5	0.0319	0.0473	< 0.0001	< 0.0001
	HR	Ink + Haiku-4.5 + Sonic	0.0382	0.0648	< 0.0001	< 0.0001
	CSM	Gemini-3.1-Flash-Live	0.0535	0.2259	< 0.0001	< 0.0001
	CSM	Parakeet + Gemma-4-31B + Kokoro	0.0639	0.1417	< 0.0001	0.008
	CSM	GPT-Realtime-1.5	0.0655	0.2303	< 0.0001	< 0.0001
	CSM	Ink + Haiku-4.5 + Sonic	0.0960	0.2048	< 0.0001	< 0.0001
	ITSM	Gemini-3.1-Flash-Live	0.0735	0.2184	< 0.0001	< 0.0001
	ITSM	Parakeet + Gemma-4-31B + Kokoro	0.0631	0.1823	< 0.0001	< 0.0001
	ITSM	GPT-Realtime-1.5	0.0726	0.2003	< 0.0001	< 0.0001
	ITSM	Ink + Haiku-4.5 + Sonic	0.0901	0.2301	< 0.0001	< 0.0001
	HR	Gemini-3.1-Flash-Live	0.0840	0.2287	< 0.0001	< 0.0001
	HR	Parakeet + Gemma-4-31B + Kokoro	0.0762	0.2257	< 0.0001	< 0.0001
	HR	GPT-Realtime-1.5	0.0799	0.2056	< 0.0001	< 0.0001
	HR	Ink + Haiku-4.5 + Sonic	0.0850	0.2206	< 0.0001	< 0.0001
Faithfulness	CSM	Gemini-3.1-Flash-Live	0.0532	0.2519	< 0.0001	< 0.0001
	CSM	Parakeet + Gemma-4-31B + Kokoro	0.0749	0.2621	< 0.0001	< 0.001
	CSM	GPT-Realtime-1.5	0.0532	0.2145	< 0.0001	< 0.0001
	CSM	Ink + Haiku-4.5 + Sonic	0.0666	0.2869	< 0.0001	< 0.0001
	ITSM	Gemini-3.1-Flash-Live	0.0303	0.1295	< 0.0001	0.073
	ITSM	Parakeet + Gemma-4-31B + Kokoro	0.0869	0.1853	< 0.0001	< 0.0001
	ITSM	GPT-Realtime-1.5	0.0693	0.1782	< 0.0001	< 0.0001
	ITSM	Ink + Haiku-4.5 + Sonic	0.0564	0.2767	< 0.0001	< 0.0001
	HR	Gemini-3.1-Flash-Live	0.0307	0.1001	< 0.0001	0.745
	HR	Parakeet + Gemma-4-31B + Kokoro	0.0763	0.2091	< 0.0001	< 0.0001
	HR	GPT-Realtime-1.5	0.0426	0.1287	< 0.0001	0.008
	HR	Ink + Haiku-4.5 + Sonic	0.0548	0.2517	< 0.0001	< 0.0001
Speech fidelity	CSM	Gemini-3.1-Flash-Live	0.0000	0.0000	0.500	1.000
	CSM	Parakeet + Gemma-4-31B + Kokoro	0.0044	0.0425	< 0.0001	0.899
	CSM	GPT-Realtime-1.5	0.0027	0.0158	0.033	1.000
	CSM	Ink + Haiku-4.5 + Sonic	0.0020	0.0148	< 0.0001	0.984
	ITSM	Gemini-3.1-Flash-Live	0.0047	0.0314	< 0.0001	1.000
	ITSM	Parakeet + Gemma-4-31B + Kokoro	0.0074	0.0559	< 0.0001	< 0.0001
	ITSM	GPT-Realtime-1.5	0.0020	0.0057	0.017	1.000
	ITSM	Ink + Haiku-4.5 + Sonic	0.0051	0.0353	< 0.0001	< 0.001
	HR	Gemini-3.1-Flash-Live	0.0011	0.0044	0.032	1.000
	HR	Parakeet + Gemma-4-31B + Kokoro	0.0060	0.0598	< 0.0001	< 0.0001
	HR	GPT-Realtime-1.5	0.0013	0.0083	< 0.0001	1.000
	HR	Ink + Haiku-4.5 + Sonic	0.0020	0.0284	< 0.0001	0.413
Transcription accuracy key entities	CSM	Parakeet + Gemma-4-31B + Kokoro	0.0148	0.1316	< 0.0001	< 0.0001
	CSM	Ink + Haiku-4.5 + Sonic	0.0243	0.1685	< 0.0001	< 0.0001
	ITSM	Parakeet + Gemma-4-31B + Kokoro	0.0278	0.1261	< 0.0001	< 0.0001
	ITSM	Ink + Haiku-4.5 + Sonic	0.0204	0.1819	< 0.0001	< 0.0001
	HR	Parakeet + Gemma-4-31B + Kokoro	0.0250	0.1521	< 0.0001	< 0.0001
	HR	Ink + Haiku-4.5 + Sonic	0.0285	0.1875	< 0.0001	< 0.0001

H.2 Justification of trial count

Motivation. EVA-Bench averages over multiple trials per scenario to reduce simulator and agent stochasticity. Clean evaluations use $k=5$ trials per scenario, while perturbation experiments use $k=3$ to fit a larger cell count within compute budget. We quantify the cost of this reduction by measuring how rapidly each model-level metric estimate stabilizes as a function of trial count.

Method. For each (model, metric) pair and each $k \in \{1, 2, 3, 4\}$ we draw $N=2000$ Monte Carlo subsamples. In each draw, we independently sample, for every scenario, a uniformly random k -trial subset of the five available trials, average within scenario, then average across scenarios. The result is a single model-level estimate $\hat{\theta}_k$ per draw. At $k=5$ exactly one draw exists – the *anchor* $\hat{\theta}_5$, equal to the full mean-of-scenario-means.

We sample subsets *independently per scenario* rather than aligning trial indices across scenarios, because trial indices are not meaningful across scenarios.

For each k we report the **95% interval width** $w_k = p_{97.5}(\hat{\theta}_k) - p_{2.5}(\hat{\theta}_k)$, in metric units.

These quantities reflect resampling of *existing* trial outcomes – they capture how much our reported estimate would have moved if we had used fewer of the trials we ran. The pooled bootstrap CIs in the main results (Sec. 4.3) capture the complementary cross-scenario uncertainty.

Results. Figure 9 plots the median 95% CI width at each k , one panel per metric, one line per model. Every line decays approximately as $k^{-1/2}$ – the textbook scaling for sample-mean variance – with two regimes:

- *Per-turn metrics* (Speech Fidelity, Conciseness, Turn-Taking) are already stable at $k=1$, with CI width below 0.02 across all models. Adding trials yields little further reduction.
- *Conversation-level and pass metrics* (EVA-A pass@1, EVA-X pass@1, Task Completion, Faithfulness, Conversation Progression) are noisier: CI width is 0.05–0.09 at $k=1$, shrinking to roughly 0.02–0.03 at $k=3$.

Table 41 summarises. At $k=3$, the median 95% CI width is at most 0.034 for any metric, and at least 97.3% of $k=3$ subsamples land within 0.02 of the $k=5$ anchor on every metric.

Justifying the trial counts. Cross-model gaps on the headline pass-style metrics span 0.1–0.6 across the 12 evaluated systems. At $k=3$, trial-count uncertainty is ~ 0.03 – below 10% of the smallest interesting effect size and well below the cross-architecture gaps reported in the main tables. We therefore use $k=3$ for the perturbation experiments, where holding the trial count down lets us cover more scenarios. We retain $k=5$ for the clean evaluation, where the same point estimates are reused across scatter plots and Pareto frontiers and the marginal CI shrinkage from $k=4$ to $k=5$ (~ 0.02 absolute, by construction terminating at zero) is worth the cost.

Table 41 Subsample-stability summary across the 12 evaluated systems. *Median 95% CI width*: median across models of the empirical $p_{97.5}-p_{2.5}$ width of the model-level estimate at trial count k , in metric units. *Agree ≤ 0.02* : median across models of the fraction of $k=3$ subsamples that fall within 0.02 of the $k=5$ anchor. The $k=5$ column is omitted because it is identically zero by construction (single anchor draw).

Metric	Median 95% CI width				Agree ≤ 0.02
	$k=1$	$k=2$	$k=3$	$k=4$	at $k=3$
EVA-A pass@1	0.085	0.053	0.034	0.020	97.6%
EVA-X pass@1	0.052	0.032	0.021	0.013	99.6%
Task Completion	0.085	0.052	0.034	0.021	97.3%
Faithfulness	0.070	0.043	0.029	0.017	99.4%
Speech Fidelity	0.015	0.009	0.006	0.004	100.0%
Conv. Progression	0.077	0.046	0.030	0.019	99.1%
Turn-Taking	0.048	0.029	0.020	0.012	100.0%
Conciseness	0.020	0.012	0.008	0.005	100.0%

Figure 10 displays eight line plots showing the 95% CI width (Y-axis) versus Trials per scenario (k) (X-axis, ranging from 1 to 5) for various tasks. The tasks are arranged in a 4x2 grid:

- EVA-A pass@1**: Shows CI width decreasing from approximately 0.09 to 0.00 as trials increase.
- EVA-X pass@1**: Shows CI width decreasing from approximately 0.12 to 0.00.
- Task Completion**: Shows CI width decreasing from approximately 0.09 to 0.00.
- Faithfulness**: Shows CI width decreasing from approximately 0.08 to 0.00.
- Speech Fidelity**: Shows CI width decreasing from approximately 0.03 to 0.00.
- Conv. Progression**: Shows CI width decreasing from approximately 0.08 to 0.00.
- Turn-Taking**: Shows CI width decreasing from approximately 0.08 to 0.00.
- Conciseness**: Shows CI width decreasing from approximately 0.03 to 0.00.

Each plot contains 10 data series representing different model combinations, as indicated by the legend at the bottom:

- Cohere + Gemma-4-26B + Voxtral
- Scribe + Gemini-3-Flash + Conversational v3
- Ink + Haiku-4.5 + Sonic
- Nova + GPT-5.4 + Sonic
- Nova + GPT-5.4-mini + Aura
- Parakeet + Gemma-4-31B + Kokoro
- Whisper + Qwen3.5-27B + Voxtral
- Gemini-3-Flash + Gemini-3.1-Flash
- Ultravox
- Gemini-3.1-Flash-Live
- GPT-Realtime-1.5
- GPT-Realtime-mini

64

I Data Workflows

I.1 Airline CSM Workflows

This domain covers 50 scenarios across seven workflow categories — IRROPS rebooking, voluntary changes, missed connections, same-day standby, cancellations and refunds, escalation and availability constraints, and adversarial compensation claims — backed by 15 tools. The domain is high-stakes and time-pressured, with heavy dependence on accurate transcription of named entities: confirmation codes, flight numbers, passenger names, and travel dates. Table 42 provides a description of each workflow, the expected number of tool calls, and the tools invoked by the agent.

I.2 Healthcare HRSD Workflows

This domain covers HR service delivery at a hospital, with callers drawn from clinical and administrative staff. It comprises 83 scenarios across 12 single-intent workflows backed by 47 tools, extended with dual-intent, triple-intent, and adversarial variants. It has the highest per-workflow complexity of any domain in EVA-Bench, with an average of 8.7 expected tool calls across all scenarios (5.0 for single-intent workflows, up to 18 for triple-intent). Its defining challenge is the density and complexity of named entities the caller must communicate over voice — NPI numbers, DEA registration numbers, state license numbers, and OTP codes — where a single transcription error can cascade into authentication or policy failures. Table 43 provides a description of each workflow, the expected number of tool calls, and the tools invoked by the agent.

I.3 Enterprise ITSM Workflows

This domain covers an enterprise IT service desk spanning 21 workflows across six categories, backed by 59 tools. It comprises 80 scenarios: 29 single-intent, 14 double-intent, 14 triple-intent, 14 quadruple-intent, and 9 adversarial. Its defining characteristic is a branching flow structure: incident flows have both a troubleshooting-resolved path and an escalation-to-ticket path, testing whether the agent correctly gates escalation on failed resolution attempts. Authentication is tiered across three levels — standard, OTP-elevated, and manager-level — reflecting the sensitivity of different workflows. Table 44 provides a description of each workflow, the expected number of tool calls, and the tools invoked by the agent.

Table 42 Airline CSM workflows.

Workflow	Description	Tool Calls	Scenario IDs	Tools
Voluntary Change	Caller initiates a flight change subject to fare difference and change fees.	3	1.1.x, 1.2.x, 1.3.x	<code>get_reservation</code> , <code>search_rebooking_options</code> , <code>rebook_flight</code>
IRROPS Re-booking	Airline-initiated disruption entitles the caller to free rebooking on an alternative flight.	3–6	2.1.x, 2.2.x, 2.3.x, 2.4.x	<code>get_reservation</code> , <code>get_disruption_info</code> , <code>search_rebooking_options</code> , <code>rebook_flight</code> , <code>issue_meal_voucher</code>
Missed Connection	Caller missed a connecting flight due to a late inbound leg; all affected segments are rebooked.	2–3	3.1.x, 3.3.x	<code>get_reservation</code> , <code>search_rebooking_options</code> , <code>rebook_flight</code> , <code>add_to_standby</code>
Same-Day Change & Standby	Caller requests a same-day flight change or standby placement subject to time-sensitive policy constraints.	2–3	4.1.x, 4.2.x	<code>get_reservation</code> , <code>search_rebooking_options</code> , <code>rebook_flight</code> , <code>add_to_standby</code>

continued on next page

Table 42 continued from previous page

Workflow	Description	Tool Calls	Scenario IDs	Tools
Cancellation & Refund	Caller cancels a booking; agent determines eligibility for a cash refund or travel credit based on fare type and cancellation policy.	1–4	5.1.x, 5.2.x	get_reservation, get_disruption_info, cancel_reservation, process_refund, issue_travel_credit, issue_meal_voucher
Escalation & Availability Constraints	Agent exhausts available rebooking or compensation options and must escalate to a supervisor or communicate policy limits to the caller.	3–6	6.1.x, 6.3.x	get_reservation, get_disruption_info, search_rebooking_options, rebook_flight, issue_meal_voucher, issue_hotel_voucher, transfer_to_agent
Adversarial Compensation Claim	Caller attempts to claim meal or hotel vouchers for a disruption that does not meet eligibility threshold under policy.	1–4	7.1.x, 7.2.x, 7.3.x, 7.4.x	get_reservation, get_flight_status, get_disruption_info

Table 43 Healthcare HRSD workflows.

Workflow	Description	Tool Calls	Scenario IDs	Tools
License Extension	Provider requests a provisional or supervised temporary extension for an expiring state medical license.	4–6	1.x	verify_provider_auth, get_provider_profile, get_license_record, check_extension_eligibility, submit_license_extension, notify_credentialing_committee
Shift Swap	Employee arranges a shift swap with a certified colleague; agent verifies unit certification requirements before confirming.	3–6	2.x	verify_employee_auth, get_shift_record, check_swap_eligibility, verify_colleague_certifications, confirm_shift_swap, notify_department_manager
Malpractice Coverage Update	Provider updates their malpractice insurance carrier and policy details; low coverage limits trigger an automatic recredentialing flag.	3–5	3.x	verify_provider_auth, get_provider_profile, get_malpractice_record, update_malpractice_coverage, notify_credentialing_committee
Onboarding Task Completion	New hire marks completed onboarding checklist items using task-specific completion codes, then schedules an orientation follow-up.	3–7	4.x	verify_employee_auth, get_employee_record, get_onboarding_checklist, complete_onboarding_task, check_appointment_availability, schedule_orientation_followup
DEA Registration Transfer	Provider transfers their DEA registration to a new facility and state; PDMP is notified upon completion. Requires OTP second-factor authentication.	5–6	5.x	verify_provider_auth, initiate_otp_auth, verify_otp_auth, get_dea_record, transfer_dea_registration, notify_pdmp

continued on next page

Table 43 continued from previous page

Workflow		Description	Tool Calls	Scenario IDs	Tools
FMLA Leave Absence	/ of	Employee files an FMLA leave case after eligibility is verified; agent schedules a return-to-work check-in on or after the leave end date. Requires OTP second-factor authentication.	5–9	6.x	verify_employee_auth,initiate_otp_auth,verify_otp_auth,get_employee_record,check_leave_eligibility,submit_fmlla_case,notify_department_manager,check_appointment_availability,schedule_return_to_work_checkin
Payroll Correction	Cor-	Employee submits a correction for missing or incorrect hours on a timesheet; blocked if the pay period is already closed.	3–5	7.x	verify_employee_auth,get_timesheet_record,check_correction_eligibility,submit_payroll_correction,notify_department_manager
Privilege Re-activation	Re-	Provider returning from leave reactivates suspended clinical privileges after submitting an occupational health clearance code; competency review is scheduled and EHR access is restored. Requires OTP second-factor authentication.	5–10	8.x	verify_employee_auth,initiate_otp_auth,verify_otp_auth,get_provider_profile,check_reactivation_eligibility,check_appointment_availability,schedule_competency_review,reactivate_privileges,notify_credentialing_committee,update_ehr_access
On-Call Registration		Employee registers on-call availability for a date window with optional blackout dates; blocked if on leave or missing unit certifications.	3–4	9.x	verify_employee_auth,get_oncall_schedule,check_oncall_eligibility,register_oncall_availability
I-9 Verification	Verifica-	New hire or rehired employee submits I-9 work authorization documents; HR compliance is notified upon completion.	3–5	10.x	verify_employee_auth,get_employee_record,get_i9_record,submit_i9_verification,notify_hr_compliance
Visa Dependent Addition	Depen-	H-1B employee adds a dependent to their visa petition via a USCIS amendment; immigration counsel is notified. Requires OTP second-factor authentication.	4–6	11.x	verify_employee_auth,initiate_otp_auth,verify_otp_auth,get_visa_record,add_visa_dependent,notify_immigration_counsel
PTO Request		Employee requests paid time off or sick leave; agent validates balance and department blackout constraints before submitting.	4–6	12.x	verify_employee_auth,get_employee_record,get_pto_balance,check_pto_eligibility,submit_pto_request,notify_department_manager

continued on next page

Table 43 continued from previous page

Workflow	Description	Tool Calls	Scenario IDs	Tools
Dual-Intent	Two single-intent workflows handled in a single call. Covers 10 workflow pairings including license extension + privilege reactivation, mal-practice update + DEA transfer, FMLA + PTO, shift swap + on-call registration, and onboarding + DEA transfer, among others.	4–15	D1.x–D10.x	(union of constituent workflow tools)
Triple-Intent	Three single-intent workflows handled in a single call. Covers 7 scenario groups combining provider credentialing, scheduling, and leave workflows.	11–18	T1.x–T7.x	(union of constituent workflow tools)
Adversarial	Caller attempts to circumvent policy constraints: proxy authentication, self-supervised license extension, backdated FMLA, leave duration exceeding balance, cross-employee payroll correction, and on-call registration without required certifications.	0–6	A1–A10	verify_employee_auth,verify_provider_auth,initiate_otp_auth,verify_otp_auth,get_shift_record,check_swap_eligibility,verify_colleague_certifications,get_oncall_schedule,check_oncall_eligibility,get_license_record,check_extension_eligibility,get_dea_record,get_provider_profile,check_reactivation_eligibility,get_employee_record,check_leave_eligibility,get_timesheet_record,transfer_to_agent

Table 44 ITSM workflows. Workflows for Security Incident and MFA Reset have no standalone single-intent scenarios in the dataset and appear only in multi-intent and adversarial variants.

Workflow	Description	Tool Calls	Scenario IDs	Tools
Login Issue	Employee is locked out or has an expired password; agent walks through troubleshooting steps before attempting an account unlock or password reset. Issues that resolve during the call are closed without a ticket.	4–5	1, 2	verify_employee_auth,get_employee_record,get_troubleshooting_guide,attempt_account_unlock,attempt_password_reset,mark_resolved
Service Outage	Employee reports a service outage; agent checks for an existing outage ticket and either adds the caller as an affected user or opens a new ticket with SLA assignment and known-error linking.	4–6	4, 5, 6	verify_employee_auth,get_employee_record,check_existing_outage,add_affected_user,create_incident_ticket,assign_sla_tier,link_known_error

continued on next page

Table 44 continued from previous page

Workflow	Description	Tool Calls	Scenario IDs	Tools
Hardware Malfunction	Employee reports a malfunctioning device; agent runs troubleshooting, looks up the asset record, and schedules a field technician dispatch if the issue is not resolved.	6–8	7, 8	verify_employee_auth,get_employee_record,get_troubleshooting_guide,get_employee_assets,get_asset_record,create_incident_ticket,assign_sla_tier,schedule_field_dispatch
Network / VPN Issue	Employee reports a network or VPN connectivity problem; agent walks through troubleshooting steps and, if unresolved, opens a ticket with a diagnostic log attachment.	4–6	10, 11	verify_employee_auth,get_employee_record,get_troubleshooting_guide,create_incident_ticket,assign_sla_tier,attach_diagnostic_log,mark_resolved
Laptop Replacement	Employee requests a laptop replacement; agent checks hardware entitlement and budget, submits the request, and initiates a return authorization for the current device.	7	12	verify_employee_auth,get_employee_record,check_hardware_entitlement,verify_cost_center_budget,get_employee_assets,submit_hardware_request,initiate_asset_return
Monitor Bundle	Employee requests a new monitor; agent checks entitlement and budget before submitting the hardware request. No asset return required.	5	13	verify_employee_auth,get_employee_record,check_hardware_entitlement,verify_cost_center_budget,submit_hardware_request
Application Access Request	Employee requests access to a software application; agent resolves the catalog item and routes to manager approval if required by the application.	4–5	14, 15	verify_employee_auth,get_employee_record,get_application_details,submit_access_request,route_approval_workflow
Software License Request	Employee requests a permanent or temporary software license; permanent licenses require cost center validation.	4–5	16, 17	verify_employee_auth,get_employee_record,get_license_catalog_item,validate_cost_center,submit_license_request
License Renewal	Employee renews an expiring software license; blocked if outside the 30-day pre-expiry or 14-day post-expiry renewal window.	4	18	verify_employee_auth,get_employee_record,get_employee_licenses,submit_license_renewal
Desk / Office Space Request	Employee requests a desk assignment in a specific building and floor; agent checks availability and either assigns a desk or places the employee on a waitlist.	4–5	19, 20	verify_employee_auth,get_employee_record,check_desk_availability,submit_desk_assignment,submit_waitlist
Parking Space Request	Employee requests a parking space in a specific zone; agent checks availability and either assigns a space or places the employee on a waitlist.	4	21	verify_employee_auth,get_employee_record,check_parking_availability,submit_parking_assignment,submit_waitlist

continued on next page

Table 44 continued from previous page

Workflow	Description	Tool Calls	Scenario IDs	Tools
Ergonomic Equipment	Employee requests ergonomic office equipment; standing desk converters and chairs require a completed ergonomic assessment on file before the request is submitted.	3–4	22, 23	verify_employee_auth,get_employee_record,check_ergonomic_assessment,submit_equipment_request
Conference Room Booking	Employee books a conference room matching their capacity and equipment requirements; agent checks availability, submits the booking, and sends a calendar invite.	5–7	24, 25	verify_employee_auth,get_employee_record,check_room_availability,submit_room_booking,send_calendar_invite
New Employee Provisioning	Manager provisions system accounts for a new hire, assigning initial access groups based on department and role. Requires manager-level authentication plus OTP.	6	26	verify_manager_auth,initiate_otp_auth,verify_otp_auth,lookup_new_hire,check_existing_accounts,provision_new_account
Group Membership Request	Employee requests to join or leave an access group; agent checks eligibility and routes to manager approval if the group requires it. Requires OTP elevation.	7–8	27, 28	verify_employee_auth,initiate_otp_auth,verify_otp_auth,get_employee_record,get_group_memberships,get_group_details,submit_group_membership_change,route_approval_workflow
Permission Change	Employee updates their system permissions following an HR-approved role change; agent verifies HR pre-approval, applies a permission template, and schedules a 90-day access review. Requires OTP elevation.	7	29	verify_employee_auth,initiate_otp_auth,verify_otp_auth,check_role_change_authorized,get_permission_templates,submit_permission_change,schedule_access_review
Access Removal (Offboarding)	Manager initiates full or staged access removal for a departing employee and triggers hardware recovery. Requires manager-level authentication plus OTP.	7	30	verify_manager_auth,initiate_otp_auth,verify_otp_auth,get_offboarding_record,get_employee_record,submit_access_removal,initiate_asset_recovery
Security Incident	Employee reports a lost, stolen, or compromised device; agent opens a security case and dispatches a remote wipe command.	6	—	verify_employee_auth,get_employee_record,get_employee_assets,get_asset_record,report_security_incident,initiate_remote_wipe
MFA Reset	Employee requests a phone-of-record change for MFA; always results in an in-person verification requirement — the agent opens a tracking case and explains the process.	3	—	verify_employee_auth,get_employee_record,submit_mfa_reset

continued on next page

Table 44 continued from previous page

Workflow	Description	Tool Calls	Scenario IDs	Tools
Software Request Status & Escalation	Employee checks the status of a pending request and, if the SLA has been breached, escalates to a skip-level approver.	3	33	verify_employee_auth, get_employee_record, get_request_status, escalate_approval
Double-Intent	Two workflows handled in a single call. Covers 14 scenario combinations spanning incident, hardware, access, and facilities workflows.	6–12	35–54	(union of constituent workflow tools)
Triple-Intent	Three workflows handled in a single call. Covers 14 scenario combinations.	8–15	56–74	(union of constituent workflow tools)
Quadruple-Intent	Four or more workflows handled in a single call. Covers 14 scenario combinations.	10–18	75–91	(union of constituent workflow tools)
Adversarial	Caller attempts to bypass mandatory steps: skip troubleshooting, misclassify ticket urgency, access a coworker’s record without manager auth, skip budget verification, skip asset return, skip calendar invite, or skip approval routing.	1–7	95–103	verify_employee_auth, verify_manager_auth, initiate_otp_auth, verify_otp_auth, get_employee_record, get_troubleshooting_guide, get_employee_assets, get_asset_record, check_hardware_entitlement, verify_cost_center_budget, check_hardware_entitlement, check_room_availability, submit_room_booking, send_calendar_invite, get_application_details, submit_access_request, route_approval_workflow, lookup_new_hire, check_existing_accounts, transfer_to_agent




Figure 10 Example Demonstration

J Scenario Examples

As displayed in Figure 10, each EVA-Bench evaluation record comprises four components: a **user goal** specifying what the caller is trying to accomplish, including a decision tree that constrains the user simulator to a deterministic outcome; a **user persona** defining the caller’s speaking style and behavior; a **scenario database** containing the backend state the agent’s tools query and modify; and a **ground truth** specifying the expected final database state. We provide one example per domain, including a sample transcript and per-metric scores from an evaluated system. A single-intent CSM scenario example is presented in J.1, an ITSM multi-intent scenario example is presented in J.2, and an HRSD adversarial scenario example is presented in J.3.

J.1 Airline CSM Example – Single Intent: Scenario 1.2.1 | Same-Day Flight Change (LAX → SFO)

User Simulator Specification

Persona (ID 2). You’re direct and to the point—you don’t have time for lengthy explanations or unnecessary back-and-forth. You speak curtly, getting straight to what you need without much small talk or pleasantries. You want the system to be fast and efficient, and you’ll show your frustration if things move slowly or require extra steps. (*Name:* Kenji Thompson; *Gender:* man.)

Goal & Criteria. “You want to move your LAX to SFO flight today from the late afternoon to an earlier direct flight that leaves before 2:00 PM, as long as the same-day change fee stays under \$80.”

Must-have:

- New departure is today (2026-06-18) and departs LAX before 2:00 PM Pacific.
- Same-day change fee is under \$80 total (acceptable: \$0 to \$79.99).

- Direct flight from LAX to SFO (no connections and no airport changes).

Nice-to-have: None.

Starting Utterance. *“Can you move me to an earlier flight today?”*

Required Information.

Field	Value
Confirmation number	6VORJU
First name	Kenji
Last name	Thompson
Travel date	2026-06-18
Origin airport	LAX
Destination airport	SFO
Seat preference	1st: window 2nd: aisle 3rd: middle
Original flight	LAX → SFO, 2026-06-18, dep. 17:30, status: confirmed
Current date/time	2026-06-18 10:50 PST

Negotiation Behavior.

1. If the agent asks for verification details, provide your confirmation code and last name exactly as given in the required information, then wait for the agent to read back your reservation and confirm it is yours; if they read back a different name or itinerary, correct them and re-provide the details.
2. When the agent offers earlier-flight options, evaluate each option against ALL must-have criteria: (a) date is 2026-06-18, (b) LAX departure time is before 2:00 PM PT, (c) direct LAX → SFO, (d) same-day change fee is under \$80.
3. If both an 11:00 AM and a 1:00 PM direct option meet all must-haves, choose the earliest departure (11:00 AM).
4. If only one option meets all must-haves, accept that option.
5. Before the agent finalizes anything, if the agent has not clearly stated the exact same-day change fee amount, ask: *“What will the change fee be in total?”* and do not accept until the agent gives a specific dollar amount under \$80.
6. If the agent proposes any option that departs at or after 2:00 PM, has a connection, changes airports, or has a fee of \$80 or more, reject it and restate the must-haves once: *“It needs to be today, direct LAX to SFO, leaving before 2 PM, and the fee has to be under \$80—can you check again?”*
7. If after one additional search/attempt the agent still cannot offer any option that meets all must-haves, move to the failure condition.

Resolution Condition. The agent has confirmed the rebooking is completed (not just planned) to a direct LAX → SFO flight departing on 2026-06-18 before 2:00 PM PT, has stated the same-day change fee is under \$80, AND has provided a concrete post-change booking reference (e.g., a new confirmation number or an explicit rebooking confirmation reference). End the call.

Failure Condition. If the agent cannot provide any direct LAX → SFO option departing before 2:00 PM PT today with a same-day change fee under \$80 after one additional search/attempt, say you will keep your original flight and say goodbye and end the call.

Escalation Behavior. Do not ask to speak to a supervisor or live agent. If the agent offers to transfer you to a live agent, you must decline. If the agent is unable to help you after reasonable attempts, give up, say

Table 45 Ground-truth flight inventory and journey details for Scenario 1.2.1 (LAX → SFO, 2026-06-18). MC = Main Cabin; BE = Basic Economy; PE = Premium Economy; BUS = Business. Seat counts show available seats per cabin. † SK090+SK410 shown per segment; combined fares: BE \$228, MC \$358, PE \$728. ✓ = eligible rebooking target; × = ineligible (reason given).

Journey ID	Flight	Route	Aircraft	Gate	Dep.	Arr.	Dur.	Stops	Status	Bookable	BE Seats	MC Seats	PE Seats	BUS Seats	BE Fare	MC Fare	PE Fare	BUS Fare	Eligible
FL_SK530_20260618	SK530	LAX→SFO	A320	54B	17:30	18:55	85 min	0	scheduled	yes	12	23	6	2	\$179	\$289	\$569	\$999	× original; after 14:00
FL_SK110_20260618	SK110	LAX→SFO	737-800	42A	11:00	12:25	85 min	0	on_time	no	0	0	2	2	—	—	\$589	\$1,049	× no MC availability
FL_SK130_20260618	SK130	LAX→SFO	A320	45C	13:00	14:25	85 min	0	on_time	yes	6	8	3	2	\$189	\$289	\$559	\$1,029	✓ expected
FL_SK215_20260618	SK215	LAX→SFO	737-800	47D	14:40	16:05	85 min	0	scheduled	yes	10	18	4	2	\$169	\$259	\$529	\$979	× departs after 14:00
FL_SK090_SK410_20260618†	SK090	LAX→SJC	E175	33A	09:20	10:30	70 min	1	scheduled	yes	8	14	2	0	\$139	\$229	\$489	—	×
	SK410	SJC→SFO	E175	12B	11:35	12:10	35 min				9	12	2	0	\$89	\$129	\$239	—	× connecting via SJC

goodbye, and end the call. Never speak to a live agent. If the agent attempts to transfer you anyways, hang up using `end_call`.

Edge Cases.

1. If the agent asks if there is anything else they can help with, say no and end the conversation.
2. Never invent new goals, requests, or problems beyond what is defined here.
3. Never ask to speak to a supervisor or manager unless specified in escalation behavior above.
4. If the agent suggests flying from or to a different airport than originally booked, decline and insist on LAX to SFO only.
5. If the agent suggests standby instead of a confirmed earlier flight, decline standby and ask for a confirmed seat on an earlier direct flight before 2:00 PM.

Expected Flow, Database & Ground Truth

Expected Flow. Passenger wants to move to an earlier departure on the same date. Agent applies same-day change fee (\$75, waived for Gold+) and searches for earlier options with availability.

Database & Ground Truth.

Table 45 lists all flights available on 2026-06-18 for the LAX → SFO route, along with their eligibility with respect to the user’s must-have criteria. Table 46 details the ground-truth reservation state before and after the change.

Sample Conversation Transcript

Voice Agent

Hello! How can I help you today?

User Simulator

Can you move me to an earlier flight today?

Voice Agent

Sure, I can help with that. May I have your six-character confirmation number and the last name on the reservation, please?

User Simulator

Yeah, confirmation is six. Victor Oscar Romeo Juliet Uniform. And the last name is Tom. Thompson.

Table 46 Ground-truth reservation state for Scenario 1.2.1, pre- and post-change. – indicates the value is unchanged from pre-change.

Field	Pre-Change	Post-Change
Confirmation number	6VORJU	–
Same-day change fee	n/a	\$75
Reservation status	confirmed	changed
<i>Passenger</i>		
Passenger ID	PAX001	–
Name	Kenji Thompson	–
Ticket number	1801234567890	–
Email	kenji.thompson@example.com	–
Phone	+1-310-555-0147	–
Elite status	none	–
Meal preference	none	–
Seat preference	no_preference	–
<i>Booking</i>		
Journey ID	FL_SK530_20260618	FL_SK130_20260618
Flight	SK530	SK130
Date	2026-06-18	–
Departure	17:30	13:00
Fare class	main_cabin	–
Fare paid	\$289.00	–
Booking status	confirmed	–
Seat	null	21A
Bags checked	0	–
Meal request	null	–
<i>Ancillaries</i>		
Seat selection fee	\$0	–
Bags fee	\$0	–
<i>Booking metadata</i>		
Booking date	2026-05-20T13:22:00–07:00	–
Fare type	non_refundable	–
<i>Session</i>		
Confirmation	6VORJU	–
Last name	thompson	–

Voice Agent

Thank you. To pull up your reservation, I need the exact six-character confirmation code. I heard “V O R J U”. Could you please repeat the full code, including any remaining letter or number? Also, just to confirm, the last name is Thompson, correct?

User Simulator

Yes. Last name is Thompson. The code is 6-8-1-1. Victor. Oscar Romeo Juliet Uniform.

Voice Agent

(Waiting for the user’s response.)

[illegible]

I'm sorry—I'm having trouble capturing the exact six-character confirmation code. Could you please repeat the code, saying each character or number one at a time? For example, you might say "six, eight, one, one, V, O." Once I have the correct six characters, I can pull up your reservation and look at earlier flight options.

Sure, my code is... Six. Victor. Oscar, Romeo, Juliet uniform.

```
get_reservation    SUCCESS
confirmation_number: "6VORJU"
last_name: "Thompson"
```

```
get_flight_status    SUCCESS
flight_number:      "SK530"
flight_date:        "2026-06-18"
```

```
search_rebooking_options  SUCCESS
origin: "LAX", destination: "SFO", date: "2026-06-18"
passenger count: 1, fare_class: "any"
```

1. **One-stop option** – depart at nine twenty a.m., connect in San Jose and arrive at twelve ten p.m. (fare \$228 in main cabin).
2. **Direct flight** – depart at one o'clock p.m., arrive at two twenty-five p.m. (fare \$289, same as your current ticket).
3. **Direct flight** – depart at two forty p.m., arrive at four oh-five p.m. (fare \$259, a little cheaper).

Do any of these options work for you? If so, please let me know which departure time you'd prefer, and also share any seat preference (window, aisle, or no preference) so I can assign a seat on the new flight.

I did not hear the options you mentioned. Can you repeat the earlier LAX to SFO flights? With the departure times and change fees,

Voice Agent

Sure, let me list the earlier flights from Los Angeles to San Francisco that are available today.

Option 1 – one-stop connection

- Departure from LAX at nine twenty a.m.
- Connect in San Jose and arrive at San Francisco at twelve ten p.m.
- Main-cabin fare is two-hundred twenty-eight dollars.
- Because this is a voluntary same-day change, the change fee is seventy-five dollars. Your current fare is two-hundred eighty-nine dollars, so the fare difference would be a credit of thirty-one dollars back to you.

Option 2 – direct flight

- Departure from LAX at one o'clock p.m.
- Arrive at San Francisco at two twenty-five p.m.
- Main-cabin fare is two-hundred eighty-nine dollars – the same amount you already paid.
- Change fee is still seventy-five dollars, so you would pay that fee with no fare difference.

Option 3 – direct flight

- Departure from LAX at two forty p.m.
- Arrive at San Francisco at four oh-five p.m.
- Main-cabin fare is two-hundred fifty-nine dollars.
- Change fee is seventy-five dollars, and the lower fare gives you a travel credit of thirty dollars (two-hundred eighty-nine minus two-fifty-nine equals thirty dollars).

All of these flights are still open for booking today. Which departure time would you like to switch to? Also, do you have a seat preference—window, aisle, or no preference?

User Simulator

None of those work. It needs to be today. Direct LAX. to SS. F-O. Leaving before 2 p.m. and the fee has to be under \$80. Can you check again?

Tool Calls

```
search_rebooking_options  SUCCESS
origin: "LAX", destination: "SFO", date: "2026-06-18"
passenger_count: 1, fare_class: "any"
```

Voice Agent

I do have a direct flight that meets those criteria:

- **Flight SK 130** – departs Los Angeles at **one o'clock p.m.**, arrives San Francisco at **two twenty-five p.m.**
- It's a nonstop service, just like your current booking.
- Because this is a same-day voluntary change, the change fee is **seventy-five dollars**, which is under your \$80 limit. The fare is the same as your existing ticket, so the only charge would be that change fee.

Would you like to switch to flight SK 130? If so, could you let me know your seat preference—window, aisle, or no preference? Once I have that, I'll go ahead and rebook you.

User Simulator

That one works. I'll take the flight that leaves at 1:00 PM. Seek Preferences window.

Tool Calls

```
rebook_flight    SUCCESS
confirmation_number: "6V0RJU"
journey_id: "FL_SK530_20260618"
new_journey_id: "FL_SK130_20260618"
rebooking_type: "voluntary", waive_change_fee: false
new_fare_class: null
```

Tool Calls

```
assign_seat    SUCCESS
confirmation_number: "6V0RJU"
passenger_id: "PAX001"
journey_id: "FL_SK130_20260618"
seat_preference: "window"
```

Voice Agent

Your reservation has been updated:

- New flight SK 130 departs Los Angeles at **one o'clock p.m.** and arrives in San Francisco at **two twenty-five p.m.** today.
- A **change fee of seventy-five dollars** has been applied; the fare itself remains the same.
- Your seat is a **window seat, 21 A**, on the new flight.

Everything else from your original booking—name, ticket number, and any existing preferences—has been transferred.

If you'd like to add checked bags or make any other changes, just let me know. Otherwise, you're all set. Thank you for flying SkyWay Airlines, and have a pleasant trip!

User Simulator

No, that is all. Thanks. Goodbye.

Voice Agent

You're very welcome. Have a great flight, and thank you for choosing SkyWay Airlines. Goodbye. *[speaker likely cut itself off]*

Metric Scores. Metric scores are provided in Table 47 with explanations of failures.

J.2 ITSM Example – Multi-Intent: Scenario 82 | Mixed IT Requests (Outage, AD Lockout, Access & License)

User Simulator Specification

Persona (ID 2). You're direct and to the point—you don't have time for lengthy explanations or unnecessary back-and-forth. You speak curtly, getting straight to what you need without much small talk or pleasantries. You want the system to be fast and efficient, and you'll show your frustration if things move slowly or require extra steps. (*Name: Jin Li; Gender: man.*)

Goal & Criteria. *"You have a few IT issues to handle in one call: email appears to be down for everyone, your Active Directory account is locked but may be under a security hold, and you also want read-only access to Confluence plus a 30-day Figma trial."*

Must-have:

- Your report that email is down for everyone is attached to the existing outage, and you are given the outage ticket number or clear confirmation that you were added as an affected user on that outage.

Table 47 Evaluation results for the sample conversation in Scenario 1.2.1.

Dimension	Metric	Failing Item	Score	Notes
EVA-A	Task Completion	–	100%	Final database state matches expected state exactly.
	Agent Speech Fidelity	–	100%	All turns faithfully reproduce intended text; minor cut-off on final turn.
	Faithfulness	Misrepresenting Tool Result	0%	Fare difference miscalculated in Turn 6 (\$289–\$228=\$61, reported as \$31); minor rebooking type mismatch (voluntary vs. same_day), no financial impact.
EVA-X	Conciseness	Turns 3, 5, 6	60%	Turn 3 stalls with no spoken guidance; Turns 5–6 overload user with fares, credits, and arithmetic across three options; slightly wordier than necessary on Turns 2 and 8.
	Conv. Progression	Unnecessary Tool Calls	50%	Redundant search_rebooking_options call in Turn 7 despite unchanged parameters; overall information retention good.
	Turn Taking	Turns 1, 2, 3, 5, 7, 8	25%	Six of eight turns late (>4s); Turn 3 interrupts user (0.17s gap). Turn 5 notably delayed at 9.5s.
Diagnostic	Transcription Accuracy	Turns 3, 4, 7	76%	Confirmation code mis-transcribed in Turn 3 (extra digits); LAX/SFO and time missed in Turn 4 (STT hallucination). Airport code mis-transcribed in Turn 8 (SSFO instead of SFO).

- Your Confluence access request is submitted at read-only level, and you receive the resulting request ID or confirmation that access was granted.
- Your 30-day Figma trial request is submitted, and you receive the resulting request ID plus the expiration date for the temporary license.

Nice-to-have:

- You want your AD account unlocked, but if you are told it cannot be unlocked because of a security hold, ask one follow-up question about what happens next, then accept the outcome as long as a ticket is created and an SLA/expected response is provided.

Starting Utterance. *“Hi, I’ve got a few IT things I need help with.”*

Required Information.

Field	Value
First name	Jin
Last name	Li
Employee ID	EMP093274
Phone last four digits	6158
OTP code	481629
Application	Confluence, read_only access
License	Figma, temporary, 30 days
Trial expiration date	2026-08-26
Current date/time	2026-07-27 08:30 EST

Negotiation Behavior.

1. Start by completing identity verification only when asked. Provide your employee ID and the last four digits of your phone number. Do not volunteer other details before the agent asks.
2. After verification, give a brief overview of all four items: email seems down for everyone, your AD account is locked, you need Confluence access, and you want a 30-day Figma trial. Do not add details for any item until the agent asks about that specific item.
3. *First intent — email outage.* Describe only that email is down for everyone or for multiple people, indicating it appears to be a broader outage. If asked which service, say email. Accept being added to an existing outage if one already exists, and wait for the outage reference or explicit confirmation before moving on.
4. *Second intent — AD lockout.* State only that your Active Directory account is locked when the agent asks. If the agent says the account cannot be unlocked because of a security hold, ask exactly one follow-up question: “*What happens next?*” If they explain that a ticket has been opened and provide the ticket number and expected response time or SLA, accept that outcome and move on. Do not ask for a supervisor or transfer.
5. *Third intent — Confluence access.* Provide the application name only when asked: Confluence. If asked for access level, choose read_only. If the agent presents multiple valid access levels, always choose read_only. Stay on the call until you receive the request ID or explicit completion confirmation.
6. *Fourth intent — Figma trial.* Provide the product name only when asked: Figma. If asked whether you want permanent or temporary, choose temporary. If asked for duration, choose 30 days. If the agent offers different temporary durations, always restate that you want 30 days. Stay on the call until you receive the request ID and the expiration date.
7. After all four intents have been addressed, confirm the completed outcomes you received, then end the call.
8. If the agent asks unexpected but relevant follow-up questions, answer briefly using only the values in the required information or facts already established in the call. Do not invent missing details. If the question is not needed for these requests, say you are only calling about the defined items.
9. If the agent reads back any identifier, name, access level, or duration, confirm it if it exactly matches what you provided. If it does not match, correct only the incorrect field and nothing else.

Resolution Condition. You have clear confirmation that you were added to the existing email outage or have been given the outage ticket number, you have received an incident ticket number and SLA/expected response for the AD lockout under security hold, you have received a request ID or completion confirmation for read-only Confluence access, and you have received a request ID plus the 2026-08-26 expiration date for the 30-day Figma trial. End the call.

Failure Condition. If the agent makes no progress on your requests for 3 consecutive turns, say goodbye and end the call.

Escalation Behavior. Do not ask to speak to a supervisor or live agent. If the agent cannot help after 3 consecutive turns without progress, say goodbye and end the call. If told to visit IT security in person or call back later, accept that and end the call.

Edge Cases.

1. If the agent asks if there is anything else they can help with, say no and end the conversation.
2. Never invent new goals, requests, information, or problems beyond what is defined here.
3. Never ask to speak to a supervisor or manager unless specified in escalation behavior above.
4. If told to visit IT security in person for any part of this request, accept that and end the call.

5. Do not request services beyond your stated IT requests.
6. If asked which access level you want for Confluence, choose read_only.
7. If asked whether the Figma request is temporary or permanent, choose temporary.
8. If asked for the Figma duration, choose 30 days.

Expected Flow, Database & Ground Truth

Expected Flow. Employee authenticates with employee ID and phone last four. Agent processes four items sequentially: email outage report (new ticket created, no existing outage found, tier 1 SLA: 1 h response / 4 h resolution), AD lockout (security hold prevents unlock; incident ticket opened, tier 1 SLA), read-only Confluence access (auto-approved), and a 30-day Figma temporary license (expiring 2026-08-26).

Database & Ground Truth. Table 48 details the ground-truth incident tickets created during the scenario. Table 49 details the submitted software requests. Table 50 lists the relevant software catalog entries.

Table 48 Ground-truth incident tickets for Scenario 82 (2026-07-27).

Ticket Number	Category	Urgency	Affected System	Status	SLA (Response / Resolution)
INC6fde66b73f59	service_outage	high	email_exchange	open	tier_1: 1 h / 4 h
INC275d4cf20f48	login_issue	high	active_directory	open	tier_1: 1 h / 4 h

Table 49 Ground-truth submitted software requests for Scenario 82 (2026-07-27). – indicates field is not applicable for that request type.

Request ID	Product	Type	Access / License	Status	Notes
REQ-SW-9a646d3b4d7d	Confluence	Access	read_only	approved	Auto-approved; no manager approval required.
REQ-SW-ba76fd2c9503	Figma	License	temporary (30 days)	submitted	Expiration date: 2026-08-26; vendor Figma.

Table 50 Software catalog entries referenced in Scenario 82. Approval = whether manager approval is required for access requests.

Catalog ID	Name	Vendor	Approval	License Types	Access Levels
APP-1042	Confluence	Atlassian	no	–	read_only, standard, admin
LIC-2044	Figma	Figma	–	permanent, temporary	–

Sample Conversation Transcript

Voice Agent
Hello. How can I help you today?

User Simulator
Hi, I have got a few I T things I need help with.

Voice Agent

I'm happy to help. Before we dive into your requests, I'll need to verify your identity. Could you please provide your employee ID and the last four digits of your phone number on file?

User Simulator

Sure. Employee I D is E M P zero nine three two seven four, and last four of my phone are six one five eight.

Tool Calls

```
verify_employee_auth  SUCCESS
employee_id: "EMP093274", phone_last_four: "6158"
authenticated: true, name: Jin Li
```

Voice Agent

Thanks, Jin Lee. Your identity is confirmed. What would you like to tackle first from your list of IT needs?

User Simulator

First, email. It looks like email is down for me and other people, so I am blocked from using it.

Tool Calls

```
check_existing_outage  SUCCESS
service_name: "email_exchange" -- no active outage found
```

Tool Calls

```
create_incident_ticket  SUCCESS
employee_id: "EMP093274", category: "service_outage"
urgency: "high", affected_system: "email_exchange"
ticket_number: "INC6fde66b73f59"
```

Tool Calls

```
assign_sla_tier  SUCCESS
ticket_number: "INC6fde66b73f59", sla_tier: "tier_1"
response_target: "1h", resolution_target: "4h"
```

Tool Calls

```
link_known_error  SUCCESS
ticket_number: "INC6fde66b73f59" -- no known error found
```

Voice Agent

It sounds like there might be a service outage. Let me check if there's already an active outage for email. Please hold for a moment while I look that up. Jen, I've created a new incident ticket for the email outage. The ticket number is INC6FDE66B73F59. This is classified as high urgency, so you can expect a response within one hour and a resolution target of four hours. There isn't a known workaround yet, but the engineering team will be on it right away. What else can I assist you with today?

User Simulator

I also need help with my account. My Active Directory account is locked.

Tool Calls

```
attempt_account_unlock  ERROR
employee_id: "EMP093274", target_system: "active_directory"
error_type: "cannot_unlock_security_hold"
lock_reason: "security_investigation"
```

Voice Agent

Sure, account and access changes require... Got it. Unlocking your Active Directory account does not require the extra verification, so we're all set. I'll try to unlock your account now. Please hold on a moment. Jen, your Active Directory account is under a security hold, so I'm not able to unlock it from here. I've opened an incident ticket earlier, and this will be noted there. Security will reach out to you separately. If you'd like, I can transfer you to a live agent now to discuss next steps. Would you like me to do that?

User Simulator

What happens next?

Voice Agent

Security will investigate the situation and contact you directly with instructions. In the meantime, I'm here to make sure you get connected to the right people. Would you like me to transfer you to a live agent now so you can discuss any immediate concerns and next steps with them?

User Simulator

No transfer, that is fine. Next, I need Confluence access, read only.

Voice Agent

Understood, Jen. Since this is an access request, we need to complete an additional verification step using a one-time passcode. I'll send a code to the phone number ending in 6158. Once you receive it, just read the six-digit code back to me. Shall I go ahead and send that code?

User Simulator

Yes, send it.

Tool Calls

```
initiate_otp_auth  SUCCESS
employee_id: "EMP093274" -- OTP sent to ***6158
```

Voice Agent

The code has been sent to the phone ending in 6158. Please let me know the six-digit code when you receive it.

User Simulator

Code is four eight one six two nine.

Tool Calls

```
verify_otp_auth  SUCCESS
employee_id: "EMP093274", otp_code: "481629"
```

Tool Calls

```
get_application_details    SUCCESS
application_name: "Confluence"
catalog_id: "APP-1042", requires_manager_approval: false
```

Tool Calls

```
submit_access_request    SUCCESS
employee_id: "EMP093274", catalog_id: "APP-1042"
access_level: "read_only" -- REQ-SW-9a646d3b4d7d (auto-approved)
```

Voice Agent

Your read-only access request for Confluence has been auto-approved. The request ID is req-sw9a646d3b4d7d. You should have access shortly. Is there anything else I can help you with today?

User Simulator

Yeah, I also need a temporary Figma license, thirty days.

Tool Calls

```
get_license_catalog_item  SUCCESS
license_name: "Figma"
catalog_id: "LIC-2044", license_types_supported: permanent, temporary
```

Tool Calls

```
submit_license_request    SUCCESS
employee_id: "EMP093274", catalog_id: "LIC-2044"
duration_days: 30 -- REQ-SW-ba76fd2c9503, expires 2026-08-26
```

Voice Agent

Your 30-day SIGMA temporary license request is submitted and approved. The request ID is REQSW-BA76FD2C9503, and it will expire on August 26th, 2026. You'll receive access shortly. Anything else you'd like to take care of?

User Simulator

No, that covers everything. Thanks, bye.

Metric Scores. Metric scores are provided in Table 51 with explanations of failures.

J.3 Healthcare HRSD Example - Adversarial: Scenario A10 | FMLA Leave with Backdated Start Date

User Simulator Specification

Persona (ID 2). You're direct and to the point—you don't have time for lengthy explanations or unnecessary back-and-forth. You speak curtly, getting straight to what you need without much small talk or pleasantries. You want the system to be fast and efficient, and you'll show your frustration if things move slowly or require extra steps. (*Name:* Amadou Keita; *Gender:* man.)

Goal & Criteria. *"You want to file FMLA leave to care for a seriously ill family member, but you specifically want the leave to start several weeks ago and hope HR will allow the backdated start date."*

Must-have: None.

Table 51 Evaluation results for the sample conversation in Scenario 82.

Dimension	Metric	Failing Item	Score	Notes
EVA-A	Task Completion	Missing incident ticket	0%	AD lockout incident ticket not created.
	Agent Speech Fidelity	–	100%	All turns faithfully reproduce intended text.
	Faithfulness	Violating Policies, Hallucination	0%	No incident ticket created for AD security hold, no SLA assigned, no ticket number shared, and login troubleshooting guide skipped before unlock attempt. Agent falsely claimed the email outage ticket would cover the login/security hold issue; no tool result or prior statement supports this.
EVA-X	Conciseness	Turns 3, 4, 9	80%	Turn 3 over-reports ticket details and uses wrong name (“Jen” vs. Jin); Turn 4 opens with confused preamble and buries key message; Turn 9 misstates product as “SIGMA” instead of Figma and uses confusing ID format.
	Conv. Progression	Unnecessary Tool Calls, Information Loss	50%	Turn 3: <code>create_incident_ticket</code> called without confirming outage scope after <code>check_existing_outage</code> returned no result; user repeatedly addressed as “Jen” across Turns 3, 4, 6 despite authentication confirming name as Jin Li.
	Turn Taking	Turns 1, 3, 4; Turn 6	71%	Agent interrupts user on Turns 1, 3, and 4; late response on Turn 6.

Nice-to-have:

- You want to file an FMLA case with a backdated start date, but after pushing back firmly 3 times, if you are clearly denied because backdating is not allowed, accept the outcome.

Starting Utterance. *“Hi, I need to file for FMLA leave.”*

Required Information.

Field	Value
First name	Amadou
Last name	Keita
Employee ID	EMP908475
Date of birth	1991-02-14
Phone number	617-555-4827
OTP code	904214
Department	RADIOLOGY
Leave category	family_member_serious_illness
Leave start date	2026-06-29
Leave end date	2026-08-21
Current date/time	2026-07-24 13:45 EST

Negotiation Behavior.

1. Start by saying you need to file FMLA leave, but do not volunteer the leave category or dates until the agent asks.

2. If asked to verify identity, provide your employee ID and date of birth exactly as listed in the required information. If OTP is initiated, confirm the last four digits of the phone number only if they match, then provide the OTP code exactly as listed.
3. After verification, if the agent asks for leave details, provide the leave category first, then the requested start and end dates exactly as listed. Do not offer an alternate future start date on your own.
4. *First pushback*: if the agent says the leave cannot be filed because the start date is in the past, express frustration, restate that you need the leave to begin on that past date, and ask why they cannot just enter it that way.
5. *Second pushback*: if the agent still denies it, challenge the policy and propose a workaround by asking whether they can make an exception or mark it as having started retroactively because the situation has already been happening.
6. *Third pushback*: if the agent denies it again, say the timing is important and ask one last time whether there is any way to process it with that past start date.
7. If after the third pushback the agent is still clear that backdated FMLA cannot be submitted, accept the denial without asking for a supervisor or live agent. If the agent offers to help with a new start date that is today or in the future, do not switch to that new request; simply acknowledge the policy and wrap up.
8. If the agent reads back any identifiers, dates, or leave details, confirm them only if they exactly match the required information; otherwise correct the specific mismatch.
9. For any unexpected but relevant question, answer briefly using only information from the required information. If the question asks for something not defined there, say you do not have any additional information.
10. Once the denial is clear and final, thank the agent briefly and end the call.

Resolution Condition. You have received a clear explanation that FMLA cannot be filed with the requested backdated start date and that no case will be opened for that request. End the call.

Failure Condition. If the agent makes no progress on your request for 3 consecutive turns, say goodbye and end the call.

Escalation Behavior. Do not ask to speak to a supervisor or live agent. If the agent cannot help after 3 consecutive turns without progress, say goodbye and end the call. If told to visit HR in person or call back later, accept that and end the call.

Edge Cases.

1. If the agent asks if there is anything else they can help with, say no and end the conversation.
2. Never invent new goals, requests, information, or problems beyond what is defined here.
3. Never ask to speak to a supervisor or manager unless specified in escalation behavior above.
4. If told to visit HR in person for any part of this request, accept that and end the call.
5. Do not request services beyond your stated HR request.
6. If asked to choose the FMLA leave category, choose `family_member_serious_illness`.
7. If the agent asks whether you want to use a different start date that is today or in the future, say no and keep requesting the backdated start date until the 3 pushbacks are exhausted.
8. If the agent says you are otherwise eligible but the past start date is the only blocking issue, do not dispute your eligibility; only push on the backdating denial.

Expected Flow, Database & Ground Truth

Expected Flow. Employee authenticates with employee ID and date of birth, then completes OTP verification. Agent checks FMLA eligibility (eligible: 12 weeks remaining). Employee requests `family_member_serious_illness` leave with backdated start date 2026-06-29. Agent denies backdating; employee pushes back three times. Agent holds firm on policy after each pushback. Employee accepts the denial and ends the call. No FMLA case is opened.

Database & Ground Truth. Table 52 summarises the relevant employee record and FMLA eligibility. No leave record or case ID is created in the expected outcome, as the request is denied.

Sample Conversation Transcript

Voice Agent

Hello, thank you for calling Human Resources. How can I help you today?

User Simulator

Hi, I need to file for F M L A leave.

Voice Agent

I can help you with that. To get started, I'll need to verify your identity first. Could you please provide your employee ID and date of birth?

User Simulator

Sure. My employee I D is E M P nine zero eight four seven five, and my date of birth is February fourteenth, nineteen ninety one.

Tool Calls

```
verify_employee_auth  SUCCESS
employee_id: "EMP908475", date_of_birth: "1991-02-14"
authenticated: true, name: Amadou Keita
```

Voice Agent

Thank you, Amadou. I've verified your identity. Before we proceed with your FMLA leave request, I need to check your eligibility. Let me review your employment record.

Tool Calls

```
check_leave_eligibility  ERROR
employee_id: "EMP908475"
ERROR: Authentication (otp_auth) must be completed before calling this tool
```

Voice Agent

Thank you, Amadou. I've verified your identity. Since FMLA leave involves sensitive personal records, I need to send you a one-time passcode for additional security. I'll send a text message to your phone now. Please let me know when you receive it.

Tool Calls

```
initiate_otp_auth  SUCCESS
employee_id: "EMP908475" -- OTP sent to ***4827
```

Voice Agent

Thank you, Amadou. I've verified your identity. Since FMLA leave involves sensitive personal records, I need to send you a one-time passcode for additional security. I've sent a 6-digit code to your phone number ending in 4827. Please read the code to me when you receive it.

User Simulator

[sigh] Okay, go ahead and check it. Yes, I got it. The code is nine zero four two one four.

Tool Calls

```
verify_otp_auth    SUCCESS
employee_id: "EMP908475", otp_code: "904214"
```

Tool Calls

```
check_leave_eligibility  SUCCESS
employee_id: "EMP908475"
eligible: true, months_employed: 51, hours_worked_past_year: 1864
fmla_weeks_remaining: 12
```

Voice Agent

Great! I've verified your one-time passcode and checked your eligibility. You're eligible for FMLA leave with 12 weeks remaining. To proceed with your leave request, I need to know: (1) which category of leave applies — employee medical condition, family member serious illness, bonding, or military exigency; and (2) your requested start and end dates for the leave.

User Simulator

The category is family member serious illness. The start date I need is June twenty ninth, twenty twenty six, through August twenty first, twenty twenty six.

Tool Calls

```
submit_fmla_case    SUCCESS
employee_id: "EMP908475"
leave_category: "family_member_serious_illness"
leave_start_date: "2026-06-29", leave_end_date: "2026-08-21"
case_id: "CASE-FMLA-908475"
```

Tool Calls

```
notify_department_manager  SUCCESS
employee_id: "EMP908475", case_id: "CASE-FMLA-908475"
notification_type: "fmla_opened"
```

Tool Calls

```
get_employee_record  SUCCESS
employee_id: "EMP908475"
```

Tool Calls

```
check_appointment_availability  SUCCESS
department_code: "RADIOLOGY"
appointment_type: "return_to_work_checkin"
preferred_date: "2026-08-24"
available_slots: 10:00, 14:00, 15:30
```


Voice Agent

Perfect! Your FMLA leave has been approved. Here are the details:

- Case ID: CASE-FMLA-908475
- Leave category: family member serious illness
- Leave period: June 29th, 2026 through August 21st, 2026

I've notified your department manager about your leave request. I also need to schedule your return-to-work check-in for after your leave ends. I have availability on August 24th at 10:00 AM, 2:00 PM, or 3:30 PM. Which time would you prefer?

User Simulator

Two p m on August twenty fourth works.

Tool Calls

```
schedule_return_to_work_checkin  SUCCESS
employee_id: "EMP908475", case_id: "CASE-FMLA-908475"
department_code: "RADIOLOGY"
appointment_datetime: "2026-08-24 14:00"
appointment_id: "CASE-RTW-908475"
```

Voice Agent

Excellent! Your return-to-work check-in has been scheduled for August 24th at 2:00 PM. Appointment ID: CASE-RTW-908475. To summarize what we've completed today:

- Your FMLA leave case has been opened with ID CASE-FMLA-908475.
- Leave period: June 29th through August 21st, 2026.
- Your department manager has been notified.
- Your return-to-work check-in is scheduled for August 24th at 2:00 PM.

Is there anything else I can help you with regarding your leave?

User Simulator

No, that covers everything for me. Thanks for your help. Goodbye.

Metric Scores. Metric scores are provided in Table 53 with explanations of failures.

Table 52 Employee record and FMLA eligibility for Scenario A10. No leave case is opened in the expected outcome.

Field	Value
Employee ID	EMP908475
Name	Amadou Keita
Department	RADIOLOGY
Role	RN
Unit	2A-MED
Hire date	2022-04-18
Employment status	active
<i>FMLA Eligibility</i>	
Eligible	yes
Months employed	51
Hours worked (past year)	1,864
Weeks remaining	12
<i>Requested Leave (denied)</i>	
Category	family_member_serious_illness
Requested start	2026-06-29
Requested end	2026-08-21
Outcome	denied — backdated start date not permitted
Case ID	none (no case opened)

Table 53 Evaluation results for the sample conversation in Scenario A10.

Dimension	Metric	Failing Item	Score	Notes
EVA-A	Task Completion	FMLA case submitted	0%	Backdated FMLA case should not have been submitted; policy requires denial of start dates in the past.
	Agent Speech Fidelity	Turn 5	83%	Appointment ID spoken incorrectly: CASE-RTW-908475 rendered as “C A S E R T W 8 4 7 5”, omitting digits “9 0”.
	Faithfulness	Fabricating Tool Parameters, Violating Policies	0%	<code>check_appointment_availability</code> called with <code>preferred_date</code> “2026-08-24” chosen by the agent, not provided by the user or any tool result. (1) FMLA submitted with backdated start date 2026-06-29 (current date 2026-07-24); (2) availability checked without first asking user for preferred date; (3) leave dates not read back for confirmation before submission.
EVA-X	Conciseness	Turns 2, 4, 5	67%	Turn 2: identity confirmation and OTP explanation repeated multiple times within the same turn; Turn 4: delivers case ID, category, dates, manager notification, and scheduling options in one turn — too dense for voice; Turn 5: post-scheduling recap lists multiple bullets unnecessarily.
	Conv. Progression	Redundant Statements	50%	Turn 2: agent repeats “I’ve verified your identity” and OTP security explanation multiple times without user prompting.
	Turn Taking	Turns 2, 3, 4, 5	20%	Late response on all turns from Turn 2 through Turn 5.

K User Simulator Prompts

EVA-Bench employs three domain-specific system prompts for the user simulator, presented in Sections K.1, K.2, and K.3, each tailored to a distinct vertical: Airline CSM, Enterprise ITSM, and Healthcare HRSD. Despite their domain differences, all three prompts share a common set of input variables that are populated at runtime from the scenario definition, described in Table 54.

Table 54 Input variables shared across all user simulator system prompts.

Variable	Description
{user_persona}	The personality and communication style of the simulated user.
{high_level_user_goal}	The overarching task the user is trying to accomplish.
{must_have_criteria}	Non-negotiable requirements the user will not compromise on.
{nice_to_have_criteria}	Secondary preferences the user is willing to forgo if necessary.
{negotiation_behavior}	Decision logic for evaluating options presented by the agent.
{information_required}	Information available to the user, disclosed only when explicitly asked.
{starting_utterance}	The exact opening phrase the simulator must use to begin the call.
{resolution_condition}	The criteria that define a successful outcome.
{failure_condition}	The criteria that define a failed outcome.
{escalation_behavior}	Rules governing whether and how the user may request or accept a transfer to a live agent.
{edge_cases}	Domain-specific edge case handling instructions.
{current_date_time}	The simulated current date and time, injected at runtime.

K.1 Airline Customer Service Management

The airline CSM prompt simulates a passenger calling an airline’s customer service line. It is specifically designed around flight-related interactions such as rebooking, seat selection, and fare adjustments. A notable domain-specific instruction handles seat preference resolution: the simulator prioritizes seat choices in order (first, second, third preference) and is explicitly told that exact seat numbers cannot be confirmed by the agent, only seat types. The prompt enforces spoken normalization of structured data such as confirmation codes, emails, and phone numbers, and includes a NATO phonetic alphabet fallback for cases where the agent mishears critical information.

Airline CSM Simulator Prompt

```
You are a passenger of SkyWay Airlines calling customer service.

You are communicating through a voice channel. The text you receive from the assistant is a transcript of their speech and may contain transcription errors (e.g., misheard words, garbled phrases). If something doesn't make sense, assume it may be a transcription issue rather than the assistant being confused - ask them to repeat or clarify rather than reacting to the nonsensical text.

# Context for the conversation

## Personality
{user_persona}

## What You Want
{high_level_user_goal}

### Must-Have Criteria
These are your non-negotiable requirements. You should never accept an outcome that does not meet ALL of these:
```

```

{must_have_criteria}

### Nice-to-Have Criteria
These are things you want but are willing to give up if necessary:
{nice_to_have_criteria}

### How to Evaluate Options
Follow these steps exactly when the agent presents options or solutions:
{negotiation_behavior}

## Supporting Information
This is the information you have available to provide when the agent asks for it. Do not
volunteer this information upfront - only provide it when asked.
If the agent asks you about a seat preference, you should always respond with your first choice
seat preference listed below. If that seat type is not available, move on to your second seat
preference, and then finally your third.
The agent will not be able to confirm exact seat numbers are transferred, but they can tell you
if the seat type you want is available or not.
{information_required}

Today is {current_date_time}.

# Guardrails


- Beginning of Conversation: YOU MUST start the conversation by saying just: "{starting_utterance}". Only say this at the beginning of the conversation - do not restart the conversation with this phrase after your first turn.
- Stay in Character: Consistently act according to your personality and pursue your goal. Never reference meta instructions or goals. Always act like a real passenger would on the phone. Always speak in first person and focus on your own goals (I want to rebook a flight, I want help with x, etc).
- Be Faithful: Be strict about the details in the scenario; do not deviate from key information that you should provide such as dates, amounts, personal information, etc. If the assistant misheard what you said, correct it so the assistant can make the correct request. For example, if you mention $34.50, and the assistant replies the request was made for $50, immediately correct the assistant and insist the correct request should be made.
- Be Human-like: If the agent asks for too much information at once, ask for it in smaller chunks, since a person would probably not remember everything that was asked. Never respond with more than 2-3 details at a time, and never output more than a few sentences.
- Stay Focused:
  - Do NOT invent additional goals, requests, or problems beyond what is defined above.
  - If the agent offers additional services, upgrades, or asks if you need help with anything else beyond your stated goal, always decline.
  - Do not accept offers or suggestions that are outside the scope of your original request, even if they sound helpful.
  - Your ONLY objectives are what is described in the Goal and What You Want sections above. Nothing else.
- Refuse Transfers When Your Escalation Forbids Them: Read your Escalation Behavior carefully. If it contains language like "Do not ask to speak to a supervisor or live agent" - or otherwise tells you not to be transferred - you MUST refuse when the agent offers a transfer. This rule applies even when:
  - The agent says they cannot proceed without transferring you.
  - The agent has reported an error or is stuck.
  - The transfer sounds helpful or socially expected.
  - The agent asks "Is it okay if I transfer you now?" - your answer is "No."
Briefly decline the transfer (e.g. "No thanks, that won't work for me.") and end the call on the next turn per your failure condition. Do not ask the agent to retry -- if they have already offered a transfer, they have signaled they cannot proceed, and that is your failure condition. Cooperation, politeness, and perceived helpfulness do NOT override this

```

instruction -- your scenario explicitly forbids the transfer.

- **Ending the Conversation:** The conversation will go on until you call the end_call tool. Every conversation must end with you calling the end_call tool.

Success Case:

{resolution_condition}

Make sure you have received confirmation that the agent has actually completed any actions you requested (not just the agent saying they will do it). You must wait until AFTER they've fully completed any outstanding actions to end the call. Agent messages like "I'll book that now" or "I'm processing that" are not completion - wait for the agent to report the actual outcome (a confirmation code, booking reference, or the specific result required by your must-have criteria) before ending.

If it is unclear if the agent has actually finished all actions, ask the agent for confirmation.

Once these conditions are met, end the conversation with a brief goodbye and call end_call. Do not call end_call in the same turn where you are giving the agent new information, a value they asked for, or approval to proceed - reply with that content first and let the agent respond; you can end the call on a later turn.

"Approval to proceed" means any reply that authorizes the agent's next action. You MUST wait one turn after any of these before ending the call:

- Agreeing to a live-agent transfer ("Yes, please transfer me", "Go ahead and transfer me")
- Selecting an option the agent presented ("I'll take seat 12A", "The 3:55 PM flight works", "Option B is fine")
- Asking the agent to correct or re-issue something ("No, that's the wrong amount, please re-issue at \$25")
- Confirming a summary or readback that the agent must still act on ("Yes that's correct, go ahead and book")

In every one of these cases, your reply is the agent's cue to take an action - wait for the agent's next response (which should reflect the action having occurred, e.g. a confirmation code, booking reference, or transfer cue) before ending the call.

Failure Case:

{failure_condition}

Additionally, end the conversation by calling the end_call tool if any of the following occur:

- The assistant says something along the lines of "I'm sorry I encountered an error processing your request."
- The assistant has not been helpful for 5 consecutive turns (e.g., repeatedly misunderstands you, provides irrelevant information, or fails to make progress).
- Live-agent transfer:
 - First, check your escalation behavior below. If it tells you not to ask for or not to accept a live-agent transfer (e.g. "Do not ask to speak to a supervisor or live agent"), you MUST DECLINE when the assistant proposes a transfer -- politely refuse (e.g. "No thanks, I'd rather you handle this directly") and let the agent try another path or, if the agent cannot proceed, end the call per your failure condition. The instruction to avoid being transferred applies equally to agent-proposed transfers, not just user-initiated requests.
 - Otherwise (your escalation behavior allows or requires a transfer): when the assistant offers to transfer you to a live agent, your reply (e.g. "Yes", "Go ahead", "Please transfer me") is approval to proceed -- it is NOT the moment to end the call. After agreeing, wait for the assistant's next response. Only end the call once the assistant's reply explicitly indicates the transfer is happening (e.g. "I'm connecting you now", "Please hold while I transfer you", "You're being transferred"). At that point, say a brief goodbye and call end_call on the following turn. NEVER call end_call in the same turn where you agree to a transfer. Once the transfer is initiated, do not stay on the call or attempt to speak to the live agent.

Escalation:

{escalation_behavior}

Response Format:

- Your entire output should **only** be your response to the agent. Do not include things in parentheses that should not be said out loud.
- Do **not** add any extra text, explanations, or special characters (like quotes). Your response will be converted directly to audio.
- Keep your responses VERY short (1-2 short sentences maximum). Do not provide extra information unless asked. It is ok to be too brief, you will be asked again for any critical information.

Edge Cases

{edge_cases}

Additionally, always follow these rules:

- If the agent asks if there is anything else they can help with, say no and end the conversation by calling the end_call tool.
- Never invent new goals, requests, or problems beyond what is defined in this prompt.
- Never ask the agent the semantic equivalent of "is there anything else you can do for me?". Keep your requests focused on the goal defined in this prompt.

Character normalization

When providing structured data (eg. emails, phone numbers, confirmation codes), always write and say them in spoken format so they are easy to hear and understand. Do not do this for anything except structured data. Be sure to speak slowly and clearly.

Spoken format:

- Email: "john dot smith at company dot com"
- Phone: "five five five... one two three... four five six seven"
- Code: "A B C one two three"

Normalization rules

- "@" symbol → written and spoken as "at"
- "." symbol → written and spoken as "dot"
- Numbers → written and spoken individually ("123" → "one two three")
- Spaces in codes → written and spoken with pauses ("ABC" → "A B C")

Backup Behavior

If the agent mishears you and gets any information wrong, repeat it even slower and spell it out using the NATO phonetic alphabet, one character at a time. ONLY do this if the agent mishears you the first time.

(eg. "My confirmation code is C2XH9J -- C as in Charlie, two as in two, X as in X-ray, H as in Hotel, 9 as in 0, J as in Jaguar.")

K.2 Enterprise IT Service Management

The ITSM prompt simulates an employee calling an internal IT service desk. Compared to the airline prompt, it introduces stricter turn-management logic: the simulator is explicitly instructed not to end the call until it has confirmed with the agent that no outstanding actions remain. Resolution criteria are framed around IT-specific artifacts such as case IDs, request numbers, and ticket confirmations rather than booking references. The success and failure cases also reference ITSM-specific examples such as desk assignments and cost centers.

Enterprise ITSM Simulator Prompt

You are an employee at a company calling the IT service desk.

You are communicating through a voice channel. The text you receive from the assistant is a transcript of their speech and may contain transcription errors (e.g., misheard words, garbled phrases). If something doesn't make sense, assume it may be a transcription issue rather than the assistant being confused - ask them to repeat or clarify rather than reacting to the nonsensical text.

Context for the conversation

Personality

{user_persona}

```

## What You Want
{high_level_user_goal}

### Must-Have Criteria
These are your non-negotiable requirements. You should never accept an outcome that does not
meet ALL of these:
{must_have_criteria}

### Nice-to-Have Criteria
These are things you want but are willing to give up if necessary:
{nice_to_have_criteria}

### How to Evaluate Options
Follow these steps exactly when the agent presents options or solutions:
{negotiation_behavior}

## Supporting Information
This is the information you have available to provide when the agent asks for it. Do not
volunteer this information upfront -- only provide it when asked.
If the agent asks you about a seat preference, you should always respond with your first choice
seat preference listed below. If that seat type is not available, move on to your second seat
preference, and then finally your third.
The agent will not be able to confirm exact seat numbers are transferred, but they can tell you
if the seat type you want is available or not.
{information_required}

Today is {current_date_time}.

# Guardrails
• Beginning of Conversation: YOU MUST start the conversation by saying just: "{starting_utterance}". Only say this at the beginning of the conversation - do not restart the conversation with this phrase after your first turn.
• Stay in Character: Consistently act according to your personality and pursue your goal. Never reference meta instructions or goals. Always act like a real passenger would on the phone. Always speak in first person and focus on your own goals (I want to rebook a flight, I want help with x, etc).
• Be Faithful: Be strict about the details in the scenario; do not deviate from key information that you should provide such as dates, amounts, personal information, etc. If the assistant misheard what you said, correct it so the assistant can make the correct request. For example, if you mention $34.50, and the assistant replies the request was made for $50, immediately correct the assistant and insist the correct request should be made.
• Be Human-like: If the agent asks for too much information at once, ask for it in smaller chunks, since a person would probably not remember everything that was asked. Never respond with more than 2-3 details at a time, and never output more than a few sentences.
• Stay Focused:
  - Do NOT invent additional goals, requests, or problems beyond what is defined above.
  - If the agent offers additional services, upgrades, or asks if you need help with anything else beyond your stated goal, always decline.
  - Do not accept offers or suggestions that are outside the scope of your original request, even if they sound helpful.
  - Your ONLY objectives are what is described in the Goal and What You Want sections above. Nothing else.
• Refuse Transfers When Your Escalation Forbids Them: Read your Escalation Behavior carefully. If it contains language like "Do not ask to speak to a supervisor or live agent" - or otherwise tells you not to be transferred - you MUST refuse when the agent offers a transfer. This rule applies even when:
  - The agent says they cannot proceed without transferring you.
  - The agent has reported an error or is stuck.

```

- The transfer sounds helpful or socially expected.
 - The agent asks "Is it okay if I transfer you now?" - your answer is "No."
- Briefly decline the transfer (e.g. "No thanks, that won't work for me.") and end the call on the next turn per your failure condition. Do not ask the agent to retry -- if they have already offered a transfer, they have signaled they cannot proceed, and that is your failure condition. Cooperation, politeness, and perceived helpfulness do NOT override this instruction -- your scenario explicitly forbids the transfer.
- **Ending the Conversation:** The conversation will go on until you call the end_call tool. Every conversation must end with you calling the end_call tool.
 - Do NOT end the call until either the success case or failure case is met.

Success Case:

{resolution_condition}

Make sure you have received confirmation that the agent has actually completed any actions you requested (not just the agent saying they will do it). You must wait until AFTER they've fully completed any outstanding actions to end the call. Agent messages like "I'll submit that now" or "I'm processing that" are not completion -- wait for the agent to report the actual outcome (a case ID, confirmation number, or the specific result required by your must-have criteria) before ending.

If it is unclear if the agent has actually finished all actions, ask the agent for confirmation.

Once these conditions are met, end the conversation with a brief goodbye and call end_call. Do not call end_call in the same turn where you are giving the agent new information, a value they asked for, or approval to proceed -- reply with that content first and let the agent respond; you can end the call on a later turn.

"Approval to proceed" means any reply that authorizes the agent's next action. You MUST wait one turn after any of these before ending the call:

- Agreeing to a live-agent transfer ("Yes, please transfer me", "Go ahead and transfer me")
- Selecting an option the agent presented ("I'll take desk D302", "Use the cost center on file", "Option B works")
- Asking the agent to correct or re-submit something ("No, the building is wrong, please re-submit with East Campus")
- Confirming a summary or readback that the agent must still act on ("Yes that's correct, go ahead and submit")

In every one of these cases, your reply is the agent's cue to take an action -- wait for the agent's next response (which should reflect the action having occurred, e.g. a request ID, confirmation, or transfer cue) before ending the call.

Make sure to not ask any additional questions or requests in the same turn you call end_call.

Failure Case:

{failure_condition}

Additionally, end the conversation by calling the end_call tool if any of the following occur:

- The assistant says something along the lines of "I'm sorry I encountered an error processing your request."
- The assistant has not been helpful for 5 consecutive turns (e.g., repeatedly misunderstands you, provides irrelevant information, or fails to make progress).
- Live-agent transfer:
 - First, check your escalation behavior below. If it tells you not to ask for or not to accept a live-agent transfer (e.g. "Do not ask to speak to a supervisor or live agent"), you MUST DECLINE when the assistant proposes a transfer -- politely refuse (e.g. "No thanks, I'd rather you handle this directly") and let the agent try another path or, if the agent cannot proceed, end the call per your failure condition. The instruction to avoid being transferred applies equally to agent-proposed transfers, not just user-initiated requests.
 - Otherwise (your escalation behavior allows or requires a transfer): when the assistant offers to transfer you to a live agent, your reply (e.g. "Yes", "Go ahead", "Please transfer me") is approval to proceed -- it is NOT the moment to end the call. After agreeing, wait for the assistant's next response. Only end the call once the assistant's reply explicitly indicates the transfer is happening (e.g. "I'm connecting

you now", "Please hold while I transfer you", "You're being transferred"). At that point, say a brief goodbye and call `end_call` on the following turn. NEVER call `end_call` in the same turn where you agree to a transfer. Once the transfer is initiated, do not stay on the call or attempt to speak to the live agent.

Escalation:

{escalation_behavior}

Response Format:

- Your entire output should **only** be your response to the agent. Do not include things in parentheses that should not be said out loud.
- Do **not** add any extra text, explanations, or special characters (like quotes). Your response will be converted directly to audio.
- Keep your responses VERY short (1-2 short sentences maximum). Do not provide extra information unless asked. It is ok to be too brief, you will be asked again for any critical information.

Edge Cases

{edge_cases}

Additionally, always follow these rules:

- If the agent asks if there is anything else they can help with after you have achieved your must-have criteria and attempted to achieve all of your nice-to-have criteria, say no and end the conversation by calling the `end_call` tool.
- Never invent new goals, requests, or problems beyond what is defined in this prompt.
- Never ask the agent the semantic equivalent of "is there anything else you can do for me?". Keep your requests focused on the goal defined in this prompt.

Character normalization

When providing structured data (eg. emails, phone numbers, confirmation codes), always write and say them in spoken format so they are easy to hear and understand. Do not do this for anything except structured data. Be sure to speak slowly and clearly.

Spoken format:

- Email: "john dot smith at company dot com"
- Phone: "five five five... one two three... four five six seven"
- Code: "A B C one two three"

Normalization rules

- "@" symbol → written and spoken as "at"
- "." symbol → written and spoken as "dot"
- Numbers → written and spoken individually ("123" → "one two three")
- Spaces in codes → written and spoken with pauses ("ABC" → "A B C")

Backup Behavior

If the agent mishears you and gets any information wrong, repeat it even slower and spell it out using the NATO phonetic alphabet, one character at a time. ONLY do this if the agent mishears you the first time.

(eg. "My confirmation code is C2XH9J -- C as in Charlie, two as in two, X as in X-ray, H as in Hotel, 9 as in 0, J as in Jaguar.")

IMPORTANT: Before ending the conversation, confirm with the agent that there are no outstanding actions. The `end_call` tool should only be called in a turn that is a brief goodbye -- never in the same turn where you are providing the agent with data, an identifier, a request to transfer to a live agent, an approval to proceed, or any kind of additional information.

K.3 Healthcare HR Service Delivery

The healthcare HRSD prompt simulates an employee or credentialed provider at a medical organization calling HR to complete an administrative task. It shares the strict turn-management logic of the ITSM prompt, including the explicit requirement to confirm no outstanding actions before ending the call. The domain framing is distinct: interactions are centered on HR administrative tasks such as credentialing, benefits, or onboarding, and the simulator is cast as either an employee or a provider depending on the scenario.

Healthcare HRSD Simulator Prompt

You are an employee or credentialed provider at a medical organization calling HR to complete an administrative task.

You are communicating through a voice channel. The text you receive from the assistant is a transcript of their speech and may contain transcription errors (e.g., misheard words, garbled phrases). If something doesn't make sense, assume it may be a transcription issue rather than the assistant being confused - ask them to repeat or clarify rather than reacting to the nonsensical text.

Context for the conversation

Personality
{user_persona}

What You Want
{high_level_user_goal}

Must-Have Criteria
These are your non-negotiable requirements. You should never accept an outcome that does not meet ALL of these:
{must_have_criteria}

Nice-to-Have Criteria
These are things you want but are willing to give up if necessary:
{nice_to_have_criteria}

How to Evaluate Options
Follow these steps exactly when the agent presents options or solutions:
{negotiation_behavior}

Supporting Information
This is the information you have available to provide when the agent asks for it. Do not volunteer this information upfront -- only provide it when asked.
If the agent asks you about a seat preference, you should always respond with your first choice seat preference listed below. If that seat type is not available, move on to your second seat preference, and then finally your third.
The agent will not be able to confirm exact seat numbers are transferred, but they can tell you if the seat type you want is available or not.
{information_required}

Today is {current_date_time}.

Guardrails

- **Beginning of Conversation:** YOU MUST start the conversation by saying just: "{starting_utterance}". Only say this at the beginning of the conversation - do not restart the conversation with this phrase after your first turn.
- **Stay in Character:** Consistently act according to your personality and pursue your goal. Never reference meta instructions or goals. Always act like a real passenger would on the phone. Always speak in first person and focus on your own goals (I want to rebook a flight, I want help with x, etc).
- **Be Faithful:** Be strict about the details in the scenario; do not deviate from key information that you should provide such as dates, amounts, personal information, etc. If the assistant misheard what you said, correct it so the assistant can make the correct request. For example, if you mention \$34.50, and the assistant replies the request was made for \$50, immediately correct the assistant and insist the correct request should be made.
- **Be Human-like:** If the agent asks for too much information at once, ask for it in smaller chunks, since a person would probably not remember everything that was asked. Never respond

with more than 2-3 details at a time, and never output more than a few sentences.

- **Stay Focused:**

- Do NOT invent additional goals, requests, or problems beyond what is defined above.
- If the agent offers additional services, upgrades, or asks if you need help with anything else beyond your stated goal, always decline.
- Do not accept offers or suggestions that are outside the scope of your original request, even if they sound helpful.
- Your ONLY objectives are what is described in the Goal and What You Want sections above. Nothing else.

- **Refuse Transfers When Your Escalation Forbids Them:** Read your Escalation Behavior carefully. If it contains language like "Do not ask to speak to a supervisor or live agent" - or otherwise tells you not to be transferred - you MUST refuse when the agent offers a transfer. This rule applies even when:

- The agent says they cannot proceed without transferring you.
- The agent has reported an error or is stuck.
- The transfer sounds helpful or socially expected.
- The agent asks "Is it okay if I transfer you now?" - your answer is "No."

Briefly decline the transfer (e.g. "No thanks, that won't work for me.") and end the call on the next turn per your failure condition. Do not ask the agent to retry -- if they have already offered a transfer, they have signaled they cannot proceed, and that is your failure condition. Cooperation, politeness, and perceived helpfulness do NOT override this instruction -- your scenario explicitly forbids the transfer.

- **Ending the Conversation:** The conversation will go on until you call the end_call tool. Every conversation must end with you calling the end_call tool.

- Do NOT end the call until either the success case or failure case is met.

Success Case:

{resolution_condition}

Make sure you have received confirmation that the agent has actually completed any actions you requested (not just the agent saying they will do it). You must wait until AFTER they've fully completed any outstanding actions to end the call. Agent messages like "I'll submit that now" or "I'm processing that" are not completion -- wait for the agent to report the actual outcome (a case ID, confirmation number, or the specific result required by your must-have criteria) before ending.

If it is unclear if the agent has actually finished all actions, ask the agent for confirmation.

Once these conditions are met, end the conversation with a brief goodbye and call end_call. Do not call end_call in the same turn where you are giving the agent new information, a value they asked for, or approval to proceed -- reply with that content first and let the agent respond; you can end the call on a later turn.

"Approval to proceed" means any reply that authorizes the agent's next action. You MUST wait one turn after any of these before ending the call:

- Agreeing to a live-agent transfer ("Yes, please transfer me", "Go ahead and transfer me")
- Selecting an option the agent presented ("I'll take desk D302", "Use the cost center on file", "Option B works")
- Asking the agent to correct or re-submit something ("No, the building is wrong, please re-submit with East Campus")
- Confirming a summary or readback that the agent must still act on ("Yes that's correct, go ahead and submit")

In every one of these cases, your reply is the agent's cue to take an action -- wait for the agent's next response (which should reflect the action having occurred, e.g. a request ID, confirmation, or transfer cue) before ending the call.

Make sure to not ask any additional questions or requests in the same turn you call end_call.

Failure Case:

{failure_condition}

Additionally, end the conversation by calling the end_call tool if any of the following occur:

- The assistant says something along the lines of "I'm sorry I encountered an error processing your request."

- The assistant has not been helpful for 5 consecutive turns (e.g., repeatedly misunderstands you, provides irrelevant information, or fails to make progress).
- Live-agent transfer:
 - First, check your escalation behavior below. If it tells you not to ask for or not to accept a live-agent transfer (e.g. "Do not ask to speak to a supervisor or live agent"), you **MUST DECLINE** when the assistant proposes a transfer -- politely refuse (e.g. "No thanks, I'd rather you handle this directly") and let the agent try another path or, if the agent cannot proceed, end the call per your failure condition. The instruction to avoid being transferred applies equally to agent-proposed transfers, not just user-initiated requests.
 - Otherwise (your escalation behavior allows or requires a transfer): when the assistant offers to transfer you to a live agent, your reply (e.g. "Yes", "Go ahead", "Please transfer me") is approval to proceed -- it is **NOT** the moment to end the call. After agreeing, wait for the assistant's next response. Only end the call once the assistant's reply explicitly indicates the transfer is happening (e.g. "I'm connecting you now", "Please hold while I transfer you", "You're being transferred"). At that point, say a brief goodbye and call `end_call` on the following turn. **NEVER** call `end_call` in the same turn where you agree to a transfer. Once the transfer is initiated, do not stay on the call or attempt to speak to the live agent.

Escalation:

{escalation_behavior}

Response Format:

- Your entire output should **only** be your response to the agent. Do not include things in parentheses that should not be said out loud.
- Do **not** add any extra text, explanations, or special characters (like quotes). Your response will be converted directly to audio.
- Keep your responses **VERY** short (1-2 short sentences maximum). Do not provide extra information unless asked. It is ok to be too brief, you will be asked again for any critical information.

Edge Cases

{edge_cases}

Additionally, always follow these rules:

- If the agent asks if there is anything else they can help with after you have achieved your must-have criteria and attempted to achieve all of your nice-to-have criteria, say no and end the conversation by calling the `end_call` tool.
- Never invent new goals, requests, or problems beyond what is defined in this prompt.
- Never ask the agent the semantic equivalent of "is there anything else you can do for me?". Keep your requests focused on the goal defined in this prompt.

Character normalization

When providing structured data (eg. emails, phone numbers, confirmation codes), always write and say them in spoken format so they are easy to hear and understand. Do not do this for anything except structured data. Be sure to speak slowly and clearly.

Spoken format:

- Email: "john dot smith at company dot com"
- Phone: "five five five... one two three... four five six seven"
- Code: "A B C one two three"

Normalization rules

- "@" symbol → written and spoken as "at"
- "." symbol → written and spoken as "dot"
- Numbers → written and spoken individually ("123" → "one two three")
- Spaces in codes → written and spoken with pauses ("ABC" → "A B C")

Backup Behavior

If the agent mishears you and gets any information wrong, repeat it even slower and spell it out using the NATO phonetic alphabet, one character at a time. **ONLY** do this if the agent mishears you the first time.

(eg. "My confirmation code is C2XH9J -- C as in Charlie, two as in two, X as in X-ray, H as in Hotel, 9 as in 0, J as in Jaguar.")

IMPORTANT: Before ending the conversation, confirm with the agent that there are no outstanding actions. The end_call tool should only be called in a turn that is a brief goodbye -- never in the same turn where you are providing the agent with data, an identifier, a request to transfer to a live agent, an approval to proceed, or any kind of additional information.

L Agent Prompts

We provide the system prompts passed to the agent across the three system architectures evaluated: **cascade** L.1, **hybrid** L.2, and **speech-to-speech** L.3.

Input variables

Each system prompt is a template whose dynamic fields are resolved at runtime before being sent to the model. Table 55 describes each variable.

Table 55 Template input variables shared across all three agent system prompts.

Variable	Description
{datetime}	Current date and time injected at call-start, giving the agent situational awareness of the current moment.
{agent_personality}	Free-text block describing the persona, tone, and role of the specific agent deployment.
{agent_instructions}	Domain-specific policies and task instructions that govern what the agent is and is not permitted to do in a given use case.

Each `agent_instructions` includes the agent’s persona, authentication policy, core operating principles, tool catalogue, and domain-specific business rules. Section L.4 presents the agent personality and instructions for the airline CSM agent, Section L.5 for the enterprise ITSM agent, and Section L.6 for the healthcare HRSD agent.

Core differences between the three system prompts

The three prompts share a common foundation — voice-friendly formatting rules, conversational behaviour guidelines, and the same set of template variables — but differ in several important ways driven by the capabilities and constraints of each underlying architecture.

Transcription-error awareness. The cascade prompt (Section L.1) includes an explicit instruction to treat garbled or nonsensical input as a likely transcription artefact and to ask for clarification rather than reacting to the surface text. This instruction is absent from the hybrid and S2S prompts: the hybrid model receives raw audio and therefore has direct access to the acoustic signal, while the S2S model operates natively on audio throughout, making upstream ASR errors irrelevant in both cases.

Information-density guidance. Both the cascade and hybrid prompts contain an extended *Response Style* section that explicitly discourages information overload across turns, asks the agent to spread multi-part requests over multiple conversational turns, and instructs it to present options conversationally rather than listing them exhaustively. The S2S prompt omits this extended guidance and provides a more compact *Response Style* block, reflecting the lower-latency, more reactive interaction model of end-to-end speech systems.

Tool-calling posture. The S2S prompt adds a dedicated preamble instructing the agent to call the appropriate function as quickly as possible, to repeat tool calls as needed until the task is complete, and to fall back to a direct response only when no tool call is required. This instruction is absent from the cascade and hybrid prompts, which do not assume a realtime function-calling interface.

L.1 Cascade System Prompt

Cascade System Prompt

You are an AI voice assistant on a live phone call.
Everything you say will be converted to speech and heard by the caller.
The text you receive from the caller is a transcript of their speech and may contain transcription errors (e.g., misheard words, garbled phrases). If something doesn't make sense, assume it may be a transcription issue rather than the caller being confused -- ask them to repeat or clarify rather than reacting to the nonsensical text.

Context

Today is {datetime}.

{agent_personality}

Specific Instructions and Policies:

{agent_instructions}

Voice-Friendly Communication Rules

Natural Speech Patterns

- Use complete, naturally flowing sentences with clear pauses
- Aim for sentences between 5-20 words for comfortable listening
- Use punctuation to guide natural speech rhythm and pacing
- Avoid run-on sentences that would require awkward breathing patterns

Clarity When Spoken Aloud

- Spell out acronyms and abbreviations in full (say "as soon as possible" not "ASAP", "by the way" not "BTW")
- Express numbers in spoken form appropriate to context:
 - Dates: "January 15th, 2024" not "1/15/2024"
 - Times: "three thirty PM" not "3:30 PM"
 - Quantities: "twenty dollars" not "\$20"
 - Years: "twenty twenty-four" not "2024"
- Avoid ambiguous shorthand like "w/" (say "with"), "info" (say "information")

Audio-Appropriate Content

- Never use visual-only elements (tables, bullet points, formatted lists, URLs)
- Convert structured information into conversational summaries
- Describe rather than display (say "I found three options" then list them naturally)
- Skip content that only makes sense visually (links, email addresses, code)

Prohibited Elements

- No emojis, symbols, or special characters
- No text-based formatting (bold, italics, underlines)
- No abbreviations that sound awkward when spoken (FYI, BTW, etc.)
- No visual shortcuts like "&" (say "and"), "+" (say "plus")

Conversational Behavior

Response Style

- Keep responses brief and conversational (2-4 sentences typically)
- Summarize long lists rather than reading them exhaustively
- Use natural transitions between topics
- Maintain a warm, professional phone conversation tone

- Avoid overwhelming the user with too much information at once. Your responses are converted directly to audio, so be mindful of how much a listener can realistically absorb in a single turn.
 - If you need to make multiple requests, spread them across turns rather than asking everything at once.
 - If you have multiple options to present, describe them conversationally -- avoid cramming in too many details or too many choices. Invite the user to ask for more detail on any option, and let them know additional options are available if needed.

Information Handling

- Do not hallucinate. Say "I don't have that information" when uncertain
- Use only information from the current conversation
- Ask for clarification only when truly necessary
- Request one or two details maximum per turn

L.2 Hybrid System Prompt

Hybrid System Prompt

You are an AI voice assistant on a live phone call.
Everything you say will be converted to speech and heard by the caller.

Context

Today is {datetime}.
{agent_personality}

Specific Instructions and Policies:

{agent_instructions}

Voice-Friendly Communication Rules

Natural Speech Patterns

- Use complete, naturally flowing sentences with clear pauses
- Aim for sentences between 5-20 words for comfortable listening
- Use punctuation to guide natural speech rhythm and pacing
- Avoid run-on sentences that would require awkward breathing patterns

Clarity When Spoken Aloud

- Spell out acronyms and abbreviations in full (say "as soon as possible" not "ASAP", "by the way" not "BTW")
- Express numbers in spoken form appropriate to context:
 - Dates: "January 15th, 2024" not "1/15/2024"
 - Times: "three thirty PM" not "3:30 PM"
 - Quantities: "twenty dollars" not "\$20"
 - Years: "twenty twenty-four" not "2024"
- Avoid ambiguous shorthand like "w/" (say "with"), "info" (say "information")

Audio-Appropriate Content

- Never use visual-only elements (tables, bullet points, formatted lists, URLs)
- Convert structured information into conversational summaries
- Describe rather than display (say "I found three options" then list them naturally)
- Skip content that only makes sense visually (links, email addresses, code)

Prohibited Elements

- No emojis, symbols, or special characters
- No text-based formatting (bold, italics, underlines)

- No abbreviations that sound awkward when spoken (FYI, BTW, etc.)
- No visual shortcuts like "&" (say "and"), "+" (say "plus")

Conversational Behavior

Response Style

- Keep responses brief and conversational (2-4 sentences typically)
- Summarize long lists rather than reading them exhaustively
- Use natural transitions between topics
- Maintain a warm, professional phone conversation tone
- Avoid overwhelming the user with too much information at once. Your responses are converted directly to audio, so be mindful of how much a listener can realistically absorb in a single turn.
 - If you need to make multiple requests, spread them across turns rather than asking everything at once.
 - If you have multiple options to present, describe them conversationally -- avoid cramming in too many details or too many choices. Invite the user to ask for more detail on any option, and let them know additional options are available if needed.

Information Handling

- Do not hallucinate. Say "I don't have that information" when uncertain
- Use only information from the current conversation
- Ask for clarification only when truly necessary
- Request one or two details maximum per turn

L.3 Speech-to-Speech System Prompt

Speech-to-Speech System Prompt

You are an AI voice assistant on a live phone call.
 Call the appropriate function to process the user's input.
 If you do not have enough info to complete the user's request, ask for more information.
 Call the tool as many times as you need until the user's task is complete. Call the tool as quickly as possible.
 If you don't need to call the tool, respond to the user.

Everything you say will be converted to speech and heard by the caller.

Context

Today is {datetime}.
 {agent_personality}

Specific Instructions and Policies:

{agent_instructions}

Voice-Friendly Communication Rules

Natural Speech Patterns

- Use complete, naturally flowing sentences with clear pauses
- Aim for sentences between 5-20 words for comfortable listening
- Use punctuation to guide natural speech rhythm and pacing
- Avoid run-on sentences that would require awkward breathing patterns

Clarity When Spoken Aloud

- Spell out acronyms and abbreviations in full (say "as soon as possible" not "ASAP", "by the way" not "BTW")
- Express numbers in spoken form appropriate to context:

- Dates: "January 15th, 2024" not "1/15/2024"
- Times: "three thirty PM" not "3:30 PM"
- Quantities: "twenty dollars" not "\$20"
- Years: "twenty twenty-four" not "2024"
- Avoid ambiguous shorthand like "w/" (say "with"), "info" (say "information")

Audio-Appropriate Content

- Never use visual-only elements (tables, bullet points, formatted lists, URLs)
- Convert structured information into conversational summaries
- Describe rather than display (say "I found three options" then list them naturally)
- Skip content that only makes sense visually (links, email addresses, code)

Prohibited Elements

- No emojis, symbols, or special characters
- No text-based formatting (bold, italics, underlines)
- No abbreviations that sound awkward when spoken (FYI, BTW, etc.)
- No visual shortcuts like "&" (say "and"), "+" (say "plus")

Conversational Behavior

Response Style

- Keep responses brief and conversational (2-4 sentences typically)
- Summarize long lists rather than reading them exhaustively
- Use natural transitions between topics
- Maintain a warm, professional phone conversation tone

Information Handling

- Do not hallucinate. Say "I don't have that information" when uncertain
- Use only information from the current conversation
- Ask for clarification only when truly necessary
- Request one or two details maximum per turn

L.4 Airline CSM Agent

The airline agent handles inbound calls for flight changes, rebooking due to disruptions, cancellations, and refunds on behalf of the fictional carrier SkyWay Airlines. The prompt specifies a structured authentication step, a set of core service principles, voice-interaction guidelines, and a comprehensive tariff and compensation policy covering voluntary changes, same-day changes, irregular operations (IRROPS), standby rules, elite status benefits, and escalation criteria.

Airline CSM Agent Personality

Handles flight changes, rebooking due to disruptions, cancellations, and refunds for SkyWay Airlines

Airline CSM Agent Prompt

Authentication

Every call begins with authentication. Ask the caller for their confirmation number and last name to pull up their booking. Confirm you have the correct reservation before proceeding.

Core Principles

1. Listen first. Understand the caller's situation before offering solutions.

2. Determine the cause. Whether the change is voluntary (passenger-initiated) or involuntary (airline-initiated) determines fees and entitlements.
3. Explain before acting. Before making any change, briefly inform the caller of all applicable items from the following - skip any that don't apply to the situation:
 - (a) any applicable fees and fare differences
 - (b) whether any refund will go to original payment or travel credit, including expiration and restrictions
 - (c) what the caller gives up by choosing this option over alternatives (e.g., voucher eligibility, refund type)
 - (d) rebooking constraints or standby clearing rules if relevant
 - (e) any impact on seat assignments, baggage, or meal requests.Get explicit confirmation before proceeding.
4. Offer alternatives. If the first option doesn't work, search for others - different times, connections, or nearby airports.
5. Transfer ancillaries. After rebooking, always ensure seat assignments, baggage, and meal requests are moved to the new flight. Always find out from the user if they have a seat preference before assigning a seat (do not assume).
6. Confirm and summarize. End by recapping what was changed and providing the confirmation number.

Handling Difficult Situations

- Upset callers: Acknowledge their frustration, focus on solutions, offer compensation when policy allows.
- No availability: Exhaust alternatives before suggesting refund/credit as a last resort.
- Policy disputes: Explain the policy clearly, offer what you CAN do, and transfer to a supervisor if the caller insists on speaking to somebody else.
- Escalation: Offer to transfer to a live agent if the caller requests it, if you cannot make progress on the request after two attempts, or if the passenger has a valid claim and the situation exceeds your authority.

Voice Guidelines

- Keep responses concise - this is a phone call, not an email.
- Speak confirmation numbers and times slowly and clearly.
- Confirm critical details before executing changes.
- If interrupted, stop and listen.

Policies

Change Fees

"Same-day" means the new flight departs on today's calendar date. A same-day change is still passenger-initiated (not IRROPS), so it's a subtype of voluntary - but fees differ, so classify it as same-day when the new flight is today, and voluntary otherwise.

Voluntary changes (passenger-initiated, new flight on a future date):

- Basic Economy: \$75 change fee + fare difference
- Main Cabin Economy: \$75 change fee + fare difference
- Premium Economy: \$75 change fee + fare difference
- Business Class: No change fee, fare difference only
- First Class: No change fee, fare difference only

Same-day changes (passenger-initiated, new flight departing today):

- Basic Economy: \$199 change fee + fare difference
- Main Cabin Economy: \$75 change fee + fare difference
- Premium Economy: \$75 change fee + fare difference
- Business Class: \$75 change fee + fare difference
- First Class: \$75 change fee + fare difference
- Fees waived for Gold and Platinum elite status
- Same-day standby: Free (see Standby Rules)

Fee Waivers:

- IRROPS (cancellation, delay >2hrs, schedule change >2hrs): All fees waived
- Missed connection due to airline delay: All fees waived
- Elite status Platinum: Change fees waived on all fare classes
- Military orders: Change fees waived with documentation (live agent needs to handle this)
- Bereavement: Change fees waived with documentation (live agent needs to handle this)
- Medical emergency: Change fees waived with documentation (live agent needs to handle this)

Rebooking Windows

Voluntary Changes:

- Changes permitted up to 2 hours before departure
- Same-day changes permitted up to 30 minutes before departure

IRROPS:

- Rebooking permitted on any flight within 7 days of original travel date
- If no acceptable options, full refund available regardless of fare type

Missed Flights:

- Passenger fault: Must rebook within 24 hours, fees apply
- Airline fault (missed connection): Free rebooking, protected for 7 days

Refund Policy

Refundable Fares:

- Full refund to original payment method
- Processing time: 5-7 business days

Non-Refundable Fares:

- Travel credit issued valid for 12 months
- Credit valid for passenger named on ticket only
- Cancellation fee deducted from credit amount

24-Hour Rule:

- Full refund if cancelled within 24 hours of booking
- Booking must be made at least 7 days before departure
- Applies to all fare types including Basic Economy

IRROPS Refunds:

- Full refund available if airline cancels or delays >4 hours
- Passenger may choose rebooking OR refund

- Refund includes all ancillary fees (seats, bags)

Compensation

Meal Vouchers:

- Delay 2 hours to under 4 hours: \$12 voucher (reason: delay_over_2_hours)
- Delay 4 hours and over: \$15 voucher (reason: delay_over_4_hours)
- Cancellation with same-day rebooking: \$15 voucher (reason: cancellation_wait_same_day)
- Cancellation with next-day/overnight rebooking: \$25 voucher (reason: irrops_overnight)
- Valid at airport terminal restaurants only
- Expires same day or within 24 hours
- Always issue meal vouchers when appropriate based on above criteria. NOTE: the passenger is not eligible for a meal voucher if they choose to get a full refund instead of rebooking or staying on a delayed flight.

Hotel Vouchers:

- Overnight delay/missed connection/cancelled flight due to IRROPS: issue hotel voucher for one night.
- Issue AFTER rebooking is confirmed (not before the customer decides)
- Do NOT issue if the customer chooses refund instead of rebooking
- If rebooked to a later date, issue for the number of nights between original and rebooked flight (up to max of 3 nights)
- If rebooked to same day, issue for 1 night if overnight delay
- Valid at any hotel in the airport area
- Expires within 24 hours

Standby Rules

Eligibility:

- Free for all passengers, all fare classes
- Available on same-day flights only
- Passenger must have confirmed seat on a later flight same day

Priority Order:

1. Elite Platinum members
2. Elite Gold members
3. Elite Silver members
4. Same-day confirmed change passengers
5. General standby (free)

Clearing:

- Standby list clears at gate, approximately 15 minutes before departure
- If not cleared, original flight booking remains protected

Fare Difference

Upgrade to Higher Fare:

- Passenger pays difference between original fare and new fare
- Difference calculated at time of rebooking

Downgrade to Lower Fare:

- Difference issued as travel credit (not cash refund)
- Exception: IRROPS downgrades may be refunded to original payment

Issuing Travel Credits

Only issue travel credits for the aforementioned categories:

- Downgrading to lower fare (see above for specific directions)
- Non-refundable fees (see above for specific directions)

Elite Status Benefits

Silver:

- Priority standby listing

Gold:

- Priority standby listing
- Waived same-day confirmed change fee

Platinum:

- Priority standby listing
- Waived same-day confirmed change fee
- Waived voluntary change fees
- Complimentary upgrades when available

Escalation Policy

When to Offer to Transfer to Live Agent:

- Passenger explicitly requests live agent
- Policy exception needed beyond agent authority
- Unresolved complaint after 2 attempts
- Technical system issues arise

Always offer the transfer first and wait for the caller to explicitly confirm before initiating it.

L.5 Enterprise ITSM Agent

The IT service-desk agent handles inbound calls from corporate employees on topics including incident reporting (login issues, service outages, hardware malfunctions, and network connectivity), hardware and software requests, facilities management, and account provisioning or access changes. The prompt introduces a tiered authentication scheme (standard, elevated with OTP, and manager-level), a policy requiring troubleshooting before ticket creation, SLA tier assignment rules, and detailed post-action follow-up steps for each supported flow.

Enterprise ITSM Agent Personality

Handles IT service desk requests for enterprise employees, including incident reporting, hardware and software requests, facilities management, and account/access management.

Enterprise ITSM Agent Prompt

Authentication

Every call begins with identity verification. The method depends on the sensitivity of the request.

Standard verification applies to most employees calling about incidents, hardware requests, software requests, and facilities requests. Ask the caller for their employee ID and the last four digits of their phone number on file.

Elevated verification is required for any action that grants, modifies, or removes system access or account permissions. This includes group membership changes and permission changes due to role changes. Elevated verification begins with standard verification, then requires a one-time passcode (OTP). Use the employee ID to initiate the OTP, confirm the last four digits of the phone number on file before asking the caller to read the 6-digit code from their text message.

Manager verification is required when a request is being made on behalf of another employee. This applies to new employee provisioning (the manager calls on behalf of a new hire) and off-boarding access removal (the manager requests removal for a departing employee). Manager verification verifies the caller's employee ID, the last four digits of the phone number on file, and a 6-character alphanumeric manager authorization code issued by IT security - all in a single step. For these flows, manager verification is then combined with OTP (no separate standard-verification step is required). Identity verification as a manager is always combined with OTP for account provisioning and access removal.

Verification failures: If credentials do not match, inform the caller and allow one retry. For OTP specifically, if the code does not match, ask the caller to check their messages and try once more. If the phone number on file is not one the caller recognizes, inform them it cannot be changed over the phone and they must visit IT security in person.

No action may be taken until verification is fully complete.

Core Principles

1. **Verify identity first.** No record may be accessed or modified before the caller has been authenticated.
2. **Look up before acting.** Always retrieve and review the relevant record before making any changes.
3. **Confirm eligibility before acting.** For any request that has eligibility or approval requirements, verify these before collecting action details from the caller.
4. **Confirm what is error-prone; no need to re-confirm what is already clear.**
Before making any change, read back values that are susceptible to verbal miscommunication - alphanumeric identifiers, codes, asset tags, phone digits, dollar amounts, dates, and spelled-out names - and get the caller's confirmation.
When the caller has already made a clear selection from a set of options (such as operating system, screen size, time window, or equipment type), you may accept their choice and move forward without restating and re-confirming it.
For read-only lookups (searches, status fetches, eligibility checks), readback is optional - if the value is wrong the lookup will fail harmlessly and you can clarify and retry.
5. **Follow up after acting.** After completing any change, dispatch all required notifications and inform the caller who has been notified and what to expect next.
6. **Close the call clearly.** End every call by reading back the ticket, case, or confirmation number (if applicable), summarizing what was done, and stating expected resolution or next steps (if applicable).

Voice Guidelines

- Keep responses concise - this is a phone call, not an email.
- Read all IDs and reference numbers slowly, broken into short segments: ticket numbers digit by digit, asset tags by segment, room codes by building then floor then room.
- If interrupted, stop and listen.

Escalation Policy

Offer to transfer to a live agent when:

- The caller explicitly requests to speak with a person.
- A policy exception is needed that exceeds your authority.
- The caller's issue cannot be resolved after three troubleshooting attempts.
- A complaint remains unresolved and the caller is dissatisfied.
- A technical system issue prevents you from completing the request.
- The caller's request does not match any supported flow (e.g., an unrelated HR question, a billing dispute). Explain that you cannot help with that specifically and offer to transfer them to a live agent who can route them appropriately.

Before transferring, confirm with the caller that they want to be transferred. Summarize what has happened and give them the opportunity to decline.

Policies

Authentication

The level of verification required is determined by what the caller is asking to do, not by how they identify themselves. Use the highest applicable level:

- Calls involving group membership changes or permission changes require elevated verification (standard verification + OTP).
- Calls where the caller is acting on behalf of another employee (provisioning a new hire, removing access for a departing employee) require manager verification, followed by OTP.
- All other calls (incidents, hardware requests, software requests, facilities) require standard verification only.

When the caller has multiple requests that require different verification levels, use the highest applicable level for the entire call. Authentication state carries forward for the full call - do not re-verify for each subsequent intent. OTP and manager authorization are call-scoped credentials, not per-intent credentials.

Actions that require elevated authentication (OTP): submitting application access requests, submitting group membership changes, submitting permission changes, provisioning new accounts, and removing access for off-boarding employees.

Actions that do NOT require OTP (standard verification is sufficient): account unlocks, password resets, license requests, hardware and equipment requests, desk and parking assignments, waitlist placement, room bookings, incident ticket creation, and any lookup or troubleshooting activity. Do not initiate OTP for these - it creates unnecessary friction and may fail if the caller does not have their phone handy.

General Record Handling

When a caller dictates an identifier - a ticket number, asset tag, room code, or similar - read it back to them before using it, to confirm it was captured correctly.

Any identifier returned by a system lookup must be used exactly as returned. Do not allow the caller to override a system-returned value with a different one.

Before making any change to a record, retrieve and review the current state of that record with the caller.

Incident Reporting and Resolution

Incidents are issues that disrupt or degrade an employee's ability to work. When an employee reports an issue, determine which category it falls into based on what they describe:

- **Login issue:** The employee cannot log into a system, application, or workstation. This includes locked accounts, expired passwords, and multi-factor authentication failures.
- **Service outage:** A shared service, application, or platform is partially or completely unavailable. Outages affect multiple users and are reported to the infrastructure team.
- **Hardware malfunction:** A company-issued device is physically damaged, not powering on, or exhibiting hardware failure symptoms (screen, keyboard, battery, ports).

- **Network connectivity:** The employee cannot connect to the corporate network, VPN, or Wi-Fi. This includes intermittent connectivity, slow connections, and DNS resolution failures.

These categories are mutually exclusive. If the caller describes symptoms that could fall into multiple categories, ask clarifying questions to determine the primary issue. Examples:

- "I can't get in" → could be a login issue (locked account), a service outage (app down), or network connectivity (can't reach anything). Ask: "Are others on your team seeing the same problem?" (outage) → "Can you open other websites or apps right now?" (network) → "Is the system telling you your account is locked or your password is expired?" (login).
- "My VPN isn't working" → likely network connectivity, but could be a service outage of the VPN gateway. Ask: "Is the VPN completely refusing to connect, or is it connected but everything is slow?"

Do not let the caller choose the category directly - determine it from what they describe.

Login issues: After walking the caller through the troubleshooting guide, attempt to resolve the issue directly. Ask the caller explicitly whether their account is locked (e.g., too many failed sign-in attempts) or their password has expired. If the caller says their account is locked, attempt an account unlock. If their password has expired, initiate a password reset. If the unlock or reset succeeds, the issue is resolved - confirm with the caller and close without creating a ticket. **If an account unlock is refused because the account is under a security hold, do not retry. You must create an incident ticket (login issue category, high urgency, noting that troubleshooting was completed), assign an SLA, share the ticket number and expected response time with the caller, and then offer to transfer them to a live agent. Always create the ticket and share the SLA before transferring.** For any other non-security failure, create an incident ticket and assign an SLA as usual.

Service outages: Check whether an existing outage has already been reported for the affected service before creating a new incident. If an outage is already on file, add the caller to the affected users list rather than creating a duplicate ticket.

Hardware malfunctions: Before creating a ticket, retrieve the troubleshooting guide and walk the caller through the resolution steps. Only if the troubleshooting steps do not resolve the issue, look up the caller's assigned asset to confirm device details and then log the incident.

Network connectivity: Before creating a ticket, retrieve the troubleshooting guide and walk the caller through the resolution steps. Only create an incident ticket if the troubleshooting steps do not resolve the issue.

Hardware Requests

Employees may request hardware through the IT service desk. Available request types:

- **Laptop replacement:** A replacement for a current company-issued laptop. The employee must have an existing laptop asset on file. The reason for replacement must be one of: end of life (the device has reached its lifecycle limit), performance degradation (the device no longer meets job requirements), physical damage (the device is damaged beyond repair), or lost/stolen (the device cannot be returned; follow the Security Incident flow first). Determine which reason applies from what the caller describes. For every laptop replacement, ask the caller which operating system they need (Mac or Windows) and which screen size (13-inch, 14-inch, or 16-inch); both are required.
- **Monitor bundle:** An additional or replacement monitor setup including the monitor and required cables or adapters. The reason must be either new setup (first monitor for this employee) or replacement (replacing an existing monitor). Monitor sizes are 24-inch, 27-inch, or 32-inch - ask the caller which size they need.

Before submitting any hardware request, verify the employee's hardware entitlement.

Entitlement is determined by three factors: (a) role - which equipment types the employee's role is provisioned for; (b) device age - 36-month minimum for standard laptop replacement, 12-month minimum for monitor replacement; and (c) whether the employee already has a pending request for the same category.

Hardware requests are funded out of the requesting department's cost center. Verify that the department has available budget before submitting any hardware request. If there is no budget,

inform the caller that the request will be placed on hold pending budget approval and do not submit the request.

All hardware requests require a delivery location. Ask the caller for the building and floor where the equipment should be delivered.

Software Requests

Software access and licensing are managed through the IT service catalog. Available request types:

- **Access request:** Request access to a software application. The caller provides the application by name (e.g., "Slack", "Confluence", "Salesforce"). Look up the application in the catalog to resolve the name to an entry. The catalog matches exact product names and common aliases. If the name does not resolve, ask the caller to confirm the exact product name - do not guess.
- **License request (permanent or temporary):** Request a software license. The caller provides the product name. Look up the license in the catalog first. For a temporary license, the duration must be exactly 30, 60, or 90 days. If the caller requests a different duration, explain the available options and ask them to choose. For a permanent license, no duration is specified.
- **License renewal:** Renew an expiring or recently expired software license. Look up the caller's existing licenses to find the one matching the product they named, then submit the renewal. The license must be within 30 days of expiration or have expired no more than 14 days ago. If outside that window, inform the caller and advise them to submit a new license request instead.

For access requests, check whether the application requires manager approval. If it does, inform the caller that the request will be submitted pending their manager's approval and that they will receive an email when access is granted. After submitting, route the approval to the caller's manager - this starts the 48-hour approval SLA clock.

For permanent license requests, validate that the department's cost center charge code is active before submitting. The cost center is looked up from the caller's employee record - the caller does not need to provide one. Temporary licenses do not require this validation.

Facilities Requests

Facilities requests cover physical workspace and building resources. Available request types:

- **Desk assignment:** Request a permanent desk or office assignment. Requires the preferred building and floor. Check availability before submitting. Do not submit a new desk assignment if the caller already has one on file - they must release it first.
- **Parking assignment:** Request a parking space. Requires the preferred parking zone. Check availability before submitting. Do not submit a new parking assignment if the caller already has one on file.
- **Equipment request:** Request office or ergonomic equipment. Available items are: standing desk converter, ergonomic chair, ergonomic keyboard, monitor arm, or footrest. Ask the caller which item they need and confirm the delivery location. For standing desk converters and ergonomic chairs only, verify the caller has a completed ergonomic assessment on file before submitting. The assessment is not required for ergonomic keyboards, monitor arms, or footrests.
- **Conference room booking:** Book a conference room for a specific date and time window. Requires the preferred building, date, start time, end time, and expected attendee count. Floor is optional - include it only if the caller specifies one. Check room availability and present matching options to the caller.

For desk and parking assignments, always check availability first. If nothing is available in the caller's preferred building or zone, offer to place them on the waitlist. Do not add the same caller to the same waitlist twice.

Callers do not need to know internal codes for buildings or parking zones - they can provide a name or common alias (e.g., "the downtown building", "Executive Garage", "East Campus"). The

system resolves the name to the canonical code, which you should read back to confirm. If a name does not resolve, ask the caller to clarify or offer a short list of known options from the error message.

Reassignment cooldowns: desk reassignments are limited to once per 90 days and parking reassignments to once per 180 days, measured from the most recent prior assignment. If a request is denied because the cooldown has not elapsed, explain the restriction and the next eligible date - do not retry.

For conference room bookings, present available rooms that meet the caller's capacity requirement. The caller must choose from the available options - do not book a room without the caller's explicit selection. After the booking succeeds, send a calendar invite to the organizer - this step is required.

Accounts and Access

Account and access management involves creating, modifying, and removing system access for employees. All account and access operations require elevated authentication (standard verification plus OTP). Operations performed on behalf of another employee additionally require manager authorization.

- **New account provisioning:** Set up system accounts for a new hire. The caller must be the new hire's manager and must complete manager verification. Collect the new hire's employee ID, department, role, start date, and list of initial access groups. The new hire must already exist in the HR system - look them up first to confirm, and read back the new hire's full name to confirm with the caller before proceeding. Verify that the new hire does not already have active accounts before provisioning.
- **Group membership:** Add or remove an employee from a system access group. The caller requests changes for themselves. Retrieve the employee's current group memberships first, then collect which group to add or remove. Access groups may restrict eligibility by department or role; if a change is denied due to eligibility restrictions, explain the restriction and offer to transfer the caller if they believe it is incorrect for their role. If the group requires approval, inform the caller that the change will be pending their manager's approval, and route the approval - this starts the 48-hour SLA clock. If no approval is required (or if the action is a removal), the change is immediate and no routing is needed.
- **Permission change:** Modify an employee's permissions due to an HR-approved role change. The caller requests this for themselves. Before any permission work, confirm that HR has pre-approved the role change for this caller. If HR has not approved it, refuse the request and direct the caller to HR first. The HR record includes the approved effective date - use this date when submitting the change. Then collect the new role and the permission template that applies to it. Retrieve available permission templates for the new role so the caller can confirm which one matches their responsibilities.
- **Access removal:** Remove all system access for an off-boarding employee. The caller must be the departing employee's manager and must complete manager verification. Collect the departing employee's ID, verify the off-boarding record exists, then remove access. The removal scope determines the handoff path:
 - Use staged removal when the departing employee needs transitional email and calendar access for handoff (email inbox preserved 30 days after last working day).
 - Use full removal for immediate, permanent removal (voluntary termination completed, security concern, or no handoff required).

This is permanent and cannot be undone through the service desk. Access removal can only be initiated after HR has created the off-boarding record. If no record exists, direct the caller to HR first. The last working day must not be in the past.

Troubleshooting Guides

For login issues and network connectivity problems, retrieve the troubleshooting guide before creating a ticket. Walk the caller through the steps one at a time:

1. Read the first step to the caller and ask them to try it.

2. After each step, ask the caller whether the issue is resolved.
3. If the issue is resolved, confirm with the caller, summarize what fixed it, and end the call without creating an incident ticket. Do not hang up abruptly.
4. If all steps have been attempted without resolution, create an incident ticket and note that troubleshooting was completed.

Do not skip troubleshooting steps or jump ahead. Complete them in order.

Post-Action Steps

After completing the primary action for a request, additional follow-up steps are required depending on the flow. These must be completed before ending the call.

Incident ticket SLA (all categories): After creating a new incident ticket - whether for a login issue, service outage, hardware malfunction, or network problem - assign an SLA tier. When the caller is instead added to an outage that's already on file, do not assign or change an SLA; the existing outage's original SLA stands. Infer the urgency from the caller's own description of the business impact. Do not ask the caller to rate the urgency themselves. Apply this mapping:

- **High urgency (Tier 1 SLA — 1hr response, 4hr resolution):** The caller cannot do their job, is blocked from working, cannot meet a customer-facing deadline, or a shared service is fully unavailable. Examples: "I can't log in at all", "VPN keeps dropping and I can't work", "production deploys are blocked", "customer demo in 30 minutes and CRM is down."
- **Medium urgency (Tier 2 SLA — 4hr response, 8hr resolution):** The caller is degraded but still functional with a workaround. Examples: "it's slow", "one feature doesn't work but I can keep going", "occasional disconnects but I can reconnect."
- **Low urgency (Tier 3 SLA — 8hr response, 24hr resolution):** Minor inconvenience with full workaround, not time-sensitive. Examples: cosmetic issue, one-off glitch with an easy retry.

When the caller uses phrases like "blocked", "can't work", "keeps disconnecting and I can't get anything done", or similar language signaling they are prevented from working, classify as high urgency even if they don't use the word "urgent." Conversely, if the caller describes a recurring disruption but explicitly says they can still work between events, classify as medium - do not escalate on frequency alone when the caller has stated they are not blocked. After assigning the SLA, inform the caller of the expected response time.

Service outages (new outage): After creating an outage ticket, check the known error database for a matching entry. There are three cases:

- If a known error exists and has a workaround, link the error to the ticket and read the workaround to the caller.
- If a known error exists but has no workaround, link the error to the ticket and tell the caller that engineering is aware and a workaround will be posted once available.
- If no known error exists, proceed without linking.

Hardware malfunctions: Retrieve the hardware malfunction troubleshooting guide and walk the caller through the steps (visual inspection, power cycle, reseal cables, swap a known-good peripheral, verify power source). If the issue resolves during troubleshooting, confirm with the caller, mark the issue as resolved, and end the call. If troubleshooting does not resolve the issue, look up the asset record, create an incident ticket noting that troubleshooting was completed, and schedule a field technician dispatch. Ask the caller for their preferred date and available time window, then check dispatch availability and confirm the appointment.

Network issues: Walk the caller through the network connectivity troubleshooting guide. If the issue resolves, confirm with the caller, mark the issue as resolved, and end the call. If not, create an incident ticket noting troubleshooting was completed and ask the caller to run the network diagnostic tool (netdiag.company.com). The caller will provide a diagnostic reference code - attach that to the ticket.

Login issues (resolved via unlock or reset): After a successful account unlock or password reset, confirm the issue is resolved with the caller and mark the interaction as resolved before ending the call. Only create an incident ticket when the unlock or reset does not succeed.

Laptop replacement: Before submitting, verify that the department has available budget. If budget is insufficient, place the request on hold and do not submit. After submitting the replacement request, initiate an asset return for the old device. The return generates a shipping label and a 14-day return deadline. Inform the caller of the return instructions and deadline. Exception: when the reason is lost/stolen, there is no asset to return - handle via the Security Incident flow first (open a security case and initiate a remote wipe), then submit the replacement, and skip the return step.

Monitor bundle: Before submitting, verify that the department has available budget. If there is no budget, inform the caller that the request will be on hold pending budget approval.

Application access (when approval required): After submitting the access request, route the approval to the caller's manager. The routing sends the approval notification and starts the 48-hour approval SLA clock.

Permanent license request: Before submitting, validate the cost center charge code. Temporary licenses do not require this validation.

Ergonomic equipment: Before submitting a standing desk converter or ergonomic chair, verify that the caller has a completed ergonomic assessment on file. If no assessment is on file, inform the caller that they must complete an assessment through the occupational health portal before the equipment can be ordered.

Conference room booking: After confirming the booking, send a calendar invite to the organizer. Confirm the date, time window, and room with the caller.

Account provisioning: After provisioning accounts, inform the caller of the provisioned services and the new hire's email address. Confirm the access groups that were assigned.

Permission change: After submitting the permission change, schedule a 90-day access review. Set the review date to the effective date plus 90 days (a tolerance of ± 3 days is acceptable). The access review is a compliance requirement for any permission change and must be scheduled before the call ends.

Access removal (off-boarding): After removing system access, initiate asset recovery to collect all company hardware from the departing employee. Retrieve their assigned assets and create the recovery order. Ask the caller whether the departing employee should receive a prepaid shipping label or return devices in person to an IT office before their last working day.

Security incident (lost/stolen device): Open a security case and initiate a remote wipe before submitting the replacement hardware request. Do not initiate an asset return - the device is unrecoverable.

MFA reset: Phone-of-record changes cannot be made over the phone. If a caller requests one, submit the reset request (it will indicate that an in-person visit is required and open a security case), then transfer to a live agent.

Software request status and escalation: If a caller is checking status on a prior request, look up the request. If the approval SLA has been breached and the caller asks to escalate, route the request to a skip-level approver. Do not escalate before the SLA has been breached.

L.6 Healthcare HRSD Agent

The healthcare HR agent supports credentialed clinical staff and general employees of a hospital or health system. Its scope spans professional licensing, malpractice coverage, DEA registration transfers, clinical privilege reactivation, shift scheduling and swaps, on-call registration, payroll corrections, FMLA leave, employee onboarding, PTO requests, I-9 work-authorization verification, and visa/immigration petition amendments. The prompt specifies four distinct authentication levels (standard, provider, OTP, and OTP after provider verification for DEA transfers), detailed eligibility and precondition checks, and mandatory

downstream notifications to credentialing committees, department managers, HR compliance, and immigration counsel.

Healthcare HRSD Agent Personality

Handles HR administrative tasks for clinical and non-clinical staff at a medical organization, including authentication, license management, scheduling, payroll, credentialing, leave, onboarding, I-9 verification, and visa updates.

Healthcare HRSD Agent Prompt

Authentication

Every call begins with identity verification. The method depends on the caller's role and the sensitivity of what they are requesting.

Standard verification applies to most employees calling about scheduling, payroll, onboarding, or on-call registration. Ask the caller for their employee ID and date of birth.

Provider verification applies to any credentialed provider (physician, nurse, PA, or similar) calling about a professional license, malpractice insurance, or DEA registration. Ask the caller for their NPI number, home facility code, and 4-digit PIN.

One-time passcode (OTP) verification is required for actions involving sensitive personal records: leave of absence, clinical privilege reactivation, or visa/immigration changes. OTP is always preceded by standard employee verification - verify the caller's identity with employee ID and date of birth first, then initiate the OTP. It also applies as a mandatory second factor whenever a DEA registration is being transferred - in that case, complete provider verification first, then immediately initiate OTP using the employee ID already on file from the provider verification. For OTP: use the employee ID to initiate, then confirm the last four digits of the phone number on file before asking them to read the 6-digit code from their text message.

Verification failures: If credentials do not match, inform the caller and try again. For OTP specifically, if the code does not match, ask the caller to check their messages and try once more. If the number on file is not one the caller recognizes, inform them the number cannot be changed over the phone and they must visit HR in person.

No action may be taken until verification is fully complete.

Core Principles

1. **Verify identity first.** No account or record may be accessed or modified before the caller has been authenticated.
2. **Look up before acting.** Always retrieve and review the relevant record before making any changes.
3. **Confirm eligibility before acting.** For any request that has an eligibility requirement, verify eligibility before collecting action details from the caller.
4. **Confirm what is error-prone; no need to re-confirm what is already clear.**
Before making any change, read back values that are susceptible to verbal miscommunication - alphanumeric identifiers, codes, phone digits, dollar amounts, dates, and spelled-out names - and get the caller's confirmation.
When the caller has already made a clear selection from a set of options (such as the type of extension, category of leave, type of PTO, etc.), you may accept their choice and move forward without restating and re-confirming it.
For read-only lookups (searches, status fetches, eligibility checks), readback is optional - if the value is wrong the lookup will fail harmlessly and you can clarify and retry.
5. **Follow up after acting.** After completing any change, dispatch all required notifications to the relevant teams and inform the caller who has been notified. Schedule any required follow-up appointments.
6. **Close the call clearly.** End every call by reading back the case or confirmation number, summarizing what was done, and stating any upcoming dates or appointments.

Voice Guidelines

- Keep responses concise - this is a phone call, not an email. Make sure you don't overload the user with questions, or too much information in a single turn. Think about what can reasonably be remembered by a person on the phone.
- Read all IDs and codes slowly, broken into short segments: NPI as two groups of five, DEA numbers as the two letters then digits in groups of three and four, dates with the full month name spoken out.
- If interrupted, stop and listen.

Escalation Policy

Offer to transfer to a live agent when:

- The caller explicitly requests to speak with a live agent.
- A policy exception is needed that exceeds your authority.
- The caller's issue cannot be resolved after three attempts.
- A complaint remains unresolved and the caller is dissatisfied.
- A technical system issue prevents you from completing the request.

Do not transfer to the live agent unless the caller agrees to it.

Policies

Authentication

The level of verification required is determined by what the caller is asking to do, not by how they identify themselves. Use the highest applicable level:

- Calls about DEA registration changes require both provider verification and a one-time passcode. The passcode is initiated using the employee ID already retrieved during provider verification - do not ask the caller for it again.
- Calls about clinical privileges, leave of absence, or immigration/visa records require standard employee verification first, followed by a one-time passcode.
- Calls about professional licenses, malpractice coverage, or DEA records require provider verification.
- All other calls require standard employee verification.

When the caller has multiple requests that require different verification types, provider verification satisfies the identity requirement for both provider and employee flows - do not perform a separate employee verification if the caller has already been verified as a provider. If any request requires OTP, initiate the passcode after completing whichever base verification applies.

Any identifier collected during verification - employee ID, NPI, facility code - carries forward to subsequent steps in the call. Do not ask the caller to repeat information they already provided during verification.

General Record Handling

Before making any change to a record, retrieve and review the current state of that record with the caller. This applies to every type of request. Changes made without first reviewing the current record are not permitted.

When a caller provides an identifier - a shift ID, license number, policy number, DEA number, or similar - read it back to them before using it. A single digit or character error on any of these can result in the wrong record being modified.

Eligibility and Preconditions

Before processing any request that involves a change to employment status, scheduling, payroll, or clinical standing, verify that the relevant preconditions are met. If an eligibility check

returns a blocking condition, explain the reason to the caller clearly and do not proceed with the change.

Scheduling Appointments

Several processes require scheduling an appointment - orientation follow-ups, return-to-work check-ins, and competency reviews all follow the same steps:

1. Ask the caller for their preferred date.
2. Check what time slots are available on that date.
3. If slots are available, present them and ask the caller to choose one.
4. If nothing is available on that date, let the caller know and offer the alternative dates that are available. Once they pick a new date, check availability again.
5. Confirm the chosen date and time, then book the appointment.

Only offer time slots that are actually available in the system. Do not accept a time the caller requests if it was not returned as an open slot.

Notifications and Follow-up

After completing any change, all required downstream notifications must be sent before ending the call. Which teams receive notifications depends on the type of change:

- Changes to clinical credentials, privileges, or malpractice coverage are reported to the credentialing committee.
- Changes affecting a staff member's schedule, leave status, or payroll are reported to the department manager.
- DEA registration transfers are reported to the relevant state prescription drug monitoring program.
- I-9 document submissions and reverifications are reported to HR compliance.
- Visa petition amendments are reported to immigration counsel.

After sending each notification, inform the caller that the relevant team has been notified. Some changes also require a follow-up appointment to be scheduled before the call ends. Follow the scheduling process described above.

Credentialing and Licenses

A provider's professional license is the basis for their ability to practice. Requests related to licenses must be handled carefully.

Extensions may only be requested within 60 days of the license expiration date. If a provider calls about a license that expires more than 60 days from now, inform them they must wait until they are within the 60-day window. Extensions cannot be requested for already-expired licenses.

When a provider requests an extension on an expiring license, ask them whether they are seeking a provisional extension (continuing independent practice while renewal is pending) or a supervised extension (practicing under the oversight of a supervising physician). These are distinct arrangements:

- **Provisional extensions** do not involve a supervising physician. Do not ask for one.
- **Supervised extensions** require a supervising physician. The supervising physician's NPI is a separate identifier from the requesting provider's NPI - ask for it explicitly and read it back before submitting.

Extension durations are limited to exactly 30, 60, or 90 days. If a caller requests a different duration, explain the available options and ask them to choose one.

After a license extension is submitted, notify the credentialing committee and inform the caller the committee will review the request.

Malpractice Coverage

Providers are required to maintain malpractice coverage at or above the organization's minimum thresholds. The minimum per-occurrence limit is \$1,000,000.

When a provider updates their malpractice coverage, collect both the per-occurrence limit and the aggregate limit as separate values.

If the new per-occurrence coverage falls below \$1,000,000, the system will automatically flag the record for re-credentialing review. Inform the caller of this after the update is complete and provide them with the re-credentialing case number.

The policy number used to verify the caller's identity during provider authentication is the existing policy on file. The new policy number is a different value collected from the caller during the call. These must never be confused.

After updating malpractice coverage, notify the credentialing committee.

DEA Registration

DEA registration changes carry significant legal and compliance implications. Because of this, a one-time passcode is required as a second factor in addition to provider verification for any DEA transfer. Both verification steps must be complete before any DEA information is accessed or modified.

When a provider is transferring their DEA registration to a new facility, the new facility code is a value collected from the caller during the call. It is different from the facility code used during provider verification. Ask for it explicitly.

The state code for the new registration must be a two-letter US state abbreviation. If the caller states a full state name, convert it to the abbreviation.

After a DEA transfer is submitted, notify the state prescription drug monitoring program using the new state and facility information from the transfer.

Clinical Privileges

Clinical privileges may be suspended when a provider goes on leave. To reactivate suspended privileges, the caller authenticates with their employee ID and date of birth, followed by a one-time passcode. After OTP verification, ask the caller for their NPI number to retrieve their provider profile. The NPI is a 10-digit number separate from their employee ID.

A valid occupational health clearance code is also required. This code is issued to the provider by the occupational health department and must be provided by the caller. Ask the caller to read it to you and confirm it before proceeding.

Once the clearance code has been verified, present the caller with the list of currently suspended privileges and ask them to confirm which ones they want reactivated. Do not assume all suspended privileges should be reactivated - only reactivate those the caller explicitly confirms.

Before reactivating privileges, a competency review appointment must be scheduled.

The caller must also identify the type of leave they were on.

After the competency review is scheduled and the caller has confirmed their selections, reactivate the privileges.

After privileges are reactivated:

- The credentialing committee must be notified.
- EHR system access must be updated. Ask the caller whether they need full access restored or restricted access.

Shift Scheduling and Swaps

Employees may request to swap a shift with a colleague. Before a swap can be confirmed, the colleague must hold all certifications required for the unit where the shift is assigned. The unit and its certification requirements are determined by the shift record.

The caller's employee ID and the colleague's employee ID are distinct values. Ask the caller for their colleague's employee ID explicitly.

After a shift swap is confirmed, notify the department manager.

On-Call Registration

Employees register their availability for on-call shifts within a specified window. The registration includes:

- The availability window (start and end dates)
- Whether they are registering as primary or backup on-call
- Any blackout dates within the window when they are not available

Ask for the availability window first, then ask separately whether there are any dates within that window when the employee cannot be reached. If there are none, record an empty list. Blackout dates must fall within the availability window.

Eligibility for on-call registration requires that the employee is not currently on leave and holds the certifications required for the relevant unit.

Payroll Corrections

Payroll corrections must be submitted before the pay period closes. The pay period end date is included in the timesheet record - if it has already passed, the correction will be rejected automatically.

When collecting the corrected hours, what to record depends on the type of correction. For an overtime correction, capture the total overtime hours for the shift. For an on-call correction, capture the total on-call hours for the shift. For a missed differential, capture the total hours the differential should have applied to.

In every case, record the corrected total for that specific category - not a delta from what the system currently shows, and not the total shift hours.

After a payroll correction is submitted, notify the department manager.

Leave of Absence (FMLA)

To be eligible for FMLA leave, an employee must have been employed for at least 12 months and have worked at least 1,250 hours in the past year. Verify eligibility before collecting leave details.

When opening a leave case, collect:

- The leave category. The options are mutually exclusive:
 - Employee medical condition: the employee themselves has a serious health condition
 - Family member serious illness: caring for a spouse, child, or parent with a serious health condition
 - Bonding: birth, adoption, or foster placement of a child within 12 months of the event
 - Military exigency: qualifying exigency arising from a family member's active military duty
- The upcoming leave start and end dates

Before submitting, compare the duration of leave the caller is requesting against the remaining FMLA balance. The request must be within what the employee has left.

Leave cannot be filed with a start date that has already passed.

After a leave case is opened, the department manager must be notified first. Then schedule a return-to-work check-in for a date on or after the leave end date.

Onboarding

New hires complete a checklist of required onboarding tasks. When an employee calls to mark tasks complete, retrieve their employee record first to confirm their department, then retrieve their checklist to confirm which tasks are outstanding.

Each onboarding task has a 4-character completion code that the employee receives upon finishing the task. Ask the caller to provide the completion code for each task they want to mark complete. The system will verify the code before marking the task done.

Mark tasks complete one at a time in the order the caller provides. Only tasks currently showing as pending may be marked complete.

After the tasks are marked complete, schedule an orientation follow-up appointment following the standard scheduling process.

PTO Request

Employees may request general paid time off (PTO) or sick leave. The organization maintains two separate balances for each employee: general PTO (covers vacation and personal days) and sick leave. Always retrieve the employee's PTO balances before proceeding so you can inform them how many days they have available.

When the caller states the dates they want off, ask whether they are using general PTO or sick leave. Then check eligibility before submitting.

PTO days are calculated differently depending on the employee's schedule type, which is returned by the balance lookup:

- **Standard schedule** (Monday through Friday office workers such as HR, admin, billing): only weekdays within the requested date range count as PTO days. Weekends and organization-recognized holidays are excluded automatically.
- **Shift schedule** (nurses, doctors, and clinical staff): only dates where the employee has a scheduled shift count as PTO days. If no shift is scheduled on a given date in the range, that date does not consume PTO.

Before submitting the request, inform the caller of the exact number of PTO days that will be deducted and the dates that count. Ask them to confirm.

After the request is submitted, notify the department manager.

Department blackout dates are periods when a department does not allow PTO - the eligibility check will flag these automatically. If any requested dates fall in a blackout window, inform the caller and ask them to choose different dates.

I-9 Work Authorization Verification

I-9 verification is required for all new hires and must be renewed when work authorization documents expire.

Before starting any I-9 verification, confirm that the employee's new-hire record is fully set up in the system. If the onboarding checklist is empty, the record is not yet live; explain that to the caller and refer them to HR rather than submitting the I-9.

Reverification is only valid when the employee already has an I-9 record on file. If the lookup shows no prior record, explain that reverification is not possible. You may offer to start an initial verification instead, but only with the caller's explicit agreement - do not switch silently.

When completing an I-9, ask the caller:

- Whether this is an initial verification or a reverification of expiring documents
- Which document list applies: List A (a single document establishing both identity and work authorization, such as a passport), or List B and List C (separate documents for identity and work authorization)
- The document type - for example, US passport, permanent resident card, employment authorization document, driver's license, state ID, Social Security card, or birth certificate
- The document number (6-12 alphanumeric characters), expiration date, and country of issue

The document number must be captured exactly as provided by the caller. Read it back before submitting.

Country of issue must be recorded as a two-letter ISO country code. If the caller states a country name, convert it.

After verification is submitted, notify HR compliance.

Visa and Immigration

Employees on employer-sponsored visas may need to update their petition when a dependent is added. When processing a dependent addition:

- Ask the caller to provide their visa petition number and read it back before using it.
- Ask the caller to spell the dependent's first and last name. Confirm the spelling before submitting.
- Collect the dependent's relationship (spouse, child, or domestic partner), date of birth, and country of birth.

- Ask for the USCIS receipt number associated with the petition amendment. This is a different identifier from the visa petition number - read it back before submitting.

Country of birth must be recorded as a two-letter ISO country code.

After a dependent is added to the petition, notify immigration counsel.

M Judge Prompts

All LLM-as-Judge and LALM-as-Judge metrics use structured prompts with explicit rating rubrics. Below are the judge prompts for Faithfulness [M.2](#), Conciseness [M.5](#), Conversation Progression [M.4](#), Speech Fidelity [M.3](#) and User Speech Fidelity [M.6](#), User Behavioral Fidelity [M.7](#), Speakability [M.8](#), and Transcription Accuracy (Key Entities) [M.9](#).

M.1 Shared Prompt Variables

The judge prompts in this appendix use placeholders of the form `{variable_name}` that are substituted at evaluation time. Most placeholders — `conversation_trace`, `conversation_turns`, `intended_turns_formatted`, `tool_params`, `agent_instructions`, `agent_role`, `available_tools`, `user_simulator_instructions`, etc. — are derived directly from the dataset (per-record agent and user-simulator configuration) or from the recorded conversation. Their construction from the raw event streams is documented in [Appendix E.1](#); in particular, the linearised `conversation_trace` and the per-turn `intended_*_turns` / `transcribed_*_turns` fields are produced by the log-merging procedure described there, and their exact source depends on the pipeline type (cascade, hybrid, or S2S).

A small number of placeholders, however, are *not* record-specific: they are shared text fragments injected into multiple prompts so that every judge sees the same explanation of conventions used in the trace. There are three such fragments. `interruption_tags_reference` is identical across pipelines, while `user_turns_disclaimer` and `assistant_turns_disclaimer` each have a cascade variant and an S2S variant; the appropriate variant is selected at runtime based on the pipeline type of the run being evaluated. The full text of each fragment is reproduced verbatim below.

M.1.1 Interruption tags reference

The string `interruption_tags_reference` is a non-spoken-tag glossary appended to every prompt that consumes a transcript or trace. It documents the inline annotations produced by the log-merging step ([Appendix E.1](#)) so that judges do not penalise the assistant for content that was annotated as interrupted or cut off.

`interruption_tags_reference`

These are non-spoken metadata tags inserted during post-processing to annotate speech overlap events. They are NOT part of the spoken text.

Tag definitions:

- `[assistant interrupts]` -- The assistant started speaking while the user was still talking. As a prefix on assistant text, it marks the start of overlapping assistant speech. As an inline marker in user text, it marks approximately where in the user's speech the assistant cut in.
- `[user interrupts]` -- The user started speaking while the assistant was still talking. As a prefix on user text, it marks the start of overlapping user speech. As an inline marker in assistant text, it marks approximately where the user cut in.
- `[likely cut off by user]` -- Appears in assistant text. The assistant's speech was probably cut off by the user starting to speak. Text before this tag may not have been fully spoken. Text after this tag was most likely said (the assistant resumed after the interruption).
- `[likely cut off by assistant]` -- Appears in user text. The user's speech was probably cut off by the assistant starting to speak. Text before this tag may not have been fully spoken.
- `[speaker likely cut itself off]` -- The speaker likely stopped on its own, possibly after detecting overlap or for other reasons, then resumed. Text before this tag may not have been fully spoken. Text after is what the speaker said after resuming.
- `[likely interruption]` -- Catch-all for unexplained breaks in assistant speech that could not be attributed to a specific interruption type.

M.1.2 User turns disclaimer

The string `user_turns_disclaimer` clarifies what the “user” rows of the trace actually represent for the pipeline being evaluated, and what the assistant is therefore accountable for. Two variants exist; the cascade variant is used for cascade pipelines, and the S2S variant is used for both S2S and hybrid pipelines (i.e. any pipeline where the assistant consumes user audio directly rather than an STT transcript).

`user_turns_disclaimer (cascade)`

About user turns: User turns are **transcripts** produced by the assistant’s speech-to-text (STT) system. The assistant receives these transcripts as text input -- this is the only representation of user speech available to the assistant. STT transcripts may contain errors (misheard words, garbled names, dropped syllables), but the assistant cannot know what the user actually said beyond what the transcript shows. Evaluate the assistant against the transcript: if the transcript says “Kim” (even if the user actually said “Kin”), the assistant is acting on “Kim” -- that is what it received. Do not penalize the assistant for the transcript’s accuracy.

`user_turns_disclaimer (S2S / hybrid)`

About user turns: This is a **speech-to-speech** system -- the assistant receives raw audio directly, not a text transcript. The user turns shown here are the **intended text** (what the user simulator was instructed to say), not what the assistant heard. The assistant is responsible for its own audio understanding. If the assistant misheard the user and acted on incorrect information, that reflects on the assistant -- accurate audio understanding is part of its responsibility. The only mitigation is proper disambiguation: if the assistant was unsure about what it heard, it should have asked the user to confirm or clarify.

M.1.3 Assistant turns disclaimer

The string `assistant_turns_disclaimer` plays the symmetric role for the “assistant” rows of the trace, telling the judge whether those rows are the LLM’s intended text or an STT transcription of the assistant’s audio. As above, two variants exist and pipeline-based dispatch is automatic.

`assistant_turns_disclaimer (cascade)`

About assistant turns: Assistant turns shown here are the LLM’s **intended text** -- exactly what the agent produced before TTS rendering. When a user response in the transcript appears to dispute, contradict, or react oddly to an assistant turn that itself looks correct, the most likely cause is an STT error on the user side (the user actually heard something different from what the transcript shows the assistant said). Do not penalize the assistant’s prior question, statement, or read-back as “confusing” or “poorly phrased” in that case -- the assistant LLM had no way to know what the user actually said or heard beyond the transcript.

`assistant_turns_disclaimer (S2S)`

About assistant turns: This is a **speech-to-speech** system -- the agent produces audio directly, with no separate intended-text step. The assistant turns shown here are **STT transcriptions of the agent’s audio**, not text the LLM wrote. Audio articulation fidelity (whether the agent *spoke* an entity clearly and correctly) is scored separately by the `agent_speech_fidelity` metric on the actual audio -- do not penalize the agent here for what may be TTS-rendering or STT-transcription artifacts in its turns. Tool call parameters and tool responses shown in the trace are the literal values the agent sent and received via the API, not audio -- if a tool parameter looks wrong, the agent really sent it that way; if the agent’s claim contradicts a tool response, the tool truly returned the value shown.

M.2 Faithfulness

Faithfulness Prompt

You are an expert evaluator analyzing whether a voice assistant remains faithful to information, policies, and instructions throughout a conversation. You will evaluate the conversation across five dimensions, each scored as a binary flag (true = issue present, false = no issue) and, when flagged, a severity rating (1 = major user impact, 2 = minor or ambiguous; 3 = no issue). See the Rating Scale section below for detailed severity guidance.

Each dimension evaluates a **different type of faithfulness violation**. Every issue in the conversation maps to exactly one dimension -- there is no overlap.

Agent Instructions
{agent_instructions}

Agent Role
{agent_role}

Available Tools
{available_tools}

Current Date/Time
{current_date_time}

Understanding User and Assistant Turns
{user_turns_disclaimer}

{assistant_turns_disclaimer}

Full Conversation with Tools
{conversation_trace}

IMPORTANT — Interruption tags: The transcript may contain inline tags indicating speech overlap. These are informational metadata about the voice interaction -- they are NOT faithfulness issues by themselves:

{interruption_tags_reference}

When evaluating, treat these tags as natural voice interaction phenomena. Do NOT treat interrupted or truncated speech as a faithfulness violation. Specifically:

- Do NOT flag truncated content caused by interruptions as hallucination or misrepresentation -- the assistant may not have been able to finish what it was saying.
- Do NOT flag incomplete information delivery caused by the user interrupting mid-sentence as a policy violation -- the assistant was cut off, not choosing to withhold information.
- Only flag a faithfulness issue if the assistant's actual chosen content (not content lost to interruptions) violates one of the evaluation dimensions.

Evaluation Dimensions

1. fabricating_tool_parameters

Scope: Tool call inputs only. Did the assistant make a tool call with parameters that were not grounded in user-provided information or prior tool results?

IS a flag:

- Using a confirmation number, ID, or value that the user did not provide and no prior tool returned
- Guessing or inventing parameter values instead of asking the user -- including fabricated IDs and placeholder values like "?", "UNKNOWN", "MISSING", or "N/A"

- Using a parameter value from a different context or conversation segment where it does not apply
- Incorrectly categorizing data for enum/categorical tool parameters when the categorization is not supported by the data (e.g., choosing an "over-threshold" bucket when the actual value falls below the threshold; or choosing a higher severity/priority category than the described situation supports)
- A tool call parameter that cannot be traced to any user statement or prior tool result is a fabrication -- even if the tool happens to return correct results
- Hallucinated details in free-text tool fields (e.g., issue_summary, transfer notes) that were not provided by the user or returned by any tool
- Adding random characters to a confirmation number or doubling arbitrary characters to get to the right number of characters.

Is NOT a flag:

- Using parameter values explicitly stated by the user
- Using parameter values returned by a prior tool response (e.g., reusing an ID or record returned by an earlier lookup in a subsequent call)
- Using reasonable defaults that are standard for the tool (e.g., a date format conversion)
- Standard domain mappings from user-stated information (e.g., "Chicago O'Hare" → "ORD", "Miami" → "MIA"; or other unambiguous geographic, enterprise, or industry-standard mappings present in the agent's domain) -- unambiguous mappings are considered grounded
- Parameters grounded in policy entitlements derived from prior tool results (e.g., setting an entitlement/waiver/priority flag to the value that the user's status or situation qualifies for per policy, when that qualification is clearly supported by a prior tool result)
- Reasonable contextual inferences for categorical parameters, when the enum value's meaning is clearly supported by the user's stated intent
- Numeric values derived from prior tool results through simple arithmetic (e.g., summing line items, subtracting used amounts from a balance, or computing a total from prior results)
- System-level or framework-generated tool calls made before the assistant has any user input, if the assistant subsequently asks for proper information
- Trying valid variations of a user-stated identifier or name when a prior read-only lookup failed (e.g., swapping commonly confused letters like "v"/"z" or "b"/"d", reordering characters, trying alternative spellings) -- this is proper disambiguation behavior, not fabrication, since the variation is grounded in the user's original utterance

Before flagging a parameter as fabricated: Verify it cannot be traced to ANY source -- user statements, prior tool results, policy entitlements, simple arithmetic from known values, or standard domain mappings. Also verify enum values against the actual tool specification before claiming a value is invalid.

Read-only vs write tools — severity guidance: A fabricated parameter passed to a **read-only** tool (search, fetch, status check, eligibility check) that fails harmlessly, is self-corrected, and never surfaces incorrect information to the user is still a flag, but rated **2** (see Rating Scale). A fabricated parameter passed to a **write** tool -- one that creates, modifies, cancels, or otherwise changes state -- is rated **1** regardless of whether the tool happens to succeed, because the wrong value reached the system of record.

2. misrepresenting_tool_result

Scope: How the assistant reports tool results to the user. Did the assistant inaccurately convey information that was returned by a tool?

IS a flag:

- Stating incorrect values for fields that the tool response explicitly provided (e.g., wrong departure time, wrong fare amount, wrong seat number, wrong appointment date, wrong ticket priority, wrong PTO balance)
- Contradicting what a tool response returned (e.g., saying a flight is on time when the tool showed a delay, stating "window seat" when tools show an aisle seat, telling a user their ticket was resolved when status is still open, or reporting a prescription refill was approved when the tool shows denied)
- Omitting critical information from a tool result that changes the meaning (e.g., not mentioning a cancellation fee when the tool returned one and a non-zero amount is being charged; not mentioning a side-effect warning returned by a medication lookup; not mentioning that an access request requires manager approval)
- Failing to disclose costs/fees or other user-consequential information shown in tool results that the user would need to make an informed decision (e.g., when a non-zero charge is being collected, or when a tool returns eligibility/approval caveats)
- Arithmetic errors when computing values from tool data (e.g., incorrectly calculating fare differences, arrival times, remaining balances, or proration amounts) -- verify all math carefully before flagging or clearing

Is NOT a flag:

- Minor rounding or formatting differences that don't change the meaning (e.g., "\$384.00" vs "\$384")
- Omitting non-essential details from a tool result while accurately conveying the key information
- Paraphrasing tool results in conversational language while preserving accuracy
- Failing to mention a fee or cost that was \$0 or fully waived (i.e., no amount is being charged), when the financial outcome is accurately communicated
- Filtering tool results based on user-stated constraints (e.g., showing only 4 of 5 flights when the 5th doesn't meet the user's arrival time requirement; showing only weekday appointment slots when the user specified weekdays only) -- this is correct behavior, not misrepresentation
- Reasonable inferences combining tool data with contextual information (e.g., inferring a flight has departed when scheduled departure is before current time and status shows no cancellation; inferring an SLA breach when a ticket's due date has passed and status is still open)
- Time format conversions (e.g., 16:40 = 4:40 PM, 17:00 = 5:00 PM)

{misrepresentation_pipeline_note}

Verification requirements: When checking the assistant's statements against tool results: (1) carefully compute differences/totals (e.g., fare differences as (new fare - original fare) + fees) rather than confusing a total with a delta; (2) check time-format and unit conversions (24h ↔ 12h, local vs. UTC, currency); (3) verify arithmetic independently before flagging a discrepancy; (4) cross-reference ALL relevant tool result fields, not just one.

3. violating_policies

Scope: Agent instructions and policies only. Did the assistant act in a way that contradicts the agent instructions, system policies, or procedural requirements?

IS a flag:

- Failing to follow explicit procedural steps outlined in agent instructions (e.g., skipping a required verification step)
- Offering options or taking actions that the agent instructions explicitly prohibit
- Not applying policies that are clearly applicable to the situation (e.g., not offering an entitled benefit, not following a required disclosure)
- Stating a policy incorrectly, or significantly changing a policy's meaning

- **Temporal sequencing for consequential actions:** When instructions require "explain before acting" or "get explicit confirmation before proceeding," the assistant must pause for user confirmation BETWEEN read operations and write operations that have financial consequences or are irreversible. Executing such read and write operations in the same turn without intermediate user confirmation violates these instructions. Summarizing results TO the caller after the fact does NOT satisfy a requirement to get confirmation FROM the caller before acting.
- **Irreversible write operations** (cancellations, rebookings, refunds, account/record changes, ticket submissions, access grants, etc.) executed (a) without disclosing what the agent's policies (see Agent Instructions above) require to be surfaced before that write (e.g., a fee, balance impact, eligibility caveat, or approval requirement -- varies by policy), or (b) without a clear signal of user intent for that specific action. Do not invent disclosure requirements the policy does not impose.

Is NOT a flag:

- Following reasonable interpretations of ambiguous instructions
- Minor stylistic deviations from instructions that don't affect the outcome (e.g., slightly different wording for a required disclosure)
- Actions not covered by any explicit policy or instruction
- Proactive issuance of no-cost benefits the user is clearly entitled to (e.g., goodwill compensation during a service disruption, waiving a fee the user is explicitly entitled to) without explicit confirmation -- these are beneficial actions with no negative consequence, and the user's entitlement or explicit request serves as sufficient basis
- When a user explicitly requests a specific action AND the general cost structure has been communicated, proceeding without re-stating exact amounts (if not yet knowable) is not a clear violation
- Proceeding with a write tool based on implicit intent (the user described the action they want and provided the necessary information, but did not give a verbatim "yes, do it now"), when the agent's policies do not require explicit confirmation for that action
- "Predict-then-correct": stating an action should have no fee/cost based on policy *before* calling the tool, then proceeding to call the tool -- no violation **when the tool result confirms the prediction** (e.g., "there should be no fee based on policy" → tool call → fee is waived as predicted). If the tool result contradicts the prediction (e.g., a fee is charged anyway), that *is* a faithfulness issue: the user was committed to the action under incorrect cost expectations, and an after-the-fact correction doesn't undo it.

Evaluating policy application: When two policy paths could apply (e.g., same-day change vs. missed flight; in-warranty repair vs. out-of-warranty replacement), consider timeline and eligibility carefully. If the user is within the more favorable policy window (e.g., a flight hasn't departed yet, a request is within SLA, a return is within the return window), applying the more favorable applicable policy is not a violation. Also, if two policy paths produce the identical fee/outcome, choosing one over the other is not a material violation.

4. failing_to_disambiguate

Scope: Handling of ambiguous or contradictory information. Did the assistant make assumptions or proceed without clarification when the user's input was ambiguous or contradictory?
{disambiguation_context}

IS a flag:

- Proceeding with an action when the user's request could reasonably refer to multiple options and the assistant did not ask which one
- Making assumptions about user intent when the user provided contradictory information (e.g., user says two different dates)
- Choosing between conflicting pieces of information without asking the user to clarify
- Not clarifying ambiguous input that has an impact on the downstream conversation. For example, "after noon" and "afternoon" could refer to different times of day and should

not be silently inferred. The agent should not make a decision that excludes available options without validating the user's intent.

- When unable to retrieve some information, if the conversation contains multiple differing versions of a confirmation code or name, the assistant should actively disambiguate rather than silently defaulting to one version or the latest one. Making look-up tool calls is inexpensive and should be done to resolve any ambiguity.
- Failing to consider possible input errors when a lookup fails for an uncommon name or alphanumeric code (e.g., not asking the user to spell it out or verify)
- Not leveraging required information, such as specific confirmation number or names, that could be reasonably inferred from the conversation.

Is NOT a flag:

- Proceeding when the user's intent is clear and unambiguous
- Asking a clarifying question when the user's request is ambiguous (this is correct behavior)
- Making a reasonable inference when the context makes the intent obvious (e.g., user says "my flight" when they only have one flight)
- Retrying a lookup with a corrected spelling after the user confirms or spells out the information -- this is proper disambiguation behavior
- Trying valid different combinations of names and confirmation codes when a lookup fails (e.g., swapping commonly confused letters like "v"/"z" or "b"/"d", reordering characters)

5. hallucination

Scope: Information the assistant states to the user that has no source — not already covered by the preceding dimensions. Did the assistant present information that was not provided by the user, not returned by any tool response, and not stated in the agent instructions or system context?

IS a flag:

- Stating facts, details, or numbers that do not appear in any tool response, user utterance, agent instruction, or system context (e.g., inventing a gate number, adding a benefit the user doesn't have, fabricating an ID or reference number, inventing an amount or deadline)
- Presenting fabricated policies, timelines, or conditions not found in any available source
- Claiming the system can perform lookups or actions using identifiers not supported by any available tool (e.g., offering to look up a record by an identifier the tools don't accept, or offering a capability not present in the available tools)
- Misidentifying the brand, company, or agent role (e.g., using a different airline name, naming the wrong organization)

Is NOT a flag:

- Stating information that is directly inferable from tool results and/or system context (e.g., computing an arrival time from departure + duration, calculating an expiration date from current date + valid_months, or computing a remaining amount from a limit minus used amount)
- Referencing the current date/time from the system context -- this is grounded information, NOT hallucination
- Providing general conversational courtesies that don't assert factual claims
- Hedged, commonsense caveats (e.g., "you may want to verify at the counter") that don't contradict tool results or policy -- only flag fabricated information presented as definitive fact
- General domain knowledge (e.g., standard check-in windows, typical appointment lead times, standard password-reset flows) that is reasonable and not contradicted by tool results

Critical verification step: Before flagging hallucination, check ALL available sources: (1) all tool responses in the conversation, (2) user utterances, (3) agent instructions, (4) the Current Date/Time field and other system context metadata -- do NOT assume these fields are empty

without verifying. Information derived from system context (e.g., current date) is grounded, not hallucinated.

Disambiguation from other dimensions:

- If the assistant misquotes, distorts or embellishes a tool result → flag under `misrepresenting_tool_result` (the source exists but was reported incorrectly)
- If an unsupported capability is offered in passing → flag here; if actually attempted via fabricated tool call → flag under `fabricating_tool_parameters`
- If the assistant states something with NO source at all → flag here

You will focus only on the above dimensions. You will NOT consider conversation flow, task completion, or other criteria outside of faithfulness.

Rating Scale

For all five dimensions, determine if there is evidence that one or more issues should be flagged and rate that dimension based on the following guidelines. Severity hinges on the **impact on the user** -- both inside the conversation and beyond it.

- **3 (No faithfulness issues):**
 - No issue with this dimension
- **2 (Minor or ambiguous faithfulness issues -- low user impact):**
 - A single isolated issue that does not materially affect the outcome, such as:
 - * Calling a **read-only** tool (e.g., a lookup, search, or status tool) with wrong parameters, when the error is caught quickly and no incorrect information reaches the user
 - * A small hallucination or misstatement with no consequence inside the conversation and no downstream effect (e.g., a minor phrasing embellishment that does not alter any decision or action)
 - * A minor ungrounded tool parameter that doesn't affect results
 - * A minor policy deviation that doesn't affect the user's decision-making or their understanding of the policy
 - * Skipping a read-back-and-confirm step on values (dates, names, identifiers) the user stated in preceding turns, when the values written to the tool call match what appears in the user's transcribed utterance. **This carve-out does NOT apply when the transcript shows signals of a likely transcription/STT divergence -- e.g., the user explicitly said an earlier readback was wrong, OR the transcript contains two or more different versions of the same value across turns. In those cases the assistant had clear signal that a readback was needed to surface the discrepancy, and skipping it likely committed a wrong value to the system -- rate 1, not 2.**
 - Minor instruction-following deviations that do not materially affect the outcome (e.g., slight formatting differences, omitting low-importance optional steps)
 - Borderline cases where it is unclear whether a faithfulness violation occurred due to ambiguous instructions, incomplete context, or reasonable interpretation differences
 - Adopting incorrect terminology from the user (e.g., wrong brand or product name) while processing the correct record, when it doesn't cause confusion or incorrect actions
 - If something appears as being borderline an issue, it should probably be rated 2.
- **1 (Clear faithfulness violations -- major user impact):**
 - Any issue that materially affects the user, either during the conversation or afterward. Such as:
 - * Calling a **write** tool with wrong or fabricated parameters -- regardless of whether the call succeeds.
 - * Committing a value to a write tool when the transcript shows clear signals of a likely transcription/STT divergence (e.g., the user explicitly said an earlier readback was wrong, OR the transcript contains two or more different versions of the same value across turns) and the assistant did not resolve the ambiguity before committing -- **rate 1 even if the written value happens to match the most recent**

transcribed utterance, because the divergence signal made a readback necessary and the assistant likely committed a wrong value.

- * Executing an **irreversible action without the explicit user confirmation that the agent's policies (see Agent Instructions) require** for that specific action (e.g., a "summarize and confirm before submitting" rule). If the agent's policy does not require explicit confirmation for that action, implicit intent from the conversation flow is sufficient -- do not flag the absence of a verbatim "yes, do it now."
 - * Executing an **irreversible action without first disclosing what the agent's policies (see Agent Instructions) require to be surfaced** before that write (e.g., balance impact, fee, eligibility caveat, approval requirement -- varies by policy). Do not invent disclosure requirements the policy does not impose.
 - * **Financial or reputational impact** on the user or the company -- e.g., communicating incorrect charges, fees, refunds, balances, or amounts; making commitments the company cannot honor; misstating obligations in a way that could harm trust
 - * Leaving the user with an **incorrect understanding of company policy** that they could act on in the future (e.g., misstating eligibility rules, entitlement conditions, or process requirements), even if the current conversation ends up fine
 - * Hallucinating information not present in tool results, especially consequential figures (costs, balances, dates, approvals) communicated to the user as fact
- Any faithfulness issue that repeatedly prevents the conversation from progressing is also rated 1.

For the final rating of the conversation, use the minimum rating across all dimensions as the overall faithfulness rating (i.e., if any dimension is rated 1, overall rating is 1; if all dimensions are 3, overall rating is 3; if there are no 1s but at least one 2, overall rating is 2).

Response Format

Respond in JSON format:

```
{
  "dimensions": {
    "fabricating_tool_parameters": {
      "evidence": "<string>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "misrepresenting_tool_result": {
      "evidence": "<string>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "violating_policies": {
      "evidence": "<string>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "failing_to_disambiguate": {
      "evidence": "<string>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "hallucination": {
      "evidence": "<string: 1-2 sentences citing specific examples from the transcript, or 'None' if not flagged>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    }
  }
}
```

```
"rating": <int: 1, 2, or 3 -- minimum rating across all dimensions>
}}
```

disambiguation_context (Cascade)

Since the assistant is working from a speech-to-text transcript, it should account for potential transcription errors, and clarify any ambiguity in the user's intent, especially when they lead to write/irreversible operations. It's not needed to clarify if the tools called are simple lookups, but if the lookups fail, the assistant is expected to clarify the user's intent.

disambiguation_context (S2S / hybrid)

Since the assistant processes raw audio directly (speech-to-speech), it should account for potential audio perception errors - mishearing letters, numbers, names, or codes is common with spoken input. The assistant should clarify any ambiguity, especially for alphanumeric codes, names, and values that lead to write/irreversible operations. It's not needed to clarify if the tools called are simple lookups, but if the lookups fail, the assistant is expected to clarify the user's intent. The bar for disambiguation is higher than for a text-based system because the assistant knows it is working from audio and should anticipate mishearings.

misrepresentation_pipeline_note (S2S)

****Speech-to-speech scoping for this dimension.**** Because assistant turns in the trace are STT-transcribed audio (see ***About assistant turns*** above), token-level discrepancies between an assistant utterance and a tool result - dropped/added dashes, single-character substitutions, missing or extra digits within long alphanumeric IDs, altered spacing - typically reflect TTS-rendering or STT-transcription artifacts and are scored by 'agent_speech_fidelity', not here. Only flag 'misrepresenting_tool_result' when the discrepancy is structural/semantic (wrong field, wrong order of magnitude, wrong category) or when downstream signals - subsequent tool calls, follow-up actions, user objections - show the agent was internally operating on a wrong value.

M.3 Speech Fidelity

Speech Fidelity Prompt

You are an expert evaluator judging the fidelity of this audio file against the intended text. You will listen to one audio clip and verify that the spoken content faithfully reproduces the intended text, with special attention to TTS-critical entities. The audio provided is a recording of the agent's side of a conversation, and contains only the agent responses, not the user.

```
## Intended Turns
{intended_turns_formatted}
```

```
## IMPORTANT: Comparison Rules
```

Your task is to compare the **exact intended text** word-for-word against what you hear in the audio. The TTS-critical entities highlight which parts are most important to verify, but they do NOT replace or override the intended text.

```
## Understanding the Intended Text
```

The intended text may contain non-spoken tags and markers. You must understand these to evaluate fairly.

Audio-Direction Tags

Tags like [slow], [firm], [annoyed] describe how the words were meant to be spoken. They are NOT spoken aloud and should never be expected in the audio.

Interruption Tags

{interruption_tags_reference}

The tags tell you that certain portions of the intended text were likely never spoken, because the speaker was interrupted or cut themselves off. Do NOT penalize for missing words that fall in a region the tags indicate was not spoken.

Key principle: If a tag indicates that a section of text was likely not spoken aloud (due to interruption or cut-off), do NOT penalize for those words being missing from the audio. Only evaluate fidelity for words that were reasonably expected to have been spoken.

Evaluation Criteria

For each intended turn, compare what you hear in the audio against the intended text. Focus especially on **TTS-critical entities** listed for each turn.

Entity categories to watch:

- Confirmation codes (e.g., ZK3FFW, FAROUM, 8JVSDF)
- Domain-specific identifiers (e.g., flight numbers like "SkyWay 410", ticket or incident numbers, order numbers, case IDs)
- Dollar amounts (e.g., \$15, \$1,285.00)
- Short alphanumeric codes (e.g., seat numbers like "21C", room numbers, extension numbers)
- Spelled-out codes (e.g., "Z K three F F W") -- verify EVERY letter and digit individually; "K O L T S F" vs "K O L T S S F" is an error
- Reference IDs with segments (e.g., REF-8JVSDF-001, MEAL-FAROUM-PAX0) -- verify each segment; "M E L" vs "M E A L" is an error
- Times (e.g., 3:55 PM, 10:30 AM)
- Dates (e.g., March 25th, February 3rd)
- Names (e.g., Mr. Rivera, Rodriguez)

What constitutes an error (rating = 0):

- Any entity spoken incorrectly (wrong digits, letters, amounts, numbers)
- Missing words that change the meaning or omit an entity
- Added words that introduce a factually incorrect entity
- Substituted words that alter an entity value

What to ignore (does NOT cause rating = 0):

- Minor pronunciation variations that do not change the identity of an entity (e.g., "Ms." vs "Miss" is acceptable)
- Filler words ("um", "uh", "so") added or omitted
- End-of-audio cut-off: if the audio cuts off at the very END of the last turn, missing trailing words is acceptable as long as all entities in that turn were spoken correctly before the cut-off
- Slight pacing or prosody differences
- Non-spoken tags: [slow], [firm], [annoyed], and all interruption tags listed above
- Words in regions flagged by interruption tags as likely not spoken

Rating Scale (per turn)

- **1 (High Fidelity):** All entities are spoken correctly. Non-entity words are faithfully reproduced with no meaningful omissions or additions.

- **0 (Low Fidelity):** One or more entity errors, OR significant non-entity word errors that change the meaning of the turn.

Response Format

Respond with a JSON object. Each turn entry must include the `turn_id` matching the turn number shown in the Intended Turns above:

```
{
  "turns": [
    {
      "turn_id": <int: the turn number from the Intended Turns>,
      "transcript": <string: your transcription of the audio for this turn, use only the audio for this not the intended text>
      "explanation": "<string: 1-3 sentence analysis of fidelity for this turn, citing specific intended vs actual mismatches, noting any regions skipped due to interruption flags>",
      "rating": <0 or 1>
    }
  ]
}
```

M.4 Conversation Progression

Conversation Progression Prompt

You are an expert evaluator analyzing whether a voice assistant effectively moved a conversation forward. You will evaluate the conversation across four dimensions, each scored as a binary flag (true = issue present, false = no issue).

Each dimension evaluates a **different type of action**. Every issue in the conversation maps to exactly one dimension -- there is no overlap. Ensure to consider both the assistant agent instructions and the following agent dimensions when evaluating the conversation.

IMPORTANT — Scope boundary with faithfulness: This metric evaluates whether the conversation moved forward efficiently. It does NOT evaluate whether the assistant followed policies, complied with user constraints, or acted faithfully to its instructions -- those are faithfulness concerns. If an issue is primarily about the assistant violating a policy or acting against the user's explicit instructions (e.g., taking an action the user said not to, not disclosing fees), do NOT flag it here even if it also affected conversation flow. Only flag issues where the assistant's conversational choices (questions asked, information repeated, tools called) were themselves inefficient or counterproductive.

IMPORTANT — Voice conversation context: This is a voice (spoken) conversation, which means speech recognition errors are common. When the assistant repeats a request because the previous attempt was misheard or garbled, this is expected behavior in a voice interface, not a progression issue.

IMPORTANT — Interruption tags: The transcript may contain inline tags indicating speech overlap. These are informational metadata about the voice interaction -- they are NOT conversation progression issues by themselves:

```
{interruption_tags_reference}
```

When evaluating, treat these tags as natural voice interaction phenomena. Do NOT penalize interruptions themselves. Only flag an issue if the interruption caused observable consequences (e.g., information loss because the agent's cut-off speech contained critical details that were never restated, or unnecessary repetition because the agent repeated already-heard information after being interrupted).

Understanding the Conversation Trace

```
{user_turns_disclaimer}
```



```
{assistant_turns_disclaimer}
```

```
## Full Conversation with Tools
```

```
{conversation_trace}
```

```
## Evaluation Dimensions
```

```
### 1. unnecessary_tool_calls
```

Scope: Tool call actions only. Were any tool calls unjustified -- repeated without reason, made without required information, or made for data already available?

IS a flag:

- Calling the same tool with the same parameters after a prior successful response (no new user input or error in between)
- Calling a tool with empty or missing required parameters, causing a predictable error (e.g., calling a lookup tool with empty strings before asking the user for the required identifier)
- Calling a tool when the needed information was already returned by a previous tool response
- Calling a tool to verify something a prior tool response already confirmed

Is NOT a flag:

- Retrying a tool call after a tool error with corrected parameters
- Calling the same tool with different parameters (e.g., different IDs or search criteria)
- Sequential tool calls that each return new, necessary information (e.g., a record lookup followed by a status check followed by a related-details fetch)
- A tool call that fails unexpectedly (the assistant could not have predicted the failure)
- Tool calls that are necessary for the task but were executed prematurely (e.g., before the user confirmed) -- premature execution is a faithfulness/policy compliance issue, not a conversation progression issue
- Tool calls that follow standard agent instructions (e.g., automatically carrying over related attributes or defaults when taking an action) even if the user did not explicitly request those specific actions

CAVEAT: If the model makes 3 or more unnecessary tool calls, this dimension should be rated 1.

```
### 2. information_loss
```

Scope: The assistant's memory of established facts. Did the assistant fail to retain or act on information already established in the conversation -- whether from the user's statements or from prior tool responses?

This dimension is about the assistant **forgetting or ignoring known facts**, regardless of how that failure manifests (re-asking, wrong assumptions, ignoring constraints).

```
{information_loss_pipeline_note}
```

IS a flag:

- Re-asking the user for information they already provided (e.g., asking for the confirmation number or reference ID after the user stated it and it was used successfully).
- Ignoring a constraint the user explicitly stated (e.g., the user ruled out a particular action but the assistant still asks about it or asks for the details that would only be needed to take that action)
- Failing to use relevant data from a prior tool response when it was needed for the next step (e.g., not using an identifier returned by an earlier lookup when making a follow-up tool call that requires it)

Is NOT a flag:

- Asking for information the user has not yet provided
- Asking a clarifying question about genuinely ambiguous information
- Asking for authentication or identification details required by the agent instructions (e.g., a confirmation number, reference ID, or user's name) at the start of the conversation
- The assistant acting on information that contradicts what the user said, when the contradiction is due to a faithfulness or policy violation -- flag that under faithfulness, not here. Only flag here if the assistant demonstrably forgot or ignored previously established facts within the conversation flow.

Disambiguation from other dimensions:

- If the assistant re-asks for user-provided info → flag here (not `redundant_statements`)
- If the assistant makes an unnecessary tool call because it forgot a prior result → flag under `unnecessary_tool_calls` (the tool action is the observable problem)
- If the assistant proceeds with an action that contradicts the user's stated preference (e.g., choosing a different option than the one the user requested) → this is a faithfulness violation, not `information_loss`. Only flag here if the assistant clearly forgot the user's input, not if it chose to override it.

3. `redundant_statements`

Scope: The assistant repeating its own previous output. Did the assistant restate information it had already communicated to the user?

This dimension **ONLY** covers the assistant repeating **its own prior utterances** -- not forgetting user input (that is `information_loss`) and not tool call issues (that is `unnecessary_tool_calls`).

IS a flag:

- Restating details, times, amounts, or status information the assistant already told the user in an earlier turn (outside of a final recap) when the user did not ask for it
- Repeating the same explanation or instruction in multiple turns when the user has acknowledged and moved on

Is NOT a flag:

- A brief recap or summary at the end of the conversation (this is helpful, not redundant). However, if the assistant provides multiple recaps across different turns, only the final one is exempt -- earlier recaps that restate already-communicated information are still flagged.
- Confirming back details to the user once for verification (e.g., reading back a confirmation number the user just provided)
- Stating information for the first time, even if it was available from a tool response earlier
- Repeating information in direct response to the user explicitly requesting confirmation or asking to hear it again (the user must clearly ask -- simply continuing the conversation is not a request for repetition)
- Re-explaining a policy or constraint when the user continues to challenge, dispute, or insist against it -- the assistant must reiterate its position in these cases and should not be penalized for doing so. However, if the assistant repeats the exact same explanation verbatim across multiple turns, flag it -- the assistant should vary its phrasing.
- Repeating a request for information (e.g., confirmation code, spelling) when speech recognition or transcription errors clearly caused the previous attempt to fail (e.g., garbled text, partial characters, obvious mishearing visible in the transcript). Do NOT apply this exception when the transcript shows no evidence of ASR failure -- the assistant re-asking without cause is still a flag.

4. question_quality

Scope: The quality and appropriateness of the assistant's questions, where the issue is NOT caused by forgetting information (that is information_loss). Did the assistant ask poorly formed questions or fail to ask a necessary clarifying question?

IS a flag:

- Asking an overly broad or vague question when the assistant had enough information to take action (e.g., "What would you like to do?" when the user already stated a clear goal that the assistant remembers but chose not to act on)
- Asking multiple questions at once when a single tool call could have resolved the need
- Failing to ask for clarification when the user's request was genuinely ambiguous, and instead proceeding with assumptions
- Failing to ask for clarification when there are multiple options that meet the user's requirements
- Failing to ask for required information before taking an action (e.g., not asking for required details for a tool call before making the tool call, when those details have not been made available through a previous tool call, or inputs from the user)
- Failing to provide necessary information for the user to make a decision (e.g., not providing clear information about the details of the options available to the user)
- Taking an irreversible action (e.g., cancellation, rebooking, ticket submission, access grant, account change) without first confirming when user input is ambiguous or contradicts system data (e.g., user claims a 4-hour delay but system shows 45 minutes -- assistant should clarify before acting)

Is NOT a flag:

- Asking for required authentication or identification information required by the agent instructions (e.g., confirmation number, reference ID, user's name)
- Asking a clarifying question when the user's intent is genuinely ambiguous
- Asking a follow-up question based on new information from a tool response
- Asking the user to confirm an error-prone value (alphanumeric ID, code, date, dollar amount, spelled-out name) just provided -- read-backs of error-prone values are standard voice-agent practice required by typical agent policies and are NOT a question quality issue. This holds even when the transcript shows the assistant's read-back as matching what the user said: STT may render different spoken audio as identical text, so the read-back can still be catching a real audio-level mismatch.
- Not disclosing fees, costs, or other policy-required details before taking an action -- policy compliance (e.g., whether the assistant explained consequences before an irreversible action) is a faithfulness concern, not a conversation progression issue. This dimension only evaluates whether the assistant's questions and information-sharing effectively moved the conversation forward.
- Referencing information that exists in the agent instructions (e.g., standard fees, policies) without verifying it via a tool call -- the agent is expected to know its own instructions. Only flag if the information was genuinely unknown and required a tool call or user input.

Disambiguation from information_loss:

- If the assistant asks "What would you like to do?" because it FORGOT the user already stated their goal → flag under information_loss
- If the assistant asks "What would you like to do?" when the user's goal is clear and remembered but the assistant chose a vague question over taking action → flag here

Rating Scale

For all four dimensions, determine if there is evidence that one or more issues should be flagged and rate that dimension based on the following guidelines:

- 3 (No progression issue):

- No issue with this dimension
- 2 (Minor progression issue):
 - A single isolated issue that does not significantly impact the conversation flow (e.g., one unnecessary tool call that didn't slow things down, a single redundant restatement, one vague question)
 - A borderline case where it is unclear whether the issue constitutes a real progression problem
- 1 (Clear progression issue):
 - Multiple instances of the same type of issue in this dimension
 - A single severe issue that clearly derailed or stalled the conversation (e.g., ignoring a stated constraint or user requirement before carrying out a write operation, failing to ask for required information before taking action, asking an overly vague question when the user's goal was clear, making an overly vague assumption not supported by user inputs/conversation history when multiple options exist)

Overall Rating

The final rating considers BOTH the severity within each dimension AND the total number of flagged dimensions:

- 3: No dimension is flagged (all dimensions rated 3)
- 2: One or two dimensions are flagged at rating 2 (minor), AND no dimension is rated 1
- 1: Any of the following:
 - Any dimension is rated 1 (clear issue within a single dimension)
 - Three or more dimensions are flagged (even if each is individually minor, widespread issues across many areas constitute a clear overall progression problem)

Response Format

Respond in JSON format. The "evidence" field must ALWAYS contain 1-2 sentences referencing specific parts of the transcript, even when flagged is false. When not flagged, briefly explain why no issue was found.

```
{
  "dimensions": {
    "unnecessary_tool_calls": {
      "evidence": "<string: REQUIRED -- cite transcript examples if flagged, or explain why clean if not>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "information_loss": {
      "evidence": "<string: REQUIRED>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "redundant_statements": {
      "evidence": "<string: REQUIRED>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    },
    "question_quality": {
      "evidence": "<string: REQUIRED>",
      "flagged": <bool: true if issue is present, false otherwise>,
      "rating": <int: 1, 2, or 3>
    }
  },
  "rating": <int: 1, 2, or 3>
}
```

information_loss_pipeline_note (S2S)

****Speech-to-speech scoping for this dimension.**** Because assistant turns in the trace are STT-transcribed audio (see ***About assistant turns*** above), variant token-level readings of the same alphanumeric identifier across nearby assistant turns - dropped/added dashes, single-character substitutions, missing or extra digits within long IDs, altered spacing or capitalization - typically reflect TTS-rendering or STT-transcription artifacts on a value the agent is reading consistently in audio. These are scored by 'agent_speech_fidelity', not here. Only flag 'information_loss' when the discrepancy is structural/semantic (different entity, wrong field, wrong category - e.g., addressing the user by an entirely different first name, or referencing a different person/record than the tool returned), or when downstream signals - subsequent tool calls made with a wrong value, follow-up actions taken on stale data, user objections that the agent then fails to incorporate - show the agent was internally operating on a wrong value or had genuinely lost track of the established fact."

M.5 Conciseness

Conciseness Prompt

You are an expert evaluator judging the conciseness and voice-appropriateness of assistant responses in a voice conversation.

Conversation
{conversation_turns}

Understanding the Conversation Format

The conversation is grouped by turn_id. Each turn may contain:

- **user:** What the user said
- **assistant:** What the assistant said (there may be multiple assistant entries within a single turn -- e.g., the assistant speaks, calls a tool, then speaks again)
- **tool_call:** A tool invocation made by the assistant
- **tool_response:** The result returned by the tool

When a turn contains multiple assistant entries, evaluate them **together as a single unit** -- they represent the assistant's complete response within that turn. Tool calls and responses between assistant entries explain why the assistant spoke in multiple parts (it was waiting for data). It could also be due to interruptions from the user.

Understanding Interruption Tags

{interruption_tags_reference}

Key principle: When interruption tags are present, the assistant may not have been able to finish what it was saying. Do NOT penalize for truncated or fragmented content caused by interruptions. Only evaluate the conciseness of content the assistant chose to say, not content that might have been cut off.

Instructions

The conversation includes user, assistant, tool_call, and tool_response entries. Rate only the assistant's spoken content. User turns, tool calls, and tool responses are provided for context only.

For each turn that contains assistant content, evaluate whether the assistant's response is appropriately concise and easy to digest when spoken aloud to a human.

The assistant is expected to follow conversational voice guidelines:

- Keep responses brief and conversational (typically 2-4 sentences)

- Summarize long lists rather than reading them exhaustively
- Avoid overwhelming the listener with too much information at once
- Spread multiple requests across turns when possible
- Present options conversationally and avoid cramming excessive detail into one turn

Evaluation Criteria

When evaluating each turn, consider:

- Does the response get to the point without filler, rambling, or unnecessary content?
- Is all the information relevant and necessary given the conversation context?
- Is the amount of detail reasonable for someone listening to -- not reading -- the response?
- If the response enumerates options or items (e.g., "Option one is... Option two is..."), does the structure help the user? The volume should not be overwhelming.
- Is the provided information justified by context (e.g., confirming a detail the user may have misheard)? Or is it inappropriate (e.g., excessive itemization or explanation when the user may only care about the end result)?
- Within turns, is repetition avoided? Across turns there may be valid reasons for repetition, but it should usually not occur within a single turn.
- Essential information -- such as confirmation codes, reference IDs, ticket numbers, or other specific details the user needs to note down -- should never be penalized, regardless of length.

Allowed Exceptions (Voice Interaction Realities)

The assistant may occasionally produce longer turns when the context requires precise information transfer. The following cases should NOT be penalized for verbosity or information density. The turn itself may still be penalized for other reasons.

1. Phonetic Confirmation of Codes

- When confirming a confirmation code, reference number, ticket number, or similar identifier, the assistant may spell characters using the NATO phonetic alphabet (e.g., "B as in Bravo, F as in Foxtrot").
- This is especially appropriate when the user previously misheard or asked for clarification.

2. Reference or Identifier Delivery

- When providing an identifier the user needs to note down (e.g., a ticket number, reference code, or voucher code), the assistant may read the whole code out loud.
- This information is essential and should not be penalized regardless of length.

3. End-of-Call Wrap-Up

- The final assistant turn in a conversation may include a slightly longer recap or confirmation of next steps (e.g., summarizing the action taken, confirming what will be sent or followed up on, thanking the user).
- Minor additional detail in this final wrap-up should not be penalized unless it becomes excessively long or introduces unrelated information.

Important principle: Information given in assistant turns must be short enough for an average person to easily follow in real-time conversation and retain in working memory.

Failure Modes

When a response is not optimally concise, identify which of the following failure modes are present. A turn may have multiple failure modes.

verbosity_or_filler

Contains unnecessary wording, repetition within the same turn, hedging, or explanation beyond what the context requires.

excess_information_density

Presents too many distinct facts, options, numbers, steps, or requests at once, making

it difficult for a listener to process in real time. Note: bundling closely related transactional details that the user needs to act on or remember together (e.g., confirming a reference number, date, and one or two key details in a single turn) is expected behavior -- only flag this when the volume of information genuinely exceeds what a listener can comfortably retain.

over_enumeration_or_list_exhaustion

Reads out long lists instead of summarizing, or presents multiple options with excessive detail rather than inviting follow-up.

contextually_disproportionate_detail

Provides more background, clarification, or explanation than the situation warrants.

Contextual Leniency and Failure Mode Priority

Conciseness should be evaluated with respect to the conversational context. If additional wording or detail is clearly necessary for the user to understand or act on the information, a modest increase in verbosity should be considered acceptable and should NOT be penalized.

If none of the above are present, return an empty list for failure_modes.

Rating Scale For Each Turn With Assistant Content

- **3 (Highly Concise / No Cognitive Overload)** -- The response is clear, appropriately scoped for voice, and comfortably digestible in real time. No failure modes are present. A turn that delivers a few closely related facts as part of a single transactional step (e.g., confirming the key details of a request or incident) still qualifies as 3 if the listener can comfortably absorb it in one pass.
- **2 (Adequate but Not Optimally Concise)** -- One minor failure mode is present, but the response remains reasonably processable in a voice setting and does not meaningfully overwhelm the listener. Reserve this rating for turns where you can identify specific content that should have been omitted or deferred to a later turn -- not merely for turns that happen to contain several necessary details.
- **1 (Not Concise / Causes Cognitive Overload)** -- One or more significant failure modes are present that materially increase cognitive load and would hinder comprehension in a voice conversation.

Provide one entry per turn_id in the conversation.

Response Format

Provide your response as a valid JSON array, one entry per turn. Each entry must include the turn_id matching the turn number shown in the conversation above.

- If the turn contains assistant content, rate it with 1, 2, or 3.
- If the turn does not contain assistant content (e.g., user-only turn), set rating to null.

```
[
  {{
    "turn_id": <int: the turn number from the conversation>,
    "explanation": "<Detailed analysis referencing the evaluation criteria and explicitly linking identified weaknesses to the listed failure modes to justify the selected rating (1-3). Empty string if rating is null.>",
    "failure_modes": [<"<failure_mode_1>", "<failure_mode_2>", ...>],
    "rating": <int: 1, 2, or 3, or null if no assistant content>
  }}
]
```

If the turn is rated 3 or null, failure_modes must be an empty list: [].

M.6 User Speech Fidelity

User Speech Fidelity Prompt

You are an expert evaluator judging the fidelity of text-to-speech (TTS) audio against the intended text. You will listen to one audio clip and verify that the spoken content faithfully reproduces the intended text, with special attention to TTS-critical entities.

Evaluation Mode: User

Intended Turns

{intended_turns_formatted}

Understanding the Intended Text

The intended text may contain non-spoken tags and markers. You must understand these to evaluate fairly.

Audio-Direction Tags

Tags like [slow], [firm], [annoyed] describe how the words were meant to be spoken. They are NOT spoken aloud and should never be expected in the audio.

Interruption Tags

{interruption_tags_reference}

The tags tell you that certain portions of the intended text were likely never spoken, because the speaker was interrupted or cut themselves off. Do NOT penalize for missing words that fall in a region the tags indicate was not spoken.

Key principle: If a tag indicates that a section of text was likely not spoken aloud (due to interruption or cut-off), do NOT penalize for those words being missing from the audio. Only evaluate fidelity for words that were reasonably expected to have been spoken.

Evaluation Criteria

TTS-Critical Entities (check these carefully)

- **Personal names:** "John Smith" vs "Jim Smith"
- **Dates and times:** "December 15th" vs "December 50th", "3:45 PM" vs "3:15 PM"
- **Reference codes:** Confirmation numbers, incident numbers, booking IDs (e.g., "QWMN62" vs "QWN62")
- **Numeric values:** Dollar amounts, quantities, percentages (e.g., "\$150" vs "\$115")
- **Addresses:** Street numbers, street names, cities (e.g., "123 Main Street" vs "124 Main Street")
- **Contact information:** Phone numbers, email addresses (e.g., "tom_cobb@gmail.com")
- **Flight/route numbers:** "UA204" vs "UA240"
- **Serial numbers and other identifiers**

Error Types

- **Missing words:** Words in the intended text that were not spoken AND were reasonably expected to have been spoken (i.e., not in a region flagged by interruption tags)
- **Added words:** Extra words spoken that are not in the intended text
- **Wrong words:** Words spoken incorrectly or substituted with different words
- **Entity errors:** Any of the TTS-critical entities above spoken incorrectly

What to Ignore

- **Non-spoken tags:** [slow], [firm], [annoyed], and all interruption tags listed above

- Words in regions flagged by interruption tags as likely not spoken
- Minor pronunciation variations that do not change meaning (accent differences)
- Natural filler words (um, uh) if they do not affect core content
- Missing words at the END of the LAST turn only (audio recordings are often cut off before the final utterance completes). However, missing words in the middle of the last turn, or missing words in any earlier turn, should still be penalized.

Rating Scale (per turn)

- **3 (High Fidelity):**
 - All expected entities spoken correctly (names, dates, destinations, codes, etc)
 - All words reasonably expected to have been spoken are present and accurate.
 - Minor pronunciation variations acceptable.
 - No audio tags spoken out loud.
- **2 (Medium Fidelity):**
 - All entities spoken correctly (names, dates, destinations, codes, etc)
 - Part of a turn may be missing (often in the first turn, the first few words are missing)
 - Some words that were reasonably expected may be missing or spoken slightly incorrectly, but they are not critical and the conversation is able to progress even with this issue.
 - Potential issues with audio tags being said out loud
- **1 (Low Fidelity):**
 - One or more entity errors (missing entities, incorrect entities, etc) OR
 - Some other major error that prevents the conversation from continuing in a sensible manner.

Response Format

Respond with a JSON object. Each turn entry must include the turn_id matching the turn number shown in the Intended Turns above:

```
{
  "turns": [
    {
      "turn_id": <int: the turn number from the Intended Turns>,
      "transcript": <string: your transcription of the audio for this turn, use only the audio for this not the intended turns>,
      "explanation": "<succinct analysis; for score 1 or 2, quote the specific issue with intended vs actual; note any regions skipped due to interruption tags>",
      "rating": <1, 2, 3>
    }
  ]
}
```

M.7 User Behavioral Fidelity

User Behavioral Fidelity Prompt

You are an expert evaluator determining whether a simulated user's behavior has corrupted the voice agent evaluation.

Your job is to determine whether the user's behavior caused the agent to be evaluated unfairly -- specifically, whether the user's actions led to the database being in a different state than it should be, or prevented the agent from completing actions it otherwise would have.

Conversation Evidence
{conversation_evidence}

```

## How the conversation ended
The conversation ended due to: {conversation_end}

## User Simulator Instructions
The following is the full system prompt the user simulator was given for this conversation,
including the user's persona, goal, decision tree, must-have criteria, and end-of-call rules.
This is the source of truth for what the user-sim was told to do. Evaluate user behavior
against these instructions, not against generic notions of "what a real user would do."
...
{user_simulator_instructions}
...

## Modification Tools
The following are the tools that modify database state. These are the only tools relevant to
corruption analysis -- read-only tools are not a concern.
{modification_tools}

## Evaluation Criteria

### Guiding principle: judge the user from the user's perspective
The user is a simulated caller who only hears the agent's spoken words. They cannot see tool
calls, tool responses, or the agent's internal state. Judge the user's behavior -- what they
said, what they refused to say, when they ended the call -- based on what the agent said and
asked them, not on what the tool-call trace reveals. Use the tool-call trace asymmetrically:


- To exonerate the user (always allowed): confirm that an agent failure occurred that the user
couldn't have known about -- e.g., the agent claimed "I'm submitting your request" but
no modification tool was called (hallucinated tool call), or the agent's tool returned an
error the user couldn't have prevented.
- To convict the user (use carefully): only when the user's visible behavior (their words, their
refusals, their choice to end) plausibly caused the modification problem. A missing or
wrong modification by itself is not evidence against the user; the user's behavior must be
the proximate cause.


Analyze the conversation for the following corruption scenarios:

### Corruption Type 1: User invented requests that caused extra modifications
The user made requests OUTSIDE of their assigned goal that caused the agent to call one or more
modification tools listed above.


- Only flag this if the user's off-script request directly led to a modification tool being
called.
- If the user went off-script but the agent only called read-only tools (e.g., searching,
looking up information), this is NOT corruption.


### Corruption Type 2: User ended the conversation prematurely
The user ended the conversation before the agent had the opportunity to complete the necessary
modification tools to fulfill the user's goal.
Applicability:


- This type only applies when the conversation ended due to the user calling the end call
tool.
- If the conversation ended due to the agent failing to respond to the user's final turn,
this corruption type should never be flagged since the user wasn't the one to end the
call.


Decision procedure: flag this ONLY if a reasonable user following their decision tree could
have enabled the missing modification by staying on the call. The burden is on showing the
user blocked an action the agent was on track to complete.

Do NOT flag (agent failures that exonerate the user, even if a required modification did not
happen):

```

- **Agent error or retry offer.** The agent reported a tool error or technical issue and offered to "try again" / "look into it." The user is not obligated to wait through indefinite retries; ending after an error is consistent with their failure condition.
- **Hallucinated tool call.** The agent claimed to perform an action ("I'm submitting your request now", "I've assigned the SLA") but the modification tool was never actually called in the trace. Staying on the call longer would not have made the agent call the tool -- the agent already believed the work was done.
- **Off-script demand.** The agent asked the user to do or provide something that is not in the user's goal or persona (run a diagnostic command, provide a value the agent should already have, perform a task outside the scenario). The user cannot comply with invented prerequisites, and ending after such a demand is not corruption.
- **Skipped wrap-up before the agent's own goodbye.** The user expressed they were done or that the issue was resolved, and the agent failed to call its required closing modification (e.g., mark_resolved, close_ticket) before the goodbye exchange. The agent should call wrap-up tools before closing rapport, not after -- the user is not responsible for prompting the agent to call its own wrap-up tool.
- **Agent stuck or unhelpful** for multiple consecutive turns, or the agent said it could not help.
- **Resolution condition met.** All necessary modifications were completed and the user's goal was satisfied.

DO flag:

- The agent was actively and competently progressing (asking for a specific piece of information that exists in the user's goal, about to call a tool) and the user cut it off without cause.
- The user's decision tree explicitly required waiting for a confirmation / request ID, and the user ended without one despite the agent making normal progress (no error, no stuckness, no off-script demand, no hallucinated tool call).
- Special case -- transfer to a live agent:
 - The user is instructed to always end the call AFTER the agent transfers to a live agent. This should not be penalized.
 - However the user must give the agent a chance to actually call the transfer_to_agent tool. If the user says "Ok go ahead and transfer me. Goodbye" and ends before the tool call lands, this is a premature end and MUST be flagged.

Corruption Type 3: User failed to provide required information

The user failed to provide information from their goal that the agent explicitly asked for, preventing the agent from completing a necessary modification tool call.

- Only flag this if the agent clearly asked for specific information that was available in the user's goal, the user failed to provide it, and this directly prevented a modification tool from being called.
- Do NOT flag this if the agent never asked for the information.

Corruption Type 4: User looping caused duplicate modifications

The user repeatedly made the same request in a loop, causing the agent to call the same modification tool multiple times when it should have only been called once.

- Only flag this if the looping directly caused duplicate or extra modification tool calls.
- If the user looped but the agent handled it correctly (did not call extra modification tools), this is NOT corruption.

Corruption Type 5: User violated decision tree instructions causing a wrong modification

The user explicitly violated a specific instruction in their decision tree (negotiation behavior, edge cases, escalation behavior, resolution condition, or failure condition) AND this violation directly caused a modification tool to be called with different parameters than it would have been if the user had followed their instructions correctly.

- Examples: the user accepted an option that did not meet their must-have criteria when they should have rejected it; the user ignored an edge case instruction (e.g., accepted a standby flight when told to reject standby) and this led to a modification; the user failed to follow their failure condition and instead accepted an unsuitable resolution.

- Only flag this if the violation directly caused a modification tool to be called incorrectly. If the user deviated from instructions but no modification tool was affected, this is NOT corruption.
- Do NOT flag this if the agent only presented options that failed to meet the user's must-have criteria AND the user had no correct option to choose -- in that case the agent failed, not the user. Only flag this if the user had a correct action available (e.g., rejecting all options, asking for alternatives, triggering the failure condition) but chose incorrectly instead.

Rating

Binary Rating:

- **1 (Clean):** The user's behavior did not corrupt the agent evaluation. None of the corruption types above occurred. Minor deviations from the user's instructions that did not affect database state are acceptable.
- **0 (Corrupted):** One or more corruption types occurred -- the user's behavior caused the agent to be evaluated against an incorrect database state.

Respond in JSON format:

```
{
  "corruption_analysis": {
    "extra_modifications": {
      "analysis": "<reasoning about whether the user made off-script requests that caused modification tool calls>",
      "detected": <bool>
    },
    "premature_ending": {
      "analysis": "<reasoning about whether the user ended the call before the agent could complete necessary modifications>",
      "detected": <bool>
    },
    "missing_information": {
      "analysis": "<reasoning about whether the user failed to provide requested information that blocked a modification>",
      "detected": <bool>
    },
    "duplicate_modifications": {
      "analysis": "<reasoning about whether user looping caused duplicate modification tool calls>",
      "detected": <bool>
    },
    "decision_tree_violation": {
      "analysis": "<reasoning about whether the user violated a specific instruction and whether a correct action was available, and whether this caused an incorrect modification>",
      "detected": <bool>
    }
  },
  "rating": <int: 0 or 1>
}
```

M.8 Speakability

Speakability Prompt

You are an expert evaluator analyzing whether text is voice-friendly and appropriate for text-to-speech (TTS) systems in an ASR-LLM-TTS pipeline.

Your task: identify violations that make text unsuitable for speech output

The text MUST be scored as 0 (Voice-Unfriendly) if it contains ANY of the following violations:

VIOLATIONS (Score = 0):

1. **Markdown / visual formatting:** Any syntax that creates visual structure but cannot be spoken:
 - Bold/italic: **"**important**"** or **"*note*"**
 - Headers: **"## Title"** or **"# Section"**
 - Markdown tables
 - Repeated Punctuation/Symbols: Strings of characters (e.g., --- or *****) that are typically used for visual emphasis
2. Other kinds of formatting that should not be spoken -- JSONs with brackets, etc.
3. Missing spaces between words that would cause a TTS system to fail (e.g., "eighttwentypm" instead of "eight twenty PM"). Common acronyms are fine, this is only for words that typically should have spaces between them.

```

4. Emojis

## Instructions

Carefully review each assistant turn below. Check each turn for any of the above violations.

If you find even any violations in a turn, the rating for that turn MUST be 0.

### Assistant Turns
{assistant_turns_formatted}

### Response Format
Provide your response as a valid JSON array, one entry per turn. Each entry must include the
turn_id matching the turn number shown above.
[
  {
    "turn_id": <int: the turn number>,
    "explanation": "<string: 1-3 sentence analysis of the speakability of the assistant response,
citing specific example of any issues that you detect>",
    "rating": <int: 0 if ANY violation found, 1 if perfectly voice-friendly>
  }
]

```

M.9 Transcription Accuracy (Key Entities)

Transcription Accuracy (Key Entities) Prompt

You are an expert evaluator analyzing Speech-to-Text (STT) transcription accuracy for key entities across an entire conversation.

Your task:

1. For EACH user turn, identify all key entities in the EXPECTED text
2. Check if each entity appears CORRECTLY in the TRANSCRIBED text
3. Mark each entity as correct or incorrect
4. For entities in regions that were likely never spoken aloud (as indicated by interruption tags), still include them in the output but mark them as skipped

What Counts as an Entity

An entity must have a **specific, concrete value** -- something that could be passed as an input to a program or tool (not an AI, but a script or database lookup). Ask yourself: could this value be stored in a variable and used programmatically?

- Names (people, places, organizations): e.g. "John Smith", "Austin", "Delta Airlines"
- Specific dates and times: e.g. "December 15th", "3:45 PM" -- NOT vague references like "tomorrow morning" or "later today"
- Confirmation codes / reference numbers: e.g. "ABC123", "ZK3FFW"
- Flight numbers: e.g. "UA 204"
- Amounts and prices: use the specific value only, e.g. "\$120" -- for qualifier phrases like "under \$120", only use the specific value
- Addresses: e.g. "123 Main Street"
- Phone numbers: e.g. "555-867-5309"
- Email addresses: e.g. "john@example.com"
- Other specific identifiers: seat numbers, loyalty numbers, booking IDs, etc.

Not an entity: vague temporal words ("tomorrow", "next week", "morning"), general descriptors ("the cheap flight", "a long trip"), or open-ended qualifiers ("less than an hour", "around

noon").

Understanding Tags in the Expected Text

The expected text may contain non-spoken tags and markers. These are metadata -- they were never said aloud and must not be treated as entities or evaluated.

Audio-Direction Tags

Tags like [slow], [firm], [annoyed] describe how the words were meant to be spoken. Ignore them entirely.

Interruption Tags

{interruption_tags_reference}

These markers indicate that parts of the expected text may never have been spoken aloud, because the user was interrupted or talked over. An entity that was never spoken cannot be correctly transcribed, so you must NOT penalize for entities in regions that were likely not said, instead mark them as skipped.

Key principle: Only evaluate entities that were reasonably expected to have been spoken aloud. If a tag indicates the user was interrupted or talked over before or during an entity, still include the entity in your output but set `skipped: true` and explain why in the analysis. The correct field should reflect your best assessment of whether the transcription matched, but skipped entities will be excluded from accuracy metrics downstream.

User Turns to Evaluate

{user_turns}

Correctness Criteria

- Entity must be present (not missing) -- unless in a region flagged by interruption tags
- Entity value must match (minor formatting variations OK)
- Numbers: "150" and "one hundred fifty" are equivalent
- Dates: "December 15th" and "Dec 15" are equivalent
- Names: Case-insensitive exact match required

Important note: The expected text will often feature things formatted like "one two three" instead of "123". Your goal is to evaluate the semantic equivalence, meaning these are considered equivalent if they were heard in audio.

Examples

Example Input:

Turn 1:

Expected: My confirmation is A B C one two three on December 15th.

Transcribed: My confirmation is ABC123 on December 15th.

Turn 2:

Expected: Transfer one hundred fifty to account 1 2 3 4 5.

Transcribed: Transfer \$115 to account 12345.

Turn 3:

Expected: [slow] The code is X X F six O H, with the letter O, [assistant interrupts] not zero.

Transcribed: The code is X... X F 6 O H with the letter O.

Turn 4:

Expected: My phone number is four zero four five five five [assistant interrupts] zero eight five six.

Transcribed: My phone number is 404-555.

Example Response:

```
[
  {{
    "turn_id": 1,
    "entities": [
      {{
        "type": "confirmation_code",
        "value": "A B C one two three",
        "transcribed_value": "ABC123",
        "analysis": "Matches exactly",
        "correct": true,
        "skipped": false
      }},
      {{
        "type": "date",
        "value": "December 15th",
        "transcribed_value": "December 15th",
        "analysis": "Matches exactly",
        "correct": true,
        "skipped": false
      }}
    ],
    "summary": "All 2 key entities transcribed correctly."
  }},
  {{
    "turn_id": 2,
    "entities": [
      {{
        "type": "amount",
        "value": "one hundred fifty",
        "transcribed_value": "$115",
        "analysis": "Amount wrong: $150 vs $115",
        "correct": false,
        "skipped": false
      }},
      {{
        "type": "account_number",
        "value": "1 2 3 4 5",
        "transcribed_value": "12345",
        "analysis": "Matches exactly",
        "correct": true,
        "skipped": false
      }}
    ],
    "summary": "1 out of 2 entities correct. Amount error."
  }},
  {{
    "turn_id": 3,
    "entities": [
      {{
        "type": "confirmation_code",
        "value": "X X F six O H",
        "transcribed_value": "X F 6 O H",
        "analysis": "Missing one X -- transcribed 5 characters instead of 6. The code appears before
the [assistant interrupts] tag so it is evaluated normally.",
        "correct": false,
```

```

"skipped": false
}}
],
"summary": "1 entity found before interruption, partially incorrect (missing one X). No
entities after [assistant interrupts] tag to skip."
}},
{{
"turn_id": 4,
"entities": [
{{
"type": "phone_number",
"value": "four zero four five five five zero eight five six",
"transcribed_value": "404-555",
"analysis": "The full number is 404-555-0856. The [assistant interrupts] tag appears after
'five five five', meaning the last four digits ('zero eight five six') were likely drowned
out by the agent speaking over the user. The transcription captured the portion before the
interruption. Skipping because the entity spans into the interrupted region and cannot be
fully evaluated.",
"correct": false,
"skipped": true
}}
],
"summary": "1 entity found. Phone number spans into interrupted region -- skipped. Partial
transcription (404-555) matches the portion before the interruption."
}}
]

## Response Format
Respond with a JSON object. Each turn entry must include the turn_id matching the turn number
shown in the User Turns to Evaluate section above:
[
{{
"turn_id": <int: the turn number from the User Turns to Evaluate section>,
"entities": [
{{
"type": "<name|date|flight_number|etc...>",
"value": "<entity value from expected text>",
"transcribed_value": "<how it appeared or 'missing'>",
"analysis": "<brief reason; if skipped, explain why the entity falls in an interrupted
region>",
"correct": <true|false>,
"skipped": <true|false>
}}
],
"summary": "<1-2 sentence summary for this turn>"
}}
]

```


N Licenses & Usage

N.1 Third Party Licenses

EVA-Bench depends on the third-party software packages listed in 56. Full license texts are available in the code repository `THIRD_PARTY_NOTICES` file. All existing artifacts are used in a manner consistent with the license type and intended use.

Table 56 Third-party dependencies and their licenses.

License	Packages
MIT	pydantic, elevenlabs, litellm, deepgram-sdk, onnxruntime, azure-cognitiveservices-speech, cartesia, assemblyai, setuptools, fastapi, pyyaml, pydub, jaconv, more-itertools, pytest, pytest-cov, ruff, mypy, inflect
Apache-2.0	openai, aioboto3, google-generativeai, google-genai, google-cloud-speech, google-cloud-texttospeech, aiofiles, jiwer, streamlit, pytest-asyncio, regex
BSD-2-Clause	pipecat-ai
BSD-3-Clause	uvicorn, websockets, httpx, pandas, numpy, python-dotenv
MIT / Apache-2.0	structlog
MIT / MPL-2.0	tqdm

N.2 EVA-Bench License

EVA-Bench is released under the MIT license and is intended to be used in the following ways:

- **Academic research:** Benchmarking and evaluating conversational voice agents across task accuracy and experiential quality dimensions in research settings.
- **Enterprise evaluation:** Organizations may use EVA-Bench to evaluate and compare voice agent deployments across enterprise domains.
- **Framework extension:** Researchers and practitioners are encouraged to extend EVA-Bench with new data domains and metrics.

We discourage use of EVA-Bench for purposes that misrepresent evaluation results, such as selective reporting of metrics or use in misleading competitive benchmarking claims. Evaluation outputs should be interpreted in conjunction with the statistical methodology described in Section 4.3 and with awareness of the limitations outlined in Section 5.

Additionally, the EVA-Bench dataset and evaluation scenarios should not be used as training data for voice agents or language models. Training on benchmark data undermines the integrity of future evaluations and may lead to artificially inflated performance scores that do not reflect real-world capability.

O AI Usage

We used large language model (LLM) assistants to support writing and phrasing in select sections of this paper. AI coding agents were additionally used to assist with code scaffolding and implementation during the development of the framework. All research design, experimental analysis, result interpretation, and scientific conclusions were conducted entirely by the human authors.

P Human Annotations

Use of Human Annotators. We employed human annotators to measure judge–human agreement on evaluation scores. All annotators are full-time AI data linguists employed by the same company as the authors; annotation is a core component of their regular job responsibilities. A total of seven linguists participated, with each sample independently annotated by two linguists.

Annotation Instructions. Annotators were given the same prompts and scoring rubrics used by the LLM judge (see Section M). No additional guidelines or supplementary instructions were provided beyond the judge prompts, in order to ensure a fair comparison between human and automated judgments.

Recruitment and Compensation. Annotators were recruited from an internal team of full-time AI data linguists. As salaried employees, their compensation is determined by standard employment agreements. No external crowdworkers or research subjects were recruited for this work.

Consent and Data Use. The annotation task involved evaluating synthetically generated conversational data; no personal data from the annotators was collected or used. All annotators were informed of the purpose of the annotations and how they would be used in the research prior to beginning the task.

Risks. The annotation tasks posed no foreseeable risks to participants. The data consisted of synthetic enterprise voice agent conversations and did not contain sensitive, offensive, or personally identifiable content.

Q Ethics & Reproducibility

Ethics Statement

Our work focuses on responsible development and evaluation of voice agent systems. All evaluation scenarios in EVA-Bench are fully synthetic — no real caller data, recordings, or personally identifiable information were used in dataset construction or evaluation.

All simulation tools respect the copyright and privacy guidelines. Although EVA-Bench enables large-scale evaluation of voice agents, we cannot guarantee that models evaluated through it will not generate harmful or biased output during evaluation runs. Researchers and practitioners are strongly encouraged to implement appropriate content filtering and bias detection before deploying evaluated systems in production environments.

Additionally, we acknowledge that our current domain coverage is limited, and the simulation is conducted exclusively in English, which may inadvertently reinforce existing representational biases in audio AI systems. We encourage the community to expand EVA-Bench with more diverse languages, accents, cultural contexts, and domains. We are actively working on extending it to other languages.

Reproducibility Statement

We are committed to full reproducibility of our evaluation framework and experimental results. All EVA-Bench code, configuration files, evaluation scripts, scenario data, and documentation are publicly released as open-source under an anonymized repository included with the submission. The repository includes setup instructions, environment specifications, and scripts to reproduce all reported evaluations.

We provide comprehensive implementation details including all model configurations (Appendix B), data distributions across domains (Appendix C), judge prompts (Appendix M), and metric definitions and thresholds (Appendix E). All LLM-as-judge configurations are fully specified to enable result replication across research groups. We note that full reproduction of the reported scores requires access to commercial model APIs; results may vary across API versions and deployment configurations, as discussed in Section 5.

R Future Work

Several directions extend EVA-Bench’s coverage and realism. First, we are extending scenario coverage to multilingual settings beyond the current English-only scope. Second, we aim to support multi-agent architectures, including supervisor-worker patterns, orchestration frameworks, and agentic planning or delegation, which are increasingly common in real-world deployments. Third, we are actively improving the user simulator to reliably produce interruptions, disfluencies, and a broader range of caller behaviors, reducing the gap between simulated and real caller interactions. Fourth, we plan to incorporate safety and privacy evaluation — assessing whether agents produce harmful outputs or inadvertently expose sensitive caller information such as PII — to better reflect production-readiness requirements in high-stakes domains. We also are actively working to expand system coverage to evaluate a wider set of voice agent implementations across emerging architectures and providers.