

llmXive Automated Review of AutoResearchClaw: Self-Reinforcing Autonomous Research with Human-AI Collaboration

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_claim_accuracy.md

The manuscript makes several strong comparative claims against existing literature that require tighter alignment with the cited sources.

In Section 3.2 (Cross-Domain Coverage), the paper asserts that baselines like AI Scientist v2 and AIDE-ML fail on High-Energy Physics and Biology tasks specifically due to “missing software stacks.” This causal attribution is not supported by the cited works (Lu et al., 2024; Yamada et al., 2025), which generally focus on the agents’ reasoning capabilities or general ML tasks. The failure observed in this study is likely a result of the specific benchmark environment or the agents’ inability to adapt to new domains within the test constraints, rather than a fundamental lack of software stacks in the cited systems. This distinction is crucial for accurate scientific reporting.

In Section 4.1, the claim that “best systems solve fewer than 40% of tasks” aggregates results from ScienceAgentBench and MLE-bench. While the general sentiment is supported by the literature, the specific “40%” threshold should be verified against the exact numbers reported in Tian et al. (2024) and Chan et al. (2024). If the cited papers report 38% or 42%, the claim “fewer than 40%” is factually imprecise.

Additionally, Appendix B presents specific quantitative improvements in “hypothesis diversity” (-23% for K=2, +8% for K=5) without defining the metric used to calculate diversity. Since this is a novel metric introduced by the authors, the specific percentages are not supported by external citations and must be clearly defined in the text to be considered accurate claims.

Action items.

- **[writing]** Section 3.2 claims AI Scientist v2 and AIDE-ML fail on HEP/Biology due to ‘missing software stacks’ (Table 2). The cited papers (Lu et al. 2024, Yamada et al. 2025) do not explicitly state this limitation; the failure is likely due to the specific benchmark setup or environment constraints, not the models’ inherent inability to handle the stacks. Rephrase to reflect that the baselines failed *in this evaluation context* rather than attributing a general capability gap to the cited works.
- **[writing]** Section 4.1 states ‘ScienceAgentBench... show even best systems solve fewer than 40% of tasks’ citing Tian et al. 2024 and Chan et al. 2024. Verify if the 40% figure is a direct aggregate from these specific papers or a generalization. If the cited papers report different specific percentages (e.g., 35% or 42%), the claim ‘fewer than 40%’ may be inaccurate or require a more precise citation.
- **[writing]** Appendix B (Design-Space Exploration) claims K=2 results in ‘-23% hypothesis diversity’ and K=5 results in ‘+8% diversity’ over K=3. The text does not define the baseline metric for ‘diversity’ (e.g., semantic similarity, unique hypothesis count). Without this definition, the specific percentage claims are unverifiable from the provided text and

citations.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

Figure 1 effectively contrasts the ‘Full-Auto’ and ‘CoPilot’ approaches using side-by-side tables, charts, and text snippets. The visual layout clearly supports the caption’s claim of a ‘silent semantic collapse’ in the Full-Auto column versus differentiated results in the CoPilot column, with all necessary context provided within the figure itself.

Figure 2

Figure 2 is a clear and well-structured pipeline diagram that effectively visualizes the system architecture described in the caption. All phases, sub-steps, and feedback loops are legible, and the color-coding for human-in-the-loop gates is consistent with the description.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The manuscript relies heavily on custom-defined macros and domain-specific shorthand that obscures meaning for non-specialist readers. The most critical issue is the use of the custom command `\system{}` (rendering as “AutoResearchClaw”) and `\bench{}` (rendering as “ARC-Bench”) throughout the text without a clear, plain-English definition at the very first instance of use in the Abstract or Introduction. While the preamble defines these, the prose assumes the reader already knows what these terms represent.

In Section 1, the acronym “HITL” is introduced as “Human-in-the-loop (HITL)” but is subsequently used as a standalone noun (“HITL gates”, “HITL ablation”) without re-establishing the full phrase, which can be jarring for readers unfamiliar with the specific jargon. Similarly, the debate agent roles (Innovator, Pragmatist, Contrarian, Optimist, Skeptic, Methodologist) are listed in Section 1 and 3.2 but lack immediate, plain-language descriptions of their specific

functional duties, relying on the reader to infer the nuance from the titles alone.

The term “semantic collapse” in Section 5.5 and the Case Study is used to describe a specific failure mode where outputs become identical. This is not standard terminology in the field and should be replaced with a clearer description (e.g., “output homogenization” or “failure to differentiate”) or explicitly defined to ensure the reader understands the nature of the error. Finally, the abbreviations CD, CE, and RA in Section 4.1 and Table 2 are used without being spelled out as “Code Development,” “Code Execution,” and “Result Analysis” at their first appearance, forcing the reader to guess or search the text for the definitions. These instances of jargon overuse and undefined acronyms create unnecessary friction for a general scientific audience.

Action items.

- **[writing]** Define the custom macros `\\system`, `\\bench`, `\\cmark`, `\\xmark`, and `\\tmark` at their first occurrence in the Abstract or Introduction. Currently, they are defined in the preamble but used without immediate context for a general reader.
- **[writing]** Replace the acronym ‘HITL’ with ‘human-in-the-loop’ on first use in the Introduction (Section 1) and ensure it is not used as a standalone noun without prior definition in the abstract.
- **[writing]** Define the specific roles of the debate agents (Innovator, Pragmatist, Contrarian, etc.) immediately upon their first mention in Section 1 or Section 3.2, rather than assuming the reader understands the functional difference between these specific titles.
- **[writing]** Clarify the term ‘semantic collapse’ in Section 5.5 and the Case Study. While descriptive, it is not standard terminology in this context and requires a brief plain-English explanation of the failure mode (e.g., ‘producing identical outputs despite different inputs’).
- **[writing]** Define the specific metrics ‘CD’, ‘CE’, and ‘RA’ used in the experimental setup (Section 4.1) and Table 2 captions. The text mentions ‘Code Dev’, ‘Code Exec’, and ‘Result Analysis’ later, but the abbreviations should be explicitly defined at first use.

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The manuscript presents a coherent architecture for autonomous research, but several logical inconsistencies exist between the reported data and the textual claims, particularly regarding the definition of success metrics and the interpretation of ablation results.

First, the Human-in-the-Loop (HITL) ablation results in Table 4 present a logical ambiguity regarding the denominator for the “Accept” rate. The table lists “Valid” runs (e.g., 8/10 for Full-Auto) and an “Accept” percentage (25.0%). If the accept rate is calculated against the total

10 topics, 25% implies 2.5 accepted runs, which is impossible. If it is calculated against the 8 valid runs, the rate would be $2/8 = 25\%$, which is mathematically possible but implies 6 valid runs were rejected. The text does not explicitly define whether “Accept” is a subset of “Valid” or a subset of the total “Topics”. This ambiguity undermines the claim that CoPilot’s 87.5% accept rate is a direct result of targeted intervention, as the baseline for comparison is unclear.

Second, the Component Ablation in Table 3 contains a direct contradiction between the textual claim and the tabular data. The text states that removing the verification module “introduces fabrication.” However, the “Fabrication” column for the “w/o Verification” configuration explicitly marks “xmark” (indicating no fabrication), identical to the full system. If fabrication is the primary mechanism driving the inflated acceptance rate (5/10 vs 3/10), the table should reflect the presence of fabrication in the ablated condition. The current presentation suggests the acceptance increase is due to a different factor (e.g., lower quality thresholds), which contradicts the narrative that verification is the sole guard against fabrication.

Finally, while the arithmetic for the 54.7% improvement is correct, the paper conflates “relative improvement” (percentage increase) with “ratio” in the discussion of the Result Analysis metric. The text highlights a “100.4% relative improvement” for the Result Analysis score. While 0.523 is indeed a 100.4% increase over 0.261, the phrasing risks confusion with the overall score improvement (54.7%). The distinction between these metrics should be explicitly defined to ensure the magnitude of the “Result Analysis” breakthrough is not misinterpreted as a doubling of the *overall* system capability.

These issues require clarification in the text and potentially a re-evaluation of the data presentation in Tables 3 and 4 to ensure the conclusions strictly follow from the premises.

Action items.

- **[science]** Table 4 (HITL Ablation) lists ‘Valid’ runs (e.g., 8/10) and ‘Accept’ rates (e.g., 25%). The denominator for the Accept rate is ambiguous: is it out of 10 total topics or 8 valid runs? If 25% of 8 is 2, the math holds, but the text must clarify if ‘Accept’ is a subset of ‘Valid’ to support the claim that CoPilot’s 87.5% is superior.
- **[science]** Table 3 (Component Ablation) claims removing Verification ‘introduces fabrication,’ yet the ‘Fabrication’ column shows ‘xmark’ (none) for the ‘w/o Verification’ row, identical to the full system. The data contradicts the textual mechanism; clarify if fabrication was undetected or if the acceptance increase stems from a different cause.
- **[writing]** The text states Result Analysis shows a ‘100.4% relative improvement’ (0.523 vs 0.261). While mathematically correct as a percentage increase, this phrasing risks confusion with the overall score improvement (54.7%). Explicitly define ‘relative improvement’ vs ‘ratio’ to ensure the magnitude of the Result Analysis gain is not misinterpreted as a doubling of overall capability.

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The manuscript makes several strong claims that extrapolate beyond the provided evidence, particularly regarding comparative performance and the specific causes of failure in baseline systems.

First, the headline claim of a “54.7%” improvement over AI Scientist v2 (Abstract, Section 1) is mathematically derived from the aggregate scores (0.648 vs 0.419) but is phrased in a way that suggests a general capability leap. Without explicit qualification that this is a relative gain on a specific, narrow benchmark (ARC-Bench), this risks over-claiming the system’s general superiority. The text should be tempered to reflect that this is a metric-specific improvement.

Second, the assertion in Section 4.3 and Table 3 that baseline systems “fail on HEP/Biology due to missing software stacks” is unsupported. The paper does not present data showing that these baselines were run on these specific domains and failed specifically because of software stack incompatibility. It is equally plausible they failed due to a lack of domain-specific prompting or reasoning capabilities. Attributing the failure solely to “missing software stacks” without direct evidence of the baselines’ execution logs or error analysis is an overreach.

Third, the conclusion’s comparison of HITL modes (Section 6) conflates “accept rate” with overall robustness. While CoPilot has a higher accept rate (87.5%) than Step-by-Step (50%), Table 4 reveals that Step-by-Step produced a higher number of “Valid” runs (10/10 vs 8/10). The claim that CoPilot “surpasses” Step-by-Step ignores the fact that Step-by-Step is more reliable in producing valid outputs, even if the quality of accepted outputs is lower. This selective reporting overstates the efficacy of the CoPilot mode.

Finally, the claim that removing verification “introduces fabrication” (Section 4.5) is a serious scientific accusation. The data in Table 5 shows an inflated score and higher acceptance, but does not explicitly list or count instances of fabricated data or hallucinated citations. To support the claim of “fabrication,” the authors must provide specific examples or a count of false claims generated in the ablation study, rather than inferring fabrication solely from a higher acceptance score.

Action items.

- **[science]** The manuscript makes several strong claims that extrapolate beyond the provided evidence, particularly regarding comparative performance and the specific causes of failure in baseline systems. First, the headline claim of a “54.7%” improvement over AI Scientist v2 (Abstract, Section 1) is mathematically derived from the aggregate scores (0.648 vs 0.419) but is phrased in a way that suggests a general capability leap. Without explicit qualification that this is a relative gain on a specific, narr

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: `agents/prompts/paper_reviewer_safety_ethics.md`

The manuscript addresses safety and ethics primarily through technical safeguards (sandboxing, verification) and a dedicated “Ethical Considerations” section (Section 6). However, several critical gaps remain regarding the human-in-the-loop (HITL) component and the broader implications of autonomous scientific generation.

First, the HITL ablation study (Section 5.4, Table 4) involves human participants making decisions at various stages of the research pipeline. The paper does not state whether this study received approval from an Institutional Review Board (IRB) or an equivalent ethics committee. Given that human judgment is a variable in the experiment, standard ethical protocols for human subjects research must be explicitly cited or a waiver justified. This is a significant omission for a paper claiming to advance human-AI collaboration.

Second, while the “Sandbox Security Model” (Appendix D) outlines Docker-based isolation and network policies, it lacks a formal threat model. Autonomous agents executing code in a research context pose dual-use risks; if the sandbox is compromised, the agent could potentially access sensitive data or launch attacks on the host network. The authors should explicitly discuss the containment guarantees and the potential consequences of a sandbox escape, rather than simply stating the configuration.

Third, the “Verifiable Result Reporting” mechanism (Section 4.4) focuses on numeric consistency and citation validity. However, it does not address the risk of “semantic hallucination” where the agent generates scientifically plausible but factually incorrect interpretations of data that pass numeric checks. The case study in Appendix C (Topic T10) shows the system can produce “zero-bias” outputs that are technically correct but scientifically uninformative. The ethical framework should discuss how the system handles such “truthful but misleading” outputs, which could mislead downstream researchers if not flagged.

Finally, the “Cross-Run Evolution” mechanism (Section 4.5) stores lessons with time-decayed weighting. The authors should clarify the data privacy implications of this persistent memory, specifically whether any sensitive data from failed experiments or human interventions is retained in the lesson store and how it is protected.

Action items.

- **[science]** The manuscript lacks a formal statement regarding IRB or ethics board approval for the human-in-the-loop (HITL) study described in Section 5.4 and Appendix E. Since human participants provided interventions and judgments, explicit confirmation of ethical oversight or a waiver justification is required.
- **[science]** The ‘Sandbox Security Model’ (Appendix D) describes network isolation but does not address the risk of the autonomous agent generating or executing code that could inadvertently harm the host infrastructure or exfiltrate data if the sandbox is compromised. A threat model or specific containment guarantees are needed.
- **[writing]** The paper claims to prevent ‘hallucinated references’ via a four-layer pipeline (Section 4.4), but does not discuss the ethical risk of the system generating plausible-sounding but unverified scientific claims in the ‘Result Analysis’ phase that pass numeric verification but lack empirical grounding.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The manuscript presents a complex system with multiple novel components, but the scientific evidence supporting the central claims is currently insufficient due to missing statistical rigor and potential confounding variables in the experimental design.

First, the primary claim of a 54.7% relative improvement over AI Scientist v2 (Abstract, Section 4.2) is presented as a point estimate without any measure of uncertainty. With a sample size of only 25 topics, the variance in performance across different scientific domains could be high. The authors must report confidence intervals (e.g., 95% CI) or conduct a significance test (e.g., paired t-test or Wilcoxon signed-rank test) to demonstrate that this improvement is statistically robust and not an artifact of the specific topic selection.

Second, the component ablation study in Table 4 (Section 4.5) introduces a critical confounding variable. The “Full system” is evaluated using a “best-of-3” protocol, yet the text does not explicitly confirm whether the baselines (AI Scientist v2, AIDE-ML) or the other ablated configurations (e.g., “w/o Debate”) were subjected to the same multi-run selection process. If the baselines were run only once while the proposed system was optimized over three runs, the reported quality gap is inflated by compute budget rather than architectural superiority. The experimental protocol must be standardized across all conditions.

Third, the “Cross-Domain Coverage” results in Table 3 (Section 4.3) are difficult to interpret scientifically. The baselines are marked with an ‘x’ (failure) on Biology and HEP tasks. Given that these are autonomous coding agents, this likely indicates a failure to install necessary domain-specific libraries (e.g., Biopython, ROOT) rather than an inability to reason about the science. Without evidence that the baselines were provided with equivalent environment setup instructions or time, this comparison does not validly support the claim that AutoResearchClaw is superior in “scientific reasoning” across domains; it may simply be better at environment configuration.

Finally, the statistical claim in Section 4.5 regarding the debate ablation (“drops quality by 1.37, $p=0.003$ ”) is incomplete. The specific statistical test used to derive this p-value and the exact sample size (N) for this specific comparison are not stated. Given the small number of topics in the ablation (implied to be a subset of the 25), the power of the test is a concern.

To proceed, the authors must re-run the experiments with standardized compute budgets, report statistical significance for all comparative claims, and clarify the nature of the baseline failures in cross-domain tasks.

Action items.

- **[science]** The manuscript presents a complex system with multiple novel components, but the scientific evidence supporting the central claims is currently insufficient due to missing statistical rigor and potential confounding variables in the experimental design. First, the

primary claim of a 54.7% relative improvement over AI Scientist v2 (Abstract, Section 4.2) is presented as a point estimate without any measure of uncertainty. With a sample size of only 25 topics, the variance in performance across di

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The manuscript presents a complex system with multiple novel components, but the statistical rigor of the experimental evaluation is insufficient to support the strong comparative claims made.

First, the primary results in Tables 1, 2, and 3 report only point estimates (e.g., 0.648, 7.27, 87.5%) without any measure of variance (standard deviation, standard error) or confidence intervals. In Section 5.2, the claim that the system “outperforms AI Scientist v2 by 54.7%” is presented as a deterministic fact. Without reporting the variability across the 25 topics or multiple runs, it is impossible to determine if this difference is statistically significant or an artifact of a specific topic distribution. The absence of error bars or confidence intervals in the text (and the inability to see them in the figures) renders the magnitude of improvement unverifiable.

Second, the statistical test cited in Section 5.5 (“Removing Debate drops quality by 1.37 ($p=0.003$)”) lacks necessary methodological details. The authors do not specify the statistical test employed (e.g., paired t-test, Wilcoxon signed-rank), the degrees of freedom, or the exact null hypothesis. Furthermore, given that multiple ablation conditions are tested (Debate, Self-Healing, Verification), the manuscript fails to address the multiple-comparisons problem. Without a correction method (e.g., Bonferroni, Holm-Bonferroni), the reported p-value may be inflated, increasing the risk of Type I errors.

Third, the experimental design for the HITL ablation (Section 5.4, Table 2) is ambiguous regarding the unit of analysis. The table notes “10 topics” and “7 intervention regimes,” but it is unclear if the reported means (e.g., 7.27) are averages over 10 independent topics or over multiple runs per topic. If the 10 topics are the only independent samples, the statistical power to detect differences between modes (e.g., CoPilot vs. Step-by-Step) is low. The “Accept” rate (87.5%) is derived from a small integer count (7/8 valid runs), yet no binomial confidence interval is provided to contextualize this proportion.

Finally, the “Strict Judge” protocol described in Appendix B mentions re-adjudication for disagreements $|\Delta| > 0.20$, but the statistical properties of this inter-rater reliability are not quantified (e.g., Cohen’s Kappa). The reliance on a single “Strict Judge” score without reporting the distribution of scores or the reliability of the human/agent evaluation process introduces potential bias that is not statistically accounted for.

To proceed, the authors must re-analyze the data to include measures of variance, explicitly

state the statistical tests and corrections used, and clarify the sample size and unit of analysis for all reported metrics.

Action items.

- **[science]** Report confidence intervals or standard errors for all aggregate metrics in Tables 1-4. The current presentation of single-point estimates (e.g., 0.648, 7.27) without variance measures prevents assessment of statistical significance or effect stability.
- **[science]** Clarify the statistical test used to derive $p=0.003$ in Section 5.5 (Table 3). Specify the null hypothesis, test statistic, and whether corrections for multiple comparisons were applied given the multiple ablation conditions tested.
- **[science]** Define the sample size (N) and unit of analysis for the HITL ablation (Table 2). It is unclear if the 10 topics represent independent experimental units or if the metrics are averaged across multiple runs per topic, which impacts the validity of the reported means and accept rates.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript presents a compelling technical contribution, but the writing quality is currently compromised by significant structural inconsistencies and editing artifacts that hinder readability.

The most critical issue is the presence of two distinct “Introduction” sections (Section 1 in chunk e000 and Section 1 in chunk e002). The second instance appears to be a revised draft that was appended rather than merged, resulting in a disjointed narrative flow where the reader encounters the same section twice with slightly different phrasing and content. This suggests the LaTeX source was not properly consolidated before submission.

Furthermore, the tables are inconsistent. Table 1 in the main body (e000) lacks features like “Result verification” and “Sandbox security” that appear in the duplicate table in e002. This inconsistency forces the reader to guess which version is authoritative.

Additionally, the Appendix includes a section titled “Writing-Quality Audit” (e001) which lists internal defects such as “Duplicated figure file” and “Bracket-style pseudo-citations.” Including a self-critique of the manuscript’s own errors in the final paper is inappropriate and confusing; this section should be removed entirely.

Finally, there are minor grammatical ambiguities, such as in Section 3.2 where the subject of “Scores complexity” is omitted. While the scientific content is dense and interesting, these structural and editorial flaws must be resolved to ensure the paper is readable and professional.

Action items.

- **[writing]** The manuscript contains a duplicate ‘Introduction’ section (Section 1 in e000 and Section 1 in e002) with conflicting content. The second instance (e002) appears to be a

revision that was not properly merged, causing structural redundancy and confusion.

- **[writing]** In Section 3.2 (e000), the phrase ‘Scores complexity c [0,1]’ lacks a clear subject. It should specify what entity performs the scoring (e.g., ‘The system scores complexity...’).
- **[writing]** Table 1 (e000) and Table 1 (e002) present conflicting feature sets (e.g., ‘Result verification’ and ‘Sandbox security’ are missing in the first version). Ensure the final manuscript uses a single, consistent comparison table.
- **[writing]** The Appendix ‘Writing-Quality Audit’ (e001) contains meta-commentary about the paper’s own quality (e.g., ‘Duplicated figure file’, ‘Bracket-style pseudo-citations’). This section should be removed or rewritten to avoid confusing the reader about the paper’s actual state.