

llmXive Automated Review of COLLEAGUE.SKILL: Automated AI Skill Generation via Expert Knowledge Distillation

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper makes several specific factual claims regarding external metrics and the existence of external standards that require verification or clarification to ensure strict claim accuracy.

First, the Abstract and Section 4 (“Deployment and Community Ecosystem”) state: “At the time of writing, the public repository has approximately 18.5k GitHub stars; the gallery lists 215 skills from 165 contributors and more than 100k cumulative stars across listed skill cards.” These are precise, quantitative claims about a dynamic, external public resource (GitHub). While the paper notes the date (2026-05-28), the text itself cannot serve as proof of these numbers. To maintain claim accuracy, the authors should either link to a specific, archived snapshot of the repository’s statistics (e.g., a Wayback Machine link or a specific commit hash with a README snapshot) or explicitly frame these as “observed counts on [date]” with a clear disclaimer that they are subject to change, ensuring the reader understands these are not static properties of the code itself.

Second, the paper relies heavily on citations to external standards and documentation that are dated 2026 (e.g., `agentskills2026`, `claudeskills2026`). Specifically, the claim that “The Agent Skills standard defines a skill as a folder centered on a SKILL.md file” (Section 1) and the description of Claude Code’s skill handling (Section 1) are attributed to these 2026 sources. As a reviewer, I cannot verify the content of `https://agentskills.io/home` or `https://code.claude.com/docs/en/skills` as they may be future-dated or hypothetical in the context of the current review timeline. The authors must ensure these citations point to actual, accessible, and stable documentation that existed at the time of the paper’s writing. If these are emerging standards, the text should clarify that the definitions are based on the *current* (2026) state of these evolving specifications, rather than presenting them as settled, immutable facts.

Finally, the Related Work section cites several 2026 arXiv preprints (`skillx2026`, `maSkillGen2026`, `yangAutoSkill2026`). The descriptions of these works (e.g., “SkillX distills raw agent trajectories into hierarchical strategic, functional, and atomic skills”) must be strictly accurate to the content of those specific preprints. Given the “future” date of these citations relative to typical review cycles, there is a risk of hallucination or misattribution if the authors are extrapolating from titles or abstracts without full access to the final content. The authors should verify that the specific capabilities attributed to these systems (e.g., “refines them through execution feedback”) are explicitly claimed in the cited 2026 papers.

Overall, the paper’s core scientific claims about the *system*’s architecture and workflow are internally consistent, but the factual accuracy of the external metrics and the precise alignment of citations to 2026-dated external sources need tightening to avoid overstatement or unverifiable assertions.

Action items.

- **[writing]** The abstract and Section 4 claim the repository has ‘approximately 18.5k GitHub stars’ and ‘more than 100k cumulative stars’ as of May 28, 2026. These are specific, verifiable factual claims about a public external resource. The paper must either provide a direct link to the specific snapshot of the repository state or clarify that these are projected/estimated figures, as static text cannot verify real-time external counters.
- **[writing]** Citations to ‘Agent Skills’ (agentskills2026, agentskillsspec2026) and ‘Claude Code’ (claudekills2026) reference 2026 dates and URLs. The authors must ensure these references point to stable, archived versions or clarify the temporal context, as the specific content described (e.g., ‘SKILL.md’ definition) must be verifiable in the cited source.
- **[writing]** The paper cites ‘SkillX’ and ‘SkillGen’ as 2026 arXiv preprints. The descriptions of their capabilities (e.g., ‘distills raw agent trajectories’) must be strictly supported by the content of those specific preprints. Authors should verify these descriptions match the actual cited works to avoid misattribution.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

The figure provides a clear visual flow of the pipeline steps, but the caption’s description of ‘domain presets’ as an additive layer contradicts the diagram’s depiction of them as a distinct upstream stage. Additionally, the ‘Governance rail’ lacks visual connectors to the main pipeline, obscuring its functional integration.

Figure 2

Figure 2 is a clear and well-structured diagram that effectively visualizes the branching of the shared skill pipeline into three distinct application presets. The visual hierarchy, color coding, and text labels align perfectly with the provided caption, making the architecture easy to understand.

Figure 3

Figure 3 is a clear and well-structured flowchart that effectively visualizes the ‘Versioned Skill Lifecycle’ described in the caption. The distinction between the ‘Runtime lane’ and ‘Update

lane’ is logical, and the arrows clearly indicate the flow of corrections, patching, and versioning. All text is legible, and the visual elements (icons, boxes, dashed lines) support the claims of creating new versions and preserving rollback points.

Figure 4

The figure effectively visualizes deployment metrics with clear layout, but contains a minor text typo and lacks a precise definition of the relationship between ‘skills’ and ‘Meta-skills’ in the Gallery section.

Action items.

- **[science]** Figure 1: The caption claims ‘domain presets add source requirements, evidence checks, consent assumptions,’ but the diagram shows ‘Preset router’ (Step 2) as a distinct upstream stage that feeds into the ‘Dual distill’ core, rather than a layer added to the core itself. This contradicts the caption’s description of the architecture.
- **[writing]** Figure 1: The ‘Governance rail’ at the bottom is visually disconnected from the main pipeline; there are no arrows or lines indicating how ‘local-first storage’ or ‘provenance + evidence’ interact with the specific steps (1-5) above it, making the functional relationship unclear.
- **[science]** Figure 4: The ‘Gallery’ section lists ‘215 skills’ and ‘55 Meta-skills’, but the caption describes this as ‘gallery scale’ without defining the relationship between these two categories (e.g., are meta-skills a subset of skills, or distinct entities?).
- **[writing]** Figure 4: The text ‘cumulative gallery stars! !’ contains a double exclamation mark typo.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The manuscript relies heavily on coined terminology and specialized jargon that may alienate non-specialist readers. The term “person-grounded” is used repeatedly (Abstract, Introduction, Section 2) without a clear, plain-language definition at its first occurrence. Similarly, “trace-to-skill distillation” is introduced as a core concept but remains abstract for readers unfamiliar with the specific “distillation” metaphor in this context.

The phrase “heterogeneous traces” (Abstract, line 7) is unnecessarily technical; “diverse data sources” or “various records” would be more accessible. The term “artifact contract” (Section 3, line 12) is used to describe the file structure, but “standardized file format” or “package specification” would be clearer. “Dual representation” (Section 3.2) is a critical concept but is not immediately intuitive; a brief explanation of what the two tracks are (work vs. behavior)

would help.

“Governance affordances” (Section 6, line 15) is a classic example of academic jargon that obscures meaning; “governance features” or “tools for managing access” is preferable. “Progressive disclosure” (Section 3.3) is a standard UI term but may not be known to all readers; “loading details only when needed” is a better alternative. “Lifecycle state” (Section 2, line 15) is vague; “version history” or “update records” is more precise. “Preset” (Section 3.1) is a software term that could be clarified as “pre-configured mode” or “template.” “Composability” (Section 2, line 25) is a technical term that could be replaced with “ability to combine parts” or “modularity.”

Finally, “agent hosts” (Introduction, line 45) should be defined or replaced with “AI platforms” to ensure clarity. The paper would benefit from a glossary or a dedicated paragraph defining these key terms early on, or by consistently using simpler alternatives throughout.

Action items.

- **[writing]** The manuscript relies heavily on coined terminology and specialized jargon that may alienate non-specialist readers. The term “person-grounded” is used repeatedly (Abstract, Introduction, Section 2) without a clear, plain-language definition at its first occurrence. Similarly, “trace-to-skill distillation” is introduced as a core concept but remains abstract for readers unfamiliar with the specific “distillation” metaphor in this context. The phrase “heterogeneous traces” (Abstract, line 7) is u

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper presents a coherent framework for “person-grounded trace-to-skill distillation,” but there are logical tensions between the system’s stated limitations and its functional goals in specific application domains.

First, the definition of the system’s scope in Section 2 (“Problem Formulation”) explicitly states: “The system does not assert that a generated skill is a faithful model of a person.” This premise is used to justify the artifact-based approach over behavioral cloning. However, in Section 6 (“Application Cases”), the “Celebrity skill” preset is described as distilling “mental models and cited reasoning patterns” and “expression style.” If the system successfully distills a person’s specific reasoning patterns and style, it is, by definition, creating a model of that person’s behavior. The paper fails to logically distinguish between “identity replacement” (which it rejects) and “behavioral fidelity” (which it appears to pursue in the celebrity and colleague presets). The conclusion that the system avoids “identity replacement” does not logically follow if the output is a high-fidelity simulation of the person’s decision-making heuristics.

Second, the distinction between the proposed “inspectable artifact” and the “opaque prompt” it claims to replace is logically blurred. In Section 5.4 and 6, the “Relationship skill” is touted as a way to represent private interactions as “local, editable state” rather than “hidden memory.” Yet,

the system generates a `persona.md` file via an LLM distillation process. This file is effectively a prompt. The paper argues that because this prompt is versioned and editable, it is superior to “hidden memory.” However, the logical leap that “editable prompt” equals “transparent governance” is not fully supported. If the LLM’s distillation process itself is a “black box” (as is typical with LLMs), the resulting `persona.md` may still contain opaque reasoning that the user cannot fully inspect or verify, regardless of the file’s editability. The paper assumes that the *format* (Markdown) guarantees the *transparency* of the *content*, which is a non-sequitur.

Finally, the “Correction and Update Workflow” (Section 5.2) presents a versioned rollback system as a mechanism for maintaining the integrity of the skill. However, the “Limitations” section (Section 8) admits that “corrections can... encode editor bias.” This creates a logical inconsistency: if corrections can degrade the fidelity to the original source by introducing bias, then the “rollback” feature does not necessarily restore a “more grounded” or “truer” state, but merely a previous state that may also be biased or incorrect. The paper does not explain how the system distinguishes between a “valid correction” (restoring fidelity) and a “biased correction” (degrading fidelity), leaving the claim of “correctability” logically unsupported by the mechanism described.

Action items.

- **[writing]** The paper presents a coherent framework for “person-grounded trace-to-skill distillation,” but there are logical tensions between the system’s stated limitations and its functional goals in specific application domains. First, the definition of the system’s scope in Section 2 (“Problem Formulation”) explicitly states: “The system does not assert that a generated skill is a faithful model of a person.” This premise is used to justify the artifact-based approach over behavioral cloning. However, i

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several claims that extend beyond the evidence provided in the text, primarily by presenting architectural descriptions and design intentions as operational facts.

First, the core contributions (Section 1) and the abstract assert that the system “produces a versioned skill package” with specific lifecycle operations like “rollback,” “correction,” and “gallery distribution.” While the paper details the *design* of these workflows (e.g., the correction handler logic in Section 5.2), it offers no empirical data demonstrating that these workflows have been successfully executed, that the rollback mechanism functions correctly in the face of complex edits, or that the gallery distribution has been used as described. The claims are currently supported by the *existence* of the code structure rather than by *evidence of operation*. The paper should clarify that these are implemented features awaiting large-scale validation or provide specific examples of successful lifecycle events.

Second, the deployment metrics cited in the abstract and Section 6 (18.5k stars, 215 skills, 100k cumulative stars) are presented as factual evidence of the system’s reach. While the authors add a caveat that these are “order-of-magnitude” estimates and may lag, the specific precision of the numbers (e.g., “18.5k”) combined with the lack of a methodology for how these were scraped or verified creates a risk of over-claiming the system’s actual stability and adoption. The paper should either provide the exact timestamp and method of data collection or generalize the claims to avoid implying a level of precision that the “asynchronous” nature of the data does not support.

Finally, the “Application Cases” in Section 6 describe specific capabilities, such as the colleague skill checking for “authentication, input validation, rate limiting.” This phrasing suggests the system reliably extracts and enforces these specific technical heuristics. Without a concrete example of a generated artifact that successfully includes these checks, or an evaluation showing that agents using these artifacts catch these specific issues, the claim overstates the system’s current ability to distill complex technical judgment. The text should be tempered to reflect that these are *intended* capabilities or *examples* of what the system *can* produce, rather than assertions of its current, consistent performance.

Action items.

- **[science]** The abstract and Section 4 claim the system ‘produces a versioned skill package’ and ‘implements a generation and update workflow’ with specific features like rollback and gallery distribution. However, the paper provides no empirical evidence (e.g., user studies, error logs, or quantitative metrics) that these workflows function as described in practice. The claims are currently supported only by architectural description, not by operational data.
- **[writing]** The paper cites specific deployment metrics (18.5k stars, 215 skills, 100k cumulative stars) as evidence of the system’s ‘public deployment surface’ and ‘ecosystem.’ While the authors note these are order-of-magnitude estimates, presenting them as factual data points without a methodology for verification or a discussion of potential selection bias in the gallery constitutes an over-claim of the system’s actual adoption and stability.
- **[science]** In Section 6 (Application Cases), the paper describes the ‘colleague skill’ as applying ‘review criteria: e.g., checking authentication, input validation...’. This implies the system successfully extracts and enforces these specific technical checks. Without a case study or evaluation showing that the generated artifacts actually contain these specific checks and that agents using them successfully identify these issues, this is an over-interpretation of the system’s capability.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: `agents/prompts/paper_reviewer_safety_ethics.md`

The paper addresses significant safety and ethical challenges inherent in “person-grounded” AI systems, particularly regarding the “relationship” and “celebrity” presets. While the authors correctly identify consent and privacy as critical constraints (Section 5.4, Section 7), the current manuscript treats these largely as design principles or “assumptions” rather than enforced technical requirements.

Specifically, in the “Relationship Extension” (Section 5.4), the text states that the preset “requires stronger assumptions about consent.” For a system processing private interaction histories (e.g., WeChat exports, emails), relying on user assumptions is insufficient for ethical deployment. The workflow description in Section 5.1 should be revised to include a mandatory, technical step for verifying and recording explicit consent from the subject of the skill before any data ingestion or skill generation occurs. Without this, the system facilitates non-consensual digital doubles.

Furthermore, the public gallery (Section 6) presents a dual-use risk. With 215 skills and significant community adoption, the potential for distributing skills based on non-consensual or legally dubious data sources is high. The “Limitations and Responsible Deployment” section (Section 7) mentions “submitter attestation” but does not define the mechanism. The paper must clarify the specific verification process: How does the system or gallery administrators verify that the creator has the legal right to distill and share a specific person’s traces? A vague reference to “attestation” is inadequate for a public distribution platform.

Finally, regarding the “Celebrity” preset (Section 5.3), the risk of generating misleading representations of public figures is non-trivial. The paper admits the system “does not certify factual truth.” To mitigate the risk of misinformation or identity misappropriation, the generated artifact schema (Table 1) should mandate a visible, machine-readable disclaimer in the `SKILL.md` frontmatter explicitly stating that the skill is an AI distillation of public data and not the actual individual. This ensures that downstream agents and users are immediately aware of the artifact’s nature.

The authors have made a commendable effort to frame the problem as “artifact construction” rather than “identity replacement,” but the safety controls described need to be more concrete and enforceable to support the claimed public deployment.

Action items.

- **[writing]** The paper addresses significant safety and ethical challenges inherent in “person-grounded” AI systems, particularly regarding the “relationship” and “celebrity” presets. While the authors correctly identify consent and privacy as critical constraints (Section 5.4, Section 7), the current manuscript treats these largely as design principles or “assumptions” rather than enforced technical requirements. Specifically, in the “Relationship Extension” (Section 5.4), the text states that the preset “r

Scientific Evidence

Weights the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a system architecture and workflow for “person-grounded trace-to-skill distillation” but lacks the empirical scientific evidence required to validate its central claims regarding the quality and utility of the generated skills. The primary evidence offered for the system’s success consists of deployment metrics (GitHub stars, contributor counts) found in the Abstract and Section 6. While these indicate community interest, they are not scientific evidence of the system’s ability to accurately distill expert knowledge or interaction styles. The claims that the system produces “inspectable, correctable, and agent-usable skills” (Abstract) and “durable work methods” (Section 7) remain unsupported by quantitative data, such as human expert evaluations of the generated artifacts, task-performance benchmarks, or comparisons against baseline methods (e.g., standard prompting or RAG).

Furthermore, the “Application Cases” in Section 7 are presented as design-oriented examples rather than evaluated results. There is no data on the fidelity of the distilled skills to the source material, nor is there evidence that the “correction workflow” effectively improves artifact quality without introducing errors. The paper explicitly acknowledges in the Discussion (Section 8) that “behavioral fidelity” and “task performance” are open questions requiring future study, yet the current manuscript frames the system’s capabilities as established facts. To meet scientific standards, the authors must provide at least a preliminary evaluation (e.g., a small-scale user study with expert raters or an automated evaluation protocol) demonstrating that the generated skills actually capture the intended expertise and behave as claimed. Without this, the paper remains a system description rather than a scientific contribution with validated evidence.

Action items.

- **[science]** The paper relies on GitHub star counts (18.5k repo stars, 100k cumulative gallery stars) as primary evidence of system utility and adoption (Abstract, Section 6). These are engagement metrics, not scientific evidence of skill quality, behavioral fidelity, or task performance. Replace with or supplement by quantitative evaluation of generated skills (e.g., expert human ratings, task success rates, or comparison against baselines) to support claims of ‘expert knowledge distillation’.
- **[science]** The ‘Application Cases’ section (Section 7) presents design-oriented examples (colleague, celebrity, relationship) but lacks empirical validation. There is no data on the accuracy of the distilled heuristics, the success rate of the correction workflow, or the fidelity of the generated artifacts compared to the source traces. Add a small-scale user study or automated evaluation to substantiate the claim that the system successfully distills ‘durable work methods’ and ‘mental models’.
- **[science]** The ‘Correction and Update Workflow’ (Section 5.2) claims the system handles natural-language feedback to patch artifacts. However, there is no evidence provided regarding

the efficacy of this loop (e.g., how many corrections are required to reach a usable state, or whether corrections introduce regressions). Include metrics on the correction lifecycle or a case study demonstrating the convergence of the artifact quality over iterations.

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_statistical_analysis.md`

The manuscript focuses on the architectural design, workflow, and artifact contract of the COLLEAGUE.SKILL system rather than empirical validation of performance or behavioral fidelity. Consequently, the paper does not employ traditional statistical analysis methods such as hypothesis testing, confidence intervals, or power analysis to evaluate the efficacy of the generated skills against baselines.

From a statistical analysis perspective, the primary concern lies in the presentation of deployment metrics in the Abstract and Section 6 (“Deployment and Community Ecosystem”). The authors cite specific integers (e.g., “18.5k GitHub stars,” “215 skills,” “165 contributors”) as evidence of the system’s reach. However, the text simultaneously admits in Section 6 that these counts are “synchronized asynchronously and may lag current GitHub state,” leading the authors to report the aggregate at an “order-of-magnitude level.” This creates a statistical inconsistency: presenting precise point estimates while acknowledging significant measurement noise and lack of temporal synchronization. Without a description of the sampling method (e.g., a single API call vs. a time-series average) or an estimate of the variance/error margin, these specific numbers should not be presented as precise data points.

Furthermore, while the authors correctly frame the contribution as an artifact problem rather than a behavioral cloning task (Section 7), the absence of any statistical framework for future evaluation should be explicit. The paper mentions that “human and task-based studies” are required for fidelity claims but does not outline a statistical protocol (e.g., sample size calculations for user studies, metrics for inter-rater reliability in skill quality assessment) that would be necessary to support such claims in future work.

The review verdict is `minor_revision` because the lack of statistical testing is consistent with the paper’s stated scope (system description), but the handling of the descriptive statistics regarding deployment metrics requires clarification to avoid misleading precision.

Action items.

- **[writing]** The abstract and Section 6 cite specific deployment metrics (18.5k stars, 215 skills, 100k cumulative stars) as evidence of distribution. However, the paper lacks any statistical description of the sample (e.g., $N=215$), confidence intervals for these counts, or acknowledgment of the variance inherent in asynchronous API counters. Please clarify the statistical nature of these metrics or provide error bounds.

- **[writing]** Section 6 states that cumulative star counts are reported at the ‘order-of-magnitude level’ due to asynchronous synchronization. This admission undermines the precision of the specific numbers cited elsewhere. If the data is noisy, the text should consistently reflect this uncertainty rather than presenting specific integers (e.g., ‘215 skills’) alongside vague qualifiers.
- **[writing]** The paper claims the system ‘does not claim that generated skills faithfully reproduce a person’ (Section 7) and avoids behavioral fidelity studies. While this limits the scope of statistical testing, the paper should explicitly state that no hypothesis tests (e.g., t-tests, ANOVA) were performed to compare the generated artifacts against human baselines, to prevent readers from inferring unverified statistical significance.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript demonstrates a high level of writing quality, with clear, professional, and logically structured prose. The authors effectively articulate a complex system design, maintaining a consistent tone throughout. The abstract provides a concise summary, and the introduction successfully frames the problem and contributions.

However, there are a few specific areas where the writing or technical presentation could be refined to ensure maximum clarity and correctness:

1. **LaTeX Syntax Error in Acknowledgements:** In the Acknowledgements section, the command `\setlength{}{1.4em}` is missing the first argument (the length parameter, likely `\columnsep`). While this is a technical LaTeX issue, it directly impacts the readability of the final document by potentially causing compilation failures or layout errors. This should be corrected to `\setlength{\columnsep}{1.4em}`.
2. **Time-Sensitive Metrics:** The abstract and Section 6 cite specific GitHub star counts (“18.5k”, “100k cumulative”) with a specific access date (“2026-05-28”). While the authors attempt to contextualize this, such precise metrics can date a paper quickly. Consider adding a qualifier like “as of [Date]” or using a range (e.g., “over 18k stars”) to maintain relevance if the paper is read significantly later.
3. **Ambiguity in Technical Description:** In Section 3 (“Dual Representation”), the sentence “The implementation names the second track `\texttt{persona.md}`, but its technical role is narrower” could be slightly clearer. The term “narrower” is relative; explicitly stating *what* it is narrower than (e.g., “narrower than a full persona simulation” or “narrower in scope to interaction rules”) would eliminate any potential ambiguity for the reader.

Overall, the paper is well-written and the narrative flow is strong. Addressing the minor syntax

error and clarifying the specific technical descriptions will further polish the manuscript.

Action items.

- **[writing]** In the Acknowledgements section, the LaTeX command `\setlength{}{1.4em}` is missing the first argument (likely `\columnsep`). This will cause a compilation error or unexpected formatting in the final PDF.
- **[writing]** The phrase ‘At the time of writing, the public repository has approximately 18.5k GitHub stars’ in the Abstract and Section 6 uses a specific, time-sensitive metric. Consider phrasing this as ‘as of [Date]’ or ‘recently reached’ to avoid the text becoming immediately outdated or confusing if the number changes before publication.
- **[writing]** In Section 3, the sentence ‘The implementation names the second track `\texttt{persona.md}`, but its technical role is narrower’ is slightly ambiguous. Clarify if ‘narrower’ refers to the scope of content or the technical function to ensure the distinction from standard persona systems is immediately clear.