

llmXive Automated Review of GateMem: Benchmarking Memory Governance in Multi-Principal Shared-Memory Agents

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so errors, misreadings, and inaccuracies are likely. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is advisory, automated feedback for the authors. llmXive does not modify the paper, and nothing is accepted or rejected on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (8 lenses): Claim Accuracy, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The review focuses on the accuracy of factual claims and the validity of citations within the manuscript.

Citation and Source Accuracy The bibliography contains a critical mismatch. Entry c-001 in the provided metadata points to a HuggingFace JSON file regarding GGUF statistics (`.../gguf-public-stats/...`). This source does not support the claims made in the text regarding LLM backbone capabilities or the specific models evaluated. The footnote in Section 4 (Experimental Setup) lists URLs for the models, but the bibliography entry c-001 appears to be a placeholder or ingestion error that fails to link to the actual model documentation or papers. This must be corrected to ensure the claims about model availability and capabilities are verifiable.

Factual Claims Regarding Model Existence The manuscript makes specific performance claims (Table 4, Section 4) for models named GPT-5.4, Deepseek-V4-Pro, Llama-4-Maverick, and Gemini-2.5-Flash-Lite. As of the current public knowledge cutoff, these specific model versions do not exist.

- GPT-5 has not been released, let alone a 5.4 variant.
- Llama-4 has not been released.
- Deepseek-V4 and Gemini-2.5 are not current public releases.

If these are hypothetical models used for simulation, the paper must explicitly state that the results are simulated or projected, as the current text presents them as empirical results from “evaluated backbones.” If these are real models with different names (e.g., GPT-4o, Llama-3), the names in the text and tables are factually incorrect. Presenting results for non-existent models as empirical data constitutes a fundamental factual error that undermines the validity of the “Main Results” section.

Metric Definitions The definition of the Memory Governance Score (MGS) in Section 3.3 (Equation 8) is mathematically sound and consistent with the text description. The claim that “a system cannot obtain a high overall score merely by being highly useful if it leaks” is accurately supported by the multiplicative formula $U \cdot (1 - A) \cdot (1 - F)$. No factual errors were found in the metric definitions or the logical derivation of the score.

Conclusion While the methodological framework is sound, the paper relies on factual claims regarding model versions that do not exist and contains a broken bibliography entry. These issues prevent the verification of the core experimental results.

Action items.

- [writing] The bibliography entry ‘c-001’ links to a JSON file of GGUF stats, which is unrelated to the cited claim about LLM backbones. Verify and replace with the correct official model documentation URLs for GPT-5.4, Deepseek-V4-Pro, Llama-4-Maverick, and Gemini-2.5-Flash-Lite.
- [science] The paper cites ‘GPT-5.4’, ‘Deepseek-V4-Pro’, and ‘Llama-4-Maverick’ as evaluated backbones. These model versions do not currently exist in public release. If these

are hypothetical or future-dated models, the claims regarding their specific performance metrics (Table 4) are factually unsupported by real-world evidence. Clarify the provenance of these models or correct the model names to existing versions.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The manuscript exhibits a tendency to overuse specialized terminology that, while precise for the authors, creates unnecessary friction for a broader audience. The term “principals” is used repeatedly (Abstract, Introduction, Section 3) to denote users or entities. In the context of LLM agents, “users” or “requesters” is more accessible; “principals” is a security-specific term that should be defined immediately if retained.

The word “governance” is applied liberally (Abstract, Introduction, Section 3, Conclusion) often as a synonym for “management” or “control.” In the Abstract, “memory quality requires governance” could be clearer as “memory quality requires strict management.” The phrase “memory governance problem” in the Introduction is vague; “memory access control problem” is more specific.

In Section 3, “hidden checkpoints” is introduced without a clear definition of what makes them “hidden” to the agent versus the evaluator. This distinction is crucial but obscured by the jargon. Similarly, “leak-target annotations” in the Abstract is a dense compound noun; “annotations identifying sensitive data to be protected” is more transparent.

Finally, “active forgetting” is a key contribution but is introduced as a category name before its mechanism (interface-level non-recovery) is fully unpacked. Defining it explicitly as “the agent’s ability to refuse to recover deleted information” upon first mention would improve clarity. The paper would benefit from replacing these dense terms with plainer language or providing immediate, simple definitions.

Action items.

- [writing] The manuscript exhibits a tendency to overuse specialized terminology that, while precise for the authors, creates unnecessary friction for a broader audience. The term “principals” is used repeatedly (Abstract, Introduction, Section 3) to denote users or entities. In the context of LLM agents, “users” or “requesters” is more accessible; “principals” is a security-specific term that should be defined immediately if retained. The word “governance” is applied liberally (Abstract, Introduction, Sec

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper presents a logically sound framework for evaluating memory governance, with clear definitions for Utility (U), Access Control (A), and Active Forgetting (F). The multiplicative Memory Governance Score (MGS) is a coherent mechanism to enforce the requirement that a system must perform well on all three dimensions to be considered successful. The experimental results in Table 4.1 consistently support the conclusion that no single method dominates across all dimensions, and the trade-offs observed (e.g., Long-Context having high U but non-negligible A/F) are logically derived from the data.

However, there is a minor logical tension in the interpretation of the MGS metric versus the qualitative analysis. The multiplicative formula $(U * (1-A) * (1-F))$ implies that a high violation rate in A or F drastically reduces the overall score, even if U is near-perfect. The text claims that “Long-context prompting often yields the best governance score,” which is numerically true, but the qualitative discussion in Section 4.2 suggests that these scores are still “far from reliable.” The logical consistency of the paper would be strengthened by explicitly discussing whether the MGS threshold for “reliable” is met by the best-performing models (e.g., is 80% MGS in the Medical domain considered “reliable” or just “best available”?). The current text leaves this threshold ambiguous, which slightly weakens the logical link between the metric and the final conclusion about deployment readiness.

Additionally, the definition of Active Forgetting Failure (F) in Section 3.2 states that any action other than `no_memory` is a failure. Yet, in the qualitative case study (Table 4.3), a `refuse` action is labeled as a “mismatch” (`warnresp`) rather than a full failure. This creates a slight inconsistency between the quantitative metric definition and the qualitative evaluation of specific cases. Clarifying whether `refuse` is treated as a partial or full failure in the MGS calculation would resolve this logical gap.

Action items.

- [writing] The paper presents a logically sound framework for evaluating memory governance, with clear definitions for Utility (U), Access Control (A), and Active Forgetting (F). The multiplicative Memory Governance Score (MGS) is a coherent mechanism to enforce the requirement that a system must perform well on all three dimensions to be considered successful. The experimental results in Table 4.1 consistently support the conclusion that no single method dominates across all dimensions, and the trade-offs

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several strong claims regarding the state of memory governance in multi-principal agents, particularly in the Abstract and Conclusion. While the experimental data supports the existence of a trade-off between utility and safety, the language occasionally extrapolates beyond what the specific benchmark results justify.

First, the Abstract states: “Across diverse baselines and backbone models, no method simultaneously achieves strong utility, robust access control, and reliable forgetting.” This is a categorical claim that is slightly undermined by the data in Table 1. Specifically, the Long-Context baseline with GPT-5.4 achieves a Memory Governance Score (MGS) of 80.1% in the Medical domain, with 91.4% Utility and only 10.4% Access Control violations. If the authors define “strong” as >90% utility and <15% violation, this method *does* achieve the triad in that specific domain. The claim would be more accurate if qualified as “no method consistently achieves... across all domains” or if the threshold for “strong” is explicitly defined. Without this, the statement risks overgeneralizing a domain-specific success as a universal failure.

Second, the Conclusion and Abstract assert that “current memory agents remain far from reliable shared institutional deployment.” This is a significant extrapolation. The benchmark evaluates agents on synthetic, multi-turn episodes with specific attack vectors (e.g., cross-patient confusion, indirect inference). While these are realistic, they do not encompass the full complexity of “institutional deployment,” which often includes human-in-the-loop oversight, layered security architectures, and domain-specific fine-tuning not present in the evaluated baselines. The paper demonstrates that *current baseline architectures* struggle with the specific governance tasks defined in GateMem, but claiming they are unfit for *any* institutional deployment overreaches the scope of the evidence. The authors should temper this to reflect that “current baseline approaches face significant challenges in reliable shared-memory governance as defined by GateMem.”

Finally, the claim that “Long-context prompting often yields the best governance score” is well-supported by the data (Table 1), but the subsequent inference that this is the “strongest governance trade-off” (Section 4.2) ignores the efficiency metrics in Table 2. Long-Context is significantly more token-intensive (e.g., 7.61k tokens/ckpt in Office vs. ~1.8k for RAG). While the paper notes the cost, the narrative of “best trade-off” leans heavily on the MGS metric without sufficiently weighing the prohibitive cost for large-scale deployment. The overreach here is subtle; presenting a high-cost solution as the “best” trade-off without a clear cost-benefit analysis or a definition of “trade-off” that includes economic feasibility.

These issues are primarily matters of phrasing and scope definition rather than fundamental scientific flaws. The data supports the existence of the problem; the text just needs to be more precise about the magnitude of the failure and the scope of the conclusion.

Action items.

- [writing] The claim that “no method simultaneously achieves strong utility, robust access control, and reliable forgetting” (Abstract) is contradicted by Long-Context (GPT-5.4)

achieving 91.4% Utility and 80.1% MGS in Medical. Define “strong” or qualify the claim to “across all domains” to avoid overgeneralization.

- [writing] The conclusion that “current memory agents remain far from reliable shared institutional deployment” (Abstract) overreaches. The benchmark tests synthetic episodes; generalizing to all institutional deployments ignores potential human-in-the-loop safeguards or hybrid architectures not evaluated here.
- [writing] Describing Long-Context as the “best governance trade-off” (Abstract) ignores its high token cost (Table 2). The claim overreaches by prioritizing MGS without explicitly weighing the prohibitive economic cost for large-scale deployment in the trade-off definition.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

The manuscript addresses a critical safety gap: the governance of shared memory in multi-principal LLM agents. The focus on access control and active forgetting is highly relevant to preventing privacy violations in institutional settings. However, several ethical and safety documentation gaps must be addressed before the benchmark can be safely adopted or the results fully trusted.

First, the dataset construction involves generating “realistic” episodes containing sensitive information such as medical diagnoses (e.g., STI charts, Hepatitis C), household access codes, and financial details. While the authors state the data is generated, there is no explicit statement regarding Institutional Review Board (IRB) approval or the specific ethical framework used to generate this synthetic data. Given the sensitivity of the topics, the authors must clarify whether the generation process involved human annotators who were exposed to these sensitive prompts, and if so, what protections were in place. If the data is purely synthetic, a statement confirming that no real-world data was used and that the content is entirely fabricated to avoid representing actual individuals is required.

Second, the evaluation methodology relies heavily on an LLM-as-a-judge to detect “leaks” of protected information (Section 3.3, Appendix A3). There is a non-trivial safety risk that the judge model, when prompted with “leak_targets” to check for their presence, might inadvertently reproduce or “hallucinate” these targets in its own output (the `notes` field), effectively leaking the sensitive data it is supposed to be auditing. The prompt template in Appendix A3 includes a “CRITICAL SAFETY FOR THE JUDGE” instruction, but the authors should explicitly discuss the results of any self-consistency checks or red-teaming performed on the judge itself to ensure it does not become a vector for data leakage during the evaluation process.

Finally, the benchmark includes “social engineering” and “authority pressure” attack vectors (Appendix A1). While these are necessary for a robust evaluation, the paper should explicitly state that these prompts are strictly confined to the offline benchmarking environment. There must be a clear disclaimer that these prompts are not intended for, and should not be used in, live adversarial testing against deployed agents without further safety alignment, to prevent the accidental release of harmful prompt engineering techniques into the wild.

Addressing these points will ensure the benchmark is ethically sound and that the evaluation process itself does not introduce new safety risks.

Action items.

- [writing] The dataset contains sensitive simulated PII (medical diagnoses, household codes). Explicitly detail the IRB/ethics approval status or the specific de-identification/synthetic generation protocols used to ensure no real individuals are represented, as required for human-subject-adjacent benchmarks.
- [writing] The ‘Active Forgetting’ evaluation relies on LLM judges to detect ‘leaks’ of deleted data. Add a specific safety warning in the appendix regarding the risk of the judge model itself hallucinating or inadvertently reproducing the ‘leak_targets’ it is supposed to audit, potentially creating a feedback loop of exposure.
- [writing] The benchmark includes ‘social engineering’ and ‘authority pressure’ attack vectors (Appendix A1). Clarify that these are strictly simulated within the dataset generation pipeline and that the evaluation protocol does not involve real-world adversarial testing against live agents to prevent accidental deployment of harmful prompts.

Scientific Evidence

Weights the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a comprehensive benchmark for memory governance, but the scientific evidence supporting the central claims regarding model performance requires clarification on reproducibility and statistical robustness.

First, the experimental results in Table 1 rely on backbone models such as “GPT-5.4”, “Deepseek-V4-Pro”, and “Llama-4-Maverick”. These identifiers do not correspond to currently public or stable model releases. If these are hypothetical or future-dated names, the empirical claims regarding the superiority of specific backbones (e.g., Deepseek-V4-Pro vs. GPT-5.4) are currently unverifiable. For the evidence to be robust, the authors must either provide the specific model weights/API keys used, or clarify if these are aliases for existing models (e.g., GPT-4o, Llama-3). Without this, the comparative analysis of backbone capabilities remains speculative.

Second, the primary evaluation metric (MGS) is entirely dependent on an LLM-as-a-judge (GPT-4o). While the authors provide a human validation study in Appendix A4, the sample size is approximately 290 cases (roughly 13% of the total 2,218 checkpoints). In security and governance tasks, false negatives in leakage detection are critical. A 13% sample may not be sufficient to detect systematic biases in the judge’s evaluation of “soft” leaks or indirect inferences. The authors should expand the human validation to a larger, stratified random sample (e.g., 10-20% of the full dataset) or provide a sensitivity analysis demonstrating how potential judge errors would alter the ranking of the baselines.

Finally, the “Active Forgetting” metric is sensitive to the specific phrasing of the recovery query. The paper reports aggregate failure rates but does not break down the variance of these rates across different attack types (e.g., direct confirmation vs. social engineering) for the same deleted fact. If a model fails only on specific prompt templates but succeeds on others, the aggregate score may misrepresent its true robustness. A more granular analysis of failure modes per attack type would strengthen the evidence that current agents are fundamentally brittle in this domain.

Action items.

- [science] Table 1 reports results for ‘GPT-5.4’, ‘Deepseek-V4-Pro’, and ‘Llama-4-Maverick’. These model names appear hypothetical or future-dated. Without access to these specific artifacts or clarification that they are aliases for existing models, the core empirical claims regarding backbone performance cannot be independently verified or replicated.
- [science] The study relies on an LLM-as-a-judge for all primary metrics. The human validation sample (Appendix A4) covers only ~13% of the 2,218 checkpoints. Given the high stakes of leakage detection, a larger stratified human audit or a sensitivity analysis showing how judge variance impacts baseline rankings is required to rule out systematic bias.
- [science] The ‘Active Forgetting’ metric depends heavily on the specific phrasing of recovery queries. The paper does not report variance in failure rates across different attack phrasings for the same deleted fact. Without this, it is unclear if reported rates reflect true model robustness or sensitivity to specific prompt templates.

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The statistical analysis framework in GateMem is conceptually sound, defining clear metrics for Utility (U), Access Control (A), and Active Forgetting (F). The use of an LLM-as-a-judge is justified by the human validation study in Appendix A4, which reports high

agreement (Cohen's $\kappa > 0.93$). However, the presentation of results lacks necessary statistical rigor to support the strong comparative claims made in the main text.

First, Table 1 presents point estimates for all metrics across 42 experimental conditions (7 methods $\times$ 6 backbones) without any measure of uncertainty. Given that the evaluation involves 2,218 checkpoints, the variance of these estimates should be reported. For instance, a difference of 1.5% in MGS between two methods might be statistically insignificant if the standard error is high. The authors should report 95% confidence intervals (e.g., via bootstrapping over episodes or checkpoints) or standard errors for the primary metrics. This is particularly critical for the “Efficiency Trade-offs” section, where latency and token counts are averaged without variance.

Second, the human validation results in Table A4 state a “maximum absolute difference of 1.04 percentage points” but do not specify the statistical nature of this bound. Is this the maximum observed error across the sample, or a bound derived from a specific confidence level? The table lists N values for field-level agreement (e.g., N=96 for Access leakage), but the aggregate metric comparison lacks the corresponding sample sizes or the method used to aggregate the errors. Without this, the claim of “closely match” is qualitative rather than quantitative.

Finally, the multiplicative nature of the Memory Governance Score (MGS) introduces non-linearity. A method with high utility but slightly higher leakage can have a drastically lower MGS than a method with moderate utility and zero leakage. While the formula is logically defensible, the paper should acknowledge that standard parametric tests (like t-tests) may not be appropriate for comparing MGS directly due to this non-linearity and the bounded nature of the metric [0, 1]. Non-parametric tests or bootstrap-based significance testing would be more appropriate for validating the claim that “Long-Context... achieves the highest MGS.”

Action items.

- [science] Report confidence intervals or standard errors for the main metrics (U, A, F, MGS) in Table 1. The current point estimates do not convey the variance across the 2,218 checkpoints or the stability of the LLM-as-a-judge evaluation.
- [science] Clarify the statistical basis for the ‘maximum absolute difference of 1.04 percentage points’ in the judge-human validation (Table A4). Specify if this is a mean absolute error, a max deviation, or a bound derived from a specific statistical test, and provide the sample size (N) for each metric in the validation set.
- [science] The multiplicative MGS formula ($U * (1-A) * (1-F)$) creates a non-linear scale. When comparing methods, consider reporting statistical significance (e.g., bootstrap confidence intervals or paired tests) for MGS differences, as small changes in A or F can disproportionately affect the final score.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript demonstrates a high standard of academic writing, with clear, precise, and well-structured prose throughout. The abstract effectively summarizes the problem, method, and findings without unnecessary jargon. The introduction successfully establishes the motivation for the benchmark, clearly distinguishing the proposed work from existing literature. The logical flow from the problem statement to the benchmark formulation and experimental results is coherent and easy to follow.

However, a few minor issues regarding sentence structure and variable naming clarity were identified. In the Introduction, the sentence structure regarding the “critical deployment regime” is slightly convoluted, which could be smoothed out for better readability. More significantly, in the Task Formulation section, the variable naming convention for the turn tuple (specifically using ‘r_t’ for timestamp) creates a potential ambiguity given the context of roles and requesters, which might momentarily confuse a reader parsing the mathematical definitions. Additionally, a broken or mismatched URL in the bibliography footnote list, while technically a reference issue, impacts the professional polish of the document if rendered.

Overall, the writing is strong and effectively communicates the complex technical contributions. Addressing these minor points will further enhance the clarity and professionalism of the paper.

Action items.

- [writing] In Section 1 (Introduction), the phrase ‘a critical deployment regime remains insufficiently studied, namely the multi-principal shared environment’ is slightly clunky. Consider rephrasing to ‘a critical deployment regime, namely the multi-principal shared environment, remains insufficiently studied’ for better flow.
- [writing] In Section 3 (Task Formulation), Equation 4 defines the turn tuple as (p_t, r_t, z_t, u_t) where r_t is the timestamp. However, the text immediately following states ‘ r_t is the timestamp’ but the variable name ‘r’ typically suggests ‘role’ or ‘requester’ in this context, potentially causing confusion with p_t (speaker). Consider renaming r_t to t_t or timestamp_t for clarity.
- [writing] In Section 4 (Experiments), the footnote listing model URLs contains a broken link or mismatched citation for ‘c-001’ (url: [.../gguf-public-stats/...](https://openai.github.io/openai-public-stats/)). While this is a bibliography issue, it disrupts the reading flow if the PDF renders a broken link. Ensure all URLs in the bibliography are valid and point to the correct model documentation.