

llmXive Automated Review of PixWorld: Unifying 3D Scene Generation and Reconstruction in Pixel Space

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: [qwen.qwen3.5-122b](#)

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper presents a unified pixel-space diffusion framework for 3D scene generation and reconstruction. The central claims regarding the architecture (two-stream DiT), the loss functions (flow matching + geometry perception), and the general superiority over latent-space baselines are supported by the reported tables and ablation studies. The internal consistency between the method description and the experimental setup is generally high.

However, there are specific instances where the textual claims are slightly more confident or specific than the provided evidence allows, or where citations are missing for critical data points:

1. **Baseline Conditioning Protocol:** In Section 4.2, the authors state that “all baselines except LVSM... additionally receive a text condition.” While this is a crucial detail for fair comparison, the tables (e.g., `tables/1view_gen.tex`) do not explicitly annotate which baselines were text-conditioned versus pose-only. Given that Gen3R and FlashWorld are often text-capable, the blanket statement “except LVSM” requires explicit confirmation in the text or table footnotes to ensure the reader can trust the comparison conditions. Without this, the claim of a fair “apples-to-apples” comparison on text-conditioned tasks is not fully substantiated by the provided text.
2. **Missing Citation for Data:** The Appendix (Implementation Details) cites `chen2025blip3` for the “10M images” from the BLIP-3o corpus. This reference key is absent from the provided `main.bib` file. While the dataset might exist, the specific citation required to verify the source of this 10M image subset is missing, making the claim unverifiable within the context of the provided manuscript.
3. **Ranking Nuance:** In Section 4.3, the text claims PixWorld “trailing only DepthSplat on SSIM for DL3DV-10K.” The table shows PixWorld (0.681) is indeed lower than DepthSplat (0.692), but YoNoSplat (0.678) is also lower. The phrasing “trailing only DepthSplat” could be interpreted as “DepthSplat is the only one better,” which is true, but the sentence structure in the context of “matching state-of-the-art” might be slightly ambiguous regarding the exact ranking order if not read carefully. A minor clarification (e.g., “trailing DepthSplat, the sole method with higher SSIM”) would eliminate any potential confusion.

These issues are minor and fixable via text edits or adding the missing bibliography entry. They do not invalidate the core scientific contribution but are necessary for full claim accuracy.

Action items.

- **[writing]** Baseline Conditioning Protocol: In Section 4.2, the authors state that “all baselines except LVSM... additionally receive a text condition.” While this is a crucial detail for fair comparison, the tables (e.g., `tables/1view_gen.tex`) do not explicitly annotate which baselines were text-conditioned versus pose-only. Given that Gen3R and FlashWorld are often text-capable, the blanket statement “except LVSM” requires explicit confirmation in the text or table footnotes to ensure the reader can trust
- **[writing]** Missing Citation for Data: The Appendix (Implementation Details) cites `chen2025blip3` for the “10M images” from the BLIP-3o corpus. This reference key is absent

from the provided main.bib file. While the dataset might exist, the specific citation required to verify the source of this 10M image subset is missing, making the claim unverifiable within the context of the provided manuscript.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

The figure presents a qualitative ablation study with quantitative metrics, but the 3D plots are missing axis labels and units, and the Ground Truth data mentioned in the caption is not visualized in the charts.

Figure 2

The figure provides visual examples of scene generation and depth prediction but fails to visually distinguish between the ‘reconstruction’ and ‘generation’ modes mentioned in the caption. Additionally, the camera trajectory plots lack a legend to define the color coding for input versus generated views.

Figure 3

Figure 3 effectively visualizes generated scenes by pairing RGB renderings with predicted depth maps as described in the caption. The layout is clear, and the visual contrast between the color images and the grayscale depth maps allows for easy assessment of the model’s performance.

Figure 4

Figure 4 effectively illustrates the architectural differences between PixWorld and existing methods with clear visual flow and labeling. The only issue is the use of raw citation keys in the caption text, which should be formatted as standard references.

Figure 5

Figure 5 provides a clear and comprehensive overview of the PixWorld architecture, effectively illustrating the model pipeline, flow matching loss, and geometry perception loss. The diagram is well-organized, and the caption accurately describes the components shown in panels (a), (b), and (c).

Figure 6

The figure effectively visualizes the model’s flexibility in handling reconstruction and generation tasks. However, the legend’s terminology (‘noise’) and style (solid lines) conflict with the caption’s description of frustums and the visual representation of image borders, creating minor ambiguity.

Figure 7

Figure 7 presents a visual comparison of PixWorld against baselines, but the layout and labeling are confusing: the top row is labeled as the method’s output rather than the input view described in the caption, and the presence of multiple scenes contradicts the singular ‘input view’ phrasing.

Action items.

- **[science]** Figure 1: The 3D plots lack axis labels and units, making it impossible to interpret the quantitative metrics (e.g., ‘Pose accuracy’) or the scale of the curves.
- **[science]** Figure 1: The legend at the bottom defines ‘GT’ (Ground Truth), but the 3D plots only display ‘Full’ and ‘w/o Geometry’ curves, leaving the GT data unvisualized despite the caption’s claim of comparing against ground-truth poses.
- **[writing]** Figure 1: The text boxes inside the 3D plots report ‘AUC@5’ values, but the caption does not define this metric or explain its relevance to the ablation study.
- **[writing]** Figure 2: The caption states it shows ‘reconstruction and generation under varying view selections,’ but the image lacks any visual distinction (e.g., labels, borders, or grouping) to identify which rows correspond to reconstruction versus generation.
- **[writing]** Figure 2: The caption mentions ‘camera trajectories with input and generated views marked,’ but the image contains no legend or key to explain the meaning of the blue and red lines in the trajectory plots.
- **[writing]** Figure 4: The caption cites ‘VAE kingma2013auto’ and ‘RAE zheng2025diffusion’ using raw citation keys instead of formatted references (e.g., [1] or Author et al.), which is inconsistent with standard figure caption style.
- **[science]** Figure 6: The legend labels ‘Clean view (input)’ and ‘Generated view (noise)’ contradict the caption’s explanation that blue frustums are clean input and red are generated; the legend implies the *images* are noise, while the caption implies the *frustums* denote the view type. This creates ambiguity about whether the red-bordered images are noisy inputs or generated outputs.
- **[writing]** Figure 6: The legend uses solid lines to represent ‘Clean view’ and ‘Generated view’, but the actual visualization uses colored borders around images and colored frustums in the 3D plots. The legend does not visually match the representation style (frustums vs. image borders) used in the figure, making it slightly confusing.
- **[science]** Figure 7: The top row is labeled ‘PixWorld (Ours)’ but lacks the ‘Input View’ indicator described in the caption; the layout implies the top row is the input, but the label

suggests it is a generated result, creating ambiguity about which image is the ground truth input versus the model’s output.

- **[writing]** Figure 7: The caption states ‘The large view on top denotes the input view,’ but the image shows a row of five distinct scenes (living room, kitchen, mountains, cafe, dining room) labeled ‘PixWorld (Ours)’ at the top, which contradicts the singular ‘input view’ description and confuses the comparison structure.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_jargon_police.md`

The paper is generally well-structured for a technical audience, but it contains several instances where notation and acronyms are introduced without immediate, explicit definition, creating minor friction for a competent reader from an adjacent field (e.g., a computer graphics researcher reading a diffusion paper, or vice versa).

Specifically, the notation $\mathbf{c}$ in Section 3.1 is used in Equation 1 as a generic conditional input but is never explicitly defined as a vector, embedding, or specific data type. A reader from a different subfield might assume it refers to a scalar class label or a different conditioning mechanism. Similarly, the sets $\Omega_{\mathbf{c}}$ and $\Omega_{\mathbf{n}}$ are central to the method’s formulation but are only formally defined in Equation 2, appearing in the text just before without a clear “let X be...” statement.

In the Appendix, the acronym “SDPA” is used in a table without expansion. While standard in some transformer literature, it is not universal enough to be assumed known by all adjacent-field PhDs. Additionally, the model name $\hat{\mathbf{v}}_3$ is used as a proper noun without a brief expansion of what the symbol stands for (Permutation-Equivariant Visual Geometry Learning), which is helpful context for a non-specialist. Finally, the abbreviation “Re10K” appears in the training details without being explicitly linked back to the full “RealEstate10K” name in that specific section, breaking the flow for a reader who might not have the full name fresh in their mind from the abstract.

These are minor issues that can be resolved with simple parenthetical expansions or one-sentence definitions at the point of first use, significantly improving the paper’s self-containment without altering the scientific content.

Action items.

- **[writing]** The paper is generally well-structured for a technical audience, but it contains several instances where notation and acronyms are introduced without immediate, explicit definition, creating minor friction for a competent reader from an adjacent field (e.g., a computer graphics researcher reading a diffusion paper, or vice versa). Specifically, the notation $\mathbf{c}$ in Section 3.1 is used in Equation 1 as a generic conditional input but is never

explicitly defined as a vector, embedding, or specific da

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper’s argument structure is logically sound and internally consistent. The central thesis—that a pixel-space diffusion paradigm unifies 3D reconstruction and generation while avoiding the information loss of latent-space methods—is supported by a coherent chain of reasoning.

1. **Premise-Conclusion Alignment:** The introduction correctly identifies the limitations of latent-space approaches (information loss, decoupled optimization) and proposes pixel-space diffusion as the direct solution. The Method section (Sec 3) implements this by defining the flow-matching objective directly on rendered images (Eq. 4, Eq. 6), ensuring the diffusion variable and the 3D supervision target are in the same domain. The conclusion (Sec 5) accurately reflects these premises without overreaching.
2. **Consistency of Definitions and Notation:** The partitioning of views into clean (Ω_c) and noisy (Ω_n) sets is defined in Eq. 2 and used consistently throughout the Method and Experiments sections. The loss function components ($\mathcal{L}_{\text{render}}$, $\mathcal{L}_{\text{depth}}$, $\mathcal{L}_{\text{geo}}$) are defined in Sec 3.2 and 3.3 and their weights (λ) are reported consistently in the Training Details (Sec 4.1) and Appendix (Appendix A).
3. **Ablation Logic:** The ablation study (Sec 4.3, Tab. 5) logically isolates the contribution of the geometry perception loss. The text states that removing this loss degrades geometric consistency (PSNR, SSIM, AUC), which is directly supported by the data in Table 5. The claim that “2D photometric and perceptual losses keep individual frames visually plausible but leave the underlying 3D geometry unsupervised” is a valid inference from the observed drop in pose accuracy (AUC) while generation quality metrics (VBench) remain relatively stable.
4. **No Contradictions:** There are no contradictions between the abstract, main body, and appendix. The parameter count (1.044B) is consistent between the text and Table 1 in the Appendix. The scope of the experiments (RealEstate10K, DL3DV-10K) is clearly bounded in the text and matches the reported tables. The limitations section (Appendix) appropriately qualifies the claims regarding resolution and dataset diversity, preventing any conflict with the main results.

The argument proceeds without logical gaps, non-sequiturs, or internal contradictions.

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several strong claims about the scope and universality of its results that slightly exceed the specific experimental evidence provided.

First, the abstract asserts that PixWorld “consistently outperforms prior latent-space generation methods.” While the method leads on most metrics, Table 2 (2-view generation on RealEstate10K) shows that LVSM (a regression-based, non-latent baseline) achieves higher PSNR (23.61 vs. 23.54) and SSIM (0.819 vs. 0.815) than PixWorld. The claim of “consistent” superiority is therefore technically inaccurate regarding reconstruction-quality metrics (PSNR/SSIM) when compared to the full set of baselines. The phrasing should be refined to specify that the superiority holds for perceptual, geometric, and camera-control metrics, or to acknowledge the specific trade-off where regression-based methods may still win on pixel-level fidelity.

Second, the conclusion frames the work as marking a “promising paradigm toward scalable and unified 3D scene modeling.” This language suggests a broad, solved trajectory for scalability. However, the “Limitations” section (Appendix A.5) explicitly admits the model is constrained by “finite resolution and compute budgets” and that “scalability to higher-resolution multi-view settings” remains an open area for improvement. The conclusion’s confident tone regarding “scalable” modeling contradicts the explicit admission that the current implementation is not yet scalable to higher resolutions. The conclusion should be hedged to reflect that the *approach* is promising for scalability, but the *current results* are limited to the tested resolution regime.

Finally, the Introduction claims the method “naturally unifies 3D scene generation and reconstruction in a single model.” While the architecture is unified, the experimental evaluation treats these as distinct tasks with separate baselines and metrics (e.g., Gen3R is excluded from the reconstruction table because it only supports point-cloud reconstruction). The narrative implies a seamless, empirically validated unification, whereas the evidence shows a unified architecture validated on separate, non-overlapping evaluation tracks. Clarifying that the unification is architectural and that performance is validated on distinct protocols would better align the claim with the evidence.

These are primarily issues of rhetorical precision rather than fundamental scientific flaws. Narrowing the scope of the claims to match the specific metrics and limitations reported will resolve the overreach.

Action items.

- **[writing]** Abstract claims PixWorld ‘consistently outperforms prior latent-space generation methods,’ but Table 2 shows LVSM (non-latent) beats PixWorld on PSNR/SSIM for RealEstate10K. Narrow claim to ‘outperforms on perceptual and geometric metrics’ or acknowledge the PSNR/SSIM trade-off.
- **[writing]** Conclusion states results ‘mark a promising paradigm toward scalable... modeling,’

implying solved scalability. However, Limitations (Appendix A.5) admit ‘finite resolution’ and ‘scalability to higher-resolution... remains open.’ Temper conclusion to reflect current resolution limits.

- **[writing]** Introduction claims method ‘naturally unifies... generation and reconstruction,’ implying seamless empirical validation. Experiments (Section 4) evaluate tasks separately with disjoint baselines (e.g., Gen3R omitted from reconstruction). Clarify that unification is architectural, validated on distinct protocols.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

The paper presents a unified pixel-space diffusion framework for 3D scene generation and reconstruction. From a safety and ethics perspective, the work is low-risk. The methodology relies on standard, publicly available datasets (RealEstate10K, DL3DV-10K, BLIP-3o) and does not involve human subjects, sensitive personal data, or private information that would require IRB approval or specific consent disclosures. The training data consists of posed multi-view scenes and single images, which are standard benchmarks in the computer vision community.

The paper includes a dedicated “Responsible Considerations” section in the Appendix (Appendix~\ref{appendix:responsible}), which appropriately addresses limitations, broader impacts, and LLM usage. The authors acknowledge potential concerns regarding privacy-sensitive scene capture and the misuse of synthetic content, and they encourage responsible data usage and human oversight. This disclosure is sufficient for the nature of the research.

There are no indications of dual-use capabilities that lower the barrier to specific harms (e.g., automated vulnerability discovery, biological synthesis, or targeted disinformation generation) beyond the general capabilities of 3D scene generation, which is a legitimate research area. The paper does not release any operational exploits, PII, or data scraped in violation of terms of service. The use of a frozen 3D foundation model as a critic is standard practice and does not introduce new safety risks.

Consequently, no specific safety or ethics action items are required. The paper meets the necessary standards for disclosure and risk mitigation.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a compelling unified framework for 3D scene generation and reconstruction. However, the evidentiary strength of the central claims is currently undermined by a lack of statistical rigor in the experimental design and potential confounds in the ablation studies.

First, the headline results in Tables 1 and 2 (1-view and 2-view generation) are presented as single-point estimates. In generative modeling, results are notoriously sensitive to random seeds, initialization, and sampling noise. The absence of standard deviations, confidence intervals, or even a statement of the number of seeds used (e.g., “results averaged over 3 seeds”) makes it impossible to assess the stability of the reported improvements. A 0.75 dB gain in PSNR or a 0.068 increase in AUC@5 could easily be within the variance of a single run. To support the claim of “superior performance,” the authors must report results across multiple random seeds (at least 3-5) with mean $\pm$ standard deviation.

Second, the ablation study in Table 3 (Section 4.3) claims that the Geometry Perception loss is the key driver of performance, citing a 1.13 dB drop when removed. However, the experimental design for this ablation is flawed: the authors state they trained the ablated variant for only 30K steps on a 10K-sequence subset, whereas the main model was trained for 200K steps on the full dataset. This introduces a severe confound: the performance drop could be entirely due to the reduced training budget and smaller dataset, not the absence of the loss term. To isolate the contribution of the geometry loss, the ablation must be run with the exact same training schedule (200K steps) and dataset size as the full model.

Finally, the inference speed comparison in the Appendix (Table 1) is not apples-to-apples. PixWorld is compared against FlashWorld, which achieves 10s inference via distillation (4 NFE), while PixWorld uses 100 NFE. While the authors acknowledge this, the comparison fails to rule out whether PixWorld’s architecture is inherently faster or if it simply lacks the distillation step. A fairer comparison would involve reporting the speed of a distilled PixWorld or a non-distilled FlashWorld to isolate the architectural contribution from the optimization technique.

Addressing these points—specifically adding seed variance, fixing the ablation training budget, and clarifying the speed comparison—is necessary to substantiate the paper’s strong claims of superiority.

Action items.

- **[writing]** The paper presents a compelling unified framework for 3D scene generation and reconstruction. However, the evidentiary strength of the central claims is currently undermined by a lack of statistical rigor in the experimental design and potential confounds in the ablation studies. First, the headline results in Tables 1 and 2 (1-view and 2-view generation) are presented as single-point estimates. In generative modeling, results are notoriously sensitive to random seeds, initialization, and sampli

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: `agents/prompts/paper_reviewer_statistical_analysis.md`

The statistical treatment of the quantitative results in this paper is insufficient for rigorous scientific validation. While the experimental design (datasets, baselines) appears sound, the reporting of the resulting numbers lacks the necessary uncertainty quantification required to trust the claimed improvements.

Specifically, Tables 1 through 4 (e.g., `tables/1view_gen.tex`, `tables/2view_gen.tex`, `tables/recon.tex`) present single point estimates for every metric (PSNR, SSIM, LPIPS, AUC, etc.) without any accompanying standard deviation (SD), standard error (SE), or confidence interval (CI). In the context of deep learning, where training dynamics are stochastic and results can vary significantly based on random seeds, initialization, or data shuffling, a single number is merely a point estimate, not a definitive property of the model. For instance, the claim that PixWorld achieves “18.88” PSNR on RealEstate10K (Table 1) implies a level of precision that cannot be verified without knowing the variance across runs. If the standard deviation were ± 0.5 , the “significant” gap to the baseline (17.82) might not be statistically robust.

Furthermore, the ablation study in Section 4.3 (`section/4_Experiment.tex`) relies on a single training run (30K steps on a 10K-sequence subset) to conclude that the geometry perception loss is “essential.” The reported drop of 1.13 dB in PSNR is substantial, but without reporting the variance across multiple seeds, it is impossible to rule out that this improvement is due to a favorable random seed or a specific data split rather than the loss function itself. The field standard for such claims is to report mean $\pm$ SD over at least 3 independent seeds.

Finally, the reporting of metrics to two or three decimal places (e.g., “0.614” AUC@5) constitutes false precision given the absence of uncertainty measures. This suggests a stability that the data does not support.

To rectify this, the authors must re-run their main experiments and ablations with multiple random seeds (3) and report the mean and standard deviation for all quantitative results. This will allow readers to assess the stability of the proposed method and the statistical significance of the reported improvements.

Action items.

- **[science]** Tables 1-4 (e.g., `tab:single_view_gen_avg`) report single point estimates (e.g., ‘18.88’ PSNR) without any measure of uncertainty (SD, SE, or CI) across training seeds. In deep learning, single-run results are unstable; report mean $\pm$ SD over 3 seeds for all quantitative comparisons to distinguish signal from noise.
- **[science]** Section 4.3 (Ablation Study) claims the geometry perception loss is ‘essential’ based on a single 10K-sequence subset trained for 30K steps. Without reporting variance across multiple random seeds or data splits for this ablation, the observed ~ 1.13 dB drop could be an artifact of a specific initialization or data split. Report mean $\pm$ SD over 3 seeds for the ablation variants.
- **[writing]** The paper reports performance to two decimal places (e.g., ‘0.614’ AUC@5) across multiple benchmarks. Without reported standard deviations, this precision is unjustified and misleading. Round to one decimal place or report the uncertainty range to reflect the true stability of the metric.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript demonstrates a high standard of writing quality, allowing a reader to move through the argument with minimal friction. The abstract effectively summarizes the problem, the proposed pixel-space solution, the specific geometry perception loss, and the resulting performance gains. The Introduction follows a logical narrative arc: it establishes the separation of reconstruction and generation, critiques the limitations of latent-space unification (specifically information loss and indirect optimization), and clearly positions PixWorld as the solution.

Paragraphs are well-structured with clear topic sentences. For instance, the “Task formulation” and “Two-stream diffusion transformer” subsections in Section 3 introduce their respective concepts immediately, followed by the necessary mathematical formalism. The transition from the general pixel-space diffusion preliminary (Section 3.1) to the specific PixWorld framework (Section 3.2) is smooth and motivated.

The prose is precise and avoids unnecessary hedging or wordiness. Technical terms like “clean stream,” “noisy stream,” and “geometry perception loss” are introduced and used consistently throughout the text. The distinction between the rendering loss (photometric) and the geometry perception loss (structural) is articulated clearly, preventing confusion about the model’s objectives.

While the paper is dense with technical details, the sentence structures remain parseable on the first pass. There are no instances of buried main points, ambiguous pronoun references, or run-on sentences that impede comprehension. The Appendix is also well-organized, with clear signposting that guides the reader to specific implementation details and additional results without feeling disjointed from the main text. Overall, the writing is professional, clear, and effectively supports the scientific content.