

llmXive Automated Review of Dual Latent Memory in Vision-Language-Action Models for Robotic Manipulation

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_claim_accuracy.md

The paper presents a coherent argument for using latent memory to address temporal bias in VLA models. However, there are specific inconsistencies between the textual descriptions of the methodology and the reported results that require clarification to ensure the claims are fully supported by the evidence.

First, in Section 3.1 (Latent Memory Curator), the authors describe the long-term memory vault as storing “action hidden states” and explicitly state it is “not a key-value bank.” However, the subsequent paragraph on “Memory Vault Updating Strategy” claims that the *same* compression strategy (averaging adjacent keys and values based on cosine similarity) is applied to the long-term vault. This creates a logical contradiction: if the vault does not store keys and values, it cannot be compressed by averaging them. The text must clarify whether the long-term vault also utilizes a key-value representation (contradicting the earlier claim) or if a different compression mechanism is used for the long-term stream.

Second, the comparative claims in the results sections (Sections 4.1 and 4.2) require precise alignment with the tables. In Section 4.1, the text states the method surpasses “recent state-of-the-art VLAs such as... SemanticVLA.” While Table 1 confirms LaMem-VLA (73.9%) outperforms SemanticVLA (65.1%), the margin is 8.8 points, not the 16.6 points highlighted for CogACT. The phrasing could be misinterpreted as implying a similar magnitude of improvement over all listed SOTA methods. Similarly, in Section 4.2, the text reports a “1.1 point” improvement over MemoryVLA (overall average) and a “3.5 point” improvement over pi0 (first-four-suite average). While the numbers in the text match the table calculations, the lack of explicit distinction between “overall average” and “first-four-suite average” in the narrative flow risks confusing the reader about the exact baseline being compared. The text should explicitly qualify these metrics (e.g., “1.1 points in overall average” and “3.5 points in the first-four-suite average”) to ensure the claim accuracy is unambiguous.

Finally, the citation of “SemanticVLA” (ni2026semanticvla) and “MemoryVLA” (shi2025memoryvla) as 2026 and 2025 publications respectively is consistent with the bibliography, but given the preprint nature of the paper, the authors should ensure these references are indeed the final versions or clearly marked as preprints if the specific numbers (e.g., 96.5% for MemoryVLA) are derived from a specific arXiv version that may differ from the published conference version. However, since the bibliography lists them as conference papers (ICLR’26, CVPR’26), the numbers should be treated as final. The primary issue remains the internal consistency of the methodology description and the precision of the comparative claims.

Action items.

- **[writing]** The paper presents a coherent argument for using latent memory to address temporal bias in VLA models. However, there are specific inconsistencies between the

textual descriptions of the methodology and the reported results that require clarification to ensure the claims are fully supported by the evidence. First, in Section 3.1 (Latent Memory Curator), the authors describe the long-term memory vault as storing “action hidden states” and explicitly state it is “not a key-value bank.” However, th

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

The figure provides a clear visual overview of the framework architecture, but the caption is grammatically broken with missing text and empty parentheses, and the diagram itself contains a spelling error in the ‘Hidden States’ section.

Action items.

- **[writing]** Figure 1: The caption contains grammatical errors and placeholders, reading ‘The Framework of .’ and ‘The memory curator ()’/‘memory seeker ()’/‘memory condenser ()’ with empty parentheses where the model name or specific terms should be.
- **[writing]** Figure 1: The ‘Hidden States’ box contains a typo ‘Insturction token’ instead of ‘Instruction token’.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The paper is generally well-written and avoids excessive in-group slang, but it suffers from a recurring pattern of introducing mathematical notation in equations without defining the symbols in the surrounding prose. A competent reader from an adjacent field (e.g., standard NLP or Computer Vision) will likely understand the high-level flow but will stall when encountering specific dimensions like N_{s} , N_{l} , and K_{q} in Sections 3.2 and 3.3.

Specifically, in Section 3.2, the variable N_{s} appears in Equation 2 and Equation 5. The text defines N_{i} as the sequence length but leaves N_{s} (the number

of tokens per short-term memory unit) and $N_{\text{ext}\{1\}}$ (the number of tokens per long-term memory unit) undefined. These are critical for understanding the dimensionality of the retrieved evidence $\mathbf{Z}^{\text{short}}$ and $\mathbf{Z}^{\text{long}}$. Similarly, in Section 3.3, $K_{\text{ext}\{q\}}$ is used in Equation 4 without a textual definition. Finally, in Section 3.4, the notation $\mathbf{1}_{L_{\text{ext}\{s\}}}$ is used in Equation 7; while likely a vector of ones, this specific subscripted notation is not standard enough to be assumed without a brief definition.

These are not fundamental scientific errors, but they create unnecessary friction for the reader who must infer the meaning of these dimensions from context or by guessing. Defining these symbols at their first point of use (or immediately preceding the equation) would resolve these barriers. The rest of the terminology (e.g., “diffusion action expert,” “Markovian assumption,” “latent space”) is standard for the field and does not require expansion.

Action items.

- **[writing]** Section 3.2 (Latent Memory Curator): The symbol $N_{\text{ext}\{s\}}$ appears in Eq. 2 and Eq. 5 without definition. While $N_{\text{ext}\{i\}}$ is defined as sequence length, $N_{\text{ext}\{s\}}$ (likely the number of tokens per short-term memory unit) is introduced abruptly. Define $N_{\text{ext}\{s\}}$ at its first occurrence in Eq. 2.
- **[writing]** Section 3.2 (Latent Memory Curator): The symbol $N_{\text{ext}\{1\}}$ appears in Eq. 5 (long-term memory retrieval) without definition. Similar to $N_{\text{ext}\{s\}}$, the dimensionality of long-term memory units is undefined. Define $N_{\text{ext}\{1\}}$ alongside $N_{\text{ext}\{s\}}$ in the text describing Eq. 5.
- **[writing]** Section 3.3 (Latent Memory Seeker): The symbol $K_{\text{ext}\{q\}}$ is introduced in Eq. 4 as the number of learnable query slots but is never defined in the prose. Define $K_{\text{ext}\{q\}}$ explicitly when first mentioning the query builder $\mathcal{B}$.
- **[writing]** Section 3.4 (Latent Memory Weaver): The symbol $\mathbf{1}_{L_{\text{ext}\{s\}}}$ is used in Eq. 7 without definition. While the subscript suggests a vector of ones of length $L_{\text{ext}\{s\}}$, the notation is non-standard for this context. Add a brief clause defining $\mathbf{1}_{L_{\text{ext}\{s\}}}$ as a column vector of ones of dimension $L_{\text{ext}\{s\}}$.
- **[writing]** Section 3.2: The term ‘SE-bottleneck compression module’ is used with a citation to Hu et al. (2018). While ‘SE’ (Squeeze-and-Excitation) is standard in CV, the specific adaptation as a ‘bottleneck’ for token compression is a specific architectural choice. Add a brief parenthetical explaining that this module compresses the token sequence dimension before the mean-pooling step.

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper presents a coherent argument for using latent memory to address temporal bias in VLA models. However, two specific logical inconsistencies were identified in the method description and result reporting.

First, in **Section 3.3 (Latent Memory Curator)**, there is a direct contradiction regarding the long-term memory vault ($\mathcal{M}^{\text{long}}$). The text explicitly states: “Unlike the short-term vault, $\mathcal{M}^{\text{long}}$ is not a key-value bank.” However, the subsequent paragraph claims: “We apply the same memory updating strategy to the long-term memory vault.” The described strategy (Eqs. 3-4) fundamentally relies on computing cosine similarity between *keys* ($\mathbf{k}_s^i, \mathbf{k}_s^{i+1}$) to identify redundant adjacent pairs. If $\mathcal{M}^{\text{long}}$ lacks keys, the “same strategy” cannot be executed as described. The argument breaks because the premise (no keys) contradicts the mechanism (key-based similarity) required for the claimed operation. The authors must either define keys for the long-term vault or clarify that the strategy is adapted (e.g., using action hidden state similarity) rather than being identical.

Second, in **Section 4.3 (Ablation Study)**, the textual summary of results creates ambiguity. The text states: “Removing both memory streams causes the largest degradation, dropping performance to 57.3% and 92.1%.” While these numbers correspond to the “w/o Dual Memory” row in Table 1, the sentence structure fails to explicitly map 57.3% to SimplierEnv and 92.1% to LIBERO-90. Given the preceding sentences discuss single-stream ablations with values like 65.6% and 64.6%, the lack of explicit mapping makes it difficult to immediately verify the claim against the table without cross-referencing column headers. Clarifying the dataset association in the text would ensure the conclusion follows unambiguously from the evidence.

These are minor logical gaps fixable by text revision.

Action items.

- **[writing]** Section 3.3 states $\mathcal{M}^{\text{long}}$ is ‘not a key-value bank’ but later claims the ‘same memory updating strategy’ (which relies on key similarity) is applied to it. This is a logical contradiction. Define keys for $\mathcal{M}^{\text{long}}$ or clarify the strategy differs.
- **[writing]** Section 4.3 claims removing both streams drops performance to ‘57.3% and 92.1%’. The text must explicitly map these to SimplierEnv and LIBERO-90 respectively to avoid ambiguity with the single-stream ablation values in the same paragraph.

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes a strong case for latent memory in VLAs, but the rhetoric occasionally exceeds the scope of the provided evidence, particularly regarding the universality of the “superiority” claim.

The abstract states that experiments “demonstrate the superiority of our LaMem-VLA.” While the method outperforms baselines on the reported metrics in SimplerEnv-Bridge and LIBERO, the term “superiority” without qualification implies a general dominance across the domain of robotic manipulation. The evidence is strictly limited to two specific simulation suites and a specific set of baselines (e.g., MemoryVLA, CogACT, π_0). To align the claim with the evidence, the abstract should specify that the method demonstrates superiority *on these specific benchmarks* or *in these simulated settings*.

Similarly, the Introduction and Conclusion frame the results as a broad validation of the “context-native latent memory” paradigm. While the authors correctly identify the limitation of simulation-only validation in the Introduction’s final paragraph, the Conclusion re-asserts the “superior performance” without immediately qualifying the scope. This creates a slight disconnect where the body admits the boundary (simulation only) while the framing layer (Abstract/Conclusion) presents the results as a definitive step forward for the field at large. Given that the real-world evaluation is explicitly marked as future work and not included in the results tables, the claims of “superiority” must be explicitly bounded to the tested environments to avoid overreach.

The paper does not claim “solves” the problem of long-horizon tasks, which is good, but the unqualified use of “superiority” suggests a generalization that the current data (two simulators, specific tasks) does not fully support. Narrowing these claims to the specific experimental scope is necessary to maintain scientific rigor.

Action items.

- **[writing]** Abstract: Change ‘demonstrate the superiority’ to ‘demonstrate improved performance on these benchmarks’. The current claim implies universal superiority, but evidence is limited to SimplerEnv-Bridge and LIBERO suites only.
- **[writing]** Introduction (last paragraph): Ensure the conclusion explicitly reiterates that ‘superiority’ is bounded to the two simulated environments tested, as real-world generalization remains unverified in this version.
- **[writing]** Conclusion: Replace ‘demonstrate the superior performance’ with ‘demonstrate improved performance on the tested benchmarks’. The unqualified ‘superior’ suggests a universal claim not supported by the two-simulator scope.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

This work presents a novel architecture for Vision-Language-Action (VLA) models focused on improving long-horizon robotic manipulation through latent memory mechanisms. The research is conducted entirely within simulated environments (SimplerEnv and LIBERO) using standard,

publicly available datasets (Bridge v2, Open-X Embodiment).

From a safety and ethics perspective, the paper poses no foreseeable, non-trivial risk of harm. The methodology does not involve:

1. **Human Subjects:** No human data, surveys, or behavioral logs were collected; the datasets used are standard robotic benchmarks.
2. **Dual-Use Capabilities:** The system is designed for robotic manipulation in simulation. It does not generate disinformation, automate cyberattacks, or possess capabilities that meaningfully lower the barrier to physical harm in a way that requires specific mitigation beyond standard robotic safety protocols.
3. **Privacy Violations:** No Personally Identifiable Information (PII) is processed or released.
4. **Deception or Surveillance:** The system is not designed to impersonate humans or conduct covert surveillance.

The paper explicitly acknowledges its current limitation to simulation (Introduction, Section 1), noting that real-world deployment is a future step. This transparency is appropriate. There are no missing disclosures regarding consent, licensing, or conflict of interest that would prevent publication. The work is low-risk by construction.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a compelling architectural shift for VLA models, moving memory from external conditioning to native latent tokens. However, the evidentiary strength of the reported gains is currently insufficient to rule out noise or confounding factors.

First, the primary results in Table 1 (SimplerEnv) and Table 2 (LIBERO) are presented as single-point averages (e.g., 73.9% vs 71.9% for MemoryVLA) with no reported standard deviation, confidence intervals, or number of random seeds used. In robotic manipulation benchmarks, success rates can fluctuate significantly due to random initialization, task sampling, or stochasticity in the environment. A 1.1% or 2.0% margin is often indistinguishable from sampling noise without variance reporting. The authors must report results averaged over at least 3-5 independent training seeds with standard deviations to demonstrate that the improvements are robust and not artifacts of a lucky seed.

Second, the ablation study in Table 3 attributes performance gains to the “dual-scale” nature of the memory. While removing both streams causes a large drop, the individual removals of short-term or long-term memory show moderate drops (~8-9%). The paper does not rule out the alternative explanation that the gain comes simply from increasing the total number of memory tokens (context length) rather than the specific dual-scale mechanism. To isolate the

contribution of the dual-scale design, the authors should add a control experiment where a single memory stream is expanded to match the total token count of the dual-stream model (e.g., 12 tokens in one stream vs 8+4). If the single-stream model performs similarly, the “dual-scale” claim is weakened.

Finally, the comparison between “Latent-native” and “Policy-side” memory (Table 4) requires clarification on the retrieval and condensation pipeline. The text implies the baseline uses “external policy-side condition,” but it is unclear if this baseline uses the same number of retrieved units ($K=8$) and the same condenser architecture, or if it uses raw retrieved evidence. If the baseline uses raw tokens or a different retrieval budget, the comparison confounds the “latent-native” integration with the benefits of better compression or retrieval. The authors must explicitly state that the baseline uses identical retrieval and condensation logic, differing only in the injection point (input sequence vs. external conditioning), to validly isolate the architectural contribution.

Action items.

- **[writing]** The paper presents a compelling architectural shift for VLA models, moving memory from external conditioning to native latent tokens. However, the evidentiary strength of the reported gains is currently insufficient to rule out noise or confounding factors. First, the primary results in Table 1 (SimplerEnv) and Table 2 (LIBERO) are presented as single-point averages (e.g., 73.9% vs 71.9% for MemoryVLA) with no reported standard deviation, confidence intervals, or number of random seeds used. In

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The statistical reporting in this paper is currently insufficient to support the strength of the quantitative claims made. While the experimental protocols (number of trials/rollouts) are described in Sections 4.1 and 4.2, the results are presented exclusively as single point estimates (e.g., “73.9%”, “97.6%”) without any accompanying measure of uncertainty such as standard deviation (SD), standard error (SE), or confidence intervals (CI).

In robotic manipulation benchmarks like SimplerEnv and LIBERO, performance can vary significantly across random seeds or initial states. Reporting only the mean (or a single best run) obscures the stability of the method. For instance, a 16.6% gain over a baseline is impressive, but without the variance, a reader cannot determine if this gap is robust or if the distributions overlap significantly. The field standard for such benchmarks is to report mean $\pm$ SD over multiple seeds or trials.

Furthermore, the paper makes numerous comparative claims (“surpasses,” “outperforms,” “significantly better”) across a large number of baselines and ablation settings (Tables 1, 2, 3, 4). These constitute multiple hypothesis tests. The current presentation treats each comparison

as an isolated event, ignoring the multiplicity problem where the probability of at least one false positive increases with the number of comparisons. While formal p-values are not always required in ML, if the authors claim “superiority” based on these numbers, they must either provide statistical tests with appropriate corrections (e.g., Holm-Bonferroni) or temper their language to reflect that these are observed differences without statistical validation.

Finally, the ablation studies (Section 4.3) compare the full model against variants. These are paired comparisons (same seed, same task, different model). The current text simply lists the success rates. To rigorously support the claim that removing a component “causes degradation,” a statistical test (e.g., paired t-test or Wilcoxon signed-rank test) should be performed on the per-trial data, and the resulting p-values or effect sizes should be reported.

The fix is primarily a reporting adjustment: re-analyze the existing trial data to compute and report SD/CI for all means, and perform the necessary statistical tests for the ablation and baseline comparisons. No new experiments are required, but the current numbers cannot be interpreted as definitive evidence of superiority without this context.

Action items.

- **[writing]** Section 4.1 and 4.2 report single point estimates (e.g., 73.9%, 97.6%) for success rates without any measure of uncertainty (SD, SE, or CI). While 24 trials (SimplerEnv) and 50 rollouts (LIBERO) are mentioned, the variance across these trials is not reported. Report mean $\pm$ SD (or 95% CI) for all main results to allow assessment of stability and effect magnitude.
- **[writing]** Tables 1-4 and Figure 1 present comparisons across multiple baselines and ablation settings (e.g., 12+ pairwise comparisons in Table 1 alone) without correcting for multiple comparisons. The claim of ‘superiority’ or ‘significant’ gains relies on uncorrected point estimates. Apply a correction (e.g., Bonferroni or Holm) for the family of comparisons or explicitly state that p-values are uncorrected and interpret ‘significance’ cautiously.
- **[writing]** Section 4.3 reports ablation results (e.g., 73.9% vs 57.3%) as fixed values. Since these are derived from the same experimental setup, the differences should be tested for statistical significance (e.g., paired t-test or bootstrap) rather than just compared as point estimates. Report the p-value or confidence interval for the difference to support the claim that the degradation is real and not due to random seed variance.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The paper is generally well-structured and the narrative flow is strong, particularly in the Introduction where the problem of “temporal short-horizon bias” is clearly articulated and the proposed solution is logically motivated. The transition from the problem statement to the

four-component architecture is smooth and easy to follow.

However, there are a few specific instances where the prose creates minor friction or where structural references are incorrect, requiring the reader to pause and re-evaluate.

First, in the caption of Figure 1 (Section 3.2), the text references `\S\ref{sec:seeker}` for the “memory condenser.” However, the condenser is described in Section 3.3 (or the subsequent section in the actual flow), not the seeker. This mismatch forces the reader to hunt for the correct section, breaking the immersion. Correcting the LaTeX label reference is a simple fix that restores the intended flow.

Second, in Section 4.3 (Ablation Study), the paragraph discussing the number of latent memory tokens (L_s and L_l) suffers from slight redundancy and awkward phrasing. The sentence “For long-term memory, fixing $L_s=8$ and increasing long-term latent memory token number L_l provides stronger task-progress...” repeats the phrase “long-term latent memory token number” unnecessarily. Tightening this to “increasing L_l ” would improve readability without losing precision.

Third, the term “retrieval budget” is introduced in the same section to describe the parameter K . While understandable in context, the term “budget” is not explicitly defined or used consistently elsewhere (where “number of retrieved units” is used). Standardizing this terminology to “retrieval count” or explicitly defining “budget” upon first use would prevent minor ambiguity.

Finally, the abstract ends with a generic placeholder sentence: “The project page will be available at...”. This breaks the summary’s momentum and feels like an unpolished draft element. Replacing this with a definitive statement about the release status (e.g., “Code and models will be released at...”) or removing it entirely would ensure the abstract remains focused on the scientific contribution.

Overall, the writing is clear and the argument is compelling. Addressing these minor reference errors and tightening the phrasing in the ablation study will make the paper seamless to read.

Action items.

- **[writing]** Section 3.2 (Method), caption of Fig 1: The text states ‘condenser... (`\S ef{sec:seeker}`)’ but the condenser is defined in Section 3.3. Update the reference to `\S ef{sec:condenser}` (or the correct label) to prevent reader confusion.
- **[writing]** Section 4.3 (Ablation Study), paragraph on ‘The Number of Latent Memory Tokens’: The sentence ‘For long-term memory, fixing $L_s=8$ and increasing long-term latent memory token number L_l provides stronger task-progress...’ is repetitive and slightly awkward. Rewrite to: ‘For long-term memory, fixing $L_s=8$ and increasing L_l provides stronger task-progress cues, reaching up to 73.9% on SimplerEnv and 97.0% on LIBERO-90.’ (writing)
- **[writing]** Section 4.3 (Ablation Study), paragraph on ‘The Number of Retrieved Memory Units K ’: The phrase ‘retrieval budget of the memory seeker’ is used, but ‘budget’ is not defined earlier. Ensure ‘retrieval budget’ is clearly defined or replace with ‘retrieval count’ for consistency with the variable K . (writing)
- **[writing]** Abstract: The sentence ‘The project page will be available at...’ is a standard placeholder that breaks the summary flow. Replace with a concrete statement about the code

release status (e.g., ‘Code and models will be released at...’) or remove if not yet available, to maintain the abstract’s focus on scientific contribution. (writing)