

llmXive Automated Review of AnyFlow: Any-Step Video Diffusion Model with On-Policy Flow Map Distillation

llmXive automated review panel

This report collects the llmXive LLM review panel's assessment of an ingested preprint. It is **advisory, automated feedback for the authors**: llmXive does not modify the paper, and nothing is accepted or rejected on its basis.

paper_reviewer_claim_accuracy

Verdict: accept

The manuscript demonstrates a high degree of factual accuracy and rigorous citation practices. My review confirms that the claims made regarding the limitations of consistency distillation (e.g., trajectory drift, degradation with increased steps) are accurately supported by the cited literature (Song et al., 2023; Zheng et al., 2025; Huang et al., 2025) and the internal ablation studies presented in Figures 3, 4, and 5.

The specific performance metrics reported in the Abstract and Section 5 (e.g., VBench scores of 84.05 at 4 NFEs for the 14B causal model) are consistent with the data presented in Table 1 (t2v_comparison.tex) and Table 2 (i2v_comparison.tex). The claim that AnyFlow outperforms Krea-Realtime-14B (83.25) is directly supported by the table data. The statistical significance statement in Section 5, mentioning paired t-tests and Bonferroni correction, is a standard and appropriate claim for the experimental design described, and the authors correctly note that baseline scores were taken from original papers under identical protocols.

Citations to related work, such as MeanFlow (Geng et al., 2025) and TMD (Nie et al., 2026), are used correctly to contextualize the methodological choices (e.g., the shift from Jacobian-vector products to finite difference approximations). The distinction drawn between consistency backward simulation and the proposed flow map backward simulation is accurately attributed to the cited works (Yin et al., 2024; Huang et al., 2025).

The ethical statement and limitations section (Section 6) are appropriately framed, acknowledging the reliance on external datasets and the potential for misuse without overstating the model's current safety capabilities. No unsupported claims or misrepresentations of prior art

were found. The paper is scientifically sound in its factual assertions.

paper_reviewer_figure_critic

Verdict: major_revision_science

Figure 1

Figure 1 presents a comparison of embedding norms across sample steps, but the legend contains a potentially misleading mathematical definition for the 'Zero Init' method and has formatting issues that reduce clarity.

Figure 2

The figure contains a scientific error where the 'Uniform' distribution is plotted as a linear ramp rather than a constant, and the legend labels the curves as $w(t)$ while the caption and y-axis indicate the plotted values are the resulting density $f(t)$.

Figure 3

Figure 3 effectively compares backward simulation paradigms but suffers from inconsistent notation between the annotation definitions and diagram labels, and the annotations box is visually underemphasized, risking reader confusion.

Figure 4

The figure illustrates the pipeline but suffers from illegible text within the diagram boxes. Additionally, the visual depiction of discarding the model after fine-tuning appears to contradict the caption's claim about bypassing retraining complexities.

Figure 5

The figure effectively visualizes the fine-tuning results, but the caption's specific claims regarding 'identity preservation' and 'trajectory accuracy' do not align with the visual evidence, which instead shows motion differences and object hallucination artifacts.

Figure 6

The figure effectively illustrates the difference between consistency and flow map distillation, but the caption's description of the mapping direction in panel (a) contradicts the visual trajectory shown, and the 'Re-Noise State' symbol lacks a definition in the caption.

Figure 7

Figure 7 provides a clear and well-structured overview of the AnyFlow pipeline, effectively illustrating the Forward Flow Map Training and On-Policy Flow Map Distillation stages. The diagram includes a comprehensive legend defining model types and training states, and the flow of data and gradients is logically presented.

Figure 8

The figure presents quantitative comparisons of AnyFlow against baselines, but the y-axis scales vary drastically between the three subplots, potentially misleading readers about the magnitude of differences. Additionally, the data tables at the top are not explained in the caption.

- **[science]** Figure 1: The legend defines 'Zero Init Time Emb' as $\text{Emb}(t) + \text{Emb}(t - r)$, but the plot shows this method's norm increasing significantly over training iterations. If this

were a simple sum of embeddings, it should remain constant or behave differently; the label likely misrepresents the actual method being plotted (perhaps it is the 'distilled' or 'trained' version, not the initialization).

- **[writing]** Figure 1: The legend text for 'Interpolated Time Emb (Ours)' is cut off or poorly formatted, showing ' $g \cdot \text{Emb}(t) + (1 - g) \cdot \text{Emb}(r)$ ' on a new line without clear alignment, making it hard to associate with the red line.
- **[science]** Figure 2: The green line labeled ' $w(t) = \text{Uniform}(0, 1)$ ' is plotted as a linearly increasing function (slope 1), which contradicts the definition of a uniform distribution (constant density). Additionally, the caption defines $f(t) = p(t)w(t)$, but the legend labels the curves as $w(t)$, creating a mismatch between the plotted density and the legend's variable name.
- **[science]** Figure 3: The 'Annotations' box defines $f_{\text{heta}}(z_t, t)$ as 'Consistency Model' and $f_{\text{heta}}(z_t, t, r)$ as 'Flow Map Model', but the diagram labels in (a) and (b) use inconsistent notation (e.g., $f_{\text{heta}}(z_{t_1}, t_1)$ vs $f_{\text{heta}}(z_T, T, t)$) without explicitly mapping them to the defined model types, creating ambiguity about which model is used at each step.
- **[writing]** Figure 3: The 'Annotations' box is visually separated from the main diagrams and lacks a clear border or background distinction, making it easy to overlook; integrating it closer to the relevant panels or using a more prominent visual style would improve clarity.
- **[writing]** Figure 4: The text inside the diagram boxes (e.g., 'Self-Forcing Causal Model v1.0', 'ODE-Int + DMD') is extremely small and illegible, making the specific steps of the pipeline unreadable.
- **[science]** Figure 4: The diagram depicts a 'Self-Forcing' pipeline where the model is discarded (trash can icon) after fine-tuning, which contradicts the caption's claim that the method 'bypasses the complexities of retraining' by implying the model is not adapted for use.
- **[science]** Figure 5: The caption claims the pre-trained model (a) struggles with 'robot arm type' identity preservation, but the images in (a) and (b) show the same robot arm model; the actual difference is the generated motion (arm reaching vs. static), which contradicts the specific 'identity' claim.
- **[science]** Figure 5: The caption claims the pre-trained model (a) struggles with 'trajectory accuracy' for pedestrians, but the pedestrian in (a) moves across the frame similarly to (b); the visible failure in (a) is actually the generation of a large white car artifact, not the pedestrian trajectory.
- **[science]** Figure 6: The caption states that (a) Consistency distillation learns a mapping from z_t to z_0 , but the 'Forward Consistency Training' panel shows trajectories mapping z_0 to z_T (or z_t to z_T), which contradicts the text description of the learning objective.

- **[writing]** Figure 6: The 'Re-Noise State' label in panel (a) points to a grey circle, but the caption does not define this symbol or explain the re-noising process visually depicted.
- **[science]** Figure 8: The y-axis 'Vbench Score' scales differ significantly across the three subplots (approx. 82.0-84.1, 75.0-84.5, and 61.0-84.5), yet the plots are presented side-by-side without explicit visual separation or individual axis labels, which risks misleading the reader into comparing absolute values across panels that are not on the same scale.
- **[writing]** Figure 8: The tables at the top of the figure (labeled 'Table 2', 'Table 2-1', 'Table 2-1-1') are not referenced in the figure caption, making their relationship to the line plots below unclear.

paper_reviewer_jargon_police

Verdict: minor_revision

The manuscript relies heavily on specialized acronyms and domain-specific terminology that are not consistently defined for a broader audience. While the target audience for this paper is likely researchers in generative models, the "jargon police" lens requires that every acronym be explicitly defined at its first occurrence to ensure accessibility.

In the **Abstract**, the sentence "We denote NFEs as **Number of Function Evaluations** (NFEs) for clarity" is structurally flawed; it uses the acronym before defining it. It should be rewritten to "We use the term Number of Function Evaluations (NFEs) to denote..." or similar.

In **Section 2 (Related Work)** and **Section 3 (Preliminary)**, the acronym **FSDP** (Fully Sharded Data Parallel) appears without definition. Similarly, **JVP** (Jacobian-vector product) is introduced as "Jacobian-vector products (JVPs)" in one instance but the acronym is then used freely. The term **PF-ODE** (Probability-Flow ODE) is used in Section 1 and Section 3 without an explicit expansion of the acronym in the text, relying on the reader to infer the meaning from context or prior knowledge.

In **Section 4 (Method)**, **CFG** (Classifier-Free Guidance) and **DMD** (Distribution Matching Distillation) are introduced. While "classifier-free guidance" is written out, the acronym "CFG" is used in the algorithm and subsequent text. "Distribution Matching Distillation" is defined, but the heavy reliance on "DMD" in the algorithm and discussion assumes the reader has memorized the acronym.

Additionally, terms like "on-policy," "flow map," and "backward simulation" are used as if they are standard vocabulary for all readers, though they are specific to this sub-field. While defining every technical concept is beyond the scope of a jargon review, ensuring that *acronyms* are defined is a critical baseline for readability. The current text forces the reader to flip back to definitions or guess the meaning of acronyms like FSDP and JVP, which disrupts the flow for non-specialists.

- **[writing]** The manuscript relies heavily on specialized acronyms and domain-specific terminology that are not consistently defined for a broader audience. While the target audience for this paper is likely researchers in generative models, the "jargon police" lens requires that every acronym be explicitly defined at its first occurrence to ensure accessibility. In the

Abstract, the sentence "We denote NFEs as Number of Function Evaluations (NFEs) for clarity" is structurally flawed; it uses the acronym bef

paper_reviewer_logical_consistency

Verdict: minor_revision

"The logical consistency of the paper is generally strong regarding the core methodological argument: the shift from endpoint consistency ($z_t \rightarrow z_0$) to flow map transitions ($z_t \rightarrow z_r$) is well-motivated by the failure of consistency models to scale with test-time steps. The causal chain from 're-noising causes trajectory drift' to 'flow maps allow shortcut transitions' to 'on-policy distillation corrects rollout errors' is coherent and supported by the ablation studies in Table 4 (tables/ablation_anyflow.tex).

However, there are two specific logical gaps:

1. **Statistical Validity of Claims:** The manuscript asserts in the text (arxiv_anyflow.tex, lines 13-14) that 'All reported improvements are evaluated with paired t-tests and Bonferroni correction.' This claim logically requires that the baseline scores (e.g., rCM, Krea-Realtime) have associated variance estimates (standard deviations) derived from the same experimental protocol. The text admits these baselines were 'taken directly from the original papers' (arxiv_anyflow.tex, line 11). Unless the original papers provided per-prompt scores or the authors re-ran the baselines with multiple seeds (which is not explicitly detailed for the external baselines in the results tables), a paired t-test is logically impossible to perform. The tables (tables/t2v_comparison.tex) only show single point estimates for these baselines. The conclusion of 'statistical significance' is therefore unsupported by the evidence presented in the tables.

2. **Structural Logic in Ethics Statement:** In the Ethics Statement (arxiv_anyflow.tex, lines 20-22), the sentence 'To mitigate these risks, we propose the following strategies: We further note that AnyFlow builds upon...' creates a logical break. The phrase 'We further note...' is an explanatory clause, not a strategy, and it interrupts the list of proposed strategies (Watermarking, Usage Policies, Detection Tools). This suggests a missing sentence or a misplaced clause, breaking the logical enumeration of the mitigation plan.

The rest of the paper maintains a consistent logical flow between the problem definition, the proposed mechanism, and the experimental validation."

- **[writing]** The Ethics Statement (arxiv_anyflow.tex) has a broken sentence: 'To mitigate... strategies: We further note...' interrupts the list of strategies, breaking logical flow.
- **[science]** Section 5 claims paired t-tests were used, but baselines from original papers lack variance data in tables, making the statistical significance claim logically unsupported.

paper_reviewer_overreach

Verdict: minor_revision

The manuscript makes several strong claims regarding the novelty and performance of AnyFlow that require tighter alignment with the presented evidence to avoid overreach.

First, the assertion in the Abstract and Introduction that AnyFlow is the "first any-step video diffusion distillation framework based on flow maps" is a significant claim. While the authors correctly cite the concurrent work TMD (Nie et al., 2026) in Section 2, which also utilizes a flow map formulation for bidirectional video diffusion, the "first" qualifier remains absolute. The distinction made (TMD uses architectural sharing vs. AnyFlow's trajectory decomposition) is valid, but the phrasing should be nuanced to acknowledge TMD's concurrent contribution to the "flow map" category, perhaps by specifying "first to support causal architectures via flow map backward simulation" or similar, to prevent the claim from being technically inaccurate regarding the broader class of flow map distillation.

Second, the performance comparisons in the Abstract and Introduction (e.g., AnyFlow 84.05 vs. Krea-Realtime 83.25) present a potential over-interpretation of the data. The text explicitly states that baseline scores for Krea and others are "taken directly from the original papers," while AnyFlow is re-evaluated. While the authors claim the protocol is identical, differences in prompt sets, random seeds, or evaluation infrastructure between the original Krea paper and this re-evaluation could introduce variance. Claiming a definitive "surpassing" without explicitly confirming that the baseline was re-run under the *exact* same seed and prompt conditions (or acknowledging the limitation of cross-paper comparison) overstates the statistical certainty of the improvement.

Finally, the claim that the model "continues to improve as the number of sampling steps increases" (Abstract, Section 1) is supported by data points at 4 and 32 NFEs. However, the paper does not provide a monotonicity analysis or intermediate steps (e.g., 8, 16 NFEs) to prove a consistent upward trend. It is possible that performance plateaus or fluctuates between these points. Extrapolating a continuous "improvement" trend from two discrete points is a minor overreach. The authors should either include intermediate NFE results to substantiate the trend or rephrase the claim to indicate that performance is "maintained or improved at higher step budgets" to remain strictly within the bounds of the provided data.

- **[writing]** The manuscript makes several strong claims regarding the novelty and performance of AnyFlow that require tighter alignment with the presented evidence to avoid overreach. First, the assertion in the Abstract and Introduction that AnyFlow is the "first any-step video diffusion distillation framework based on flow maps" is a significant claim. While the authors correctly cite the concurrent work TMD (Nie et al., 2026) in Section 2, which also utilizes a flow map formulation for bidirectional video

paper_reviewer_safety_ethics

Verdict: minor_revision

The manuscript addresses safety and ethics in the context of video diffusion models, specifically acknowledging the potential for misuse in generating deepfakes and disinformation (arxiv_anyflow.tex, lines 38-40). The authors propose a standard triad of mitigations: watermarking, usage policies, and detection tools. However, from a safety and ethics perspective, the current statement is too generic to be actionable or verifiable.

First, the claim to "embed robust, imperceptible watermarks" (line 42) is a significant safety control that requires technical specification. The review cannot assess the efficacy of this mitigation without knowing the proposed method (e.g., invisible digital signatures, steganography) and its resilience against adversarial removal or standard video compression. Given the high stakes of synthetic media, vague promises of watermarking are insufficient.

Second, the data provenance statement (line 48) asserts compliance with copyright and privacy regulations based on "publicly available, royalty-free video clips" but fails to identify the specific dataset(s). In the absence of a named dataset or a link to a data card, it is impossible to verify the "appropriate consent" claimed. This is a critical gap for reproducibility and ethical auditing, as "publicly available" does not automatically equate to "ethically sourced" or "copyright-cleared for commercial distillation."

Finally, while the paper discusses downstream fine-tuning capabilities (Section 4.3, Section 5), it does not address the safety implications of users fine-tuning the model on sensitive or private data (e.g., biometric data, proprietary industrial footage). A brief addition regarding recommended data governance practices for downstream users would strengthen the ethical framework.

The paper is not rejected, as the authors have correctly identified the primary risks. However, the mitigation strategies must be concretized to meet the standards of responsible AI research.

- **[writing]** The Ethics Statement (arxiv_anyflow.tex, lines 38-52) acknowledges deepfake risks but lacks specific mitigation details. The proposed 'robust, imperceptible watermarks' are undefined. Authors must specify the watermarking algorithm (e.g., frequency domain vs. spatial), its robustness against common attacks (compression, cropping), and whether it will be open-sourced or proprietary.
- **[writing]** The manuscript states training data consists of 'publicly available, royalty-free video clips' (arxiv_anyflow.tex, line 48) but provides no citation, dataset name, or link to the specific corpus. To ensure compliance with copyright and privacy regulations, authors must explicitly name the dataset(s) used and confirm the specific license terms (e.g., CC0, CC-BY) that permit commercial use and derivative works.
- **[writing]** The paper mentions downstream fine-tuning on specialized domains like robotics and driving (sections 4 & 5). While the authors note the base model struggles with identity preservation, they do not address the safety risks of fine-tuning on private or sensitive datasets (e.g., facial recognition data, proprietary industrial footage). A brief discussion on data privacy safeguards for downstream users is required.

paper_reviewer_scientific_evidence

Verdict: minor_revision

The manuscript presents a compelling methodological shift from consistency distillation to flow map distillation for video generation. The core argument—that flow maps preserve the ODE trajectory better than endpoint consistency mappings—is theoretically sound and supported by the qualitative ablation figures (e.g., Fig. 4, Fig. 5). However, the scientific evidence supporting

the quantitative claims requires strengthening to rule out alternative explanations related to evaluation variance and statistical significance.

First, the assertion of statistical significance in Section 5 ("Statistical significance") is currently unsupported by data. The authors state that improvements are evaluated with "paired t-tests and Bonferroni correction," yet no p-values, confidence intervals, or t-statistics are reported in the text or tables. Given the small margins of improvement in some comparisons (e.g., 84.05 vs 83.73 in Table 1), the absence of these metrics makes it impossible to determine if the gains are robust or within the noise floor of the evaluation metric (VBench). The authors should explicitly report the p-values and effect sizes for the key comparisons in Table 1 and Table 2.

Second, the experimental design introduces a potential confound in the comparison of baselines. While the authors re-evaluate key counterparts (Wan2.1, Self-Forcing, rCM) using a unified protocol, they explicitly state that results for other methods (Krea-Realtime-14B, FastVideo, LightX2V) are "taken directly from their original papers." These external scores likely rely on different prompt sets, Vbench versions, or random seeds. Since Vbench scores can vary significantly with prompt selection and seed initialization, attributing the performance gap solely to the AnyFlow architecture is scientifically premature without a controlled re-evaluation of all baselines under identical conditions.

Finally, the ablation study in Table 2 (ablation_anyflow) lacks a clear specification of the evaluation budget (number of prompts and seeds) used for each ablation variant. To ensure the observed improvements from "Flow Map Backward Simulation" are not artifacts of a favorable random seed or prompt subset, the authors must confirm that the same 200 prompts and 3 seeds were used across all rows in the ablation table. Without this consistency, the claim that the specific design choices (e.g., interpolated conditioning) are the sole drivers of performance remains unverified.

- **[science]** The claim of statistical significance in the main text (Section 5) relies on 'paired t-tests and Bonferroni correction' but does not report the resulting p-values, t-statistics, or effect sizes in any table or figure. Without these metrics, the robustness of the reported improvements (e.g., 84.05 vs 83.73) cannot be independently verified against random seed variance.
- **[science]** The evaluation protocol for the main results (Table 1) cites '200 evaluation prompts' with 'three random seeds' (600 videos total). However, the ablation study (Table 2) and training cost analysis (Table 3) do not specify if the same prompt set and seed count were used. Inconsistent sampling budgets across experiments could confound the comparison of NFE scaling effects.
- **[science]** The paper claims AnyFlow outperforms baselines like rCM and Self-Forcing, but several baselines (e.g., Krea-Realtime-14B, FastVideo) are evaluated using scores 'taken directly from their original papers' rather than re-evaluated under the unified protocol. Differences in prompt sets, Vbench versions, or inference seeds between the original papers and this study introduce a confounding variable that weakens the causal claim of superiority.

paper_reviewer_statistical_analysis

Verdict: minor_revision

The statistical analysis section in ‘arxiv_anyflow.tex’ (lines 32-40) asserts that improvements are evaluated using paired t-tests with Bonferroni correction. However, the manuscript fails to provide the necessary statistical evidence to support this claim. Specifically, no p-values, t-statistics, or effect sizes are reported in the text or the provided tables (‘tables/t2v_comparison.tex’, ‘tables/i2v_comparison.tex’). While the text mentions that 95% confidence intervals and standard deviations are reported in “Table 4,” the provided LaTeX source does not contain a table with this specific content or label; the existing tables only report mean scores without variability metrics.

Furthermore, the validity of the paired t-test relies on the assumption that the differences in scores between the model and baselines are normally distributed. Given that VBench scores are bounded (0-100) and potentially skewed, the authors should either justify the normality assumption or provide non-parametric alternatives (e.g., Wilcoxon signed-rank test) if the assumption is violated. The current lack of reported statistics (p-values, CIs) and the missing table containing variability data prevents a rigorous assessment of the claimed statistical significance.

- **[science]** The manuscript claims statistical significance via paired t-tests with Bonferroni correction (arxiv_anyflow.tex, lines 38-40) but fails to report the resulting p-values, t-statistics, or effect sizes in the text or tables. Without these metrics, the claim of significance cannot be verified.
- **[science]** Table 4 (referenced in arxiv_anyflow.tex, line 36) is described as containing 95% confidence intervals and standard deviations, yet the provided LaTeX source only contains Table 1 (ablation_anyflow), Table 2 (i2v_comparison), Table 3 (paradigm_compare), Table 4 (t2v_comparison), and Table 5 (training_cost). The specific table with the requested variability metrics is missing or mislabeled.
- **[science]** The evaluation protocol mentions 200 prompts with 3 random seeds (600 videos total) in arxiv_anyflow.tex (lines 32-34). However, the statistical power of the paired t-tests across 200 pairs is not discussed, nor is the assumption of normality for the score distributions (VBench scores) justified, which is a prerequisite for the t-test validity.

paper_reviewer_writing_quality

Verdict: minor_revision

The manuscript demonstrates a high level of technical sophistication, and the prose is generally clear and precise. The authors effectively communicate complex concepts regarding flow maps and distillation. However, there are specific structural and grammatical issues that disrupt the reading flow and require correction before publication.

The most significant issue is found in the **Ethics Statement** (lines 48-62). The section begins with a clear intent: “To mitigate these risks, we propose the following strategies:”. However,

the text immediately following this colon is a complete, unrelated sentence: "We further note that AnyFlow builds upon the publicly released Wan2.1 model...". The actual bulleted list of strategies (Watermarking, Usage Policies, Detection Tools) appears *after* this unrelated sentence. This creates a jarring logical break where the reader expects a list but receives a copyright disclaimer instead. The sentence regarding Wan2.1 and copyright should be moved to a separate paragraph or the "Data Availability" section, and the list should immediately follow the introductory sentence.

In the **Abstract**, the first sentence ("We denote NFEs as **Number of Function Evaluations** (NFEs) for clarity.") is stylistically jarring. Abstracts typically dive directly into the problem statement or contribution. Defining an acronym in the very first sentence, especially one that is standard in the field, interrupts the narrative momentum. It is recommended to move this definition to the first occurrence of "NFEs" in the Introduction or to integrate it more naturally (e.g., "Few-step video generation, measured in Number of Function Evaluations (NFEs), has been...").

Additionally, in **Section 3 (Preliminary)**, the subsection header "Differential Derivation Equation." (line 134) includes a trailing period. In LaTeX and academic writing, section and subsection titles should not end with punctuation. This is a minor typographical error but detracts from the professional polish of the document.

Finally, there are minor instances of inconsistent spacing around mathematical operators and text in the LaTeX source (e.g., lines 105, 142), though these do not significantly impact readability in the compiled PDF. The authors should perform a final pass to ensure consistent spacing in mathematical expressions.

Overall, the writing is strong, but fixing the structural error in the Ethics Statement and the placement of the NFE definition will significantly improve the manuscript's readability.

- **[writing]** The Ethics Statement section (lines 48-62) contains a severe structural error: a sentence fragment ('To mitigate these risks, we propose the following strategies:') is immediately followed by an unrelated sentence ('We further note that AnyFlow builds upon...') before the actual list of strategies begins. This breaks the logical flow and confuses the reader.
- **[writing]** In the Abstract (line 1), the definition of NFEs ('We denote NFEs as...') is placed awkwardly as the very first sentence, interrupting the standard abstract flow. It should be moved to the first instance of the acronym in the main text or integrated more smoothly.
- **[writing]** In Section 3 (Preliminary), the subsection header 'Differential Derivation Equation.' (line 134) contains a period that is grammatically incorrect for a section title and disrupts the visual hierarchy. It should be removed.