



Training

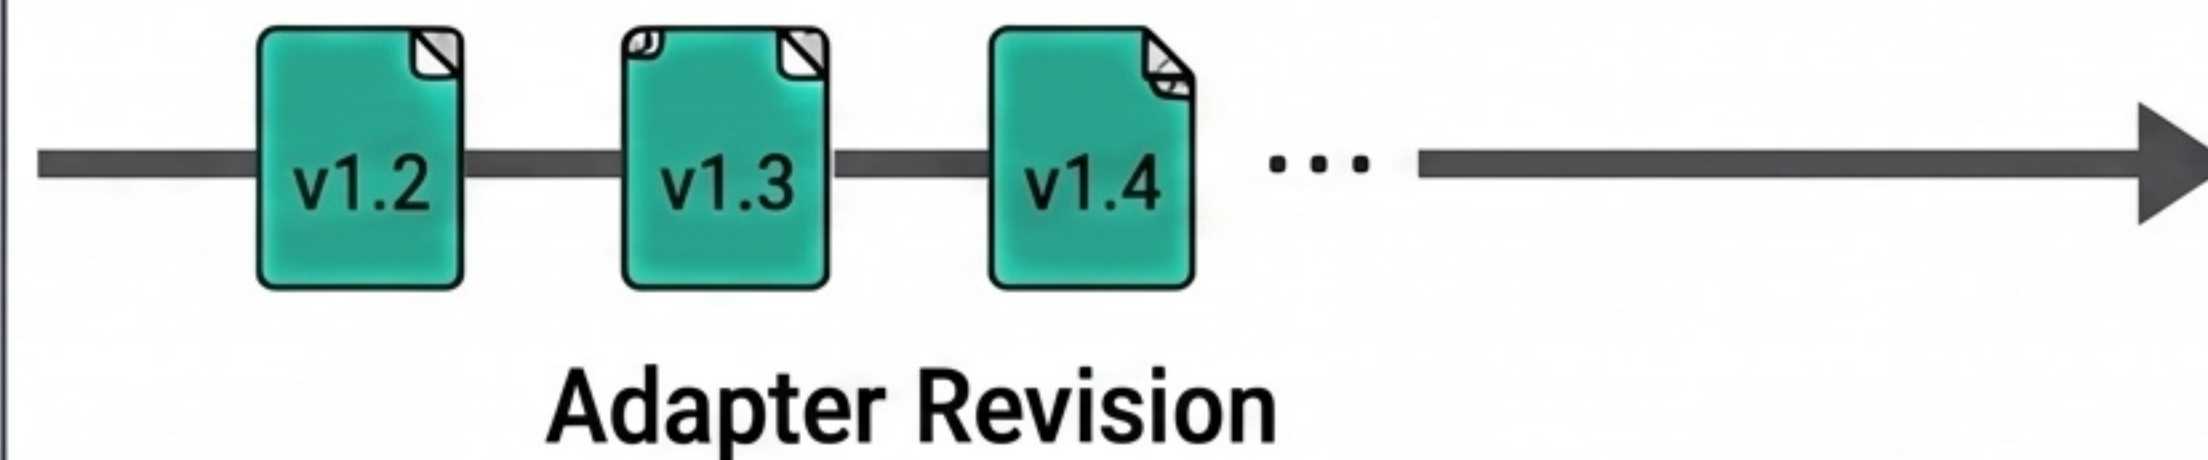
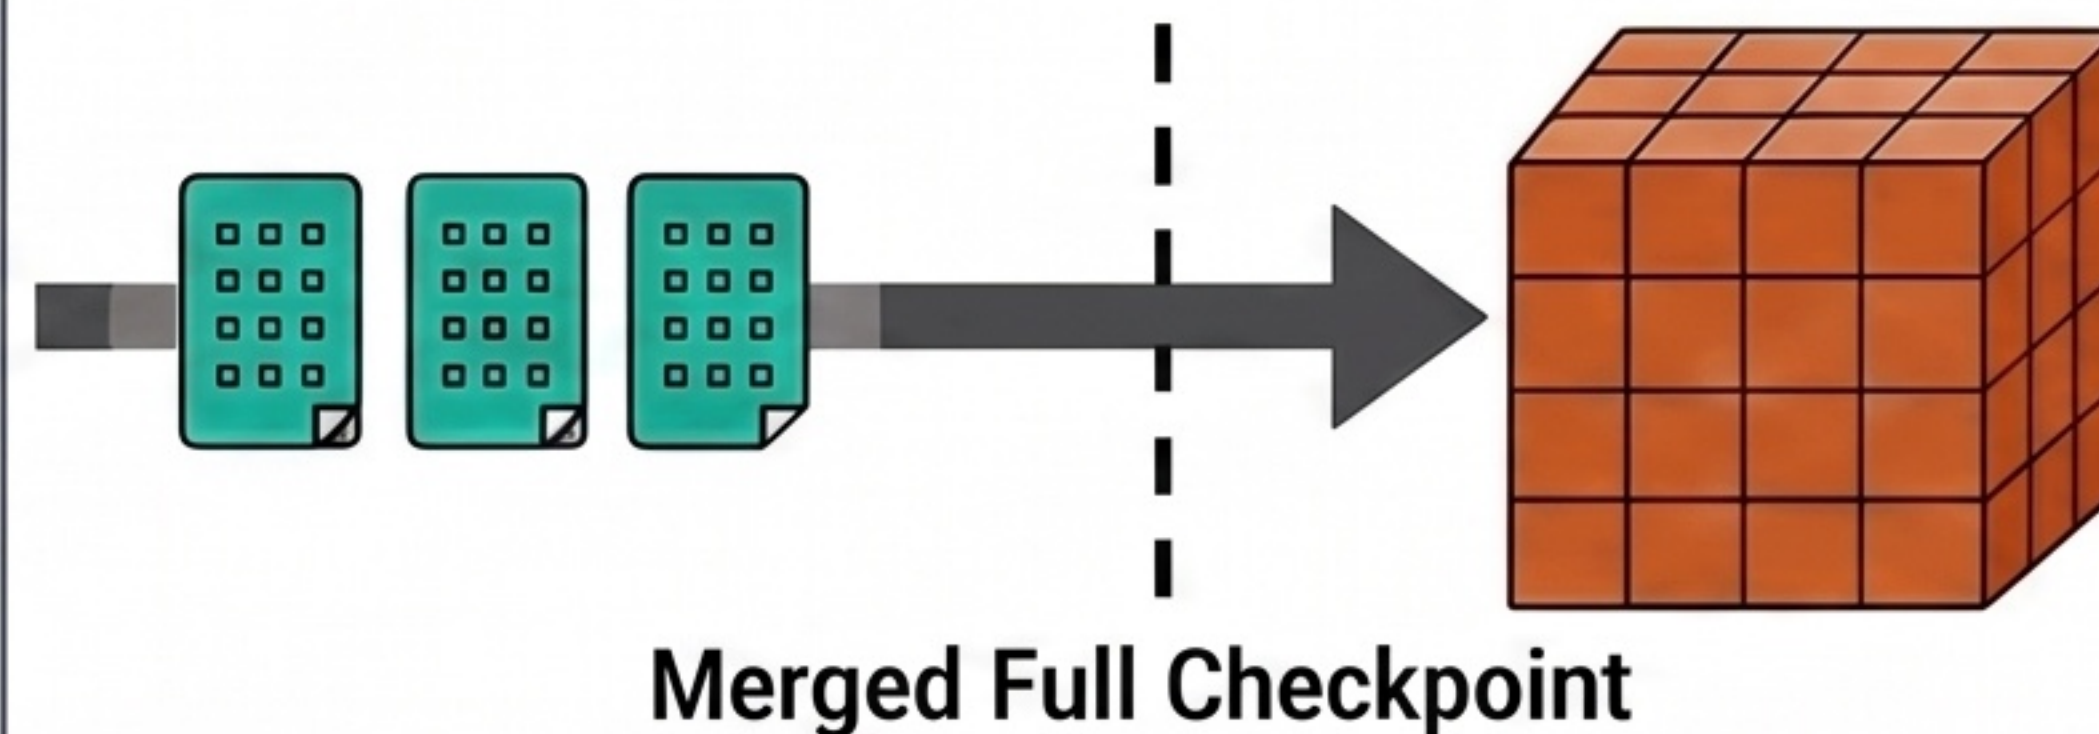
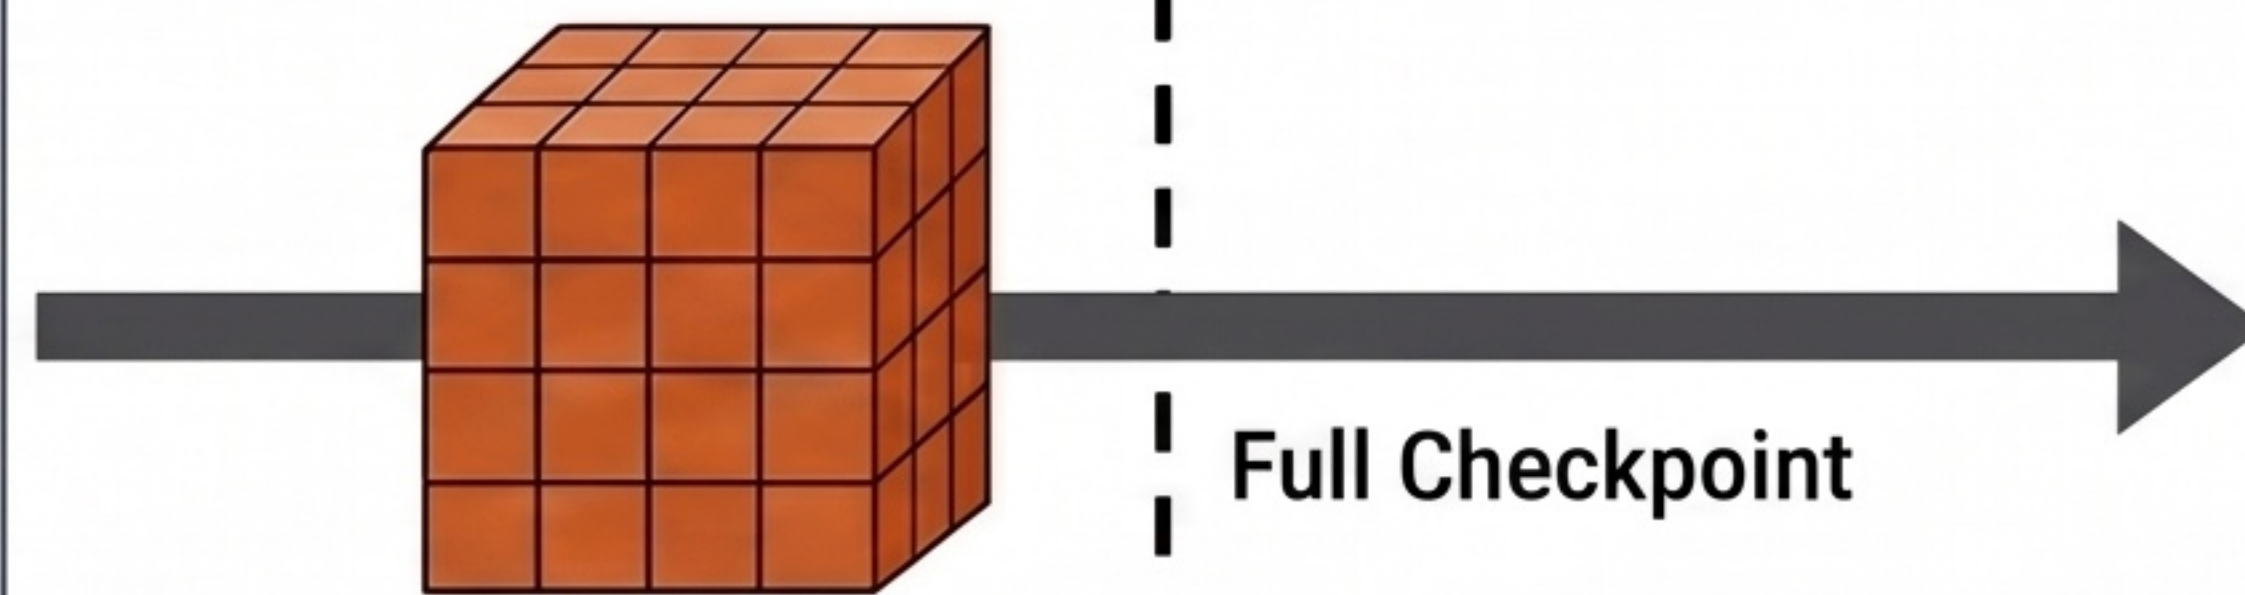
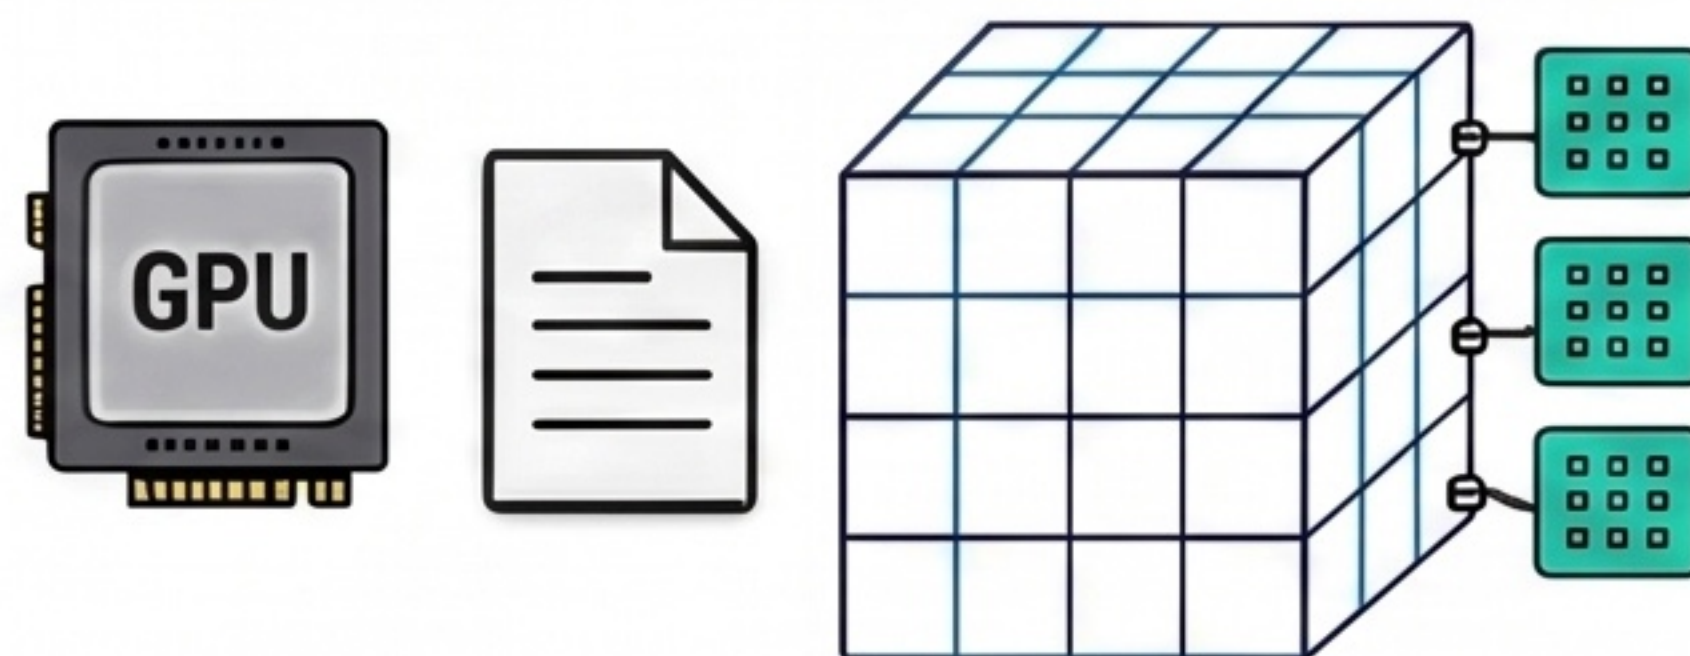
Full Fine-Tuning



Merge-Based LoRA

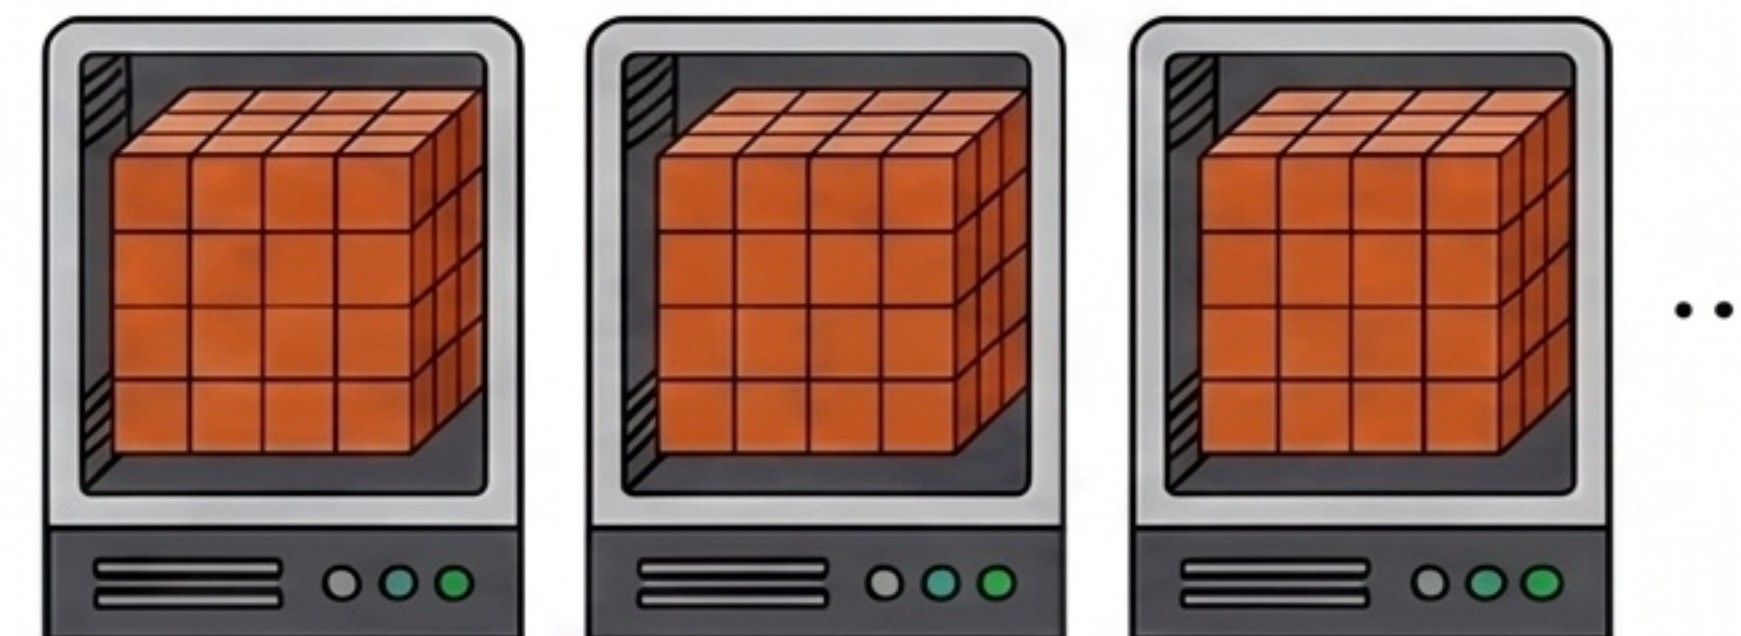


MinT / Multi-LoRA

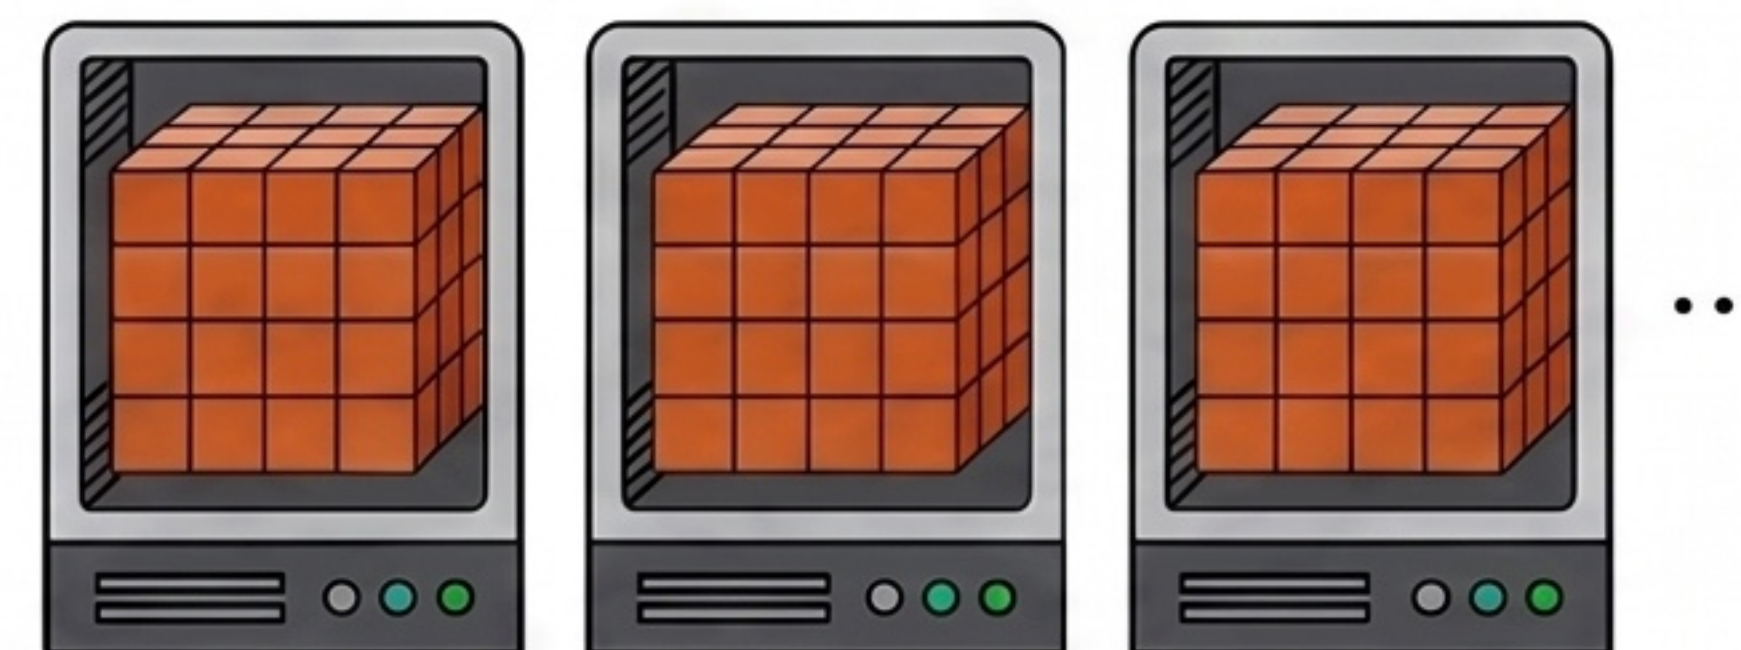


Serving

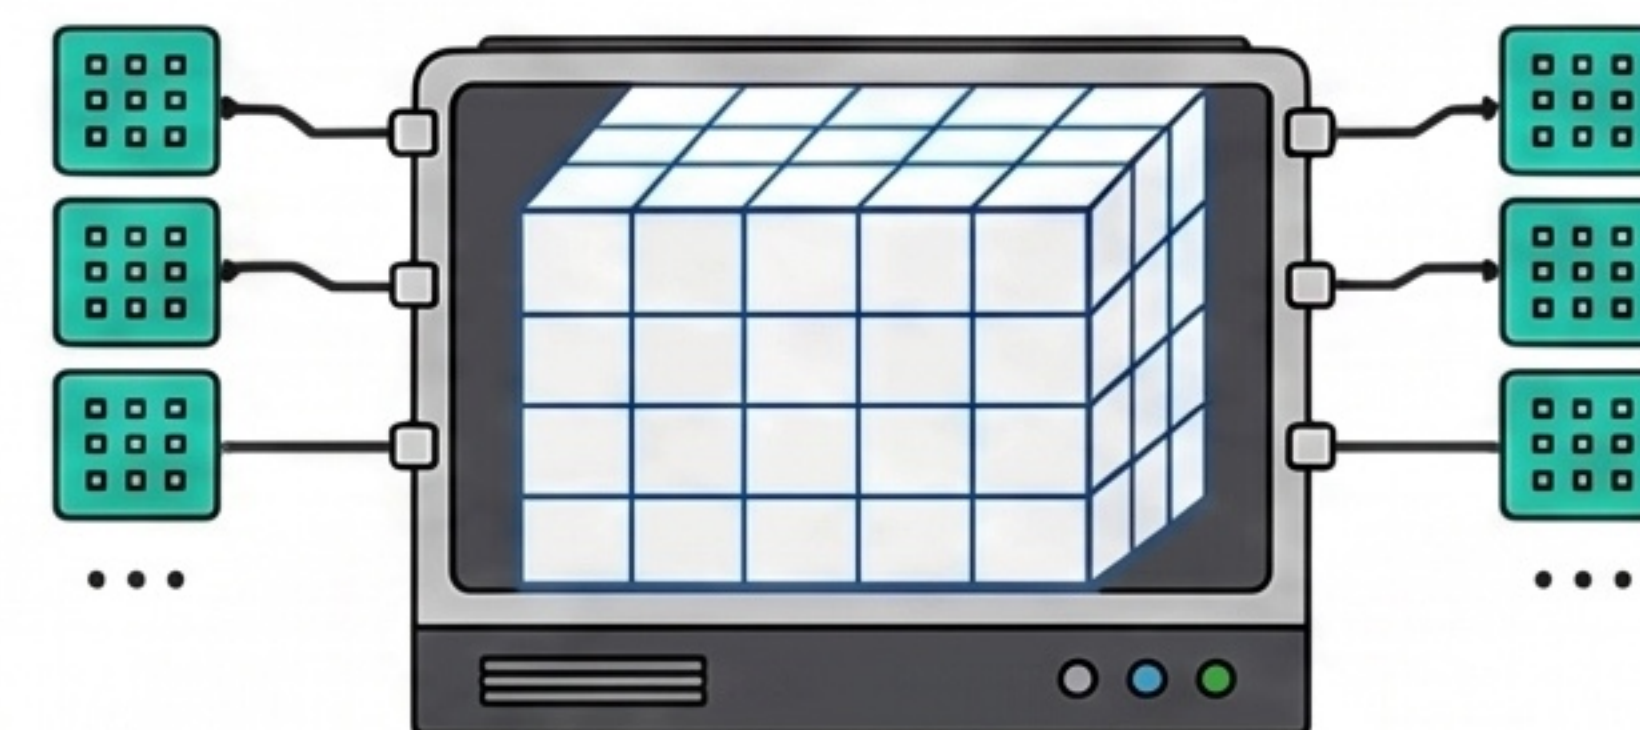
Separate Base Copies



Merged Model Copies

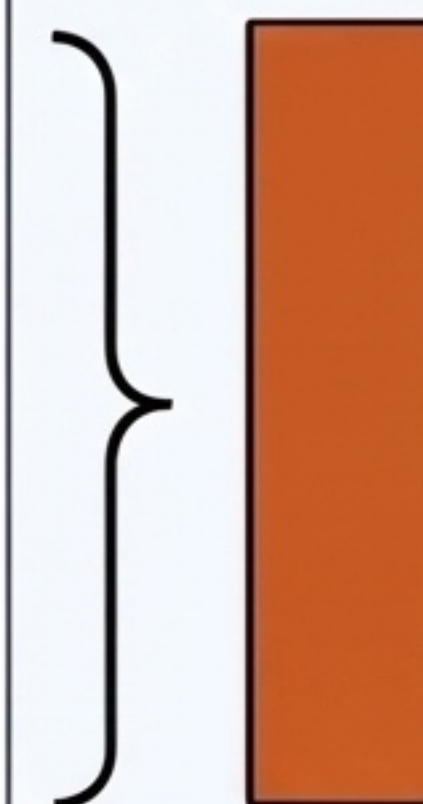


One Resident Base



Many Adapters

Memory grows with policies



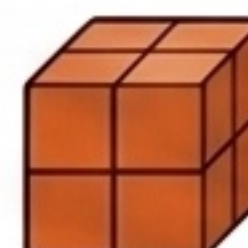
Base
memory
shared



Frozen base



Trained adapter



Materialized checkpoint



Serving actor