

Per-model failure-mode fingerprint (sorted by HR)

	HR (trustworthy)	PR (prejudice)	CR (confabulation)	IR (integration-fail)
Gemini 3 Flash	33.5	17.2	44.7	28.7
GPT-5.5	28.0	15.5	46.4	36.5
Gemini 3.1 Pro	27.4	10.8	53.2	33.4
Gemini 2.5 Pro	20.0	16.9	52.3	48.6
GPT-5.4	17.9	33.0	46.1	48.7
Gemini 2.5 Flash	16.9	28.0	43.9	59.0
Claude Opus 4.6	16.8	40.7	48.4	46.1
Qwen3.5-397B	15.9	41.5	54.3	40.9
Qwen3-VL-235B	12.8	47.0	56.4	43.8
Claude Sonnet 4.6	12.5	46.1	47.4	54.3
Claude Haiku 4.5	12.4	55.9	44.6	46.4
Qwen3-VL-30B	12.4	52.4	58.8	38.8
GPT-5.4-mini	11.3	55.4	49.4	48.2
Gemma-4-31B	11.3	29.8	75.4	39.2
o4-mini	7.8	48.2	71.7	51.7
Llama-4-Maverick	5.8	64.5	75.8	37.4
GLM-4.6V	4.6	52.7	80.7	51.2
GPT-4o	4.5	69.7	75.7	38.3
MiMo-VL-7B	3.6	56.5	84.1	42.2
Qwen3-VL-8B	2.8	62.4	85.5	47.4
Step3-VL-10B	0.9	62.3	92.9	62.1
MiniCPM-o 2.6	0.6	67.8	95.1	55.5
Qwen2.5-Omni-7B	0.5	79.9	95.8	55.6
GPT-4o-mini	0.3	87.9	95.2	45.5
Qwen2.5-VL-7B	0.1	86.3	98.9	58.9
InternVL3-8B	0.0	75.6	99.8	54.9
LLaVA-NeXT-7B	0.0	82.3	100.0	62.9