

llmXive Automated Review of SoCRATES: Towards Reliable Automated Evaluation of Proactive LLM Mediation across Domains and Socio-cognitive Variations

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so errors, misreadings, and inaccuracies are likely. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is advisory, automated feedback for the authors. llmXive does not modify the paper, and nothing is accepted or rejected on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_claim_accuracy.md

The paper presents several quantitative claims that are numerically consistent with the provided tables. The claim that the topic-localized evaluator achieves a Pearson correlation of $r = 0.82$ (trajectory) and $r = 0.80$ (outcome), “more than doubling” the ProMediate baseline, is supported by Table 2 (0.823 vs 0.372 and 0.801 vs 0.432). The calculation $0.823/0.372 \approx 2.21$ validates the “more than doubling” assertion.

The claim that the strongest mediator closes “roughly a third” of the unmediated consensus gap is supported by the reported average consensus gain of 34.4% for GPT-5.4-mini (Table 1). The breakdown of the 40 scenarios into “5 per domain” across 8 domains is mathematically consistent with the total count.

However, the claim that the five socio-cognitive axes are applied “independently” to “isolate failure modes” requires scrutiny. While the experimental design (Table 3) lists 15 non-stacked conditions, the paper does not present a statistical interaction analysis (e.g., testing for significant interaction effects between axes) to empirically prove that the effects are truly independent. The claim that the design “isolates” failure modes is a methodological intent, but presenting it as a proven result of the data without interaction tests may overstate the evidence. The text should clarify that the design *aims* to isolate these factors, rather than asserting the isolation as a confirmed statistical fact.

Additionally, the citation for the “Harvard conflict simulation framework” (fisher2011getting) is used to justify the scenario tuple $(\mathcal{B}, \mathcal{P}, \mathcal{T}, \mathcal{W})$. While Fisher’s work is foundational to negotiation, the specific 8-domain taxonomy used for scenario curation is not derived from this book. The claim that the framework follows this specific citation for the *structure* is acceptable, but the specific domain list should not be implied as coming from that source.

Finally, the comparison to ProMediate relies on a specific baseline value (0.372). It is crucial to verify that this value is explicitly reported in the cited Liu et al. (2025) paper. If the authors re-ran ProMediate to generate this number, the citation should be accompanied by a statement clarifying that the baseline was re-evaluated under the same conditions to ensure a fair comparison.

Action items.

- [writing] Section 3.2 claims ‘40 hard scenarios (5 per domain)’ across 8 domains. While mathematically consistent, the specific 8-domain taxonomy lacks an external citation, appearing to be derived solely from the authors’ agentic process. Clarify the source of this taxonomy.
- [writing] Section 4.1 states the evaluator ‘more than doubles’ ProMediate’s performance (0.82 vs 0.37). Verify that the 0.372 baseline value is explicitly reported in the cited Liu et al. (2025) paper, not a new run by the authors, to ensure the comparison is valid.

- [science] Section 3.3 claims axes are applied ‘independently’ to ‘isolate failure modes.’ Without statistical interaction tests (e.g., ANOVA), this is a design intent, not a proven result. Soften the claim to ‘designed to isolate’ to avoid overstating the evidence.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

Figure 1 provides a clear and readable visual overview of the system pipeline, but the caption contains grammatical errors and a missing noun that slightly reduce its professional quality.

Figure 2

Figure 2 presents radar charts for eight different models, but the caption’s claim of measuring ‘Mediator adaptation’ is ambiguous and potentially misleading, as it suggests a single mediator’s performance across conditions rather than a comparison of multiple mediators. The legend is redundant with the subplot labels, and the axis abbreviations are not defined in the caption.

Figure 3

The figure presents heatmaps of consensus gain shifts, but the mapping between the caption’s three abstract axes and the specific condition labels in the subplots is unclear. Additionally, inconsistent colorbar scales across subplots hinder direct visual comparison of effect magnitudes.

Figure 4

Figure 4 presents intervention effectiveness across conversation progress but fails to link the listed models to the five socio-cognitive axes claimed in the caption, and lacks clarity on the y-axis units and x-axis scale representation.

Figure 5

The figure fails to match its caption, as it displays four unlabeled heatmaps instead of the three specified subplots (a, b, c). Additionally, the axis labels do not align with the persona axes described in the text, and inconsistent colorbar scales hinder cross-panel comparison.

Figure 6

The figure appears to display ‘Intervention Effectiveness’ rather than the ‘Change in Consensus Gain’ described in the caption, creating a direct contradiction between the visual data and the text. Additionally, the caption is a verbatim copy of Figure 5’s description, failing to accurately label the specific content of Figure 6.

Figure 7

Figure 7 is a clear and legible screenshot of the annotation interface described in the caption. It effectively illustrates the pairwise comparison layout, scenario details, and evaluation criteria without any missing labels or confusing elements.

Figure 8

Figure 8 is a clear and well-annotated screenshot of the user interface for the consensus score evaluation task. The caption accurately describes the image as an example of the annotation template, and all relevant UI elements, instructions, and conversation snippets are legible.

Action items.

- [writing] Figure 1: The caption contains a grammatical error and missing noun in the opening phrase (‘Overview of : agentic scenario curation...’), likely omitting the system name ‘SoCRATES’.
- [writing] Figure 1: The caption contains a grammatical error in the phrase ‘expose where mediators fails’, which should be ‘fail’ to agree with the plural subject.
- [science] Figure 2: The legend lists 8 models, but the caption states the figure measures ‘Mediator adaptation’ (singular). The radar charts show 8 separate subplots (a-h), one per model, rather than a single comparison of one mediator’s adaptation across axes. This creates a disconnect between the caption’s claim of ‘adaptation’ and the visual presentation of static performance profiles for multiple distinct mediators.
- [writing] Figure 2: The legend uses color to distinguish models, but the subplots (a-h) also use color to distinguish models. This redundancy is unnecessary and potentially confusing if the colors in the legend do not perfectly match the fill colors in the subplots (e.g., ‘Gemini-3.1-FL’ is red in the legend and red in (a), but ‘Qwen3-30B’ is blue in the legend and blue in (h)). The legend is redundant given the subplots are already labeled.
- [writing] Figure 2: The axis labels (GEN, SA, MS, LONG, EMO, CUL) are defined in the large left chart but are not explicitly defined in the caption. While the caption mentions ‘five socio-cognitive axes’, it does not map the abbreviations (e.g., ‘SA’ for Strategic Adaptation) to the full terms, forcing the reader to infer from the large chart.
- [science] Figure 3: The caption describes three axes (strategic posture, emotional reactivity, cultural identity), but the subplots are labeled with specific condition pairs (e.g., ‘Avoiding’, ‘Accommodating’, ‘Com-Com’, ‘US-US’) without defining which axis each

column group represents, making the mapping between the caption's abstract axes and the concrete data ambiguous.

- [writing] Figure 3: The colorbars for the three subplots have inconsistent scales (e.g., -60 to 60 vs -40 to 40), which prevents direct visual comparison of the magnitude of consensus gain shifts across the different axes.
- [science] Figure 4: The caption claims to show results for 'five socio-cognitive axes' (General - 5 hard conditions), but the legend only lists 6 models (Average, Gemini 3-1 FL, GPT-5 6-m, DeepSeek V3.2, Gemma 4-25B, Nemotron 3-12B, Qwen 3-25B, Qwen 3-30B). There is no mapping between the models and the specific socio-cognitive axes (e.g., Strategic Posture, Emotional Reactivity) mentioned in the caption.
- [writing] Figure 4: The y-axis label 'Intervention Effectiveness' is present, but the unit or scale (e.g., percentage points, raw score) is not defined in the axis or caption, making the magnitude of the effect ambiguous.
- [writing] Figure 4: The x-axis label 'Conversation Progress (%)' is clear, but the caption's claim that 'turns are mapped to a 0-100% scale' is not visually represented; the plot shows discrete bins (0-20%, 20-40%, etc.) rather than a continuous scale, which may mislead readers about the data granularity.
- [science] Figure 5: The caption claims the figure shows three subplots (a, b, c) for strategic posture, emotional reactivity, and cultural identity, but the rendered image contains four distinct heatmaps with no subplot labels (a, b, c) to distinguish them.
- [science] Figure 5: The x-axis labels (e.g., 'Avoiding', 'Accommodating', 'Com-Com') do not match the categories described in the caption (strategic posture, emotional reactivity, cultural identity), making it impossible to verify which panel corresponds to which axis.
- [writing] Figure 5: The colorbars for the four heatmaps use inconsistent scales (e.g., -60 to 60 vs -40 to 40), which prevents direct visual comparison of the consensus gain shifts across the different conditions.
- [science] Figure 6: The caption claims to show 'Change in Consensus Gain' (a delta metric), but the y-axis is labeled 'Intervention Effectiveness' and the data trends (e.g., rising curves) contradict the expected behavior of a relative change metric which should center around zero or show degradation/improvement shifts.
- [writing] Figure 6: The caption text is identical to Figure 5's caption, yet the subplots (a, b, c) and data trends differ significantly; the caption fails to describe the specific axes or conditions shown in this figure.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The manuscript relies heavily on specialized terminology that creates a barrier for readers outside the immediate subfield of LLM agent evaluation. While the concepts are sound, the density of jargon without immediate definition or plain-language equivalents reduces accessibility.

First, the term “agentic” is overused as a standalone adjective (e.g., “agentic process,” “agentic deep research,” “agentic scenario curation”). In Section 1 and Section 3.2, this term is used to describe the methodology but is never defined. Does it imply autonomy, multi-agent interaction, or a specific tool-use capability? Replacing “agentic” with “agent-driven” or “multi-agent” would immediately clarify the mechanism for a broader audience.

Second, the phrase “socio-cognitive axes” and “socio-cognitive probing” are introduced in the Abstract and Introduction as if they are standard, self-explanatory categories. While the paper later lists the axes (strategy, emotion, etc.), the umbrella term “socio-cognitive” is not defined. A reader unfamiliar with the specific literature on “social cognition” in AI might not grasp that this refers to “variations in social and psychological factors.” Defining this term or using “social-psychological variations” would be more inclusive.

Third, “trajectory-aware” (Introduction) and “multi-state tracking” (Section 5.2) are technical descriptors that obscure the underlying meaning. “Trajectory-aware” simply means the evaluation considers the sequence of turns, not just the final state. “Multi-state tracking” likely refers to the difficulty of monitoring multiple parties’ hidden states simultaneously. These terms should be replaced with their functional descriptions (e.g., “sequence-aware,” “multi-party state tracking”) to ensure the specific challenge is understood without requiring domain-specific vocabulary.

Finally, the use of “rejection sampling” in Section 3.2 to describe the scenario filtering process is technically accurate but potentially confusing in a non-statistical context. Clarifying that this means “filtering out scenarios that resolve on their own” would prevent misinterpretation by readers who might expect a statistical resampling technique.

Addressing these points by defining terms at first use or substituting them with plainer language will significantly improve the paper’s readability for a general AI audience without sacrificing precision.

Action items.

- [writing] The manuscript relies heavily on specialized terminology that creates a barrier for readers outside the immediate subfield of LLM agent evaluation. While the concepts are sound, the density of jargon without immediate definition or plain-language equivalents reduces accessibility. First, the term “agentic” is overused as a standalone adjective (e.g., “agentic process,” “agentic deep research,” “agentic scenario curation”). In Section 1 and Section 3.2, this term is used to describe the methodolo

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_logical_consistency.md`

The logical flow of the SoCRATES framework is generally sound, with clear premises leading to the proposed methodology. The argument that evaluation is the bottleneck for LLM mediation is well-supported by the cited limitations of existing testbeds (e.g., lack of trajectory awareness, conflation of axes). The introduction of topic-localized evaluation logically follows the identified problem of noise in per-turn scoring.

However, there are minor logical gaps regarding the experimental design and metric definitions that require clarification:

1. **Independence of Axes:** The paper asserts that socio-cognitive axes are applied “independently” to isolate failure modes (Section 3.2). While the text states “no stacking,” the “Cultural Identity” axis generates 6 conditions (US-US, CN-CN, KR-KR, US-CN, US-KR, CN-KR). This inherently involves varying the *composition* of parties (e.g., US vs. CN) which overlaps conceptually with the “Party Composition” axis (2 vs. 3 parties). If the “Party Composition” axis adds a third party, does the “Cultural Identity” axis also apply to 3-party scenarios, or is it strictly 2-party? The current description suggests the 15 conditions are a sum of disjoint sets, but the interaction between “Cultural Identity” (which implies specific party pairings) and “Party Composition” (which changes the number of parties) needs explicit logical separation to ensure the “independent” claim holds. If a 3-party scenario is also culturally mixed, the axes are stacked, violating the isolation premise.
2. **Metric Sensitivity:** The “Consensus Gain” formula normalizes by $(1 - S^{unmed})$. Logically, if the unmediated baseline (S^{unmed}) is already high (e.g., 0.9), the denominator becomes very small (0.1). This makes the metric highly sensitive to minor absolute improvements in the mediated score, potentially exaggerating the “gain” in scenarios where the conflict was already nearly resolved without intervention. The paper does not address whether this normalization introduces a logical bias where mediators appear more effective in “easy” (high consensus) scenarios than in “hard” (low consensus) ones, or if the metric is intended to measure *relative* improvement only.
3. **Baseline Comparison Logic:** The claim that the evaluator “more than doubles” ProMediate’s performance is mathematically correct regarding the trajectory-level correlation (0.823 vs 0.372). However, the non-expert baseline (0.527) is also a significant reference point. The logic of the claim relies on the specific comparison to ProMediate. While valid, the phrasing could be slightly more precise to avoid implying a universal doubling across all comparison points, though this is a minor semantic issue.

Overall, the causal claims regarding the framework’s ability to isolate failure modes are

strong, provided the “independent axes” definition is rigorously clarified in the text to ensure no hidden confounding variables exist between the cultural and party composition axes.

Action items.

- [writing] Independence of Axes: The paper asserts that socio-cognitive axes are applied “independently” to isolate failure modes (Section 3.2). While the text states “no stacking,” the “Cultural Identity” axis generates 6 conditions (US-US, CN-CN, KR-KR, US-CN, US-KR, CN-KR). This inherently involves varying the *composition* of parties (e.g., US vs. CN) which overlaps conceptually with the “Party Composition” axis (2 vs. 3 parties). If the “Party Composition” axis adds a third party, does the “Cultural I
- [writing] Metric Sensitivity: The “Consensus Gain” formula normalizes by $(1 - S^{unmed})$. Logically, if the unmediated baseline (S^{unmed}) is already high (e.g., 0.9), the denominator becomes very small (0.1). This makes the metric highly sensitive to minor absolute improvements in the mediated score, potentially exaggerating the “gain” in scenarios where the conflict was already nearly resolved without intervention. The paper does not address whether this normalization introduces a logical bias wher

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several strong claims regarding the superiority of the SoCRATES framework and the limitations of current LLM mediators. While the methodology is robust, there are instances where the conclusions extrapolate slightly beyond the specific experimental controls presented.

First, the claim in the Introduction and Section 3.2 that the topic-localized evaluator “more than doubles” the performance of ProMediate requires careful qualification. Table 1 in the main text reports a Pearson correlation of $r = 0.823$ for SoCRATES. However, the comparison baseline for ProMediate ($r = 0.372$) is found in Table 2 of the Appendix, which explicitly notes that the evaluation was conducted using Qwen3-235B as the backbone. The main SoCRATES result likely utilizes DeepSeek-V3.2 (as indicated in Section 3.2 and Table 1 of the Appendix). The paper does not explicitly state whether the “doubling” of performance is due to the *evaluation method* (topic-localized vs. per-turn) or the *backbone model* difference. If the backbone contributes significantly to the score, claiming the method alone “doubles” the performance is an overreach. The authors should either provide a direct comparison using the same backbone or rephrase the claim to reflect the combined effect of method and backbone.

Second, the conclusion in Section 4.1 that “scale alone does not order the field” is some-

what overstated. The authors cite the case where Nemotron-3-120B (120B) underperforms Gemma-4-26B (26B) as evidence. However, the data in Table 1 shows a clear trend where larger models within the same family (e.g., Qwen3-30B vs. Qwen3-235B) perform significantly better, and proprietary models (generally larger/more capable) outperform smaller open-source ones. The Nemotron case is an outlier rather than a rule. The claim should be softened to acknowledge that while scale is a strong predictor, architectural differences and training data also play critical roles, rather than implying scale is ineffective.

Finally, the assertion in Section 3.2 that socio-cognitive axes are applied “independently” to isolate failure modes is technically inaccurate regarding the Cultural Identity axis. Table 1 (Appendix) lists six conditions for this axis, including cross-cultural pairings (e.g., US-CN, US-KR). These conditions inherently vary the attributes of *two* parties simultaneously (the mediator’s context or the disputants’ relative positions), which is a joint variation, not a single-axis perturbation. This conflation makes it difficult to attribute performance drops solely to “cultural identity” without considering the interaction effect between the two parties’ specific cultural profiles. The text should clarify that while axes are varied independently of *other* axes (e.g., not stacking culture on top of emotion), the cultural axis itself involves multi-party interaction variations.

Action items.

- [science] The claim that the evaluator ‘more than doubles’ ProMediate’s performance relies on comparing different backbones (DeepSeek vs. Qwen). Clarify if the gain holds when using the same backbone, or rephrase to avoid attributing the full gain to the method alone.
- [writing] The claim that ‘scale alone does not order the field’ is over-generalized from a single outlier (Nemotron vs. Gemma). Data shows strong scale correlation within families. Temper the claim to acknowledge scale is a strong predictor but not the sole determinant.
- [science] The claim that axes are applied ‘independently’ is contradicted by the Cultural Identity axis, which varies two parties’ identities simultaneously (e.g., US-CN). This is a joint variation, confounding the attribution of performance drops to a single ‘cultural’ factor.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

The paper addresses safety and ethics primarily through its “Ethical Considerations” section and the design of its simulation framework. However, several areas require clarification to fully satisfy safety and ethics standards for a publication involving human-subject

validation and real-world data sourcing.

First, there is a potential tension between the claim that scenarios are “synthesized and anonymized” (Ethical Considerations) and the methodology described in “Agentic Scenario Curation,” which involves agents searching the web for “real public disputes” (e.g., the detailed Mount Sinai hospital closure case in Table 5). While the authors state they use fictional names, the underlying structural details of real conflicts (e.g., specific financial losses, regulatory timelines, and stakeholder roles) are derived from real events. The manuscript must clarify the extent of this derivation to ensure that no sensitive or private information (PII) from the original real-world events is inadvertently exposed in the benchmark scenarios. A statement confirming that all real-world seeds have been sufficiently abstracted to prevent re-identification or misuse is needed.

Second, the validation of the evaluator involves human annotators (graduate students and MTurk workers) rating 1,844 dialogue snippets. The paper mentions compensation rates but does not explicitly state whether this research received Institutional Review Board (IRB) approval or an exemption. Given that the study involves human participants in a research context, a formal statement regarding IRB oversight and adherence to ethical guidelines for human-subject research is required.

Finally, the inclusion of “Cultural Identity” axes based on Hofstede’s dimensions (Section 3.3) carries a risk of reinforcing cultural stereotypes. While the authors use these profiles to test mediator adaptation, the paper should briefly discuss the safeguards in place to prevent the benchmark from generating or validating harmful generalizations about specific cultures (KR, US, CN). A sentence acknowledging this risk and the steps taken to mitigate it would strengthen the ethical standing of the work.

Overall, the paper demonstrates awareness of ethical issues, but the specific details regarding data provenance, human-subject oversight, and cultural bias mitigation need to be explicitly addressed in the text.

Action items.

- [writing] The ‘Ethical Considerations’ section states scenarios are ‘synthesized and anonymized’ but the ‘Agentic Scenario Curation’ section explicitly describes using LLM agents to ‘search the web for real public disputes’ (e.g., the Mount Sinai hospital closure case in Table 5). The manuscript must clarify the exact boundary between real-world data ingestion and synthetic generation to ensure no PII or sensitive real-world details are inadvertently leaked in the benchmark scenarios.
- [writing] The validation of the ‘Topic-localized Evaluator’ relies on 1,844 dialogue snippets rated by ‘two expert annotators’ and ‘MTurk workers’ (Section 4, Appendix). The paper lacks a formal statement regarding IRB approval or exemption for this human-subject research. Given the nature of conflict simulation, a statement confirming ethical oversight and informed consent procedures for the annotators is required.
- [writing] The benchmark includes ‘Cultural Identity’ axes using Hofstede profiles for KR, US, and CN (Section 3.3). The paper must explicitly address the risk of reinforcing cultural stereotypes through these synthetic personas. A brief discussion on how the

framework mitigates the potential for the evaluation to produce biased or harmful generalizations about specific cultures is necessary.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a comprehensive framework for evaluating LLM mediators, but the scientific evidence supporting the robustness of its central claims requires strengthening in three key areas: statistical validation of the evaluator’s superiority, experimental replication for the main benchmark, and effect size reporting.

First, the claim that the topic-localized evaluator “more than doubles” the performance of prior baselines (Section 3.2, Table 2) relies solely on point estimates of Pearson correlation ($r = 0.823$ vs $r = 0.372$). The manuscript does not provide confidence intervals for these correlations or a statistical test (such as Fisher’s z-transformation) to confirm that the difference is significant. Given that correlations are sensitive to sample distribution, the absence of a significance test for the *difference* between these dependent correlations leaves the magnitude of the improvement unverified.

Second, the main benchmark results in Section 4 suffer from a lack of replication. The text explicitly states, “Each mediator runs once on every scenario-condition pair” (Section 4). While Appendix E004 demonstrates stability for the “general condition” across three runs (Kendall’s $W = 0.929$), this robustness check is not extended to the 14 socio-cognitive conditions or the domain-specific results. Without variance estimates (standard error or confidence intervals) for the 4,800 main runs, it is impossible to determine if the reported differences between mediators (e.g., the 1.1–2.5 point gap between proprietary and open-source models) are statistically significant or merely stochastic noise inherent to LLM generation.

Finally, the analysis of socio-cognitive impacts (Section 4.2) reports raw consensus gain drops (e.g., “18.9–64.1” for Competing postures) but fails to report effect sizes. In the context of LLM benchmarks where performance can fluctuate significantly, raw percentage point differences are insufficient to support claims of “sharp” variations without normalizing against the baseline variance. The authors should calculate and report effect sizes (e.g., Cohen’s d) to quantify the practical significance of these socio-cognitive failures.

Action items.

- [science] Report confidence intervals or a significance test (e.g., Fisher’s z) for the difference between the new evaluator’s correlation (0.82) and ProMediate’s (0.37) to statistically validate the ‘doubling’ claim in Section 3.2.
- [science] Provide variance estimates (standard error or CI) for the main benchmark

results in Section 4. The current ‘single-run’ design (4,800 runs) makes it impossible to distinguish real mediator differences from stochastic noise without replication data.

- [science] Report effect sizes (e.g., Cohen’s d) for the socio-cognitive axis impacts in Section 4.2. Raw percentage point drops (e.g., 18.9–64.1) are insufficient to claim ‘sharp’ variations without normalizing against baseline variance.

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_statistical_analysis.md`

The statistical analysis in the SoCRATES paper is generally well-structured in its design, particularly the independent variation of socio-cognitive axes to isolate failure modes. However, the reporting of results lacks the necessary statistical rigor to fully support the quantitative claims made.

First, the primary benchmark results presented in Table 1 (Performance by Conflict Domain) and Table 3 (Stability Analysis) rely exclusively on point estimates (means or medians). While the authors mention running simulations multiple times (e.g., “three independent runs” in Appendix E004), the main results tables do not display confidence intervals, standard errors, or standard deviations. Without these measures of dispersion, it is impossible to determine if the reported differences between mediators (e.g., the 1.1–2.5 point advantage of proprietary models) are statistically significant or merely artifacts of stochastic variance. The large half-range reported for Nemotron-3-120B (≈ 6.8) in Table 5 suggests that some performance gaps might not be robust.

Second, the claim that the topic-localized evaluator “more than doubles” the performance of baselines (Section 3.2, Table 2) is based on a comparison of Pearson correlation coefficients ($r = 0.823$ vs $r = 0.372$). The paper does not state whether a statistical test (such as Fisher’s z -transformation) was performed to confirm that this difference is significant. Given the sample size ($N = 1,844$ snippets), the difference is likely significant, but the absence of a p -value or confidence interval for the difference in correlations weakens the claim.

Third, the analysis of socio-cognitive axes involves multiple comparisons across 15 conditions and 8 mediators. The text highlights specific performance drops (e.g., “18.9–64.1” for Competing posture) but does not mention any correction for multiple testing (e.g., Bonferroni or Benjamini-Hochberg). Without such corrections, the risk of Type I errors (false positives) in identifying specific failure modes is elevated.

Finally, while the authors report Kendall’s $W = 0.929$ to demonstrate ranking stability across runs, this metric only assesses rank agreement, not the magnitude of score variance. A more robust statistical approach would involve a mixed-effects model or ANOVA to partition the variance attributable to the mediator, the condition, and the random run

effect, thereby providing a formal test of the mediator’s main effect.

To improve the paper, the authors should supplement the point estimates with confidence intervals, perform and report significance tests for the correlation improvements and performance gaps, and address the multiple comparisons issue in the socio-cognitive analysis.

Action items.

- [science] Report confidence intervals or standard errors for the primary benchmark metrics (Consensus Gain, Timeliness, Effectiveness) in Table 1 and Table 3. Currently, only point estimates are provided, making it impossible to assess the statistical significance of the reported differences between mediators (e.g., the 1.1–2.5 point gap between proprietary and open-source models).
- [science] Clarify the statistical test used to validate the ‘doubling’ claim for the evaluator’s Pearson correlation ($r=0.82$ vs $r=0.37$). A simple comparison of correlation coefficients is insufficient; a Fisher’s z-transformation or a permutation test is required to demonstrate that the improvement is statistically significant and not due to sampling variance.
- [science] Address the multiple comparisons problem in the socio-cognitive probing analysis. With 15 conditions and 8 mediators, numerous pairwise comparisons are made. The paper reports specific drops (e.g., ‘18.9–64.1’) without indicating if these were corrected for family-wise error rate (e.g., Bonferroni) or false discovery rate (FDR).
- [science] The stability analysis reports Kendall’s $W = 0.929$ for rankings across three runs. However, the raw variance (half-range) for Consensus Gain in Table 5 is substantial for some models (e.g., Nemotron-3-120B: ± 6.8). Provide a formal test (e.g., ANOVA or mixed-effects model) to determine if the observed performance differences between mediators are robust against this run-to-run variance.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript demonstrates exceptional writing quality, characterized by clarity, logical flow, and precise academic prose. The authors effectively articulate the complexities of evaluating LLM-mediated conflict resolution, presenting the SoCRATES framework in a manner that is both accessible and rigorous.

The structure is well-organized, guiding the reader seamlessly from the introduction of the problem space through the methodology, validation, and experimental results. The introduction (Section 1) clearly establishes the motivation and contributions, while the Related Work section (Section 2) provides necessary context without becoming cluttered. The transition between the three core components of the framework—Agentic Scenario

Curation, Socio-Cognitive Probing, and Topic-Localized Evaluation—is handled smoothly, with each subsection building logically on the previous one.

Sentence-level grammar is flawless, and the authors make excellent use of technical terminology without sacrificing readability. The definitions of key metrics (e.g., Consensus Gain, Intervention Timeliness) are precise and easy to follow. The use of active voice and concise phrasing enhances the overall impact of the text. For instance, the description of the “Topic-Localized Evaluation” mechanism in Section 3.3 is particularly clear, effectively explaining how the evaluator isolates relevant turns to reduce noise.

The paper also excels in its presentation of results. The text accompanying the tables and figures (e.g., Section 4.1 and 4.2) provides insightful analysis that complements the data, avoiding mere repetition of numbers. The authors successfully convey the nuances of their findings, such as the trade-off between intervention timeliness and effectiveness, with clarity and depth.

Minor stylistic choices, such as the use of bolding for key terms and the consistent formatting of mathematical expressions, contribute to the document’s professional appearance. There are no instances of ambiguous phrasing, grammatical errors, or disjointed paragraphs that would impede understanding.

In conclusion, the writing quality of this paper is outstanding. It meets the highest standards of academic communication, making a complex technical contribution easy to understand and appreciate. No revisions are necessary regarding the writing.