

llmXive Automated Review of Evolving Agents in the Dark: Retrospective Harness Optimization via Self-Preference

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: [qwen.qwen3.5-122b](#)

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper makes strong claims about achieving a 78% pass rate on SWE-Bench Pro from a 59% baseline without external labels. While the results in Table 1 appear consistent with the text, the accuracy of the baseline claim depends on the specific model version used. The citations `openai2025codex` and `openai2026gpt55` are ambiguous; if the baseline refers to the original Codex (GPT-3.5), a 59% score is significantly higher than typical reported results, suggesting a potential mismatch between the cited model and the reported baseline. If the baseline is indeed GPT-5.5, the citation should be clarified to distinguish it from legacy Codex.

The claim of “no ground-truth labels” is technically accurate for the optimization process but risks confusion regarding the evaluation. The final pass rate is determined using the standard SWE-Bench Pro grading script, which relies on ground-truth test cases. The manuscript should explicitly state that the *optimization* is label-free while the *evaluation* uses standard benchmarks to avoid implying the 78% score is self-evaluated.

Additionally, the reliance on `liu2026webtrap` to support the risk of “adversarial content injected mid-task” requires verification. If this source is specific to web security, the generalization to a broad “harness” context may be an overstatement. Finally, the “self-validation” and “self-consistency” steps effectively generate internal labels based on the model’s own judgment. The paper should acknowledge that these internal signals act as proxy ground truth, which is a critical nuance for understanding the limitations of the “label-free” approach.

Action items.

- **[science]** Verify the baseline model for the 59% SWE-Bench Pro score. The paper cites ‘openai2025codex’ and ‘openai2026gpt55’. If the baseline is GPT-3.5 Codex, 59% is anomalously high; if GPT-5.5, clarify the citation to avoid confusion with legacy Codex.
- **[writing]** Clarify the ‘no ground-truth labels’ claim. While optimization is label-free, evaluation uses standard benchmarks. Explicitly distinguish between the optimization loop (self-preference) and the evaluation phase (ground-truth grading) to prevent misinterpretation.
- **[science]** Confirm that citation ‘liu2026webtrap’ specifically supports the claim of ‘adversarial content injected mid-task’ in a general harness context, as the title suggests a web-specific focus that may not generalize.
- **[writing]** Acknowledge that ‘self-validation’ and ‘self-consistency’ generate internal proxy labels. Clarify that the method relies on the model’s internal correctness judgment, which acts as a substitute for external ground truth.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_figure_critic.md`

Figure 1

The figure effectively visualizes the behavioral shifts and performance gains of the proposed method across three benchmarks. However, the caption contains a missing method name, and the legends in the bottom row lack explicit context for the reported percentage changes.

Figure 2

Figure 2 presents clear visualizations but suffers from incomplete caption text ('s DPP') and an undefined baseline ('Vanilla Codex'), which hinder full interpretation of the results and method attribution.

Figure 3

The figure itself is a clear, well-organized diagram of the harness artifacts, but the caption is broken with missing variable values and text placeholders that prevent understanding of the context.

Figure 4

The figure is a schematic diagram that fails to match the caption's description of a quantitative comparison between optimization methods. It lacks the necessary data, axes, or metrics to support the claims made in the text.

Figure 5

The figure effectively visualizes the pipeline described in the caption, but the caption itself contains a placeholder variable, and the embedded plots lack necessary legends and axis labels to interpret the data points and scores.

Action items.

- **[writing]** Figure 1: The caption contains a placeholder 'after .' where the method name (RHO) should be, making the sentence grammatically incomplete and unclear.
- **[science]** Figure 1: The top row y-axis label 'Fraction of tasks solved by step k' is ambiguous; it is unclear if this represents the cumulative success rate (pass@1) or the fraction of currently active tasks solved at that specific step.
- **[writing]** Figure 1: The legends in the bottom row (e.g., 'Verify +61%') lack context for the baseline; it is unclear if these percentages represent the increase over the Vanilla baseline or the absolute difference in action frequency.
- **[science]** Figure 2: The caption claims 'Difficulty or diversity alone trails even random sampling,' but the bar chart in (b) shows 'Difficulty' (0.62) and 'Coverage' (0.58) outper-

forming ‘Random’ (0.64) is false; specifically, Difficulty (0.62) is lower than Random (0.64), but Coverage (0.58) is also lower, yet the text implies both trail random, while the chart shows Difficulty is close to Random but Coverage is worse. Wait, actually $0.62 < 0.64$ and $0.58 < 0.64$, so both do trail random. However, the capt

- **[writing]** Figure 2: The caption contains an incomplete phrase ‘and ‘s DPP balancing both’ — the method name (likely ‘RHO’) is missing before ‘s DPP’, making it unclear which method uses the DPP. Additionally, the ‘Vanilla Codex: 0.59’ baseline in panel (b) is not defined in the caption or legend, leaving its purpose ambiguous.
- **[fatal]** Figure 3: The caption contains unreadable placeholders (‘=-1’, ‘on’, ‘Appendix .’) instead of the specific method name (likely ‘rho=-1’), the verb ‘produced by [method]’, and the appendix number, rendering the figure description incomplete.
- **[fatal]** Figure 4: The rendered image is a schematic diagram illustrating a workflow, but the caption describes a quantitative comparison (‘versus validation-feedback harness optimization’) that is not shown. The figure lacks axes, data points, or metrics to support the comparison claimed in the caption.
- **[science]** Figure 4: The diagram contains placeholder text (e.g., ‘=-1’) and generic icons without specific data or labels, making it impossible to verify the scientific claims regarding ‘validation-feedback’ vs ‘retrospective’ optimization methods described in the caption.
- **[writing]** Figure 5: The caption contains a placeholder ‘=-1’ at the start, likely a missing variable name (e.g., ‘RHO’ or ‘EvoAgent’).
- **[writing]** Figure 5: The ‘Trajectory Distribution’ plot lacks axis tick labels and a legend defining the red vs. pink dots, despite the caption mentioning ‘difficulty-diverse’.
- **[writing]** Figure 5: The ‘Harness Proposal’ section uses specific numerical values (-0.5, +1.4, +0.7) without a legend or axis explaining what these scores represent (e.g., reward delta, win rate).

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The manuscript relies heavily on specialized terminology and mathematical notation that creates a barrier for non-specialist readers. The term “harness” is the most critical omission; it is used throughout the Abstract, Introduction, and Method sections as if it were common knowledge, yet it is never explicitly defined in plain English (e.g., “a collection of tools, prompts, and skills”). This forces the reader to infer the meaning from context.

Additionally, the acronym “DPP” (Determinantal Point Process) appears in Section 5.1

without a full definition at first use, relying on the reader to recognize the statistical method. The concept of “self-preference” is central to the paper’s contribution but is treated as a given rather than explained in accessible terms. The text frequently defaults to dense mathematical notation (e.g., defining $\tilde{\mathbf{r}}_i$ and α) without accompanying prose that explains the *intent* of these calculations in natural language.

To improve accessibility, the authors should:

1. Define “harness” immediately upon its first introduction.
2. Spell out “Determinantal Point Process” before using the acronym DPP.
3. Provide a plain-English explanation of “self-preference” and “coreset” before diving into the algorithmic details.
4. Ensure that every mathematical variable introduced in the text is accompanied by a brief, non-technical description of its role.

These changes are necessary to ensure the paper’s novel contributions are understandable to a broader audience beyond those already familiar with specific agent optimization jargon.

Action items.

- **[writing]** Define “harness” immediately upon its first introduction.
- **[writing]** Spell out “Determinantal Point Process” before using the acronym DPP.
- **[writing]** Provide a plain-English explanation of “self-preference” and “coreset” before diving into the algorithmic details.
- **[writing]** Ensure that every mathematical variable introduced in the text is accompanied by a brief, non-technical description of its role. These changes are necessary to ensure the paper’s novel contributions are understandable to a broader audience beyond those already familiar with specific agent optimization jargon.

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper presents a logically coherent framework for retrospective harness optimization, where the core premise—that self-preference signals derived from past trajectories can substitute for ground-truth labels in optimizing agent behavior—is consistently maintained throughout the methodology and results sections. The causal chain from Coreset Selection (DPP) to Group Rollout (diagnosis) to Best-of-N Proposal (ranking) is well-structured, and the ablation studies in Table 4 (Diagnosis Contribution) effectively support the claim that both self-validation and self-consistency are necessary components for the observed performance gains.

However, a few logical gaps and ambiguities weaken the precision of the conclusions:

1. **Causal Ambiguity in “Long-Horizon” Claims:** The Abstract and Figure 4 caption claim the method “sustains accuracy in long-horizon sessions” and “shifts action mix toward verification.” While the data shows improved pass rates on benchmarks known for long horizons (SWE-Bench), the paper does not explicitly present data correlating the *duration* or *step count* of sessions with the optimized harness. The logical leap from “higher success on long tasks” to “sustains accuracy in long sessions” (implying stability over time or duration) is not fully supported by the provided metrics. The term “sustains” is ambiguous: does it mean the agent doesn’t degrade over time, or that it can complete longer tasks? This needs clarification to ensure the conclusion follows strictly from the evidence.
2. **Evaluation vs. Optimization Logic:** The paper repeatedly emphasizes “no ground-truth labels used” (Abstract, Fig 1). While this is true for the *optimization signal* (the self-preference mechanism), the *evaluation* of the final harness relies entirely on ground-truth test suites (SWE-Bench, etc.). The logical distinction between the *training signal* (label-free) and the *evaluation metric* (label-dependent) is blurred. A reader might incorrectly infer that the method is entirely unsupervised in a way that precludes any external validation, which is not the case. The text should explicitly state that while the *harness evolution* is label-free, the *performance assessment* still requires standard benchmarks with ground truth.
3. **Threshold Logic in Best-of-N:** The acceptance condition $S_j > 0$ is clearly defined, but the text in Appendix e002 mentions “mean zero is rejected.” Given the ranking scale of $[-10, 10]$, a mean of 0 represents a tie between the candidate and the baseline. The logic that a tie should be rejected is sound, but the paper should explicitly address the probability of ties and whether the “strictly positive” condition is robust against statistical noise in the ranking scores. If the ranking model is uncertain, a score of 0.1 might be indistinguishable from 0, yet the logic treats them as distinct. A brief discussion on the robustness of this threshold would strengthen the logical consistency of the acceptance mechanism.

Overall, the internal logic of the proposed method is sound, but the interpretation of results regarding “long-horizon” stability and the distinction between optimization and evaluation signals requires tighter phrasing to avoid overclaiming.

Action items.

- **[writing]** Causal Ambiguity in “Long-Horizon” Claims: The Abstract and Figure 4 caption claim the method “sustains accuracy in long-horizon sessions” and “shifts action mix toward verification.” While the data shows improved pass rates on benchmarks known for long horizons (SWE-Bench), the paper does not explicitly present data correlating the *duration* or *step count* of sessions with the optimized harness. The logical leap from “higher success on long tasks” to “sustains accuracy in long sessions” (impl
- **[writing]** Evaluation vs. Optimization Logic: The paper repeatedly emphasizes “no ground-truth labels used” (Abstract, Fig 1). While this is true for the *optimization signal* (the self-preference mechanism), the *evaluation* of the final harness relies entirely on ground-truth test suites (SWE-Bench, etc.). The logical distinction between the *training signal* (label-free) and the *evaluation metric* (label-dependent) is blurred. A reader might incorrectly infer that

the method is entirely unsupervised in

- **[writing]** Threshold Logic in Best-of-N: The acceptance condition $S_j > 0$ is clearly defined, but the text in Appendix e002 mentions “mean zero is rejected.” Given the ranking scale of $[-10, 10]$, a mean of 0 represents a tie between the candidate and the baseline. The logic that a tie should be rejected is sound, but the paper should explicitly address the probability of ties and whether the “strictly positive” condition is robust against statistical noise in the ranking scores. If the ranking model is u

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several claims that extend beyond the immediate evidence provided, particularly regarding the nature of the “label-free” setting and the comparative superiority over validation-based methods.

First, the Abstract and Introduction repeatedly emphasize that the method operates “without ground-truth feedback” and is “label-free.” While technically true that no human-annotated labels are used, the method relies heavily on an LLM judge to perform “self-validation” and “self-consistency” checks (Sec 3.2, Algorithm 1). These checks generate scores and diagnostic instructions that function as a reward signal. By not explicitly framing this as “proxy-labeling” or “self-supervised reward modeling,” the paper risks misleading readers into thinking the method is entirely free of supervision signals. The limitations section mentions “trusts past trajectories” but does not adequately address the risk of the LLM judge hallucinating preferences or reinforcing its own biases, which is a critical limitation of any self-preference system.

Second, the claim in the Contributions and Section 5 that the method “outperforms... validation-feedback under comparable budget” is an overreach. Table 2 compares the proposed method (1 round) against Meta-Harness (1 round and 10 rounds). While the 1-round comparison favors the proposed method (0.78 vs 0.62), the 10-round Meta-Harness result (0.80) actually exceeds the proposed method’s performance. The paper frames the 1-round comparison as the definitive result, implying that validation-feedback is inherently less efficient or effective. However, the data suggests that validation-feedback is simply more compute-intensive to reach a higher ceiling. The claim should be nuanced to reflect that the method achieves *competitive* performance with *fewer* optimization rounds, rather than broadly “outperforming” the category.

Finally, the behavioral analysis in Section 5.1 and Figure 5 claims that the method “alters behavior patterns” and “sustains accuracy.” While the figures show trends, the text lacks statistical rigor (e.g., significance tests) to confirm that these shifts are not due to random variance, especially given the small coreset size ($k = 10$) used for optimization. The claim of “sustained accuracy” is particularly strong given the lack of long-term stability testing across multiple optimization rounds or different seed trajectories.

Action items.

- **[science]** The ‘label-free’ claim overreaches by ignoring that LLM judges act as implicit, uncalibrated reward signals. The paper must explicitly acknowledge this proxy-labeling risk and the potential for hallucinated preferences.
- **[science]** The claim of ‘outperforming validation-feedback’ is unsupported; Meta-Harness (10 rounds) achieves 0.80 vs. 0.78. Qualify the claim to ‘outperforms single-round baselines’ or ‘achieves comparable performance with fewer rounds’.
- **[science]** Claims of ‘altered behavior patterns’ lack statistical significance testing (p-values/CIs) for the observed shifts in action mix, risking over-interpretation of variance as robust change.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

The paper addresses a critical area of AI safety: the potential for autonomous agents to optimize their own behavior in ways that amplify existing biases or introduce unsafe procedures without external oversight. The authors correctly identify the dual-use risk in the “Ethics Statement” (Section e000), noting that modifying behavior based on model-generated judgments risks “amplification of mistaken preferences or unsafe procedures.” However, the current mitigation strategies proposed—specifically “audit logs, human approval for sensitive edits, and domain-specific safety checks”—are too vague to be actionable or reassuring in a high-stakes deployment context.

The primary safety concern lies in the core methodology: “Retrospective Harness Optimization” relies entirely on the agent’s internal “self-preference” and “self-consistency” signals (Section 5.2) to guide optimization. Without ground-truth labels or external safety validators, there is a significant risk that the agent could converge on a “safe” local optimum that actually encodes harmful heuristics or exploits loopholes in the task environment (reward hacking). For instance, if the training trajectories contain subtle adversarial injections (as noted in the Limitations citing Liu et al., 2026), the self-reinforcing loop could entrench these behaviors rather than filter them. The paper currently lacks a dedicated analysis of these failure modes or a discussion on how the “self-validation” step specifically guards against safety violations versus mere task failure.

Furthermore, the “Ethics Statement” mentions “domain-specific safety checks” but does not define what these entail. Given that the method optimizes tools and scripts (e.g., **repair-verify** in Appendix e001), there is a risk of generating code that could be misused or that inadvertently compromises system integrity. The review requires a more concrete description of the proposed human-in-the-loop safeguards. Specifically, what constitutes a “sensitive edit”? How are audit

logs structured to detect drift toward unsafe behaviors?

Finally, the use of “judge-relaxation switches” in the GAIA-2 dataset (Appendix e002) raises questions about the robustness of the safety evaluation. If these switches were active, the reported improvements might be artifacts of relaxed constraints rather than genuine safety-aligned optimization. The authors must clarify the experimental configuration regarding these switches and discuss their implications for the safety claims made in the paper.

Action items.

- **[writing]** The Ethics Statement (Section e000) acknowledges risks of amplifying unsafe procedures but lacks a concrete mitigation strategy for the ‘adversarial content’ risk cited from Liu et al. (2026). Explicitly detail the proposed ‘audit logs’ and ‘human approval’ mechanisms (e.g., specific triggers, human-in-the-loop checkpoints) to satisfy safety review standards.
- **[science]** The method relies on ‘self-preference’ and ‘self-validation’ to optimize agent behavior without ground truth. This creates a high risk of reward hacking or reinforcing latent biases in the training trajectories. The paper must include a dedicated subsection analyzing potential failure modes where the agent optimizes for ‘plausible but unsafe’ behaviors that pass internal consistency checks.
- **[writing]** The dataset specifications (Appendix e002) mention ‘judge-relaxation switches’ for GAIA-2. The Ethics Statement must clarify if these switches were used during the reported experiments and, if so, how this impacts the validity of the safety claims regarding ‘unsafe procedures’ or ‘adversarial content’ injection.

Scientific Evidence

Weights the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a compelling methodology for retrospective harness optimization, but the scientific evidence supporting the magnitude and robustness of the reported gains requires strengthening.

Sample Size and Statistical Rigor: The primary claim is a 19% absolute improvement on SWE-Bench Pro (0.59 to 0.78) based on a held-out test set of 100 tasks (Appendix **app:datasets**). While the sample size (N=100) is reasonable for a benchmark, the paper reports only point estimates. There is no mention of multiple random seeds for the optimization process or the evaluation. LLM-based agents exhibit high variance; a single run could be an outlier. The absence of standard deviations, confidence intervals, or statistical significance tests (e.g., McNemar’s test for paired proportions) makes it difficult to assess whether the observed gains are robust or potentially due to chance. The ablation study in Table **tab:diagnosis** also lacks variance metrics, making it hard to judge the significance of the drop when self-consistency is removed.

Experimental Controls and Replication: The experimental design relies on a single round of optimization. The comparison with Meta-Harness (Table [tab:meta-harness](#)) is interesting but potentially confounded by the “matched budget” definition. Meta-Harness is run for 10 rounds to reach 0.80, while `\method{}` is a single pass. While the authors argue this is a fair comparison of “single-pass” efficiency, the lack of a multi-round baseline for `\method{}` (if computationally feasible) or a clearer breakdown of the variance in the 10-round Meta-Harness results would strengthen the claim that the single-pass approach is superior. Furthermore, the coreset selection uses a DPP based on difficulty scores from an LLM judge. The stability of these difficulty scores across different judge prompts or seeds is not discussed. If the coreset is unstable, the resulting harness optimization might be overfitting to a specific subset of “hard” tasks that happened to be selected.

Alternative Explanations: The paper attributes the improvement to the specific harness artifacts (tools, skills) generated. However, the “Best-of-N” selection process involves re-solving the coreset tasks with candidate harnesses. It is possible that the improvement stems partly from the *selection* of a harness that happens to work well on the specific coreset tasks, rather than a generalizable improvement in the agent’s behavior. The paper mentions that the chosen candidate is better than the lowest, but without a distribution of scores for all N candidates or a comparison against a random selection baseline (beyond the “Mean” in Table [tab:bestofn](#)), the risk of overfitting to the coreset remains a plausible alternative explanation. The claim that the method “sustains accuracy in long-horizon sessions” (Abstract) is supported by Figure [fig:steps](#), but the figure’s data (likely showing step counts or success rates over time) is not fully described in the text with statistical backing.

In summary, while the results are promising, the lack of statistical significance testing, variance reporting, and robustness checks against overfitting to the coreset prevents a full acceptance of the central claims regarding the magnitude of the improvement.

Action items.

- **[science]** Report statistical significance (p-values or confidence intervals) for the +0.19 gain on SWE-Bench Pro and other benchmarks. With N=100 test tasks, a 19% absolute jump is large, but the variance across tasks is unknown. Without error bars or significance tests, the robustness of the claim against random fluctuation is unverified.
- **[science]** Clarify the independence of the coreset selection and evaluation sets. The paper states the first 100 tasks are used for training (coreset selection) and the next 100 for testing. However, the coreset is selected from the *past trajectories* of the training set. If the ‘past trajectories’ were generated by the same model version being optimized, ensure there is no data leakage where the test set influenced the training distribution or the difficulty scores used for DPP selection.
- **[science]** Provide effect size metrics beyond raw pass rates. The paper reports absolute improvements (e.g., +0.19) but lacks normalized effect sizes (e.g., Cohen’s d) or variance measures (standard deviation of pass rates across multiple seeds if applicable). Given the stochastic nature of LLM rollouts, a single run per method may not represent the true expected performance.

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The manuscript presents a novel framework for retrospective harness optimization, but the statistical rigor of the reported results is insufficient to support the strong claims of consistent improvement.

Lack of Variance and Significance Testing: The primary results in Table 1 (Main Results) and Table 4 (Best-of-N) report single point estimates (e.g., 0.78 vs 0.59 for SWE-Bench Pro). In LLM agent evaluation, performance is highly stochastic due to sampling variance and non-deterministic environments. The paper fails to report standard deviations, confidence intervals, or results from multiple independent runs (seeds). Without this, it is impossible to determine if the observed +0.19 improvement is statistically significant or an artifact of a single lucky run. The “Mean” and “Std” columns in Table 4 are ambiguous; it is unclear if these are calculated over the N=3 candidates in a single run or over multiple experimental seeds. If the latter, the sample size (N=3) is likely too small to estimate population variance reliably.

Hyperparameter Sensitivity: The coreset selection relies on a Determinantal Point Process (DPP) with a fixed difficulty/diversity trade-off parameter $\theta = 0.7$ (Section 5.1, Appendix Table 1). The paper asserts this balances difficulty and diversity but provides no ablation study or sensitivity analysis showing how performance varies with θ . A single hyperparameter setting without variance reporting suggests the results may not be robust.

Best-of-N Selection Bias: The “Best-of-N” strategy (Section 5.3) selects the candidate with the highest self-preference score. The paper reports that the “Chosen” candidate outperforms the “Lowest” (Table 4), but this is a tautological result of selection. The critical statistical question is whether the *selected* candidate generalizes to the held-out test set better than the baseline *on average* across multiple seeds. The current presentation conflates the selection mechanism’s internal score with external test performance without providing the necessary distributional data.

Recommendation: The authors must re-run the experiments with at least 3-5 independent random seeds for all datasets. They should report mean $\pm$ standard deviation for all pass rates in Table 1 and Table 4. A statistical test (e.g., paired t-test or bootstrap) should be performed to confirm the significance of the improvements over baselines. Additionally, a sensitivity analysis for the DPP weight θ is required to demonstrate the robustness of the coreset selection.

Action items.

- **[science]** Report statistical significance for the reported performance gains (e.g., +0.19 on SWE-Bench Pro). The paper presents point estimates in Table 1 but lacks confidence intervals, p-values, or variance estimates across multiple random seeds to confirm the gains are not due to stochasticity.
- **[science]** Clarify the statistical basis for the ‘Best-of-N’ selection. The paper reports a single

‘Chosen’ score in Table 4 against a ‘Mean’ and ‘Lowest’, but does not specify if the ‘Mean’ is derived from a distribution of $N=3$ candidates across multiple seeds or a single run. Re-run experiments with multiple seeds to provide standard deviations for all reported metrics.

- **[science]** The coreset selection uses a DPP with a fixed weight $\theta=0.7$. Provide a sensitivity analysis or justification for this specific hyperparameter choice. Without reporting performance variance across different θ values, the robustness of the coreset selection strategy remains unverified.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript exhibits significant structural and completeness issues that currently impede readability. The most critical flaw is a hard truncation in the LaTeX source within Section 4 (“Problem Setting”). The text cuts off mid-sentence at “ $\tau = \mathrm{so}$ ” (end of chunk e003), leaving the definition of the execution process and the subsequent algorithmic description incomplete. This prevents the reader from understanding the formal problem setup.

Furthermore, the document appears to be a concatenation of multiple draft versions or chunks rather than a cohesive final manuscript. There are duplicate sections, particularly in the Introduction and Experiments, where content is repeated with slight variations in phrasing and formatting (e.g., Figure captions and table structures differ between chunks e001 and e002). This fragmentation disrupts the logical flow and creates confusion regarding the definitive version of the text.

While the prose within the complete segments is generally clear and the technical terminology is used appropriately, the manuscript cannot be evaluated as a finished paper in its current state. The authors must merge the fragmented sections, remove redundancies, and ensure the text is fully complete, particularly the mathematical definitions in Section 4, before the writing quality can be considered acceptable.

Action items.

- **[writing]** The LaTeX source contains a critical truncation error in Section 4 (Problem Setting). The text ends abruptly at “ $\tau = \mathrm{so}$ ” (end of e003), cutting off the definition of the execution process and the subsequent algorithm description. The manuscript must be completed to be readable.
- **[writing]** The document structure is fragmented with duplicate content. Section 2 (Introduction) and Section 5 (Experiments and Results) appear to be repeated or split across chunks (e000, e001, e002) with inconsistent formatting and redundant figure captions. The authors should consolidate these into a single, linear narrative flow.
- **[writing]** In the ‘Problem Setting’ section (e003), the mathematical notation for the trajectory

generation is incomplete due to the truncation. Ensure all equations are fully rendered and the text following the equation is present to maintain logical flow.