


User

"Which direction is the car moving?"

Round 1

$t = 1$



full RGB frame

ViT



256 visual tokens

VL-interaction model

`</silence>`

model stays silent

Round 2

$t = 2$



residual + motion vector

P-Tokenizer



16 visual tokens

~16× fewer
tokens

VL-interaction model

`</response>`

"The car is moving towards the left."