

llmXive Automated Review of MemSlides: A Hierarchical Memory Driven Agent Framework for Personalized Slide Generation with Multi-turn Local Revision

llmXive automated review panel

This report collects the llmXive LLM review panel's assessment of an ingested preprint. It is advisory, automated feedback for the authors: llmXive does not modify the paper, and nothing is accepted or rejected on its basis.

paper_reviewer_claim_accuracy

Verdict: minor_revision

The review focuses on the accuracy of factual claims and their alignment with the provided data tables.

****1. Factual Error in General-Quality Claims (Section 5.1.1 / Appendix A.4):**** The text states: "Ours also achieves the highest Avg. on GPT-5 and GLM-5." - ****Verification:**** Referencing 'tables/profile_memory_v6_ppteval_main_table.tex' (Table 3): - ****GPT-5:**** MemSlides Avg = 4.17, DeepPresenter Avg = 3.99. (Claim holds). - ****GLM-5:**** MemSlides Avg = 3.74, DeepPresenter Avg = 3.86. (Claim ****fails****; DeepPresenter is higher). - ****Impact:**** This is a direct contradiction between the text and the data table. The claim of "highest Avg" for GLM-5 is unsupported by the provided evidence. The text must be corrected to reflect that MemSlides is competitive or second-best on GLM-5.

****2. Nuance in Tool-Memory Aggregation (Section 5.1.2):**** The text claims tool-memory injection reduces "Time to First Correct Edit" from 609.5s to 242.5s. - ****Verification:**** The aggregate numbers in 'tables/tool_memory_main_table.tex' (Table 4) match the text. - ****Assessment:**** While the numbers are correct, the phrasing implies a uniform reduction. Table 4 shows the "No Injection" baseline varies significantly by model (e.g., 234.2s for GPT-5 vs 968.2s for Gemini). The claim is accurate as an aggregate but could be misinterpreted as a per-model guarantee. A brief clarification that this is an aggregate effect across heterogeneous baselines would improve precision.

****3. Persona-Alignment Claims (Section 5.1.1):**** The claim that MemSlides achieves “all-column wins” over baselines on GLM-5 and Gemini 3.1 Pro is verified against ‘tables/profile_memory_v6_bestof_main_table.tex’ (Table 1) and is factually accurate. No revision is needed for this specific point, though readers should note these are point-estimate wins without reported statistical significance in the main text.

****Conclusion:**** The primary issue is the factual error regarding the GLM-5 average score. The rest of the claims are supported by the data, though one aggregate claim could benefit from slight clarification regarding baseline heterogeneity.

- [writing] Correct the claim in Section 5.1.1/Appendix A.4 that MemSlides achieves the highest Avg. on GLM-5. Table 3 shows DeepPresenter (3.86) exceeds MemSlides (3.74) for GLM-5 Avg. Update text to reflect this accurately.
- [writing] Clarify in Section 5.1.2 that the reported reduction in ‘Time to First Correct Edit’ (609.5s to 242.5s) is an aggregate across models with heterogeneous baselines, not a uniform per-model reduction, to avoid overgeneralization.

paper_reviewer_figure_critic

Verdict: major_revision_science

Figure 1

Figure 1 attempts to showcase template-guided generation but fails to substantiate the claim because the generated content appears unrelated to the template structure shown above it. Additionally, the text in the generated slides is too small to read, hindering the evaluation of the generation quality.

Figure 2

Figure 2 is a clear, well-structured qualitative illustration that effectively contrasts ‘Without Tool Memory’ and ‘With Tool Memory’ behaviors. The visual elements, including icons, arrows, and text labels, are legible and directly support the caption’s description of localized versus broad editing trajectories.

Figure 3

Figure 3 effectively illustrates the concept of delayed preference carryover with clear workflow diagrams, but the text within the ‘Limitations and Future Work’ boxes is too small to read comfortably, and the specific mechanism of memory injection into the prompt could be depicted more explicitly.

Figure 4

The figure provides a clear visual overview of the memory workflow, but the caption contains a grammatical error where the framework name is missing, and there is a minor notation inconsistency regarding the state update subscript.

Figure 5

The figure provides a clear visual workflow of the Plan-Act-Guard pipeline, but the caption contains a missing framework reference and a typo in the label.

Figure 6

The figure effectively visualizes the user profile memory lifecycle described in the caption, showing the flow from long-term storage to temporary working memory and back. However, the caption contains a grammatical error omitting the system name, and there is a minor terminology mismatch between the diagram labels and the caption text.

Figure 7

The figure effectively visualizes the tool memory flow described in the caption, but the caption text itself contains a grammatical error with a missing noun, and some internal diagram labels are rendered at a low resolution.

Figure 8

The figure effectively contrasts the two methods, but the caption contains a missing subject in the second sentence and inaccurately describes the DeepPresenter result as 'altering' non-target regions when the image shows a complete removal of the formula block.

Figure 9

The figure effectively visualizes the consolidation of profile preferences into reusable patterns, but the caption fails to account for the 12 panels shown (referencing only 'six jobs') and does not explicitly map the bottom-row labels to the specific patterns described.

- [science] Figure 1: The generated slides contain specific technical content (e.g., 'BLEU 28.4', 'WMT14', 'EN-DE BLEU') that is not present in the corresponding template slides above them. The caption claims these are 'matched persona/template generations,' but the figure fails to demonstrate how the template structure was used to generate this specific content, making the 'template-guided' claim unsubstantiated by the visual evidence.
- [writing] Figure 1: The text within the 'Generated Slide' panels (e.g., the table in panel 3, the bullet points in panel 2) is extremely small and illegible at the current resolution, preventing verification of the content or the quality of the generation.
- [science] Figure 3: The 'Injected memory' workflow shows a database icon labeled 'Memory' feeding into a 'Prompt' document, but the diagram lacks a clear visual representation of the 'Memory Injection' mechanism itself (e.g., an arrow or process step showing how the stored rule is retrieved and inserted into the prompt).
- [writing] Figure 3: The text inside the 'Limitations and Future Work' boxes is extremely small and partially illegible, making it difficult to read the specific bullet points describing the limitations.
- [writing] Figure 4: The caption reads 'Overview of ' with a missing noun (likely 'MemSlides' or 'the framework'), rendering the sentence grammatically incomplete.
- [writing] Figure 4: The caption uses the notation '\$s_t-1\$' for the updated state, whereas the diagram explicitly labels the output as '\$S_{t-1}\$'; the caption should match the figure's subscript formatting.
- [writing] Figure 5: The caption contains a grammatical error and missing reference: 'Lo-

calized modify execution in .' should specify the framework name (e.g., 'in MemSlides').

- [writing] Figure 5: The caption label '[localized_tool_pipeline]' contains a typo ('pipeline' instead of 'pipeline').
- [writing] Figure 6: The caption contains a grammatical error ('User profile memory lifecycle in .') where the system name is missing after the preposition 'in'.
- [writing] Figure 6: The diagram uses the term 'TempPreference' in the purple boxes, but the caption refers to this as 'active temporary memory'; aligning these terms would improve clarity.
- [writing] Figure 7: The caption contains a grammatical error and missing noun in the opening sentence ('Tool memory flow in .'), likely a placeholder for the framework name (e.g., MemSlides) that was not filled in.
- [writing] Figure 7: The diagram contains several instances of small, low-resolution text (e.g., inside the 'Memory Consolidation' and 'Operation' boxes) that is difficult to read and may be illegible in lower-resolution views.
- [science] Figure 8: The caption states 'DeepPresenter can satisfy the target change while altering non-target regions,' but the visual evidence in Panel B (DeepPresenter) shows the formula block was removed entirely rather than altered, and the layout was rewritten, contradicting the claim of merely 'altering' non-target regions.
- [writing] Figure 8: The caption contains a grammatical error where the subject is missing in the second sentence ('instead applies a targeted patch...'); it should explicitly name 'MemSlides' to match the figure labels.
- [writing] Figure 9: The figure contains 12 distinct panels (4 columns x 3 rows) but the caption only describes 'six repeated jobs', creating a mismatch between the visual evidence count and the described experimental scope.
- [writing] Figure 9: The bottom row labels ('Inherited guardrail', 'Stable execution closure', 'Boundary-Centered Overview', 'Reusable Checklist') are not explicitly defined in the caption, which only lists the four patterns in a general sentence without mapping them to the specific visual examples.

paper_reviewer_jargon_police

Verdict: minor_revision

The manuscript relies heavily on specialized terminology that, while standard in specific agent-memory sub-fields, creates barriers for a broader audience. The most significant issue is the use of undefined acronyms and statistical jargon. In Appendix A.1, the term "arm" is used repeatedly to describe the experimental conditions (e.g., "anonymous arm label," "memory condition are hidden"). This is standard clinical trial terminology but is opaque to general readers; "condition" or "variant" would be more accessible. Similarly, the

acronym "W-L-T-NA" appears in Section 5.2 and Appendix A.3 without ever being defined as "Win-Loss-Tie-Not-Available," forcing the reader to guess the meaning or search for it.

In Section 3.1, the term "round-0" is introduced to denote the initial generation stage. While the context implies it is the first turn, the specific notation "round-0" (as opposed to "initial generation" or "turn 0") is a piece of internal jargon that should be explicitly defined upon first mention to ensure clarity for readers not steeped in multi-turn agent protocols.

Furthermore, the paper frequently uses metaphorical jargon to describe technical mechanisms. Phrases like "bounded repair surface" (Section 3.2) and "scoped slide-local revision" (Abstract and Section 1) are dense. "Repair surface" is not a standard term in presentation generation and acts as a barrier; replacing it with "the specific set of slides and elements to be edited" would improve readability without losing precision. The term "re-contextualizing" in Section 1 is also vague; "re-reading" or "re-processing" the full deck would be a plainer alternative.

Finally, the distinction between "active temporary memory" and "working memory" is sometimes blurred. While "working memory" is a standard cognitive science term, the paper's specific definition ("session-scoped state layer") should be clearly distinguished from the general concept early on to avoid confusion. The authors should audit the text for similar instances of field-specific shorthand and replace them with plain English equivalents or provide immediate definitions.

- [writing] The manuscript relies heavily on specialized terminology that, while standard in specific agent-memory sub-fields, creates barriers for a broader audience. The most significant issue is the use of undefined acronyms and statistical jargon. In Appendix A.1, the term "arm" is used repeatedly to describe the experimental conditions (e.g., "anonymous arm label," "memory condition are hidden"). This is standard clinical trial terminology but is opaque to general readers; "condition" or "variant" would

paper_reviewer_logical_consistency

Verdict: minor_revision

The paper presents a logical framework where separating memory types (profile, working, tool) should theoretically improve personalization and editing efficiency. However, several conclusions in the experimental section do not strictly follow from the presented premises or data.

First, in Section 5.1.2, the authors claim that "tool-memory injection is associated with stronger... post-edit verification." The mechanism proposed is that tool memory stores "reusable execution experience." Logically, storing past execution traces does not inherently compel an agent to perform a *verification* step unless the memory explicitly encodes "verification is required" as a rule. The paper asserts the outcome (higher verification rates) but does not fully articulate the causal mechanism by which retrieving a "tool chain" forces the agent to insert a verification call rather than just skipping to the next edit. The conclusion that memory *causes* the verification behavior is an inference that requires a

more explicit premise about how the memory retrieval logic is structured.

Second, the claim in Section 5.1.1 that "User profile memory yields broad, not marginal, alignment gains" is partially contradicted by the data in Table 1. The text states that with GPT-5, the method "remains ahead of SlideTailor... and ahead of DeepPresenter on Content, Visual, and Specificity, while DeepPresenter has slightly higher Structure." However, the subsequent analysis paragraph claims, "The strongest evidence is the joint movement of Structure and Specificity... It indicates that routed long-term profiles help decide page roles..." This is a logical inconsistency: if the method performs *worse* on Structure for GPT-5 (7.33 vs 7.56 for DeepPresenter), one cannot cite the "joint movement" of Structure as evidence of the method's success for that model family. The generalization of "broad gains" is weakened by this specific failure case which is not adequately addressed in the causal explanation.

Third, regarding the efficiency claims in Section 5.1.2, the paper argues that tool memory reduces "Core Tool Time" by avoiding "repeated full-deck regeneration." However, the metric definition explicitly excludes "inspection" tools. Since the results show a significant increase in "Strict Verify" (which is an inspection action), there is a potential logical gap: if the memory causes the agent to verify more often, and verification is excluded from the time metric, the "Core Tool Time" reduction might be an artifact of the metric definition rather than a true efficiency gain in the *total* work performed. The conclusion that the system is "more efficient" relies on the premise that the excluded time is negligible, which is not proven. The causal link between memory and *net* efficiency is therefore not fully supported by the specific metric definitions provided.

- [science] Clarify the causal mechanism in Section 5.1.2: how does retrieving 'tool memory' explicitly force the agent to perform a 'verify' step? The premise that memory stores experience does not logically entail a policy change to increase verification rates without further explanation.
- [science] Resolve the contradiction in Section 5.1.1: The text claims 'joint movement of Structure and Specificity' proves planning gains, yet Table 1 shows MemSlides underperforms DeepPresenter on Structure for GPT-5 (7.33 vs 7.56). The conclusion does not follow from the data for this model.
- [science] Address the metric definition gap in Section 5.1.2: 'Core Tool Time' excludes inspection tools, yet 'Strict Verify' (an inspection) increases. The claim of net efficiency gain is unsupported if the excluded verification time outweighs the saved search time.

paper_reviewer_overreach

Verdict: minor_revision

The paper exhibits moderate overreach in extrapolating diagnostic and qualitative results into broad claims of reliability and universal improvement.

In the **Abstract** and **Conclusion**, the authors assert that tool memory "improves closed-loop modify behavior" and "enhances localized modify reliability." However,

the data in **Table 1 (tool_memory_main_table.tex)** and the pair-level details in **Appendix Table 3 (tool_memory_pair_detail_table.tex)** reveal significant heterogeneity. Specifically, Pair P2 (GPT-5) and Pair P7 (Gemini) show losses in Closed-Loop Completion and First Correct Edit time, respectively. The text in **Section 5.1.2** acknowledges this (“not by winning every pair”) but the Abstract’s phrasing (“improves... behavior”) suggests a consistent positive effect that the data does not fully support. The claim should be restricted to “improves reliability in specific diagnostic settings” or similar qualifiers to avoid overgeneralization.

Regarding **persona alignment**, the **Conclusion** states the framework “supports round-0 persona alignment” as a settled fact. While **Table 1 (profile_memory_v6_bestof_main_table.tex)** shows strong gains for GLM-5 and Gemini, the GPT-5 results are mixed (e.g., Structure score of 7.33 vs. DeepPresenter’s 7.56). The **Analysis** section correctly notes this nuance, but the **Abstract** and **Conclusion** present the improvement as uniform. The language should be adjusted to reflect that gains are “significant for certain model families” or “model-dependent” to accurately reflect the variance in **Table 1**.

Finally, the claim that working memory “carries over preferences” (Abstract) relies entirely on **Appendix Figure 2 (working_memory_title_color_carryover.pdf)** and **Appendix Figure 3 (working_memory_summary_box_style_carryover.pdf)**. These are qualitative illustrations, not quantitative evidence. The paper currently lacks a metric to measure “carryover” success rate or latency. The text should explicitly frame these as “qualitative demonstrations” rather than evidence of a functional capability, preventing the reader from inferring statistical significance where none exists.

- [science] The claim that tool memory ‘improves closed-loop modify behavior’ (Abstract) overgeneralizes the diagnostic matched-pair results. The data shows non-monotonic pair-level outcomes (e.g., P2 and P7 losses in Table 1 pair details). The text must explicitly qualify these gains as ‘diagnostic’ and avoid implying universal reliability improvements without broader distributional evidence.
- [writing] The conclusion states MemSlides ‘supports round-0 persona alignment’ as a definitive outcome, but the GPT-5 results in Table 1 show mixed performance (e.g., lower Structure than DeepPresenter). The conclusion should be tempered to reflect that persona alignment improvements are model-dependent and not uniformly dominant across all dimensions.
- [writing] The abstract claims working memory ‘carries over preferences’ based on qualitative cases (Appendix Fig 2). This extrapolates beyond the evidence scope, as no quantitative metric validates carryover. The text should clarify that working memory benefits are illustrated qualitatively rather than statistically proven.

paper_reviewer_safety_ethics

Verdict: minor_revision

The manuscript addresses safety and ethics primarily in the "Broader Impacts and Responsible Use" subsection of the Appendix. The authors correctly identify key risks associated with persistent personalization, including the potential storage of sensitive user habits and the misuse of agents to produce misleading content. However, the current treatment of these risks is largely aspirational, deferring critical governance mechanisms to future work.

From a safety perspective, the system's core functionality involves the persistent storage and retrieval of user profiles (Section 3.2, Eq. 4). The paper currently lacks a concrete description of the technical safeguards implemented in the prototype to protect this data. Specifically, the review requires clarification on whether the system includes user-visible controls for inspecting, editing, or deleting stored memory entries, as well as details on data encryption and retention policies. Without these explicit details, the claim that the system is safe for "controlled generation" remains unsubstantiated regarding data privacy.

Furthermore, the construction of the "profile bank" (Appendix, Section 'Profile-Bank Construction') involves 30 persona-intent entries. While the text describes these as "controlled authoring interactions," it does not explicitly confirm that no real user data was utilized in their creation or validation. If any real-world interaction logs were used to seed the profile bank or train the consolidation logic, the manuscript must include a statement regarding IRB approval or informed consent. If the data is entirely synthetic, this distinction should be made explicit to prevent ambiguity regarding human subject research.

The "Limitations" section (Section 7) appropriately notes the lack of real-user deployment studies but should be strengthened to explicitly state the absence of current privacy-preserving features if they are indeed missing. The current phrasing suggests these are future research directions rather than acknowledged gaps in the current system's safety posture.

- [writing] The 'Broader Impacts' section defers privacy controls to future work. Explicitly detail current technical safeguards (e.g., encryption, user-accessible memory deletion) implemented in the prototype to mitigate risks of storing sensitive user profiles.
- [writing] Clarify if the 'profile bank' (Appendix) uses any real user data. If real data was used for seeding or validation, provide an IRB exemption statement or consent protocol. If entirely synthetic, state this explicitly to avoid ambiguity regarding human subjects.

paper_reviewer_scientific_evidence

Verdict: minor_revision

The scientific evidence presented for the MemSlides framework is generally sound in its experimental design but relies on small sample sizes and LLM-based evaluation metrics that require stronger robustness checks.

The primary quantitative evidence for the **tool-memory** component comes from a diagnostic matched-pair setting involving only nine pairs (Appendix Table 2). While the authors correctly apply a sign test and report p-values (e.g., $p=0.0195$ for Strict Verify), the absolute sample size ($N=9$) is extremely low for drawing broad conclusions about "reliable localized editing." The results are highly sensitive to the specific prompts chosen; for instance, one pair (P7) shows a loss in Closed-Loop Completion. The authors appropriately label this as "diagnostic," but the central claim that tool memory enhances reliability should be tempered to reflect that this evidence is limited to a narrow, controlled subset of tasks. A larger, randomized set of modification requests would be necessary to support a stronger claim of general reliability.

For the **user profile memory** component, the evidence rests on persona-alignment judgments (Table 1) with large reported effect sizes (e.g., -3.30 points on Content for GPT-5). However, the evaluation protocol relies on LLM-as-a-judge without reporting inter-rater reliability metrics (e.g., Cohen's kappa or Krippendorff's alpha) across the three blind votes mentioned in the appendix. Without these statistics, it is difficult to assess the stability of the scores against potential judge bias or variance. Additionally, the "Specificity" metric uses distractor personas, but the analysis does not explicitly rule out the possibility that the model is simply overfitting to the specific distractor set rather than learning a generalizable persona.

A significant methodological concern lies in the **profile bank construction** (Appendix A.4). The authors describe a "seeded completion" step where missing profile fields are filled using a "role-preference registry." This means the "long-term memory" being tested is partially synthetic and derived from a static registry rather than purely from accumulated user interaction history. This introduces a confound: the performance gains might reflect the quality of the seed registry or the prompt engineering used to fill it, rather than the agent's ability to learn and retrieve preferences from a dynamic history. The paper should clarify the extent to which the results depend on this synthetic completion versus genuine learning from interaction traces.

Finally, while the **general-quality metrics** (Table 2) show that MemSlides remains competitive, the diversity metric (Vendi score) is computed on a suite level rather than per-deck, which limits the granularity of the analysis regarding visual variety. The evidence supports the hypothesis that memory separation improves specific, targeted behaviors, but the sample sizes and reliance on synthetic profile data prevent a definitive conclusion about broad, real-world efficacy.

- [science] The tool-memory evaluation relies on only nine matched pairs (Appendix Table 2). While the sign test shows significance for Strict Verify and Core Tool Time Ratio, the sample size is too small to rule out random variance or specific prompt sensitivity. The authors should explicitly frame these as 'diagnostic' findings and avoid generalizing to 'reliable localized editing' without a larger, randomized test set.
- [science] The persona-alignment results (Table 1) show large effect sizes (e.g., -3.30 on Content for GPT-5), but the evaluation relies on LLM-as-a-judge without reporting

inter-rater reliability (e.g., Cohen’s kappa) or variance across the three blind votes per dimension. Without this, the robustness of the 0-10 scale scores against judge bias is unclear.

- [science] The profile bank construction involves a ‘seeded completion’ step using a registry to fill missing fields (Appendix A.4). This introduces a potential confound where the ‘memory’ effect might partly reflect the quality of the seed registry rather than the agent’s ability to learn from history. The analysis should control for or discuss the impact of this synthetic data generation on the persona-alignment scores.

paper_reviewer_statistical_analysis

Verdict: minor_revision

The statistical analysis in the paper is generally well-structured, particularly in the use of matched-pair designs for the tool-memory evaluation (Section 5.1.2). The authors correctly identify the small sample size (nine pairs) and appropriately choose a non-parametric exact sign test (Appendix Table 5) rather than a parametric t-test, which is a sound methodological choice given the data constraints. The separation of “diagnostic” settings from general distribution claims is also a prudent statistical framing.

However, several statistical reporting gaps weaken the rigor of the conclusions. First, in Appendix Table 5 (‘tables/tool_memory_paired_robustness_table.tex’), the sign test for “Closed-Loop Completion” yields a p-value of 0.3125 (3 wins, 1 loss). While the authors label this “Directional,” the text in Section 5.2 (“Tool memory improves overall reliability...”) implies a stronger effect than the statistics support. The manuscript should explicitly state that this specific metric did not reach statistical significance at the 0.05 level to prevent readers from inferring a robust effect where the data is inconclusive.

Second, the efficiency claims rely on the “Core Tool Time Ratio” (Table 4, ‘tables/tool_memory_main_table.tex’). The authors report a geometric mean of 0.327x but fail to provide any measure of dispersion (e.g., geometric standard deviation) or a confidence interval. With only nine pairs, the geometric mean is highly sensitive to outliers. Without uncertainty bounds, the claim of “reduced core tool time” lacks statistical precision.

Finally, the primary persona-alignment results (Table 1, ‘tables/profile_memory_v6_bestof_main_table.tex’) depend on LLM-as-judge scores. The paper mentions “three blind votes” per dimension but does not report inter-rater reliability metrics (such as Krippendorff’s alpha or Cohen’s kappa). Given that the reported improvements (e.g., -2.42 points on a 0-10 scale) are derived from these judgments, the absence of reliability statistics makes it difficult to assess whether the observed differences are statistically distinguishable from judge noise. The authors should calculate and report these reliability metrics in the appendix or main text to validate the stability of the evaluation protocol.

- [writing] The sign test in Appendix Table 5 (tool_memory_paired_robustness_table.tex) reports p-values for N=9 pairs. For ‘Closed-Loop Completion’ (3 wins, 1 loss), the p-value is 0.3125. The text describes this as ‘Directional’ but does not explicitly state that the

result is not statistically significant at $\alpha=0.05$. Clarify the interpretation of non-significant p-values in the main text to avoid overclaiming robustness.

- [science] Table 4 (tool_memory_main_table.tex) reports 'Core Tool Time Ratio' as a geometric mean. The text states the ratio is 0.327x. However, the table footnote and text do not specify the confidence interval or standard error for this geometric mean, which is critical for assessing the stability of the efficiency claim given the small sample size ($N=9$ pairs). Add uncertainty estimates.
- [science] The persona-alignment scores in Table 1 (profile_memory_v6_bestof_main_table.tex) are averages of three blind votes per dimension. The paper does not report the inter-rater reliability (e.g., Krippendorff's alpha or Cohen's kappa) for these LLM-as-judge evaluations. Without this metric, the reliability of the 0-10 scale differences (e.g., -2.42) cannot be statistically validated.

paper_reviewer_writing_quality

Verdict: minor_revision

The manuscript demonstrates a high level of technical sophistication, and the writing is generally clear and professional. The logical flow between the problem formulation, the proposed hierarchical memory framework, and the experimental validation is well-structured. The abstract effectively summarizes the core contributions, and the introduction successfully motivates the need for persistent personalization in slide generation.

However, there are a few minor grammatical errors and stylistic inconsistencies that should be addressed to ensure the prose is as polished as the research. Specifically, in Section 3, the sentence regarding the organization of tool memory ("organized two execution flows") is missing a preposition, which disrupts the reading flow. Additionally, the use of "carryover" as a verb in the abstract ("carryover preferences") is non-standard; it should be the two-word phrasal verb "carry over."

There are also instances of slightly informal phrasing in otherwise formal sections. For example, in Section 3, the statement "repeatedly regenerating the full deck is undesirable" could be elevated to "suboptimal" or "inefficient" to better align with the academic tone. Similarly, the subheading in Section 5 ("The gains are planning-level persona gains") suffers from repetition that could be smoothed out for better readability.

Overall, the paper is well-written, but these minor linguistic adjustments will enhance the overall quality and professionalism of the text. No major structural or clarity issues were found that would impede understanding of the scientific content.

- [writing] In Section 3 (Problem Formulation), the phrase 'Therefore, repeatedly regenerating the full deck is undesirable' is slightly informal. Consider revising to 'Consequently, repeated full-deck regeneration is suboptimal' to better match the technical tone of the surrounding equations.
- [writing] In Section 3 (Tool Memory), the sentence 'In \methodname, tool memory is

organized two execution flows' contains a missing preposition. It should read 'organized into two execution flows' or 'organized as two execution flows'.

- [writing] In the Abstract, the phrase 'carryover preferences' is used as a verb phrase ('ability to carryover preferences'). 'Carryover' is typically a noun or adjective; the verb form is 'carry over' (two words). Please correct to 'carry over preferences'.
- [writing] In Section 5 (Analysis and Discussion), the subheading 'The gains are planning-level persona gains' is repetitive. Consider rephrasing to 'The gains reflect planning-level persona alignment' for better flow and clarity.