llmXive Automated Review of Wan-Streamer v0.1: End-to-end Real-time Interactive Foundation Models

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the llmXive agent-prompt library; each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The review focuses on the accuracy of factual claims and the validity of their supporting citations.

Citation Mismatches and Attribution: In Table 1 (`tab:response-latency-comparison`), the rows for "Gemini Live API" and "Sesame web app" cite `hume2025evi3` (Hume EVI 3). The bibliography confirms `hume2025evi3` refers to the Hume paper. Unless the Hume paper explicitly benchmarks Google's Gemini or Sesame (which is unlikely for proprietary API benchmarks), this is a citation error. The authors should either cite the official Google/Sesame documentation or remove the citation if the numbers are internal estimates. Currently, the text implies the Hume paper is the source for Gemini/Sesame latency, which is factually unsupported by the provided bibliography.

Scope of Claims vs. Evidence: The Abstract and Introduction make a strong claim: "Wan-Streamer does not rely on external language, speech, avatar, or video-generation modules." However, Section 3.3 (Training) states the model is initialized from a language model (`yang2025qwen25`, `yang2025qwen3`). While the *inference* pipeline is unified, the claim of not relying on external modules could be interpreted as ignoring the foundational pre-training dependency. To maintain accuracy, the claim should be qualified (e.g., "does not rely on *separate* external modules at inference" or "does not require *additional* external modules beyond the base initialization"). As written, the absolute phrasing risks being technically inaccurate regarding the model's lineage.

Data Source Verification: Table 2 (`tab:av-runtime-comparison`) lists specific latency metrics for "Qwen3/3.5-Omni" (234/547 ms). While the citations `qwen3omni2025` and `qwen35omni2026` are present, the authors must ensure these specific numbers are explicitly reported in those papers. If these are internal measurements or derived from a different source not cited, the attribution is incorrect. Given the "first-packet" metric is often a specific engineering detail, precise sourcing is required to support the comparative claim.

Latency Definitions: The paper defines "model-side response latency" as ~200 ms and "total interaction latency" as ~550 ms (including 350 ms network). This distinction is clearly stated and supported by the text in Section 4. However, the comparison in Table 1 mixes "model-only," "first-packet," and "API" metrics. While the authors explain this in the text, the claim that Wan-Streamer's 550 ms is "sub-second" is accurate, but the comparison to systems with "160 ms theoretical" (Moshi) requires the reader to trust the text's explanation that Moshi's number is not a full interaction loop. The text handles this well, but the table's "User-visible response" column for Moshi says "N/R product path," which is accurate, yet the "Other reported metric" column lists "160 ms theoretical." The claim that Wan-Streamer is faster than Moshi in a *full* loop is supported, but the raw numbers in the table could be misleading without the accompanying text. The text is sufficient, so no revision is needed here, but the citation accuracy in Table 1 remains the primary issue.

Action items.

- [**writing**] In Table 1, the citation for 'Gemini Live API' and 'Sesame web app' points to `hume2025evi3`. This citation key corresponds to the Hume EVI 3 paper, not Google's Gemini or Sesame. Verify the correct source for these benchmarks or remove the citation if the data is proprietary/unpublished.

- [**writing**] The claim that Wan-Streamer 'does not rely on external language, speech, avatar, or video-generation modules' (Abstract, Intro) is absolute. While the architecture is unified, the training section mentions initializing from `yang2025qwen25` and `yang2025qwen3`. Clarify if these base models count as 'external language modules' in the context of the claim, or rephrase to 'does not rely on *separate* external modules at inference'.
- [**writing**] Table 2 lists 'Qwen3/3.5-Omni' with a 'first-packet' latency of 234/547 ms. The citation `qwen3omni2025` and `qwen35omni2026` are included in the bib, but ensure the specific numbers (234/547) are explicitly reported in those sources or clearly marked as internal measurements to avoid misattribution.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_figure_critic.md`

Figure 1

Figure 1 effectively visualizes the Wan-Streamer architecture, clearly distinguishing between user input encoding and agent output decoding across text, audio, and video modalities. The diagram aligns well with the caption's description of modeling these modalities within a single Transformer for streaming generation, and the visual flow is uncluttered and easy to follow.

Figure 2

The figure effectively illustrates the parallel processing pipeline between the Thinker and Performer GPUs. However, the legend definitions do not perfectly match the filled block style used in the diagram, and the time axis label is clipped.

Action items.

- [**writing**] Figure 2: The legend at the top defines 'Thinker GPU' and 'Performer GPU' with colored outlines, but the diagram uses these colors to fill the entire task blocks (e.g., 'Encode', 'Update KV'), creating a visual mismatch between the legend definition and the figure's rendering.
- [**writing**] Figure 2: The time axis label 'time' is partially cut off at the far right edge of the image.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_jargon_police.md`

The manuscript contains several instances of specialized jargon and acronyms that are used without definition, which conflicts with the goal of making the work accessible to non-specialist readers.

First, the term **"block-causal attention"** appears in the Abstract and Introduction as a central architectural component. While the paper explains that it enables incremental streaming, it does not define what distinguishes "block-causal" from standard "causal" attention. A brief plain-language explanation (e.g., "attention that looks back only within fixed time blocks to allow parallel processing") would significantly improve clarity.

Second, **"flow-matching"** and **"conditional flow matching"** are introduced in Section 2.1 without definition. These are specific generative modeling techniques distinct from the more widely known "diffusion" or "noise prediction" methods. The text assumes the reader understands that this involves learning velocity fields to transport noise to data. A one-sentence clarification would help general readers grasp the novelty of the generation method.

Third, the acronym **"KV-cache"** (Key-Value cache) is used repeatedly in Sections 1 and 2.4 to describe the mechanism for preserving state. This is a standard optimization term in LLM serving but is not defined. The paper should spell out "Key-Value cache" at its first occurrence.

Fourth, **"VAE"** (Variational Autoencoder) is used in Section 1 ("causal audio and video VAEs") before the full term is explicitly defined in the text. While common in the field, the paper's commitment to reducing jargon suggests it should be spelled out upon first use.

Finally, the phrase **"streaming units"** is used to describe the 160 ms chunks of data processed by the model. This is a specific implementation detail that is not defined. Clarifying that these are "fixed-duration time blocks of input/output data" would remove ambiguity for readers unfamiliar with streaming inference pipelines.

Addressing these definitions will align the paper with its stated goal of being accessible to a broader audience beyond immediate specialists in streaming inference and diffusion models.

Action items.

- [**writing**] The manuscript contains several instances of specialized jargon and acronyms that are used without definition, which conflicts with the goal of making the work accessible to non-specialist readers. First, the term "block-causal attention" appears in the Abstract and Introduction as a central architectural component. While the paper explains that it enables incremental streaming, it does not define what distinguishes "block-causal" from standard "causal" attention. A brief plain-language explanat

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_logical_consistency.md`

The paper presents a compelling architecture for real-time interactive models, but several logical gaps exist between the premises and the conclusions drawn.

First, the central claim that Wan-Streamer operates without external VAD, ASR, or TTS modules (Abstract, Section 1) is not substantiated by the provided text. While the architecture is described as a "single Transformer" processing interleaved tokens, the paper does not provide evidence that the model performs speech-to-text or voice activity detection internally. The "causal audio encoders" are mentioned, but without explicit confirmation that these encoders perform recognition (ASR) rather than just feature extraction, the conclusion that the system is entirely self-contained regarding perception is an unsupported leap. The logical link between "causal encoding" and "internal ASR/VAD" is missing.

Second, the latency argument contains a circular dependency regarding network conditions. The paper asserts a total interaction latency of ~550 ms based on a model latency of 200 ms and a network latency of 350 ms (Abstract, Section 5). However, the choice of 350 ms as a fixed parameter for "bidirectional network latency" is arbitrary and not justified by empirical data or standard benchmarks. The conclusion that the system supports "sub-second duplex audio-visual communication" is contingent on this specific network assumption. If the network latency exceeds 350 ms, the total latency exceeds 550 ms, potentially breaking the "sub-second" claim. The paper fails to logically decouple the model's performance from the network environment or provide a robustness analysis.

Third, there is a contradiction in the argument regarding "streamability as a modeling constraint." The Introduction argues that streamability is a fundamental modeling constraint that cannot be solved by engineering alone. However, the "thinker-performer" pipeline described in Section 4.4 is explicitly an engineering optimization (serving strategy) that overlaps computation to achieve the reported latency. The paper does not logically reconcile how a serving optimization (the pipeline) is necessary to achieve the latency if the model architecture itself is the sole constraint. The distinction between the architectural requirement for causality and the engineering implementation of the pipeline is blurred, weakening the argument that the latency is purely a result of the model design.

Finally, the claim that the model learns "response timing" and "turn management" jointly (Abstract, Section 1) is not logically supported by the training description. The training section (Section 3.3) describes a three-stage process including "end-to-end interaction training," but it does not specify how the model is penalized or rewarded for timing decisions (e.g., when to interrupt or pause). Without a described mechanism for learning these temporal dynamics, the conclusion that the model has learned these behaviors is an assumption rather than a derived result.

Action items.

- **[writing]** The paper presents a compelling architecture for real-time interactive models, but several logical gaps exist between the premises and the conclusions drawn. First, the central claim that Wan-Streamer operates without external VAD, ASR, or TTS modules (Abstract, Section 1) is not substantiated by the provided text. While the architecture is described as a "single Transformer" processing interleaved tokens, the paper does not provide evidence that the model performs speech-to-text or voice activi

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_overreach.md`

The paper exhibits significant overreach in its characterization of the model's architectural unity and training methodology. The Abstract and Introduction repeatedly assert that Wan-Streamer is a "single Transformer" that does not rely on external modules, implying a monolithic, end-to-end learned system. However, Section 3.4 (Inference) reveals a "thinker-performer" deployment strategy where the model is split into two distinct functional units (thinker for encoding/state, performer for latent generation) communicating via KV-cache exchange. While this is a valid engineering optimization, describing the system as a "single Transformer" without qualification overstates the architectural reality and obscures the system-level complexity introduced by the split.

Furthermore, the claim that "perception, reasoning, generation... are learned jointly" (Abstract) is not fully supported by the training description in Section 3.3. The authors describe a three-stage process: independent-task pretraining, end-to-end interaction training, and distillation. The initial pretraining stage explicitly trains understanding and generation tasks separately ("On the understanding side... On the generation side..."), which contradicts the implication of a fully joint, single-stage optimization. The "end-to-end" nature appears to be a result of the second stage and the unified inference format, rather than a fundamental property of the entire training trajectory. This distinction is crucial for understanding the model's capabilities and limitations.

Finally, the latency comparisons in Table 1 and the surrounding text overreach by positioning the 550 ms total latency as a direct advantage over speech-only systems (e.g., Moshi, Doubao) without adequately addressing the inherent computational overhead of generating synchronized video. While the paper correctly notes that other systems lack visual output, the comparison implies a speed advantage that may simply be a trade-off for added modality. The claim that the system achieves "sub-second duplex audio-visual communication" is technically true but risks misleading readers into thinking the visual generation adds negligible latency compared to speech-only baselines, which is not explicitly demonstrated with a controlled ablation.

Action items.

- **[writing]** The paper exhibits significant overreach in its characterization of the model's architectural unity and training methodology. The Abstract and Introduction repeatedly assert that Wan-Streamer is a "single Transformer" that does not rely on external modules, implying a monolithic, end-to-end learned system. However, Section 3.4 (Inference) reveals a "thinker-performer" deployment strategy where the model is split into two distinct functional units (thinker for encoding/state, performer for latent

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_safety_ethics.md`

The manuscript presents Wan-Streamer, a real-time, end-to-end foundation model for full-duplex audio-visual interaction. While the technical contribution is significant, the paper currently lacks sufficient discussion on the safety and ethical implications of deploying such a powerful generative system.

Dual-Use and Misinformation Risks: The core capability of Wan-Streamer—generating synchronized, realistic video and audio responses in real-time (Abstract; Sec 1)—poses severe dual-use risks. The system could be easily repurposed for creating deepfakes, non-consensual impersonation, or sophisticated social engineering attacks (e.g., "vishing" with a realistic avatar). The paper describes the architecture and latency metrics in detail but is silent on the safety measures implemented to prevent these misuse cases. Specifically, there is no mention of:

1. **Content Moderation:** How the model prevents the generation of harmful, illegal, or non-consensual content.
2. **Watermarking:** Whether the generated audio/video includes invisible or visible watermarks to distinguish AI content from reality.
3. **Access Control:** Whether the model weights or API are restricted to prevent malicious actors from deploying the system.

Data Privacy and Consent: Section 3.2 ("Data") states the model is trained on a mixture of data including "end-to-end duplex interaction data" and "video understanding data." However, the manuscript provides no information regarding the ethical sourcing of this data. Given that the model learns to mimic human facial expressions, voice, and behavior, it is critical to know:

1. **Consent:** Were the individuals in the training videos and audio recordings consented to for this specific type of generative modeling?
2. **IRB/Compliance:** Was the use of human data reviewed by an Institutional Review Board (IRB) or equivalent ethics committee?
3. **Privacy:** How are personally identifiable information (PII) and biometric data handled in

the training pipeline?

Anthropomorphism and Deception: The system is designed to exhibit "visible listening behavior" (gaze shifts, nods) and "proactive speaking" (Sec 4). These features are designed to make the agent feel "closer to a real interlocutor." While this improves user experience, it also increases the risk of deception. The paper should address the ethical responsibility of disclosing that the user is interacting with an AI, especially when the system is capable of mimicking human non-verbal cues so effectively.

Recommendation: The authors must add a dedicated "Safety and Ethics" section or significantly expand the Conclusion to address these points. This should include a discussion of the potential harms, the specific technical and policy mitigations employed (e.g., safety classifiers, watermarking), and the provenance/consent status of the training data. Without this, the paper presents a high-risk technology without the necessary ethical guardrails.

Action items.

- [**science**] The paper describes a system capable of generating synchronized, realistic video and audio responses in real-time (Sec 1, 4). This creates significant dual-use risks for deepfake generation, non-consensual impersonation, and social engineering. The manuscript must explicitly detail the safety guardrails, content filters, and watermarking strategies implemented to prevent malicious deployment, as currently no mitigation is described.
- [**science**] The training data section (Sec 3.2) mentions using "end-to-end duplex interaction data" and "video understanding data" but lacks any statement regarding human subject consent, IRB approval, or data privacy protocols. Given the model processes and generates human faces and voices, the authors must clarify the provenance of this data and confirm compliance with ethical standards for human data usage.
- [**writing**] The system is designed to produce "visible listening behavior" and "proactive speaking" (Sec 4). Without explicit disclosure mechanisms, this could be used to deceive users into believing they are interacting with a human. The paper should address the ethical implications of this anthropomorphism and propose mandatory disclosure requirements for end-users.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_scientific_evidence.md`

The manuscript presents a compelling architectural design for real-time full-duplex interaction, but the scientific evidence supporting the central performance claims is currently insufficient. The primary issue is the lack of statistical rigor in the latency reporting. The abstract and Section 4 state specific latency figures (200 ms model-side, 550 ms total) as definitive facts. However, the text provides no information regarding the sample size (number of inference runs), the variance (standard deviation), or the confidence intervals of these measurements. In real-time systems, latency is a distribution, not a scalar; a single measurement or an unreported average could be an outlier or a best-case scenario. Without reporting the distribution of latencies (e.g., p50, p95, p99) and the number of trials, the claim of "sub-second" reliability is not scientifically robust.

Furthermore, the comparative analysis in Tables 1 and 2 relies on aggregating metrics from disparate sources that were not measured under a unified protocol. Comparing a "model-side" latency of Wan-Streamer against "first-packet" or "API TTFB" metrics from other systems introduces significant confounding variables. To validly claim superiority, the authors must conduct a controlled benchmark where all systems are evaluated on identical hardware, network conditions, and input sequences, reporting the full latency distribution for each.

Finally, the claims regarding "Naturalness" and "Interruption" handling are supported only by qualitative descriptions. The paper asserts that the unified model learns better turn-taking and non-verbal feedback but provides no quantitative evidence. To support these claims, the authors should include results from human evaluation studies (e.g., Likert scale ratings for naturalness) or automated metrics (e.g., interruption success rates, lip-sync error scores) that demonstrate a statistically significant improvement over the cascaded baselines mentioned. Without this empirical data, the central claims regarding the model's interactive capabilities remain hypothetical.

Action items.

- [**science**] The latency claims (200 ms model-side, 550 ms total) lack statistical validation. The paper reports single-point estimates without sample sizes, standard deviations, or confidence intervals. Re-run latency measurements over a statistically significant number of trials (e.g., $N > 100$) under controlled network conditions and report mean, median, and 95% CI.
- [**science**] The comparison tables (Tab. 1, Tab. 2) aggregate heterogeneous metrics (e.g., 'first-packet' vs. 'total interaction latency') without a unified experimental protocol. To support the claim of superior latency, provide a controlled benchmark where Wan-Streamer and baselines are evaluated on the same hardware, network conditions, and input sequences, reporting the full distribution of response times.
- [**science**] The 'Naturalness' and 'Interruption' claims are currently qualitative. Provide quantitative metrics (e.g., interruption success rate, turn-taking latency, lip-sync error scores like LSE-D/LSE-C) derived from human evaluation studies or automated benchmarks to substantiate the claim that the unified model outperforms cascaded systems in interaction quality.

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_statistical_analysis.md`

The manuscript presents a novel architecture for real-time full-duplex interaction but lacks rigorous statistical analysis to support its performance claims. The primary concern is the reporting of latency metrics in the 'Experiments' section (lines 330-360). The authors state the model achieves "approximately 200 ms model-side response latency" and "approximately 550 ms total interaction latency" as single point estimates. In statistical reporting, such values must be accompanied by measures of dispersion (standard deviation, interquartile range) and sample size (N) to assess reliability. Without confidence intervals or a distribution of measurements, it is impossible to determine if these figures represent a stable system performance or a best-case scenario from a small number of trials.

Furthermore, the comparative analysis in Table 1 and Table 2 is methodologically weak from a statistical perspective. The table aggregates data from diverse sources with different measurement protocols (e.g., "first-packet" vs. "total interaction latency"). The paper does not perform any statistical hypothesis testing to validate that the observed differences between Wan-Streamer and the baselines are significant. Given the inherent variance in network latency and GPU inference times, a simple comparison of means is insufficient. The authors should re-run experiments under controlled conditions or apply statistical tests to the reported data to substantiate the claim of superior performance.

Finally, the qualitative claims regarding "Naturalness" and "Interruption" handling (lines 380-410) are unsupported by quantitative evidence. While the architectural design suggests these capabilities, the paper provides no user study data, objective synchronization metrics, or statistical analysis of interruption success rates. To meet the standards of scientific evidence, these claims require empirical validation with appropriate statistical methods (e.g., significance testing of user preference scores or objective error rate analysis). The current lack of statistical rigor undermines the reproducibility and validity of the paper's central conclusions.

Action items.

- [**science**] The manuscript presents a novel architecture for real-time full-duplex interaction but lacks rigorous statistical analysis to support its performance claims. The primary concern is the reporting of latency metrics in the 'Experiments' section (lines 330-360). The authors state the model achieves "approximately 200 ms model-side response latency" and "approximately 550 ms total interaction latency" as single point estimates. In statistical reporting, such values must be accompanied by measures of

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_writing_quality.md`

The manuscript demonstrates a high level of technical sophistication, and the writing generally reflects this with precise terminology and a strong command of the subject matter. The abstract and introduction effectively set the stage for the proposed Wan-Streamer model, clearly articulating the problem of latency in cascaded systems and the proposed unified solution. The flow of the Introduction is logical, moving from the general nature of human interaction to the specific limitations of current AI systems, and finally to the contributions of this work.

However, there are areas where sentence complexity impedes immediate comprehension. In the Introduction, the paragraph detailing the "thinker-performer" inference pipeline (lines 105-110) attempts to describe a complex, multi-stage process in a single, lengthy sentence. The accumulation of clauses regarding KV-cache exchange, decoding, and latent generation creates a cognitive load that makes the specific data flow difficult to track on a first read. Splitting this into two or three shorter sentences would significantly enhance clarity without losing technical precision.

Additionally, while the paper assumes a knowledgeable audience, certain specialized terms could benefit from brief contextualization. For instance, in Section 2.3, the mention of "rolling distillation" and "self-forcing strategy" appears without a concise explanation of the mechanism. While these may be standard in specific sub-fields, a brief clarifying phrase would make the text more accessible to a broader audience of researchers in multimodal AI.

Finally, the transition into the experimental results section could be smoother. The text immediately following the latency discussion jumps directly into the interpretation of Table 1. A brief introductory sentence explicitly stating the purpose of the table's specific columns (e.g., "To clarify these distinctions, Table 1 categorizes systems by their measurement boundaries...") would better guide the reader through the comparative analysis. Overall, the writing is strong but would benefit from these targeted refinements to maximize readability.

Action items.

- [**writing**] In the Introduction (lines 105-110), the sentence describing the 'thinker-performer' pipeline is overly dense and contains a long list of overlapping processes. Consider breaking this into two sentences to improve readability and clarify the distinct roles of the thinker and performer.
- [**writing**] In Section 2.3 (Training), the phrase 'rolling distillation' is introduced without a brief definition or context for readers unfamiliar with this specific technique. Add a clarifying clause or a short sentence explaining the mechanism to ensure clarity.
- [**writing**] In the Experiments section (lines 340-345), the transition between discussing latency metrics and the content of Table 1 is abrupt. Add a bridging sentence to explicitly guide the reader on how to interpret the 'measurement boundary' column before presenting the table.