

CVPR 2022 · LMU MUNICH & IWR HEIDELBERG · RUNWAY ML

High-Resolution with Latent D

Robin Rombach · Andreas Blattmann
Run diffusion in a compact latent



0:01 / 8:27

ICLR 2022

Self-Supervised Vision · arXiv:2106.08254

BEiT

BERT Pre-Training of Image Transformers
Bringing the BERT recipe from language to vision

Hangbo Bao · Li Dong · Songhao Piao · Furu Wei
Harbin Institute of Technology · Microsoft Research

Masked Image Modeling

A self-supervised task: a Vision Transformer predicts discrete visual tokens for randomly masked patches.
No labels required.

Pre-trained on ImageNet-1K

Self-supervised, no labels · then fine-tuned

83.2% / 85.2%
Top-1 ImageNet · Base / Large

BERT-style pre-training

Not just for text —

it works for images.



Scan → Paper



Code

Microsoft Research

BEiT brings the BERT recipe from language to vision.

0:01 / 3:39

ICML 2022 · AUCTION DESIGN
Automated mechanism design with deep learning

A Context-Integrated Transformer for Auction D

CTransNet — one network that
Zhijian Deng · J...

