

# llmXive Automated Review of AlayaWorld: Long-Horizon and Playable Video World Generation

*llmXive automated review panel*

---

**Please read first — this review is fully automated and not checked by any human.** It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

## How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

**This paper's panel (9 lenses):** Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

## Claim Accuracy

*Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.*

**Reviewer model:** [qwen.qwen3.5-122b](#)

**Prompt:** `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper makes several specific factual claims regarding model versions and release timelines that require verification against the provided bibliography and public records.

First, the Introduction explicitly states that “The complete technical details, experimental results, and full codebase will be released in mid-July.” Given that the paper is ingested from an arXiv preprint (implied by the URL structure and the “third-party” context), this creates a temporal contradiction. If the paper is already published on arXiv, the code and results should either be available now (requiring a text update to “are released”) or the paper is a pre-announcement (requiring a clarification that the current submission is a proposal/abstract). The current phrasing suggests the work is not yet reproducible, which conflicts with the Abstract’s claim of releasing a “full-stack open-source framework” alongside the paper.

Second, Section 3.1 identifies the base model as “LTX-2.3”. The bibliography only contains an entry for “LTX-Video” (hacohen2024ltx) from 2024. There is no public record of an “LTX-2.3” model in the cited reference or general knowledge up to the current date. This appears to be either a hallucinated version number or a specific internal build not documented in the provided reference list. The authors must verify the exact base model name and version to ensure the claim is supported by the cited source or correct the model name to the publicly available version (e.g., LTX-Video).

Third, the bibliography includes a future-dated entry “hyworld22026” for “HY-World 2.0” (year 2026). While the paper cites this as a related work, the use of a 2026 date for a paper currently under review (or published in 2025/2026 context) raises questions about the existence of this specific version at the time of writing. If “HY-World 2.0” is a hypothetical or future system, it should not be cited as an existing baseline or related work without qualification. The authors should confirm if this is a real, published work or if the citation needs to be adjusted to reflect the actual status of the HY-World series.

These issues do not necessarily invalidate the core methodology but represent “claim citation mismatch” and potential “future-dated” hallucinations that undermine the factual accuracy of the paper’s claims regarding its foundation and availability.

**Action items.**

- **[writing]** The paper makes several specific factual claims regarding model versions and release timelines that require verification against the provided bibliography and public records. First, the Introduction explicitly states that “The complete technical details, experimental results, and full codebase will be released in mid-July.” Given that the paper is ingested from an arXiv preprint (implied by the URL structure and the “third-party” context), this creates a temporal contradiction. If the paper is a

**Figure Critic**

*Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_figure\_critic.md

### Figure 1

The figure displays a grid of video generation results with overlaid control icons, but the empty caption and lack of a legend for the icons render the visual data uninterpretable.

### Figure 2

Figure 2 is a visual collage of generated images without any caption or scientific context, rendering it impossible to assess its relevance or claims within the paper.

### Figure 3

Figure 3 is a grid of images showing video generation results but lacks any descriptive caption, making it impossible to understand what is being presented or how to interpret the visual data.

### Figure 4

The figure displays a grid of generated video frames with time-step labels, but it lacks a descriptive caption to explain the context or results, and the metadata filename contradicts the figure number.

### Figure 5

Figure 5 is a grid of gameplay images that lacks any descriptive caption, labels, or context, rendering it scientifically useless for peer review. Additionally, the filename in the caption contradicts the figure number, indicating a likely error in the manuscript's figure assembly.

### Action items.

- **[fatal]** Figure 1: The caption is empty ('no caption'), providing no context for the four rows of images, the overlaid control icons, or the specific scenes depicted.
- **[science]** Figure 1: The figure lacks a clear legend or labels explaining the overlaid 'WASD' and arrow icons, making it impossible to determine if they represent user inputs, generated actions, or navigation paths.
- **[fatal]** Figure 2: The figure contains no caption text, making it impossible to determine the claims it supports, the meaning of the visual elements (e.g., the 'First Frame' label, the icons, the grid structure), or the context of the generated scenes.
- **[science]** Figure 2: The visual content appears to be a collage of fantasy-themed image generation results (e.g., pyramids with magical effects, various landscapes) but lacks any scientific data, axes, quantitative comparisons, or methodological indicators required for a scientific preprint figure.

- **[fatal]** Figure 3: The figure has no caption text; it is labeled only as ‘(no caption) [fig6.jpg]’, making it impossible to understand the figure’s purpose, what the rows/columns represent, or how to interpret the visual content.
- **[science]** Figure 3: Without a caption, the relationship between the row labels (HY-World, LingBot-Fest, Dreax-World, Ours) and the column labels (First Frame, Turn 1, etc.) is unclear, and the meaning of the overlaid control icons is undefined.
- **[writing]** Figure 4: The figure lacks a descriptive caption explaining the content, methodology, or significance of the generated video sequences shown.
- **[writing]** Figure 4: The image filename [fig7.jpg] in the caption metadata contradicts the label ‘Figure 4’, creating potential confusion regarding file organization.
- **[fatal]** Figure 5: The figure is completely devoid of a descriptive caption, axis labels, or legends, making it impossible to determine the experimental conditions, variables, or specific claims being illustrated by the image grid.
- **[science]** Figure 5: The visual content (gameplay screenshots with UI overlays) does not match the filename ‘fig2.jpg’ referenced in the caption, suggesting a rendering or file mapping error in the preprint.

## Jargon Police

*Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.*

**Reviewer model:** `qwen.qwen3.5-122b`

**Prompt:** `agents/prompts/paper_reviewer_jargon_police.md`

The paper is generally well-structured for a technical audience, but it relies on several acronyms and specialized terms that are undefined at their first occurrence, creating minor barriers for a competent reader from an adjacent field (e.g., computer vision or robotics).

Specifically, the term “DiT” (Diffusion Transformer) is used in Section 3 without expansion, despite being central to the method. Similarly, “AdaLN” (Adaptive Layer Normalization) appears in Section 3.1 without definition. While these are standard in the specific subfield of diffusion transformers, they are not universal enough to assume knowledge from a general ML or robotics PhD.

Additionally, “Plücker ray embeddings” in Section 3.1 is a specific geometric representation that would benefit from a brief parenthetical explanation for non-graphics specialists. The term “loop-closing” in Section 3.2 is used frequently but never explicitly defined as the act of revisiting a location to check for consistency, which is a core concept in SLAM but might be opaque to a pure video-generation researcher. Finally, “KV-recache” in Section 3.4 is an in-group shorthand for a specific caching mechanism that should be expanded to “Key-Value cache recalculation” or similar upon first mention.

Addressing these five points by adding brief expansions or definitions at first use will signifi-

cantly improve the paper’s accessibility without diluting its technical precision.

#### Action items.

- **[writing]** Section 3.1 (Interaction): The term ‘AdaLN’ is used in ‘AdaLN-style camera-control module’ and ‘AdaLN-style modulation’ without definition. While common in DiT literature, it is not expanded (Adaptive Layer Normalization) for an adjacent-field reader. Add ‘(Adaptive Layer Normalization, AdaLN)’ at first use.
- **[writing]** Section 3.1 (Interaction): The acronym ‘DiT’ appears in ‘autoregressive DiT’ (Section 3) and ‘DiT features’ (Section 3.1) without being spelled out as ‘Diffusion Transformer’. Define at first occurrence in Section 3.
- **[writing]** Section 3.1 (Interaction): The term ‘Plücker ray embeddings’ is used without a brief gloss. An adjacent-field reader may not know this refers to a specific 6D representation of lines in 3D space. Add a short parenthetical explanation, e.g., ‘Plücker ray embeddings (a 6D line representation)’.
- **[writing]** Section 3.2 (Consistency): The term ‘loop-closing’ is used repeatedly (e.g., ‘loop-closing trajectories’) without definition. While standard in SLAM/robotics, it is not explicitly defined here as ‘returning to a previously visited location to verify consistency’. Add a brief definition at first use.
- **[writing]** Section 3.4 (Runtime): The acronym ‘KV-recache’ is introduced in the context of ‘LongLive’ without expansion. It likely refers to ‘Key-Value cache recalculation’. Define this acronym at first use to ensure clarity for readers unfamiliar with specific inference optimization jargon.

## Logical Consistency

*Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_logical\_consistency.md

The paper’s argument structure is generally sound, with clear definitions of the four challenges (control, consistency, stability, runtime) and corresponding methodological solutions. The logical flow from problem definition to proposed mechanism is coherent. However, there is a notable inconsistency in the temporal status of the work between the Introduction and the Experiments section.

In Section 01 (Introduction), the authors state: “The complete technical details, experimental results, and full codebase will be released in mid-July.” This phrasing positions the paper as a pre-announcement or a roadmap for future work, implying that the results are not yet available or finalized.

Conversely, Section 04 (Qualitative Results) presents specific figures (Figures 3, 6, 7) and makes definitive claims about the system’s performance, such as “AlayaWorld faithfully follows the

requested viewpoint changes” and “AlayaWorld maintains visual quality... without pronounced accumulation of artifacts.” The text treats these results as established facts supporting the current paper’s thesis.

This creates a logical dissonance: the Introduction asserts the results are future-tense (“will be released”), while the body of the paper relies on those results as present-tense evidence. While this may be a stylistic choice to manage expectations regarding code availability, it weakens the internal consistency of the argument. If the results are sufficient to support the claims in Section 04, the Introduction’s claim that they “will be released” is misleading or requires qualification.

To resolve this, the authors should align the tense and status claims. Either the Introduction should acknowledge that *some* results are presented here (e.g., “We present preliminary results..”) or the phrasing in Section 01 should be adjusted to reflect that the *full* suite of details/code is coming later, while the current paper already contains the core experimental validation. As it stands, the reader is asked to accept the results in Section 04 as proof of the system’s capabilities, while the Introduction suggests those results are not yet public/final.

#### Action items.

- **[writing]** Section 01 states results ‘will be released in mid-July,’ but Section 04 presents specific figures and definitive performance claims. This creates a logical tension between a ‘future release’ status and ‘present results.’ Clarify if results are preliminary or if the release date refers only to code, ensuring tense consistency.

## Overreach

*Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_overreach.md

The paper presents an ambitious framework for interactive world generation, but the rhetoric in the Abstract and Introduction significantly outpaces the evidence provided in the Experiments section. The primary issue is the conflation of qualitative demonstrations with general capabilities.

The Abstract asserts that AlayaWorld “enables open-ended real-time interaction” and allows users to “freely navigate.” However, Section 4 (Qualitative Results) only presents static figures (Figures 3, 4, 6, 7) showing specific, pre-chosen camera paths and prompt switches. There are no quantitative metrics regarding latency (e.g., frames per second, inference time) to substantiate the “real-time” claim, nor are there any stress tests demonstrating the system’s robustness to truly “open-ended” or adversarial user inputs. The claim of “open-ended” capability is currently supported only by cherry-picked examples, not by a demonstration of the system’s ability to handle arbitrary, unbounded interaction sequences.

Furthermore, the Abstract claims the framework captures “diverse visual appearances and physical dynamics.” While Figure 2 demonstrates impressive style transfer (Minecraft, ink painting, etc.), this is a change in visual texture, not a validation of learned “physical dynamics.” The paper does not present any experiments testing the simulation of physics (e.g., object

permanence, collision response, or fluid dynamics) beyond the visual consistency of the generated frames. The claim of capturing “physical dynamics” implies a level of world modeling that the provided visual style-transfer examples do not license.

Finally, the Abstract states the work establishes a “practical foundation” and unifies “complete development.” While the architecture is described in Section 3, the paper offers no evidence of the system’s “extensibility” or “modularity” in practice (e.g., how easily a new module could be added, or how the system performs when components are swapped). The claim of a “practical foundation” is a strong assertion of utility that requires evidence of the system’s adaptability, which is absent.

To resolve these overreach issues, the authors should:

1. Replace absolute claims like “enables open-ended real-time interaction” with more precise descriptions of the demonstrated capabilities (e.g., “demonstrates real-time interaction in controlled scenarios”).
2. Add quantitative latency metrics to support the “real-time” claim.
3. Qualify the claim about “physical dynamics” to reflect that the current evidence is limited to visual consistency and style transfer, unless specific physics benchmarks are added.
4. Remove or soften the claim of establishing a “practical foundation” until the extensibility and modularity of the framework are empirically demonstrated.

#### Action items.

- **[writing]** Abstract claims ‘open-ended real-time interaction’ but Section 4 only shows qualitative figures without latency metrics or stress tests. Replace ‘enables open-ended’ with ‘demonstrates interaction in selected scenarios’ or add FPS benchmarks.
- **[writing]** Abstract claims the framework ‘unifies complete development’ and is a ‘practical foundation,’ yet Section 4 lacks any evaluation of extensibility or modularity. Narrow to ‘presents a modular architecture’ and remove ‘practical foundation’ until adaptability is shown.
- **[writing]** Abstract claims capture of ‘physical dynamics,’ but Section 4 (Fig 2) only shows visual style transfer. No physics benchmarks (collisions, fluid) are presented. Qualify to ‘visual appearances’ or add specific physical consistency tests.

## Safety Ethics

*Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_safety\_ethics.md

The paper describes a generative video world model (AlayaWorld) capable of long-horizon, interactive, and prompt-driven video synthesis. From a safety and ethics perspective, the work



falls into the category of dual-use generative AI, which inherently carries risks of misuse (e.g., generating deceptive media, simulating harmful scenarios, or creating content for disinformation campaigns). However, the paper does not present a specific, non-trivial risk that is unmitigated or unacknowledged in a way that warrants a revision.

The methodology relies on fine-tuning an existing open-source video model (LTX-2.3) and adding architectural modules for camera control and memory. It does not introduce novel capabilities for biological/chemical synthesis, cyber-attack automation, or targeted surveillance that would require specific operational safeguards beyond standard AI safety norms. The “playable” nature of the world is framed within the context of gaming and simulation, and the paper does not claim to generate photorealistic, indistinguishable-from-reality deepfakes of specific real-world individuals or events, nor does it provide a pipeline for doing so.

The paper mentions training on “gameplay recordings and real-world videos” (Abstract) but does not detail the specific datasets used for the real-world component. While this is a data provenance issue, it is primarily a matter of reproducibility and licensing compliance (which other reviewers may address) rather than an immediate safety hazard like the release of PII or unconsented human subject data. The paper does not appear to release a dataset containing PII or re-identifiable information.

There is no evidence of human-subjects research requiring IRB approval, as the evaluation appears to be based on model generation quality and consistency metrics rather than human behavioral studies or surveys. The paper does not disclose a conflict of interest that would bias the safety assessment, though the authors are affiliated with “Alaya Lab” and “Shanda,” which is standard for industry-academia collaborations.

Given that the paper does not describe a specific, actionable vulnerability, a release of sensitive data, or a system explicitly designed for deception/surveillance without safeguards, and considering that the general risks of generative video models are well-known in the field and do not require a specific “misuse” paragraph for every incremental improvement, the paper passes the safety lens. The standard disclaimer regarding the dual-use nature of generative models is implicitly covered by the field’s norms, and no specific, paper-unique risk gap was identified.

## Scientific Evidence

*Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_scientific\_evidence.md

The paper’s central claims regarding AlayaWorld’s capabilities in long-horizon generation, consistency, and interactive control are currently unsupported by the evidence presented in Section 4. The entire experimental section relies exclusively on qualitative figures (Figs 3, 5, 6, 7) and descriptive prose, lacking any quantitative metrics, statistical variance, or rigorous baseline comparisons.

Specifically, the claim of “strong stability under purely forward exploration” (Section 4.4) is



illustrated by a single figure of a one-minute rollout. Without a quantitative metric for drift (e.g., Fréchet Video Distance over time, object identity preservation scores) or a comparison against a standard autoregressive baseline, it is impossible to determine if the observed stability is a genuine effect of the proposed error bank and training strategy or simply a result of the specific scene or random seed chosen. Similarly, the “Consistency” claim (Section 4.3) relies on visual inspection of loop-closing trajectories. The design fails to rule out that the consistency is due to the static nature of the specific test scenes rather than the proposed spatial memory mechanism, as no quantitative loop-closure error is reported.

Furthermore, the “Open-ended Action” claim (Section 4.2) is backed only by a few hand-picked examples of prompt switching. There is no systematic evaluation of the semantic latency (time from command to visual change) or the visual continuity across switches, leaving the claim that the system supports “responsive control” unverified. The absence of any numerical results, seed counts, or confidence intervals means the reported effects could plausibly arise from luck or cherry-picked examples. To support these claims, the authors must provide quantitative metrics for drift, consistency, and latency, along with comparisons to strong baselines (e.g., Yume 1.5, Matrix-Game) across multiple random seeds.

#### Action items.

- **[science]** The paper’s central claims regarding AlayaWorld’s capabilities in long-horizon generation, consistency, and interactive control are currently unsupported by the evidence presented in Section 4. The entire experimental section relies exclusively on qualitative figures (Figs 3, 5, 6, 7) and descriptive prose, lacking any quantitative metrics, statistical variance, or rigorous baseline comparisons. Specifically, the claim of “strong stability under purely forward exploration” (Section 4.4) is illus

## Statistical Analysis

*Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_statistical\_analysis.md

The paper currently presents a purely qualitative evaluation of AlayaWorld’s capabilities in camera control, consistency, and long-horizon generation (Section 4). While the figures (Figs 3, 6, 7) illustrate the system’s behavior, the text makes strong comparative claims (“faithfully follows,” “strong stability,” “visually plausible”) without any accompanying statistical evidence.

Specifically, Section 4 states that “all baselines are evaluated under the same input conditioning” but provides no numerical metrics (e.g., Fréchet Video Distance, LPIPS, or user study scores) to quantify the performance gap. In the absence of quantitative metrics, the claims of superiority are anecdotal. Furthermore, even if the authors intend to rely on qualitative demonstrations, the standard practice in this field for any comparative claim involves reporting results over multiple random seeds to demonstrate robustness. The current text implies a deterministic or single-run result (“AlayaWorld maintains...”), which is statistically insufficient

for a generative model where output variance is high.

To address this, the authors should either:

1. Include a quantitative results table with metrics (mean  $\pm$  standard deviation) computed over at least 3 independent seeds for both AlayaWorld and the baselines, allowing for a statistical assessment of the improvements.
2. If the paper is intended as a systems/demo paper with no quantitative benchmarks, the language in Section 4 must be softened to explicitly frame the results as “qualitative demonstrations” rather than evidence of “strong stability” or “faithful” performance, and the claim of “fairness” in evaluation should be qualified to reflect the lack of statistical aggregation.

Currently, the lack of uncertainty reporting (no SD, SE, or CI) and the absence of any hypothesis testing or metric aggregation renders the comparative claims statistically unsupported.

#### Action items.

- **[writing]** Section 4 reports qualitative results (Figs 3, 6, 7) with claims of ‘faithful’ control and ‘strong stability’ but provides no quantitative metrics (e.g., FVD, PSNR, LPIPS) or statistical summaries (mean  $\pm$  SD over seeds) to support these assertions. Add a results table with standard deviation across 3 seeds for key metrics or explicitly state that results are purely qualitative demonstrations without statistical validation.
- **[writing]** The paper claims AlayaWorld is ‘fine-tuned from LTX-2.3’ (Sec 4) and compares against baselines ‘under the same input conditioning’ (Sec 4), but reports no variance across random seeds or training runs for either the proposed method or baselines. Without reporting standard deviation or confidence intervals, the reported performance differences cannot be distinguished from random noise. Report mean  $\pm$  SD over at least 3 independent runs for all quantitative comparisons.

## Writing Quality

*Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.*

**Reviewer model:** qwen.qwen3.5-122b

**Prompt:** agents/prompts/paper\_reviewer\_writing\_quality.md

The paper presents a clear and ambitious vision, but the prose suffers from several structural and grammatical issues that impede smooth reading. The most significant friction points occur in the Introduction and Method sections, where complex lists and sentence fragments force the reader to pause and reconstruct meaning.

In the Introduction (Section 1), the fourth paragraph lists four challenges but fails to complete the grammatical structure for the first two. The sentence “Whether navigation is endless and whether action is arbitrary, unbounded by preset physical laws” lacks a main verb, leaving the reader to guess the intended predicate. This pattern of incomplete thought disrupts the logical

flow of the problem statement. Furthermore, the fifth paragraph attempts to summarize the entire technical contribution in a single, overloaded sentence listing six distinct modules. This “list-dump” style obscures the specific mapping between the proposed components and the four challenges just defined. A restructuring that separates the component list from their functional roles would significantly improve clarity.

In the Method section (Section 3), minor but distracting errors appear. In the “Prompt-driven Action” subsection, the sentence “Navigation is a type of basic interaction while AlayaWorld also support...” contains a subject-verb agreement error (“support” should be “supports”) and a run-on structure that could be split for better pacing. Additionally, the phrase “to achieve a real playable world” contains a typo (“achive”) and awkward phrasing (“real playable world” vs. “truly playable world”).

Finally, the definition of “Consistency” in Section 3.2 suffers from tautology: “Consistency requires spatial and temporal consistency...” This redundancy weakens the definition. The sentence should be rephrased to define the property directly (e.g., “Consistency requires that the generated world maintains...”).

Addressing these specific grammatical errors and restructuring the dense lists in the Introduction will allow the reader to move through the argument without unnecessary cognitive load. The scientific content is strong, but the delivery requires polishing to match the quality of the research.

#### **Action items.**

- **[writing]** The paper presents a clear and ambitious vision, but the prose suffers from several structural and grammatical issues that impede smooth reading. The most significant friction points occur in the Introduction and Method sections, where complex lists and sentence fragments force the reader to pause and reconstruct meaning. In the Introduction (Section 1), the fourth paragraph lists four challenges but fails to complete the grammatical structure for the first two. The sentence “Whether navigation