

llmXive Automated Review of UniClawBench: A Universal Benchmark for Proactive Agents on Real-World Tasks

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer’s exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer’s section also names the underlying model that produced it.

This paper’s panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper presents a well-structured benchmark for proactive agents, and the factual claims made are consistently supported by the provided evidence, including internal tables, figures, and citations.

Specifically:

1. **Benchmark Composition:** The claim of “400 bilingual real-world tasks” (Abstract, Section 1, Section 3) is consistent with the stated design of 5 capabilities with 40 English and 40 Chinese tasks each ($5 * 40 * 2 = 400$).
2. **Experimental Results:** The performance metrics (Pass Rate and Average Score) reported in the text (Section 4.2, 4.3) align precisely with the data presented in Tables 1, 2, 3, 4, 5, and 6. For instance, the claim that “absolute success rates remain strictly below 50%” is supported by the highest overall Pass Rate of 0.475 (Claude Opus-4.8 in Table 5) and 0.480 (Claude Opus-4.6 in Table 3).
3. **Citation Accuracy:** The citations used to support claims about existing benchmarks (e.g., WebArena, OSWorld) and agent frameworks (OpenClaw, Nanobot, EDICT) point to relevant works. While some cited works (e.g., “GPT-5.4”, “Gemini-3.1-pro”) are dated 2026, this is consistent with the paper’s own timeline (NeurIPS 2026 submission) and the bibliography entries provided. There is no evidence of hallucinated references within the context of the paper’s internal logic.
4. **Methodology Claims:** The description of the “three-role closed-loop evaluation” and the “information firewall” mechanism is detailed in Section 3.2 and Appendix C, and the case study in Appendix D provides concrete evidence that the system functions as described (e.g., the user simulator receiving only a high-level status signal).

No unsupported claims, citation mismatches, or numerical inconsistencies were found that would undermine the paper’s factual accuracy. The claims are appropriately qualified and backed by the presented data.

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_figure_critic.md`

Figure 1

The figure effectively illustrates the three-role closed-loop evaluation strategy with clear visual separation of components. However, the caption contains a typo at the end, and there are minor legibility issues with the legend and a stray label at the bottom.

Figure 2

Figure 2 effectively visualizes the diversity of UniClawBench across domains, inputs, and outputs. However, the caption lacks a precise definition of what the ‘within-category percentage’ represents (e.g., row-wise vs. column-wise normalization), which is critical for interpreting the heatmaps correctly. The colorbar is clear but could be more explicitly labeled.

Figure 3

Figure 3 presents token usage and performance data clearly in terms of data trends, but the x-axis labels in subplots (a) and (b) are rotated too steeply, causing overlap and making model names and task dimensions difficult to read.

Figure 4

Figure 4 provides a clear overview of the benchmark and evaluation strategy, but the results visualization (radar chart) is flawed by a legend-to-line mismatch and a missing axis scale definition.

Action items.

- **[writing]** Figure 1: The caption ends abruptly with ‘returned to the executor f’, which appears to be a typo or incomplete sentence.
- **[writing]** Figure 1: The legend at the bottom left is labeled ‘Legend’ but the text is extremely small and difficult to read.
- **[writing]** Figure 1: The text ‘Metric 1’ at the bottom right is unexplained and appears to be a stray label not connected to the diagram.
- **[science]** Figure 2: The heatmaps display ‘within-category percentage’ values (e.g., 45% for Media/Entertainment in Multimodal), but the column headers (Multimodal, Long Context, etc.) represent the benchmark’s core capabilities. The caption fails to clarify if these percentages represent the proportion of tasks within a domain that use a specific capability, or the proportion of a capability’s tasks that fall into a domain. Without this definition, the ‘diversity’ claim is ambiguous and the data interpret
- **[writing]** Figure 2: The colorbar on the right is labeled ‘Share (%)’, but the individual cells contain specific percentage values (e.g., ‘0%’, ‘15%’). The term ‘Share’ is slightly imprecise for ‘Percentage of tasks’ or ‘Distribution’, though acceptable. However, the colorbar scale

(0-100) matches the cell values, which is good, but the lack of a clear title for the colorbar (e.g., ‘Percentage of Tasks’) makes the visualization slightly less self-explanatory.

- **[science]** Figure 3: The x-axis labels in subplots (a) and (b) are rotated at a steep angle, causing significant overlap and illegibility for model names (e.g., ‘Claude Sonnet 4 6’, ‘Gemini 3.1 Pro’) and task dimensions.
- **[science]** Figure 3: Subplot (a) displays ‘Avg Input Tokens (M)’ and ‘Avg Output Tokens (K)’ on dual y-axes, but the x-axis labels (model names) are crowded and partially illegible due to rotation.
- **[writing]** Figure 3: The caption states the data is from the ‘OpenClaw system,’ but the figure itself lacks a title or label explicitly identifying the system name, relying solely on the caption for context.
- **[writing]** Figure 4: The radar chart legend lists nine model/framework combinations, but the chart contains ten distinct colored lines, making it impossible to map one line to a specific entry.
- **[science]** Figure 4: The radar chart lacks a unit or scale definition for the ‘Normalized Pass Rate’ axis (0.0–1.0), leaving the magnitude of performance differences ambiguous.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The paper is generally well-written and accessible to a competent reader from an adjacent field (e.g., a researcher in NLP or systems). Most acronyms (Docker, API, JSON, YAML, PDF, SVG, IIIF, CLI, GUI, PR, AS) are either standard across the discipline or defined at first use. The benchmark’s core concepts (executor, supervisor, user simulator) are clearly introduced and consistently used.

However, there are a few instances where terminology could be slightly more precise or explicitly defined to avoid any potential confusion for a reader not deeply embedded in the specific subfield of proactive agent frameworks:

1. **“Codex agents” (Section 4.1):** The phrase “we use separate Codex agents based on GPT-5.4” is slightly ambiguous. “Codex” is a specific model name (OpenAI Codex), but “Codex agents” could be interpreted as a specific type of agent architecture or framework. Clarifying that these are agents *powered by* or *based on* the Codex model (or GPT-5.4, if that’s the intended meaning, as Codex is often considered a precursor to GPT-3.5/GPT-4) would remove ambiguity. Given the context mentions GPT-5.4, it’s likely a slight conflation or a specific internal naming convention that needs a brief gloss.

2. **“Advanced Packaging Tool (apt) skill” (Section 4.1):** While “apt” is universally known in Linux contexts, the phrasing “self-written Advanced Packaging Tool (apt) skill” might momentarily confuse a non-Linux specialist into thinking “apt” is a novel tool created by the authors. A slight rephrase like “a skill wrapping the standard Linux Advanced Packaging Tool (apt)” would make it immediately clear.
3. **“Coordination friction” (Section 4.3):** This term is used descriptively to explain performance issues in the EDICT framework. While the explanation is clear, it’s not a standard, widely recognized technical term in the broader agent literature. Adding a brief parenthetical explanation or citing a source where this term is used (if applicable) would strengthen the clarity for an adjacent-field reader.

These are minor points, and the overall readability is high. The definitions provided for the benchmark’s unique components (UniClawBench, the three-role strategy, the capability taxonomy) are excellent and make the paper self-contained for its intended audience.

Action items.

- **[writing]** Section 4.1 (Environmental Setup) uses ‘Codex agents’ without defining the term. While ‘Codex’ is a known model name, the phrase ‘Codex agents’ implies a specific system or role not previously defined. Clarify if this refers to agents powered by the Codex model or a specific agent framework named Codex, e.g., ‘agents based on the Codex model’.
- **[writing]** Section 4.1 (Environmental Setup) introduces ‘Advanced Packaging Tool (apt)’ as a skill name. While ‘apt’ is standard Linux terminology, the phrasing ‘self-written Advanced Packaging Tool (apt) skill’ is slightly ambiguous. Ensure the reader understands ‘apt’ is the standard Linux package manager being wrapped as a skill, not a novel tool invented by the authors.
- **[writing]** Section 4.3 (Cross-Framework Benchmark Results) uses the term ‘coordination friction’ without a formal definition or operational metric. While the context explains the phenomenon (stalled pipelines), a brief gloss or citation to a standard definition would help an adjacent-field reader understand if this is a specific technical term or a descriptive phrase.

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper’s argument structure is generally sound: the premises (limitations of existing benchmarks, the need for capability-driven evaluation) logically support the proposed solution (UniClawBench), and the experimental design (cross-model and cross-framework comparisons) follows from the stated goal of disentangling model capabilities from framework choices. The definitions of the five capabilities are consistent throughout the text and the appendix examples.

However, there are two minor logical inconsistencies regarding data reporting and citation alignment that break the internal consistency of the results presentation:

1. **Citation Mismatch in Results Table:** In Table 5 (`Tables/5_benchmark_model_avg.tex`), the last row lists the model as “Claude Opus-4.8” but the citation key used is `\cite{claude-opus-4.6}`. The bibliography (`main.bib`) contains a distinct entry for `claude-opus-4.8`. This creates a logical disconnect where the table claims to report results for version 4.8, but the reference points to version 4.6. While likely a typo, it undermines the precision of the reported evidence.
2. **Ambiguity in Timeout Application:** Section 4.1 explicitly defines two sets of timeout limits: standard (30/20 min) and long-context (45/30 min). Section 4.2 presents “Overall” Pass Rates and Average Scores in Table 5. The text does not explicitly clarify whether the “Overall” metrics are calculated using the standard limits for all tasks (implying long-context tasks were artificially constrained) or if the specific long-context limits were applied to the relevant subset before aggregation. If the latter, the “Overall” metric conflates two different evaluation conditions without stating it; if the former, the claim that the benchmark evaluates “long-context reasoning” fairly is weakened by the lack of explicit confirmation that the extended limits were used in the final reported numbers. Clarifying this in the text is necessary to ensure the conclusion (that models struggle with long-context tasks) follows validly from the reported metrics.

These issues are fixable via text correction and do not require re-running experiments, but they currently represent small breaks in the logical chain between the experimental setup and the reported conclusions.

Action items.

- **[writing]** In Table 5, the model ‘Claude Opus-4.8’ cites ‘claude-opus-4.6’. Correct the citation key to match the bibliography entry for version 4.8 to ensure the table label and reference align.
- **[writing]** Section 4.1 defines distinct timeouts for long-context tasks (45/30 min) vs standard (30/20 min). Section 4.2 reports ‘Overall’ metrics without clarifying if these aggregate results use the extended limits for the long-context subset. Explicitly state whether the reported Overall metrics reflect the specific timeout rules per task type to ensure the conclusion follows from the setup.

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_overreach.md`

The paper makes several claims regarding the scope and novelty of UniClawBench that exceed

the evidence provided in the experiments.

First, the title and abstract describe the work as a “Universal Benchmark” for “Real-World Tasks.” The term “Universal” suggests applicability across all languages, domains, and model architectures. However, the benchmark is explicitly limited to 400 tasks in only two languages (English and Chinese) as detailed in Section 3 and the task examples in the appendix. Furthermore, the evaluation is restricted to a specific set of 10 models and 3 frameworks. While the benchmark is ambitious, the language “Universal” is an overreach that implies a generality not yet demonstrated. The claim should be narrowed to reflect the specific scope (e.g., “Comprehensive Capability-Driven Benchmark”) or the limitations regarding language and domain coverage should be explicitly highlighted in the abstract.

Second, the abstract claims UniClawBench is “the first capability-driven benchmark designed to evaluate proactive agents in dynamic, real-world settings.” While the capability-driven taxonomy is a key contribution, the Related Work section (Section 2) acknowledges prior benchmarks like ClawMark and LiveClawBench that also operate in dynamic, real-world environments with evolving backends. The novelty lies specifically in the *three-role closed-loop evaluation strategy* and the *capability-oriented taxonomy*, not necessarily in being the first to evaluate agents in dynamic settings generally. The claim of being the “first” should be qualified to reflect the specific methodological innovation (the evaluation framework) rather than the general setting.

Finally, the conclusion asserts that the benchmark “successfully pinpoints the root causes of agent failures.” While the taxonomy is designed to isolate capabilities, the results in Section 4.2 show extremely low pass rates for certain categories (e.g., Multimodal at ~7.5%). The paper demonstrates *where* failures occur but does not provide sufficient evidence to definitively distinguish between a failure due to a specific capability deficit versus a failure due to the extreme difficulty of the task or the limitations of the specific frameworks tested in those domains. Without further ablation studies isolating these variables, the claim of identifying “root causes” is stronger than the data supports. The language should be softened to “provides diagnostic evidence for failure modes” or “helps identify potential root causes.”

Action items.

- **[writing]** Title/Abstract claim ‘Universal Benchmark’ but Section 3 limits scope to 400 tasks in only 2 languages (English/Chinese) and 3 frameworks. Replace ‘Universal’ with ‘Comprehensive’ or ‘Capability-Driven’ and qualify ‘Real-World’ to reflect the specific tested domains.
- **[writing]** Abstract claims ‘first capability-driven benchmark... in dynamic, real-world settings’ (Abstract). Section 2 cites ClawMark and LiveClawBench which also use dynamic real-world settings. Qualify the claim to ‘first with a three-role closed-loop evaluation strategy’ to avoid overclaiming novelty of the setting itself.
- **[writing]** Conclusion states the benchmark ‘successfully pinpoints the root causes of agent failures’ (Section 5). Results show very low pass rates in some categories but do not ablate framework vs. model vs. task difficulty to prove root causes. Soften to ‘provides diagnostic evidence for failure modes’ or ‘helps identify potential root causes’.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

This paper presents a benchmark for evaluating proactive AI agents in real-world environments. The methodology involves running agents in isolated Docker containers with injected tasks, using a “hidden supervisor” to grade performance against private rubrics, and a “user simulator” to provide feedback.

From a safety and ethics perspective, the work is low-risk. The benchmark tasks (e.g., license audits, trip planning, code refactoring) are benign and do not involve generating harmful content, exploiting vulnerabilities, or targeting real individuals. The “real-world” aspect is carefully contained within Docker containers with specific, non-sensitive fixtures (e.g., local files, mock APIs, or public web pages like The Met or Library of Congress), avoiding the risk of agents interacting with live, sensitive production systems or private user data.

The paper explicitly addresses the risk of evaluation leakage via a “three-role” architecture with an information firewall, ensuring the grading criteria remain hidden from the agent and the user simulator. This is a robust design choice that mitigates the risk of the benchmark itself being used to train agents to game evaluations.

There are no human subjects involved; the “user simulator” is an automated agent, and the tasks do not require collecting or processing personal data from real people. The datasets and code are intended for public release, but the task descriptions and examples provided in the appendix do not contain PII or sensitive information.

No dual-use capabilities are introduced that would lower the barrier to harm (e.g., automated vulnerability discovery or disinformation generation). The research focuses on evaluation methodology, not on creating new offensive capabilities. Consequently, there are no missing disclosures or unmitigated risks requiring action.

Scientific Evidence

Weighs the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a comprehensive benchmark for proactive agents, but the experimental design supporting the central claims regarding model and framework rankings lacks sufficient statistical rigor to rule out noise or confounding variables.

First, the primary results in Tables 1, 2, 5, and 6 report single-point estimates (Pass Rate and Average Score) for 400 tasks without any measure of variance. The benchmark splits tasks into five categories with only 40 tasks each (20 English, 20 Chinese). In agent evaluation, performance

is notoriously volatile across random seeds and task ordering. A difference of 1-2% in Pass Rate (e.g., GPT-5.4 at 0.407 vs. Claude Opus at 0.388 in Table 6) is well within the margin of error for a sample size of 40 per category. Without reporting standard deviations across multiple seeds or bootstrap confidence intervals over the test set, the paper cannot substantiate claims that one model or framework is definitively “better” than another; the observed rankings could easily be artifacts of a lucky seed or specific task subset.

Second, the cross-framework comparison (Section 4.3) attributes performance differences to “framework architecture” but fails to control for hyperparameter tuning asymmetry. The text notes that Nanobot is “ultra-lightweight” and EDICT uses “multi-agent orchestration,” yet it does not disclose whether these frameworks were given equal opportunities for prompt engineering, timeout tuning, or skill configuration compared to OpenClaw. If OpenClaw received more extensive tuning or a more favorable timeout configuration (e.g., the 30-minute global timeout might be insufficient for the more verbose EDICT), the performance gap reflects resource allocation rather than architectural superiority. A fair comparison requires a controlled ablation where all frameworks are subjected to the same hyperparameter search budget and resource constraints.

Finally, the reliability study (Section 4.1) validates the supervisor against human raters on 50 trajectories but does not address potential framework bias in the evaluation rubric. If the “hidden rubrics” were developed or refined using OpenClaw traces, the supervisor might inadvertently favor the specific artifact formats or logging styles of OpenClaw over the multi-agent logs of EDICT. To ensure the benchmark is truly “universal,” the authors should demonstrate that the supervisor’s scoring consistency holds across a held-out set of trajectories generated by all three frameworks, ensuring the evaluation metric itself is not confounded by the framework being tested.

Action items.

- **[writing]** The paper presents a comprehensive benchmark for proactive agents, but the experimental design supporting the central claims regarding model and framework rankings lacks sufficient statistical rigor to rule out noise or confounding variables. First, the primary results in Tables 1, 2, 5, and 6 report single-point estimates (Pass Rate and Average Score) for 400 tasks without any measure of variance. The benchmark splits tasks into five categories with only 40 tasks each (20 English, 20 Chinese).

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The statistical treatment of the benchmark results is generally transparent regarding the metrics used (Pass Rate, Average Score) and the evaluation protocol. However, two specific reporting issues regarding uncertainty and precision need correction to ensure the numbers mean what

the paper claims.

First, in Section 4.2 (“Benchmark Evaluation Reliability Study”), the authors validate their automatic supervisor against human experts using a sample of 50 trajectories. They report a point estimate of 92.0% agreement and correlation coefficients ($r = 0.71$, $\rho = 0.68$). While these numbers suggest strong reliability, the sample size ($N = 50$) is relatively small for establishing high-precision estimates. Without confidence intervals (CIs) or standard errors, a reader cannot determine if the 92.0% agreement is stable or if the correlation could plausibly be lower in a different random sample. For instance, the 95% CI for a proportion of 0.92 with $N = 50$ is approximately [0.81, 0.97]. Reporting the CIs for the agreement rate and the correlations (or at least the standard errors) is necessary to substantiate the claim of “strong correlation” and “high agreement” quantitatively.

Second, Tables 1 through 6 report performance metrics (Pass Rate and Average Score) to three decimal places (e.g., 0.438, 0.834). The benchmark consists of 40 tasks per capability category. For a proportion $p \approx 0.5$ with $n = 40$, the standard error is $\sqrt{p(1-p)/n} \approx 0.079$. This implies that the uncertainty in the estimate is in the second decimal place. Reporting three decimal places (e.g., distinguishing 0.438 from 0.442) implies a precision that the sample size cannot support, constituting false precision. The authors should round all reported metrics to two decimal places to align with the statistical resolution of the data.

These are reporting fixes that can be applied immediately to the manuscript without re-running experiments, provided the raw per-task data is available to recalculate the rounded values or CIs.

Action items.

- **[writing]** Section 4.2 reports ‘92.0% agreement’ and correlation coefficients ($r=0.71$, $\rho=0.68$) for the reliability study based on $N=50$ trajectories. No confidence intervals, standard errors, or p-values are provided for these statistics. Given the small sample size ($N=50$), the precision of these estimates is unknown. Report 95% CIs for the agreement proportion and the correlation coefficients, or state the standard error, to allow readers to assess the stability of the reliability claim.
- **[writing]** Tables 1-6 report Pass Rates and Average Scores to three decimal places (e.g., 0.438, 0.834) for 40 tasks per category. With $N=40$, the standard error for a proportion near 0.5 is ~ 0.08 , making the third decimal place statistically meaningless (false precision). Round all reported metrics to two decimal places (e.g., 0.44, 0.83) to match the resolution supported by the sample size.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript is generally well-structured and the narrative flow is strong, guiding the reader clearly from the problem statement to the proposed solution and results. The abstract effectively summarizes the contributions, and the section transitions are logical. However, there are several recurring grammatical errors, specifically regarding subject-verb agreement and pluralization, that interrupt the reading experience and require correction.

In Section 1, the opening paragraph contains a clear grammatical slip: “assist user by directly control” should be “assist users by directly controlling.” This error appears early and sets a slightly unpolished tone. Later in Section 1, the sentence structure regarding the limitations of existing benchmarks is slightly fragmented; merging the two sentences about scenario-based taxonomies would improve the rhythm.

Section 3 is mostly clean, but the list of Docker container features lacks a serial comma, which is a minor but noticeable punctuation oversight in a technical list.

Section 4 contains the highest density of grammatical issues. In the “Environmental Setup” subsection, “each frameworks” should be “each framework.” In the “Cross-Framework Benchmark Results” subsection, the sentence “EDICT consume a huge amount of token” contains two errors: the verb should be “consumes” and “token” should be plural “tokens.” Similarly, the phrase “resulting in a notably lower overall pass rates” mixes singular and plural forms incorrectly; it should be either “a notably lower overall pass rate” or “notably lower overall pass rates.”

Finally, there is a minor inconsistency in spelling conventions. The paper generally adheres to American English (e.g., “behavior” in Appendix C), but the caption of Figure 4 uses the British spelling “categorised.” While not a critical failure, standardizing to one dialect (likely American English for NeurIPS) would improve professional polish.

These issues are purely mechanical and do not obscure the scientific meaning, but correcting them will significantly enhance the readability and professional quality of the paper.

Action items.

- **[writing]** Section 1, Paragraph 1: ‘assist user by directly control’ contains a subject-verb agreement error and missing article. Rewrite as: ‘assist users by directly controlling’.
- **[writing]** Section 1, Paragraph 2: The sentence ‘Finally, and most critically, existing benchmarks organize tasks by application scenario... This approach conflates fundamentally different abilities’ splits a single logical point into two sentences. Merge them for better flow: ‘Finally, and most critically, existing benchmarks organize tasks by application scenario (e.g., office, research), an approach that conflates fundamentally different abilities.’
- **[writing]** Section 3, Paragraph 1: The phrase ‘All tasks run inside Docker containers equipped with real software, live browsers and local file systems’ lacks a serial comma before ‘and’, creating a slight ambiguity in the list. Add a comma: ‘...live browsers, and local file systems’.
- **[writing]** Section 4, Paragraph 1: ‘installing and preparing each frameworks’ has a number agreement error. Change to ‘installing and preparing each framework’ or ‘installing and preparing the frameworks’.
- **[writing]** Section 4, Paragraph 3: ‘EDICT consume a huge amount of token’ has subject-verb agreement and number errors. Rewrite as: ‘EDICT consumes a huge amount of tokens’.

- **[writing]** Section 4, Paragraph 3: ‘resulting in a notably lower overall pass rates’ has a number agreement error. Change to ‘resulting in notably lower overall pass rates’ or ‘a notably lower overall pass rate’.
- **[writing]** Section 4, Paragraph 4: ‘categorised by model’ uses British spelling (‘categorised’) while the rest of the paper uses American spelling (‘categorized’ is implied by ‘realized’ elsewhere, though ‘categorized’ is not explicitly used, ‘behavior’ is used in Appendix C). For consistency with ‘behavior’ in Appendix C and standard NeurIPS style, ensure consistent spelling. If the paper aims for US English, change ‘categorised’ to ‘categorized’.