

llmXive Automated Review of Vidu S1: A Real-Time Interactive Video Generation Model

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: [qwen.qwen3.5-122b](#)

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper presents a compelling system for real-time interactive video generation, but several factual claims regarding hardware capabilities and baseline comparisons require precise alignment with the reported evidence.

First, the abstract and Section 3.2 repeatedly claim the model runs at “up to 42 FPS on regular consumer GPUs.” However, the experimental setup in Section 3.1 and the results in Table 1 explicitly specify that these measurements were conducted on “RTX 5090 GPUs.” The RTX 5090 is a high-end, likely unreleased or professional-grade card (depending on the exact 2026 context), and certainly not representative of a “regular consumer GPU” (e.g., RTX 3060/4070). This overstatement misleads readers about the accessibility of the real-time performance. The claim should be qualified to “high-end GPUs” or the specific hardware should be named to ensure accuracy.

Second, there is a potential citation drift or hallucination regarding baselines. The text in Section 1 and the context of Table 1 imply comparisons against leading commercial systems. While the bibliography includes entries for HeyGen and Kling Avatar 2.0, the text or table captions (if they referenced specific model versions like “GPT-5.5” or “Gemini-3.1” as suggested by the future-dated nature of the paper’s timeline) must be verified. If the authors claim to compare against models that do not exist in the public record or are not cited in the bibliography, this constitutes an unsupported comparative claim. The bibliography provided does not list GPT-5.5 or Gemini-3.1; if these were used, they must be added; if not, the comparison claim must be removed or corrected to reflect the actual baselines used.

Finally, the claim of a “100% preference rate against HeyGen and LemonSlice” in Section 3.1 is a strong absolute statement. While the paper mentions this refers to “subject controllability,” the quantitative table shows HeyGen achieving a CSIM of 0.9191, nearly identical to Vidu S1’s 0.9192. To maintain rigor, the authors should ensure the “100%” figure is clearly scoped to the specific “subject controllability” metric of the human evaluation and not conflated with overall quality, or provide the specific breakdown of the A/B test results to substantiate the absolute win rate.

These issues are primarily matters of precision and scope rather than fundamental scientific flaws, but they are critical for the credibility of the performance claims.

Action items.

- **[writing]** The paper presents a compelling system for real-time interactive video generation, but several factual claims regarding hardware capabilities and baseline comparisons require precise alignment with the reported evidence. First, the abstract and Section 3.2 repeatedly claim the model runs at “up to 42 FPS on regular consumer GPUs.” However, the experimental setup in Section 3.1 and the results in Table 1 explicitly specify that these measurements were conducted on “RTX 5090 GPUs.” The RTX 5090 is

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_figure_critic.md

Figure 1

The figure effectively illustrates the model’s capabilities through clear visual examples, but the caption is grammatically broken with a missing subject, and the timeline axis lacks clear definition regarding its scale.

Figure 2

The Sankey diagram effectively visualizes the data filtering pipeline and flow, but the caption’s mention of ‘caption generation’ is not visually represented in the diagram, which only shows the final embedding block.

Figure 3

The figure is clear and readable, but the caption contains a grammatical error omitting the model name, and the ‘Subject Controllability’ label is awkwardly split across lines.

Figure 4

The figure effectively demonstrates the qualitative differences between the two models, but the caption is poorly formatted with raw color code artifacts and missing model names, reducing its clarity and professionalism.

Action items.

- **[writing]** Figure 1: The caption contains a grammatical error and missing subject (‘Overview of . supports...’), likely due to a placeholder variable not being replaced with the model name.
- **[writing]** Figure 1: The timeline labels (e.g., ‘0:00’, ‘15:00’) are ambiguous; it is unclear if they represent absolute timestamps or relative intervals, and the jump from 30:00 to 60:00 suggests a non-linear scale that is not explained.
- **[writing]** Figure 2: The caption states the pipeline shows ‘caption generation’, but the diagram only shows a ‘Caption + Embedding’ block without illustrating the generation process or inputs.
- **[writing]** Figure 2: The ‘Duration Filter’ is shown as a separate branch in the diagram, but the caption lists ‘single-shot clipping’ as a single step, creating a slight disconnect between

the text description and the visual breakdown.

- **[writing]** Figure 3: The caption reads ‘Human preference evaluation of versus HeyGen...’, missing the subject name (likely ‘Vidu S1’) before ‘versus’.
- **[writing]** Figure 3: The row label ‘Subject Controllability’ is split across two lines, causing the text ‘Subject’ to be visually disconnected from ‘Controllability’.
- **[writing]** Figure 4: The caption contains a placeholder artifact ‘[rgb]0.55,0.0,0.0’ and ‘[rgb]0.0,0.45,0.2’ instead of natural language descriptions for the red and green boxes, and the model name is missing (e.g., ‘highlight the results of [Model Name]’).
- **[writing]** Figure 4: The caption text ‘Qualitative comparison of instruction following and visual consistency’ is incomplete or generic; it should explicitly name the compared models (Kling Avatar 2.0 vs. Vidu S1) as shown in the figure labels.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The paper is generally well-structured for a technical audience, but it relies on several subfield-specific acronyms and notation conventions that are not explicitly defined for a reader from an adjacent field (e.g., a systems researcher or a general computer vision PhD).

First, in Section 2 under “Notation,” the symbol $_j$ appears in Equation 3 and the surrounding text to denote the noise level of the historical prefix. While $\mathbf{t}_j$ is clearly defined as the diffusion timestep for the current state, $_j$ is only implicitly contrasted as a different timestep. A reader might struggle to determine if $_j$ is a fixed constant, a random variable, or a specific function of $\mathbf{t}_j$. A one-sentence definition clarifying its role as the “diffusion timestep for the historical prefix” would resolve this.

Second, the “TwinCache” strategy in Section 2.2.1 introduces the terms “noisy cache” and “clean cache.” While the mechanism is described, these terms are used as proper nouns for specific data structures without a formal definition. Explicitly stating that the “noisy cache” refers to “latent states extracted at an intermediate denoising step” and the “clean cache” refers to “the final denoised states” would make the terminology self-contained.

Finally, in Section 3.1, the evaluation metrics CSIM, Sync-D, and DOVER are introduced with citations but without operational definitions. While these are standard in the specific niche of audio-driven avatar generation, they are not universal across computer vision or machine learning. A competent adjacent-field reader would need to look up the citations to understand that CSIM measures identity similarity, Sync-D measures lip-sync distance, and DOVER measures video quality. Adding brief parenthetical explanations (e.g., “CSIM (identity similarity)”) at their first mention would significantly improve accessibility without sacrificing precision.

Action items.

- **[writing]** The paper is generally well-structured for a technical audience, but it relies on several subfield-specific acronyms and notation conventions that are not explicitly defined for a reader from an adjacent field (e.g., a systems researcher or a general computer vision PhD). First, in Section 2 under “Notation,” the symbol σ_j appears in Equation 3 and the surrounding text to denote the noise level of the historical prefix. While t_j is clearly defined as the diffusion timestep for the current stat

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper’s logical structure is generally sound, with the methodology (sliding window, distillation) providing a clear mechanism for the claimed capabilities (real-time, long-horizon). However, there are minor inconsistencies in the scope of claims versus the specific experimental conditions described.

First, the Introduction and Abstract repeatedly assert the model supports “infinite-length” generation without drift. While Section 2 explains the mechanism (sliding window with a fixed context size), the term “infinite” is used as an absolute property of the model. Logically, a model with a fixed sliding window cannot attend to the *entire* infinite history, only a finite subset. The conclusion that it supports “infinite” generation follows only if “infinite” is interpreted as “unbounded duration via sliding window,” which is a standard but slightly imprecise shorthand. To ensure strict logical consistency, the text should explicitly qualify “infinite” as “arbitrarily long duration via sliding window inference” to distinguish it from a model with unbounded attention memory, preventing a potential contradiction with the finite context mechanism described in Section 2.

Second, the Abstract and Conclusion claim the system achieves 42 FPS on “regular consumer GPUs.” Section 3 specifies this result was measured on “RTX 5090 GPUs.” While the RTX 5090 is a consumer card, it is a high-end, next-generation device. The logical jump from “RTX 5090” to “regular consumer GPUs” (which might imply mid-range or older cards) is a slight scope inflation. The claim holds if “regular” is interpreted broadly, but it would be more logically rigorous to specify “high-end consumer GPUs” or explicitly state the hardware class to ensure the performance claim is not misinterpreted as applying to lower-tier hardware.

These are primarily issues of precise phrasing and scope definition rather than fundamental breaks in the argument. The causal links between the proposed methods (TwinCache, DMD) and the results (stability, speed) are well-supported within the text.

Action items.

- **[writing]** Section 1 claims ‘infinite-length’ generation, but Section 2 uses a ‘fixed-length’ sliding window. Clarify that ‘infinite’ refers to unbounded duration via sliding window, not unbounded attention context, to avoid logical contradiction.

- **[writing]** Abstract claims 42 FPS on ‘regular consumer GPUs,’ but Section 3 specifies ‘RTX 5090.’ Ensure ‘regular consumer’ aligns with the specific high-end hardware tested to prevent scope inflation.

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper contains several instances where the rhetoric in the Abstract and Conclusion exceeds the scope of the evidence presented in the Experiments and Method sections.

First, the Abstract and Conclusion assert that the model supports “infinite-length real-time video generation without blurring, drift, or visual distortion.” However, the Introduction (Section 1) explicitly states that “small errors can accumulate over time, causing drift,” and the Method section (Section 2.2) describes the approach as a way to “mitigate error accumulation.” The evidence demonstrates a reduction in these artifacts, not their total elimination. The absolute phrasing “without drift” is contradicted by the paper’s own admission of the underlying challenge. This should be rephrased to reflect the mitigated nature of the problem (e.g., “significantly reducing drift”).

Second, the Abstract claims the system outputs videos “on regular consumer GPUs,” while the Experiments section (Section 3.1) specifies that the 42 FPS throughput was measured on “RTX 5090 GPUs.” The RTX 5090 is a high-end, enthusiast-class GPU, which does not align with the term “regular consumer GPU” (typically implying mid-range cards). This creates a misleading impression of the hardware accessibility required to achieve the stated real-time performance. The claim should be qualified to specify “high-end consumer GPUs” or explicitly name the hardware class used.

Finally, the Conclusion states the model achieves “the best overall performance.” While Table 1 shows superior scores on specific metrics (CSIM, Sync-D, DOVER) compared to the listed baselines, the “Instruction Following” and “Real-Time” columns are binary indicators. Furthermore, the human preference evaluation (Figure 2) is limited to comparisons against three specific commercial systems. The phrase “best overall” implies a universal superiority across all possible metrics and baselines, which is not fully licensed by the specific, limited set of comparisons presented. The claim should be scoped to “best performance among evaluated baselines on standard metrics” or “superior human preference against selected commercial systems.”

Action items.

- **[writing]** Abstract/Conclusion claim ‘without blurring, drift, or visual distortion,’ but Section 1 admits errors accumulate and Section 2.2 only ‘mitigates’ them. Change ‘without’ to ‘mitigating’ to match the evidence of error reduction, not elimination.

- **[writing]** Abstract claims ‘regular consumer GPUs’ support 42 FPS, but Section 3.1 specifies ‘RTX 5090’ (high-end/enthusiast). This overstates accessibility. Clarify as ‘high-end consumer GPUs’ or specify the hardware class to avoid misleading readers about typical requirements.
- **[writing]** Conclusion claims ‘best overall performance,’ but Table 1 uses binary metrics for instruction following and real-time capability, and human eval is limited to three commercial systems. Scope the claim to ‘best among evaluated baselines on standard metrics’ to match the evidence.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

This paper presents a real-time interactive video generation model (Vidu S1) designed for voice-controlled digital characters. From a safety and ethics perspective, the work does not present a foreseeable, non-trivial risk of harm that is unacknowledged or unmitigated in the text.

The primary safety consideration for this class of technology is the potential for misuse in generating deceptive content (deepfakes), impersonation, or non-consensual imagery. The paper addresses the data provenance and safety filtering explicitly in Section 2.1 (“Data Preparation”). It details a multi-stage pipeline including “content safety” checks to filter out NSFW content and “subject filtering” to ensure single-subject consistency. While the paper does not include a dedicated “Ethics” or “Broader Impacts” section (which is common for systems papers focused on efficiency and architecture), the description of the safety filtering pipeline serves as the necessary disclosure regarding the mitigation of harmful data ingestion.

The paper mentions the ability for users to “upload custom images of real people” (Abstract). While this capability carries inherent risks of misuse (e.g., creating non-consensual deepfakes), the paper describes the system as a research model with an online demo. It does not provide operational details on how the system prevents the upload of non-consensual images (e.g., face-matching against a database of protected individuals), nor does it claim to have solved this problem. However, given that this is a preprint describing a novel architecture and benchmark, the absence of a comprehensive deployment-level safety policy (which would be a product requirement, not a research paper requirement) does not constitute a fatal flaw or a missing disclosure of a specific research risk. The authors have not claimed to have solved the deepfake problem, nor have they released a dataset containing PII or unconsented human data.

There is no evidence of human-subjects research requiring IRB approval; the data is described as collected from public livestreams and films/TV, processed through automated filtering. No PII is released in the paper or its artifacts. The dual-use nature of the technology is inherent to the field of generative video, but the paper does not lower the barrier to a specific harmful capability (like automated vulnerability discovery) in a way that requires a unique mitigation

discussion beyond standard responsible AI practices, which are partially addressed via the data filtering description.

Consequently, there are no specific, nameable gaps in disclosure or mitigation that require action items for this review. The paper is accepted on safety/ethics grounds.

Scientific Evidence

Weights the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The evidence presented in Section 3 and Table 1 is insufficient to support the headline claims of “best performance” and “real-time interactivity” due to missing variance reporting, unfair hardware comparisons, and a lack of ablation studies.

First, the quantitative results in Table 1 are presented as single-point estimates without any measure of uncertainty. The reported CSIM score of 0.9192 for Vidu S1 is only 0.0001 higher than HeyGen’s 0.9191. Without standard deviations or confidence intervals derived from multiple seeds or bootstrap resampling, this difference is statistically indistinguishable from noise. Similarly, the Sync-D improvement (7.847 vs 8.037) is small; without variance, it is impossible to determine if this is a robust effect or a lucky fluctuation. The claim of “best performance” is therefore unsupported by the current data design.

Second, the real-time performance claim is confounded by hardware disparities. Vidu S1 is benchmarked on an RTX 5090, while baselines like HeyGen and Kling Avatar 2.0 are listed with “–” for throughput or measured on unspecified hardware. Since inference speed is heavily dependent on GPU architecture and memory bandwidth, the 42 FPS figure cannot be attributed solely to the model’s efficiency (TurboDiffusion/TurboServe) without a fair comparison on identical hardware. The current design fails to rule out the alternative explanation that the speedup is simply due to newer hardware.

Third, the qualitative claim of a “100% preference rate” in human evaluation (Section 3.2) is suspicious without statistical context. In a 500-sample benchmark, a perfect win rate suggests either a trivial task or a biased evaluation set. The paper does not report the raw counts of wins, losses, or ties, nor does it provide a confidence interval. This omission prevents the reader from assessing whether the result is a genuine superiority or a statistical anomaly.

Finally, the paper attributes the stability of long-horizon generation to specific architectural components (TwinCache, RoPE Repositioning) but offers no ablation studies. The observed stability could equally result from the DMD distillation process, the bidirectional teacher initialization, or the specific training data filtering. Without a control run that removes these specific components while holding the rest of the pipeline constant, the causal link between the proposed mechanisms and the claimed stability is unproven.

To resolve these issues, the authors must report variance metrics (mean $\pm$ SD) across multiple seeds, re-run baseline comparisons on the same RTX 5090 hardware, provide raw counts and

confidence intervals for human evaluations, and include ablation studies isolating the proposed caching and positional embedding techniques.

Action items.

- **[science]** The evidence presented in Section 3 and Table 1 is insufficient to support the headline claims of “best performance” and “real-time interactivity” due to missing variance reporting, unfair hardware comparisons, and a lack of ablation studies. First, the quantitative results in Table 1 are presented as single-point estimates without any measure of uncertainty. The reported CSIM score of 0.9192 for Vidu S1 is only 0.0001 higher than HeyGen’s 0.9191. Without standard deviations or confidence interval

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The statistical reporting in Section 3 and Table 1 is currently insufficient to support the precision of the claims made. The primary issue is the absence of uncertainty quantification for the quantitative metrics. Table 1 presents specific point estimates (e.g., CSIM = 0.9192, Sync-D = 7.8470) to four decimal places. In generative modeling, results vary significantly across random seeds and initialization. Reporting a single number without a standard deviation (SD), standard error (SE), or confidence interval (CI) implies a false precision that the data cannot support. For instance, the difference between Vidu S1 (0.9192) and HeyGen (0.9191) is 0.0001; without variance data, it is impossible to determine if this is a real improvement or random noise.

Additionally, the human preference evaluation in Section 3.1 and Figure 2 lacks statistical rigor. The text claims the model is “consistently preferred” and cites a “100% preference rate” in specific sub-categories. While the sample size (500) is mentioned, the results are presented as raw percentages without confidence intervals or significance testing. A 100% win rate on a small subset of instructions could be a statistical fluke if the subset size is small. The authors should report the win rates with 95% confidence intervals (e.g., using the Wilson score interval) or perform a binomial test to validate the significance of the preference.

Finally, the claim of “best performance” across multiple metrics (CSIM, Sync-D, DOVER) constitutes multiple comparisons. While the paper does not explicitly claim statistical significance for every metric, the presentation of a “best” table without acknowledging the multiplicity of tests or the lack of correction is misleading. At minimum, the authors should clarify that these are point estimates and avoid language implying statistical superiority unless a formal test is provided. The fix requires re-running the experiments with multiple seeds to compute variance or, if raw data exists, aggregating it to report mean $\pm$ SD and adding confidence intervals to the human preference results.

Action items.

- **[writing]** Table 1 and Section 3.2 report point estimates for CSIM (0.9192), Sync-D (7.847), and DOVER (0.5660) without any measure of uncertainty (e.g., standard deviation, standard error, or confidence intervals) across multiple seeds or runs. Given the stochastic nature of diffusion models, report mean $\pm$ SD over at least 3 independent runs to distinguish stable performance from lucky initialization.
- **[writing]** The claim of ‘best performance’ in Table 1 is based on a single comparison against baselines without reporting p-values or effect sizes. While the field often omits formal tests, the absolute differences (e.g., CSIM 0.9192 vs 0.9191) are trivial. Report the magnitude of improvement (absolute difference) and, if feasible, a paired statistical test (e.g., bootstrap or t-test) to confirm these differences are not noise.
- **[writing]** Section 3.1 states the benchmark contains 500 samples, but the human preference results (Figure 2) lack statistical reporting. If pairwise A/B tests were conducted, report the win rate with a confidence interval (e.g., 65% [58%, 72%]) or a binomial test p-value to substantiate the ‘consistently preferred’ claim, rather than relying on visual inspection of the bar chart.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_writing_quality.md`

The paper is generally well-structured, but several sections suffer from dense, math-heavy paragraphs that disrupt the narrative flow, forcing the reader to pause and re-parse definitions.

In the **Introduction**, the “Background and Demand” paragraph introduces a hypothetical mathematical model for demand scaling (α , β , m) in the middle of a qualitative argument. This creates a “garden-path” effect where the reader must switch mental gears to parse the algebra before returning to the main point about interactive video. This theoretical detour should be moved to a footnote or a separate motivation subsection to maintain the momentum of the introduction.

In the **Method** section, the “Notation” paragraph under “Training” is a classic example of buried information. It defines five different variables across three sentences with varying levels of detail. A reader trying to understand Equation 1 must hunt for the definition of $\text{bm}\{c\}$ or τ_j within the text. Consolidating these definitions into a bulleted list or a small table would allow the reader to grasp the notation instantly, rather than parsing a block of text.

Additionally, the transition into the **Stage 3** description is slightly jarring. The text jumps from the DMD gradient equation directly to “Nevertheless, we observe that optimizing solely...” without explicitly stating the specific failure mode of DMD in *this* autoregressive context before proposing the PCM fix. A brief bridging clause explaining the specific instability (e.g., “However, in our streaming setting, this leads to...”) would smooth the logical handoff.

Finally, in **Section 3**, the “Visualization” paragraph references Table 1 after discussing Figure 2. This breaks the standard “Quantitative -> Qualitative” or “Setup -> Results” flow. The reader is asked to visualize results, then reminded of a table, then told what the table says. Reordering to present the quantitative table results immediately after the benchmark setup, followed by the qualitative figure analysis, would create a more linear and digestible argument.

These are primarily structural and flow issues that can be resolved with reordering and consolidation, without changing the scientific content.

Action items.

- **[writing]** The paper is generally well-structured, but several sections suffer from dense, math-heavy paragraphs that disrupt the narrative flow, forcing the reader to pause and re-parse definitions. In the Introduction, the “Background and Demand” paragraph introduces a hypothetical mathematical model for demand scaling (α , β , m) in the middle of a qualitative argument. This creates a “garden-path” effect where the reader must switch mental gears to parse the algebra before returning to the mai