





(a) Peak inference VRAM (GB) across precisions



(b) Decoding latency (TPOT, ms/token)