

llmXive Automated Review of Trust Region Policy Distillation

llmXive automated review panel

Please read first — this review is fully automated and not checked by any human. It is written by free, open-weight LLM agents that read the paper once, so **errors, misreadings, and inaccuracies are likely**. Nothing here has been verified by a person. Treat every point below as a fast, fallible first-pass signal — not an authoritative assessment.

How this review was produced

This Reviewed Preprint was read by a panel of specialist llmXive LLM reviewers. Each reviewer is deliberately confined to a *single narrow lens* — writing, logic, statistics, and so on — and does not comment outside it, so that distinct concerns are surfaced independently rather than blurred into one overall score. The panel runs *once* over the paper. Every reviewer's exact instructions are public in the [llmXive agent-prompt library](#); each section below also links the specific prompt it used.

This report is **advisory, automated feedback for the authors**. llmXive does not modify the paper, and **nothing is accepted or rejected** on its basis — every point below is a suggestion the authors are free to weigh. Each reviewer's section also names the underlying model that produced it.

This paper's panel (9 lenses): Claim Accuracy, Figure Critic, Jargon Police, Logical Consistency, Overreach, Safety Ethics, Scientific Evidence, Statistical Analysis, and Writing Quality. Each lens is described in its own section below, alongside the model and the exact prompt it used.

Claim Accuracy

Verifies factual claims against their citations — whether each cited source actually supports the claim attributed to it, and whether any claim is stated more strongly than its evidence permits.

Reviewer model: `qwen.qwen3.5-122b`

Prompt: `agents/prompts/paper_reviewer_claim_accuracy.md`

The paper presents a novel method (TOP-D) with a clear theoretical framework and empirical

results. However, there are specific instances where the claims are stated with a confidence level that slightly exceeds the granularity of the provided evidence or requires clarification to avoid misinterpretation.

First, the claim regarding “breaking the on-policy data-reuse barrier” in the Introduction and Conclusion is technically imprecise. The method employs internal trust region iterations (multiple epochs on the same batch), which improves *within-batch* efficiency, but the global training loop still requires fresh on-policy rollouts for every global step (Algorithm 1, Line 4). The term “breaks the barrier” might lead readers to believe the method supports true off-policy learning (reusing old data from previous global steps), which is not the case. The claim should be qualified to specify that it breaks the *single-epoch* data-reuse limitation, not the strict on-policy nature of the global update.

Second, the reproducibility claim in the “Computational Resources” section (“fully reproduced on a single standard 8-GPU node”) is a strong assertion that is not directly supported by the experimental setup described. The primary results were generated on 32 GPUs. While it is plausible that the method scales down, the paper does not provide the necessary hyperparameter adjustments (e.g., reduced global batch size, adjusted learning rate) or a specific “8-GPU” result table to substantiate that the *reported* performance (e.g., 50.42% on AIME24) is achievable on 8 GPUs. Without this, the claim remains an unsupported load-bearing statement for the paper’s “efficiency” argument.

Finally, the abstract’s claim of “dramatically enhances... final performance” is vague compared to the specific “25.84% improvement on AIME24” cited in the Introduction. While the Introduction is precise, the abstract’s generalization could be interpreted as a global average across all benchmarks, which is not explicitly calculated or reported as a single “final performance” metric in the tables. Aligning the abstract’s language with the specific benchmark results would improve accuracy.

These issues are primarily matters of precision and qualification rather than fundamental scientific flaws, but they are necessary to ensure the reader’s trust in the specific claims made.

Action items.

- **[writing]** The paper presents a novel method (TOP-D) with a clear theoretical framework and empirical results. However, there are specific instances where the claims are stated with a confidence level that slightly exceeds the granularity of the provided evidence or requires clarification to avoid misinterpretation. First, the claim regarding “breaking the on-policy data-reuse barrier” in the Introduction and Conclusion is technically imprecise. The method employs internal trust region iterations (multiple

Figure Critic

Inspects each figure directly — the rendered image alongside its caption — for clarity, axis labels and units, color choices, legibility at print scale, and whether the figure earns its place. This lens runs on a vision-capable model.

Reviewer model: qwen.qwen3.5-122b

Prompt: `agents/prompts/paper_reviewer_figure_critic.md`

Figure 1

Figure 1 is a clear and effective conceptual diagram that visually supports the caption’s claim by contrasting the unstable ‘Naive OPD’ trajectory with the stable ‘TOP-D’ trajectory within a trust region. All symbols and paths are legible, and the visual metaphor aligns well with the text.

Figure 2

Figure 2 effectively communicates the performance comparison across AIME24, AIME25, and AIME26 benchmarks. The bar charts are clearly labeled with specific accuracy values, consistent axes, and a color scheme that aligns with the caption’s description of the methods (Base, RLVR, OPD, TOP-D).

Figure 3

The figure conceptually illustrates a normalization step but fails to visually support the caption’s claim of normalizing ‘across responses’ by showing only a single response sequence. Additionally, the mathematical indices in the formula are undefined in the context of the diagram.

Figure 4

The figure is visually clear but fails to support its caption’s claim regarding the isolation of the external proximal teacher because the corresponding legend entry for $\alpha=1.0$ is missing from the plot.

Action items.

- **[science]** Figure 3: The diagram depicts a single prompt and a single response sequence, yet the caption claims ‘token-level normalization across the responses’ (plural). The visual fails to illustrate the cross-response normalization mechanism described.
- **[writing]** Figure 3: The formula uses the symbol $ildeR_k^i$, but the diagram labels the response block with $ildeR_k^i$ (or similar) without defining the indices k and i in the caption or figure, making the normalization scope ambiguous.
- **[science]** Figure 4: The caption claims to isolate the effect of the external proximal teacher ($\alpha = 1.0$), but the legend entry ‘TOP-D ($\alpha = 1.0$)’ is missing from the plot; the pink line is visible but unlabelled in the legend, making it impossible to verify the claim.
- **[writing]** Figure 4: The caption contains a typo in the mathematical notation for the alpha parameter, written as ‘(= 1.0)’ instead of ‘($\alpha = 1.0$)’.

Jargon Police

Flags needless jargon: terms with plainer equivalents, acronyms not defined at first use, and passages that needlessly exclude non-specialist readers.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_jargon_police.md

The paper is generally well-structured for a competent reader in machine learning, with most acronyms (OPD, RLVR, SFT) defined at first use. However, there are a few instances where notation or specific subfield shorthand is introduced without immediate operational definition, which could stall a reader from an adjacent field (e.g., a pure NLP researcher less familiar with specific RL notation conventions).

Specifically, in Section 2.1, the symbol ρ_k appears in Equation 2 before it is defined in the subsequent text. While the definition is provided shortly after, the standard convention for mathematical clarity is to define symbols before or simultaneously with their first appearance in an equation. Similarly, in Section 3.2, the phrase “group-based reinforcement learning settings” is used as a justification for the notation without explicitly unpacking what “group-based” entails in this context (i.e., generating G responses per prompt). A brief explanatory clause would bridge this gap. Finally, in Theorem 1 (Section 4.1), the constant C^* is presented as a “universal mathematical constant” without a brief descriptor of its origin or value range in the theorem statement itself, forcing the reader to jump to the proof to understand its nature. These are minor accessibility barriers that can be resolved with precise, one-sentence additions.

Action items.

- **[writing]** Section 2.1, Eq 2: The symbol ρ_k is introduced in the equation without a preceding definition. While the text later defines it as the probability ratio, the symbol appears first in the equation. Define ρ_k explicitly in the sentence immediately preceding Eq 2 or within the equation’s caption.
- **[writing]** Section 2.2: The term ‘undiscounted setting ($\gamma = 1$)’ is used. While standard in RL, for a reader from a pure NLP background, explicitly stating ‘where γ is the discount factor’ at first use would prevent ambiguity.
- **[writing]** Section 3.2: The term ‘group-based reinforcement learning settings’ is used to justify the superscript i . This is in-group shorthand for specific RLVR implementations (like GRPO). Add a brief clause explaining that this refers to generating multiple candidate responses per prompt to compute advantages.
- **[writing]** Section 4.1, Theorem 1: The constant C^* is introduced as a ‘universal mathematical constant’ without definition or reference to its derivation in the proof. While the proof derives it, the theorem statement should briefly note that C^* is the maximum value of the function $f(u)$ derived in the proof, or refer to the specific equation in the proof.

Logical Consistency

Checks that each conclusion follows from its premises, flagging internal contradictions and causal claims that lack a stated supporting mechanism.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_logical_consistency.md

The paper presents a coherent argument for Trust Region Policy Distillation (TOP-D), linking the construction of a proximal teacher to variance reduction and monotonic improvement. The logical flow from the problem (unbounded rewards in OPD) to the solution (interpolated teacher) and the theoretical justification (bounded variance, convergence bounds) is generally sound.

However, there are three specific instances where the logical connection between the mathematical premises and the textual conclusions is slightly overstated or imprecise:

1. **Scope of Boundedness:** In Section 3.1, the text claims the proximal teacher makes the reward “strictly lower-bounded” and “analytically prevents variance explosion.” While the derivation correctly shows the reward is bounded below by $\log(1-\alpha)$, it remains unbounded above as the probability ratio $\rho_k \rightarrow \infty$. The variance bound in Theorem 1 relies on the specific behavior of the function $h(p,q)$ which is bounded in both directions due to the expectation over the policy, but the text’s phrasing suggests the *reward signal itself* is globally bounded. This is a minor logical gap; the variance is controlled, but the reward is not strictly bounded in the upper direction. The text should be refined to specify that the *negative* tail is bounded, which is the primary source of instability in OPD.
2. **Novelty of Off-Policy Mechanism:** The Introduction claims TOP-D “safely breaks the strict on-policy data-reuse barrier.” Section 3.2 implements this via “internal trust region iterations” using a fixed behavior policy $\pi_{\theta_{\text{old}}}$ for multiple epochs. This is a standard mechanism in algorithms like PPO and TRPO, not a novel “breaking” of a barrier specific to OPD. The logical leap from “we use off-policy iterations” to “we break the barrier” is an overstatement of the contribution’s novelty. The argument would be stronger if it framed this as *applying* established off-policy trust region methods to the distillation setting, rather than claiming to break a barrier that standard RL algorithms already navigate.
3. **Ablation Framing:** In Section 4.3, the ablation study describes setting $\alpha=1.0$ as “removing the external proximal teacher.” Mathematically, $\alpha=1.0$ reduces the TOP-D objective exactly to the standard OPD objective (Equation 5). While functionally this acts as an ablation of the *modification*, the phrasing “removing the teacher” is slightly misleading because the “proximal teacher” is defined as the interpolation $\alpha\pi^* + (1-\alpha)\pi$. When $\alpha=1$, the proximal teacher is the target teacher π^* . The text should clarify that $\alpha=1.0$ corresponds to the baseline OPD regime where the proximal smoothing is disabled, rather than implying the teacher itself is removed.

These issues do not invalidate the core argument but require precise rewording to ensure the conclusions strictly follow from the mathematical premises presented.

Action items.

- **[writing]** Scope of Boundedness: In Section 3.1, the text claims the proximal teacher makes the reward “strictly lower-bounded” and “analytically prevents variance explosion.” While the derivation correctly shows the reward is bounded below by $\log(1-\alpha)$, it remains unbounded above as the probability ratio $\rho_k \rightarrow \infty$. The variance bound in Theorem 1 relies on the specific behavior of the function $h(p,q)$ which is bounded in both directions due to the expectation over the policy, but the text’s phrasing is misleading.
- **[writing]** Novelty of Off-Policy Mechanism: The Introduction claims TOP-D “safely breaks the strict on-policy data-reuse barrier.” Section 3.2 implements this via “internal trust region iterations” using a fixed behavior policy $\pi_{\theta_{\text{old}}}$ for multiple epochs. This is a standard mechanism in algorithms like PPO and TRPO, not a novel “breaking” of a barrier specific to OPD. The logical leap from “we use off-policy iterations” to “we break the barrier” is an overstatement of the contribution’s novelty. The
- **[writing]** Ablation Framing: In Section 4.3, the ablation study describes setting $\alpha=1.0$ as “removing the external proximal teacher.” Mathematically, $\alpha=1.0$ reduces the TOP-D objective exactly to the standard OPD objective (Equation 5). While functionally this acts as an ablation of the *modification*, the phrasing “removing the teacher” is slightly misleading because the “proximal teacher” is defined as the interpolation $\alpha\pi^* + (1-\alpha)\pi$. *When $\alpha=1$, the proximal teacher is the target π^**

Overreach

Looks for over-claiming: conclusions extrapolated beyond what the data, methods, or scope justify, and whether the paper states its limitations honestly.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_overreach.md

The paper makes several strong claims regarding the universality and completeness of its solution to On-Policy Distillation (OPD) instability, which exceed the scope of the provided empirical evidence.

First, the Abstract and Conclusion assert that TOP-D “resolves” and “systematically resolves” the inherent optimization instability of standard OPD. While the ablation study (Section 5.3) demonstrates that removing the proximal teacher leads to instability, this evidence is strictly confined to mathematical reasoning benchmarks (AIME, AMC, MATH-500) using Qwen3 models. The rhetoric implies a general solution for all LLM post-training scenarios, yet the paper offers no evidence on non-mathematical domains (e.g., coding, creative writing, or instruction following) where the nature of the “capacity gap” and reward signals might differ. The claim should be narrowed to “resolves instability in mathematical reasoning tasks” or supported by cross-domain experiments.

Second, the Conclusion describes TOP-D as a “robust paradigm for aligning foundation models.” This generalization is not licensed by the experiments, which test only two student scales (1.7B and 8B) and one teacher scale (30B). The “Limitations” section explicitly admits

that scaling behavior beyond 30B parameters and performance on extended training horizons remain untested. The conclusion re-opens the scope closed by the limitations section, presenting a method validated on a narrow slice of the model size spectrum as a universal paradigm. The text should be revised to reflect the specific scales and domains tested, or the limitations section must be expanded to explicitly state that the “robustness” claim is currently unverified for larger models or different task types.

Finally, the Introduction claims the method “safely breaks the strict on-policy data-reuse barrier.” While the algorithm incorporates off-policy updates (Algorithm 1), the theoretical guarantee of monotonic improvement (Theorem 3) depends on the optimization error ϵ_k being minimized. The paper does not provide empirical bounds on ϵ_k in the off-policy setting, nor does it prove that the clipping mechanism *always* prevents divergence in the off-policy regime, only that it *can* improve sample efficiency. The phrasing “safely breaks” suggests a guarantee that the current evidence (which shows improved performance but not a formal proof of safety for arbitrary off-policy data) does not fully support. Rephrasing to “enables off-policy data reuse” would be more accurate.

Action items.

- **[writing]** Abstract and Conclusion claim TOP-D ‘resolves’ and ‘systematically resolves’ the instability of OPD. Section 5 (Ablation) shows that removing the proximal teacher ($\alpha=1.0$) causes instability, but the method is only tested on mathematical reasoning (AIME/AMC). Replace ‘resolves’ with ‘mitigates’ and scope the claim to ‘mathematical reasoning tasks’ or provide evidence on non-math domains.
- **[writing]** The Conclusion states TOP-D is a ‘robust paradigm for aligning foundation models’ generally. Experiments (Section 5) are limited to Qwen3-1.7B and 8B students on math datasets. The ‘Limitations’ section admits no testing beyond 30B parameters or non-math domains. Narrow the conclusion to ‘mathematical reasoning’ or add a specific limitation acknowledging the untested scope of general alignment.
- **[writing]** The Introduction claims TOP-D ‘safely breaks the strict on-policy data-reuse barrier.’ The method uses off-policy epochs (Algorithm 1, line 12) with a clip mechanism. However, the theoretical guarantee (Theorem 3) relies on minimizing optimization error ϵ_k , which is not proven to be zero in practice. The claim of ‘safely breaking’ the barrier is slightly overstated; rephrase to ‘enables off-policy data reuse’ and acknowledge the reliance on the clipping hyperparameter for stability.

Safety Ethics

Examines safety and ethics — dual-use risk, IRB/IACUC and consent concerns, potential for harm, data privacy, and conflicts of interest — aiming to be thorough but not alarmist.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_safety_ethics.md

This paper presents a theoretical and empirical study on “Trust Region Policy Distillation”

(TOP-D), a method designed to stabilize the training of Large Language Models (LLMs) for mathematical reasoning tasks. The methodology involves standard reinforcement learning techniques (policy gradients, trust regions) applied to distillation objectives.

From a safety and ethics perspective, the work is low-risk. The research focuses on improving the stability and sample efficiency of training algorithms for mathematical reasoning (AIME, AMC benchmarks). It does not involve:

1. **Human Subjects:** The datasets used (DAPO-Math-17k) and benchmarks (AIME, AMC) are standard, public, and do not contain personal identifiable information (PII) or require human-subject consent.
2. **Dual-Use Capabilities:** The method improves the training of models for math reasoning. While LLMs can be used for various purposes, this specific algorithmic contribution (bounding gradient variance in distillation) does not inherently lower the barrier to generating harmful content, cyberattacks, or biological threats in a way that differs from standard LLM training.
3. **Deception or Surveillance:** The paper does not propose systems designed to deceive users, impersonate humans undetectably, or conduct surveillance.
4. **Data Licensing Issues:** The paper cites standard public datasets and models (Qwen series, DAPO-Math). There is no indication of scraping data in violation of Terms of Service or releasing proprietary data.

The paper includes a “Limitations” section and discusses computational resources, which is good practice, though a specific “Broader Impacts” statement is not strictly required for this type of algorithmic optimization paper in many venues, and its absence does not constitute a safety failure here. The claims are technical and do not raise immediate ethical red flags.

Verdict: Accept. No action items required.

Scientific Evidence

Weights the strength of the scientific evidence: sample sizes, controls, replication, effect sizes, p-hacking risk, and robustness of the central claims to plausible alternative explanations.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_scientific_evidence.md

The paper presents a compelling theoretical framework for Trust Region Policy Distillation (TOP-D), but the empirical evidence provided in Tables 1 and 2 is insufficient to support the headline claims of “massive” and “definitive” improvements. The primary concern is the complete absence of variance reporting. The results are presented as single-point estimates (e.g., 50.42% vs 24.58%) without any indication of standard deviation, standard error, or the number of random seeds used. In reinforcement learning and distillation tasks, performance can vary significantly based on initialization and sampling noise. A 25-point gap is large, but without error bars or multiple runs, it is impossible to rule out that this result is a statistical outlier or a “lucky seed.” The authors must report results averaged over at least 3-5 seeds with standard

deviations to establish that the effect is robust.

Furthermore, the experimental design contains a significant confound in the comparison between TOP-D and the standard OPD baseline. The hyperparameter table reveals that the OPD baseline is trained with a mini-batch size of 512 and effectively 1 epoch per batch (no off-policy reuse), whereas TOP-D utilizes a mini-batch size of 32 with 16 epochs (off-policy data reuse). This means TOP-D is optimized with 16x more gradient updates on the same data compared to the baseline. The observed performance gain could be entirely attributed to this increased optimization effort (more training steps) rather than the proposed “proximal teacher” mechanism. To isolate the contribution of the method, the authors must run a control experiment where OPD is trained with the same off-policy data reuse (multiple epochs) and matching total training steps.

Finally, the ablation study in Figure ablation curves fails to cleanly isolate the “internal trust region iterations.” The “w/o off-policy” ablation removes both the off-policy data reuse and the specific trust region clipping/objective, reverting to a standard on-policy setup. This conflates the benefit of data efficiency (reusing data) with the benefit of the trust region constraint. A proper ablation should keep the off-policy data reuse (multiple epochs) but remove the specific trust region objective (e.g., using a standard policy gradient with the proximal reward but no clipping), to determine if the “trust region” component itself is necessary or if simple data reuse is the primary driver. Without these controls, the claim that the specific TOP-D components are responsible for the gains is not fully supported.

Action items.

- **[science]** The paper presents a compelling theoretical framework for Trust Region Policy Distillation (TOP-D), but the empirical evidence provided in Tables 1 and 2 is insufficient to support the headline claims of “massive” and “definitive” improvements. The primary concern is the complete absence of variance reporting. The results are presented as single-point estimates (e.g., 50.42% vs 24.58%) without any indication of standard deviation, standard error, or the number of random seeds used. In reinforcement

Statistical Analysis

Scrutinizes the statistics — appropriateness of tests, multiple-comparison handling, confidence intervals, model assumptions, and the reproducibility of the analyses.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_statistical_analysis.md

The statistical treatment of the empirical results in this paper is insufficient to support the strong quantitative claims made. The primary issue is the complete absence of uncertainty reporting. Tables 1 and 2 present performance metrics (e.g., “50.42%”) as exact, deterministic values derived from a single training run. In the context of stochastic deep learning training, a single point estimate is statistically meaningless; it provides no information about the variance of the method or the stability of the result. The authors must report results as mean $\pm$ standard deviation (or standard error) across multiple independent random seeds (typically 3) to establish

that the observed improvements are reproducible and not due to favorable initialization or sampling luck.

Furthermore, the paper frequently uses the term “significantly” (e.g., in the abstract and Section 5.2) to describe performance gaps. However, no formal hypothesis tests (such as paired t-tests or non-parametric equivalents) are conducted, and no p-values are reported. Without a test statistic and a corresponding p-value, or at least non-overlapping confidence intervals derived from multi-seed runs, the claim of statistical significance is unfounded. The observed differences, while large in magnitude, must be validated against the inherent variance of the training process.

Finally, the experimental design involves multiple comparisons: the proposed method is tested against four baselines across six different benchmarks, resulting in 24 distinct pairwise comparisons. The paper highlights the largest improvements without applying any correction for multiple testing (e.g., Bonferroni or Benjamini-Hochberg). This inflates the family-wise error rate, making it highly probable that at least one of the reported “significant” wins is a false positive. The authors need to either apply a standard correction to their p-values or explicitly acknowledge the multiplicity issue and refrain from claiming statistical significance for the uncorrected comparisons. The ablation study suffers from the same lack of uncertainty quantification, making it impossible to assess the robustness of the component contributions.

Action items.

- **[science]** Tables 1 and 2 report single-point accuracy metrics (e.g., 50.42%) for TOP-D and baselines without any measure of uncertainty (SD, SE, or CI) or mention of the number of random seeds used. In deep learning, single-run results are unreliable. Report mean $\pm$ SD over at least 3-5 independent training seeds for all reported numbers to distinguish signal from noise.
- **[writing]** The abstract and Section 5.2 claim TOP-D is ‘significantly better’ or ‘dominates’ baselines (e.g., +25.84% on AIME24) but provide no statistical hypothesis tests (e.g., paired t-test, Wilcoxon signed-rank) or p-values to support these inferential claims. Without variance estimates or formal tests, ‘significant’ is an unsupported qualitative descriptor.
- **[science]** Section 5.3 and Table 1 compare TOP-D against 4 baselines across 6 benchmarks (24 pairwise comparisons) and highlight the best results. No correction for multiple comparisons (e.g., Bonferroni, Holm, or FDR) is applied or mentioned. With 24 tests at $\alpha=0.05$, ~ 1 false positive is expected by chance. Apply a correction or explicitly state the uncorrected nature of the comparisons.
- **[science]** The ablation study (Section 5.3) isolates components by setting $\beta=1.0$ or disabling off-policy updates, but reports learning curves and final performance without uncertainty bands or seed counts. To validate that the observed stability and performance gains are robust and not artifacts of a specific random seed, report mean performance with error bars ($\pm$ SD) across multiple seeds for the ablated variants.

Writing Quality

Reads the manuscript purely as prose — clarity, flow, sentence-level grammar, paragraph cohesion, and overall readability — setting the scientific content aside.

Reviewer model: qwen.qwen3.5-122b

Prompt: agents/prompts/paper_reviewer_writing_quality.md

The manuscript presents a technically dense argument with a generally strong logical flow, but the prose occasionally suffers from informal phrasing, redundancy, and structural ordering that impedes the reader’s momentum.

The most significant writing issue is the inconsistent register. The paper oscillates between rigorous mathematical formalism and colloquialisms that undermine its authority. For instance, the abstract opens with “Big goals are hard to achieve all at once,” which is too conversational for a NeurIPS submission. Similarly, Section 3.1 states, “Thanks to mathematics, we do not need to explicitly construct a proximal teacher,” which reads as a casual aside rather than a formal derivation. These instances force the reader to adjust their mental model of the text’s tone repeatedly.

Structurally, the “Limitations” section (Section 5) suffers from a “burying the lead” problem. It begins by restating the paper’s main contributions (“Although TOP-D establishes a highly reliable... paradigm... yielding massive improvements...”) before finally addressing the actual limitations. This redundancy wastes the reader’s attention. The section should open directly with the constraints (e.g., model scale, training steps) to maintain the critical momentum established in the previous sections.

Additionally, there are minor issues with precision and signposting. In Section 3.2, the phrase “In the naive formulation” is ambiguous without immediate context. In Section 4.1, the use of “definitively proves” is stylistically risky; scientific writing typically prefers “strongly suggests” or “demonstrates” to avoid overclaiming. Finally, the transition between the theoretical analysis and the experimental setup could be sharper; the current setup section jumps quickly into datasets without a brief sentence orienting the reader to *why* these specific benchmarks were chosen to test the theoretical claims.

Addressing these specific phrasing and structural points will significantly improve the readability and professional polish of the manuscript without altering the underlying science.

Action items.

- **[writing]** The manuscript presents a technically dense argument with a generally strong logical flow, but the prose occasionally suffers from informal phrasing, redundancy, and structural ordering that impedes the reader’s momentum. The most significant writing issue is the inconsistent register. The paper oscillates between rigorous mathematical formalism and colloquialisms that undermine its authority. For instance, the abstract opens with “Big goals are hard to achieve all at once,” which is too conversational